

1 Physical Primitives of Mind: Modeling
2 Consciousness via Symmetries and Landauer
3 Erasure

4 Sebastian Bobadilla-Suarez
 `sebastian.bobadilla.s@gmail.com`

5 August 31, 2026

6 **Abstract**

7 We present a mathematical synthesis and formal Lean 4 verifi-
8 cation demonstrating how cognitive architectures can emerge as at-
9 tractor states from established physical primitives, without introduc-
10 ing new fundamental physics. Starting from the Poincaré group of
11 spacetime symmetries and finite phase space, we construct a frame-
12 work for predictive processing. Spontaneous symmetry breaking gen-
13 erates necessary topological boundaries, and finite phase space sub-
14 jects these boundaries to Landauer’s thermodynamic erasure. Agency
15 arises when systems evolve along paths of least action, minimizing en-
16 tropy production via internal structural resonance with external per-

17 turbations. Scaling this dynamic to macroscopic levels aligns with
18 emerging paradigms in neuroscience—specifically classical electrody-
19 namic fields (ephaptic coupling)—and Kuramoto phase synchroniza-
20 tion, culminating in a topological fixed point: a unified “Self.” The
21 formalization is reported with its scope attached: the development
22 declares no axioms and every structure carrying physical content is
23 inhabited by a witness, while the two links that remain assump-
24 tions—that the field reaches the synchronized configuration, and that
25 a global section of the probability sheaf is the unity of experience—are
26 marked as such throughout rather than absorbed into the deriva-
27 tion. As a corollary, we show why rigid von Neumann architectures
28 (e.g., GPUs) are thermodynamically disadvantaged or structurally
29 constrained from reaching these specific continuous thermodynamic
30 attractors.

31 1 Introduction

32 The pursuit of a physical theory of consciousness typically involves search-
33 ing for the biological correlates of mind or treating subjective experience as
34 an abstract computational algorithm. This paper asks a more fundamental
35 question: If we assume only basic physical primitives—continuous symme-
36 tries (the Poincaré group), geometry, and thermodynamics—is the emergence
37 of a unified, perceiving entity a natural physical attractor?

38 We propose a framework where the structural properties of consciousness

39 are derived as physical attractor states. Establishing a deductive chain from
 40 basic principles, we demonstrate how continuous symmetries break to form
 41 boundaries, how those boundaries dissipate free energy as non-equilibrium
 42 steady states, and how their evolution along paths of least action results
 43 in physical resonance with their environment. This eventually scales into
 44 the macroscopic electromagnetic unity proposed by emerging field theories
 45 of mind. To avoid critiques of “math theatre,” we delineate between literal
 46 physical mechanisms (e.g., thermodynamic erasure via Landauer’s principle,
 47 boundaries from spontaneous symmetry breaking) and theoretical mappings,
 48 utilizing Lean 4 formalization as our mathematical anchor. We adopt a per-
 49 spective aligned with Russellian Monism: objective measurements of macro-
 50 scopic electromagnetic fields constitute the third-person “outside,” while the
 51 system’s internal physical deformation to reach structural resonance consti-
 52 tutes the first-person intrinsic experience.

53 **2 Formalization in Lean 4**

54 A core contribution of this work is the verification of our deductive chain
 55 using the Lean 4 interactive theorem prover. This approach builds on recent
 56 efforts to formalize consciousness and brain semantics, such as the ‘OnlyOne’
 57 system (Scherf, 2026) and sheaf-theoretic models of neural manifolds (Inoué,
 58 2026). It is necessary to explicitly distinguish which claims have been math-
 59 ematically derived as theorems within the proof assistant, and which remain

60 irreducible physical postulates. Table 1 maps the primary claims of this
61 paper to their current formalization status.

Table 1: Mapping of physical claims to Lean 4 formalization status. The development declares **no axioms**: `#print axioms` on any result above reports only `propext`, `Classical.choice` and `Quot.sound`. Earlier drafts carried five axioms, three of them inconsistent; all five became class fields, so a postulate is an obligation on each instance rather than a global assertion (Section 2.1), and two of the five have since been eliminated entirely (Derivation 5). “Theorem” still does not mean “physically established”: the postulates live in hypotheses and instance obligations, and every such class is inhabited by a witness in `Examples.lean`, so no row is vacuous.

Conceptual Claim	Lean 4 Implementation	Status
Pointwise Vacuum Minimization	<code>pointwise_vacuum_of_global_min</code> (Phase1), potential energy defined as an integral; holds a.e., and a double-well witness proves the a.e. cannot be strengthened	Theorem
Landauer’s Principle	Derived from statistical mechanics (Phase3); the axiom-free reversibility core becomes a heat bound via the field <code>StatisticalMechanics.heat_eq</code>	Theorem + instance postulate
Symmetry Breaking implies Boundaries	Derived from pointwise potential minimization (Phase1); witnessed on a double well with a $\mathbb{Z}_2$ symmetry whose minimisation hypothesis is proved, not assumed	Theorem
Structural Resonance KL Bound	Gibbs’ inequality proved; the bound is the field <code>StructuralResonance.kl_bound</code> (Phase3_KLBound)	Theorem + instance postulate
Structural Resonance via Least Action	Gradient flow of σ in coupling space (Phase8), linked to <code>entropy_production_rate</code> ; finite substrates	Theorem
Phase-locking minimizes entropy production	For a <i>reciprocal</i> coupling kernel, the phase-locked field has the least entropy production of <i>any</i> field (<code>phase_locked_minimizes_entropy_of_symm</code> , Phase8 §2a). The mean-matching condition the earlier form assumed over the comparison class is derived, not imposed: the coupling term is anti-symmetric under exchanging its two sites, so total drift is conserved across phase fields (<code>total_drift_eq_of_symm</code>). Reciprocity is necessary — a one-way coupling on two sites strictly beats its own phase-locked state (Examples §16)	Theorem
Electrodynamic Discretization	Lebesgue-Bochner edge weights (Phase2) sitting on faces of the complex and partitioning the region’s stress-energy; mesh refinement converges, witnessed on $[0, 1)$. The scalar read off the stress-energy tensor is a marked modelling choice: the partition holds for any non-negative choice, and two instances differing only in it have different couplings (Examples §6)	Theorem + instance postulate

Conceptual Claim	Lean 4 Implementation	Status
Phase synchronization to Unity	Sheaf-theoretic gluing (Phase5) over a ThermodynamicCover , whose fields are now all local: the global section is constructed by probability_glue_unique and the former class field sync_to_section_eq is a theorem about it. Witnessed by two patches built from different mass profiles, each section constructed over its own patch, whose glued state is neither of them (Examples §14). The cover's equilibrium hypothesis is no longer assumed on every instance: ThermodynamicCover.ofConvergentTrajectory (Phase5.EquilibriumBridge) discharges it from kuramoto_tendsto_global_minimum , and Examples §17.1 is a three-patch cover whose phase field is the limit of a trajectory started <i>unsynchronized</i> . The overlap-agreement hypothesis is untouched	Theorem + instance postulate
Reflexive Topology (the Self)	Banach fixed point for a contracting self-prediction map (Phase6); metric, completeness and contraction are hypotheses, all three discharged on the three-site substrate (Examples §10), where global sections are proved to be exactly the finite measures on it, the metric is the uniform distance between their densities, and the self-prediction map contracts by exactly 1/2 without being constant. The avatar's link to the field is the predicate IsRestriction-Resonance , discharged on the same witness; under it a covering family of avatars determines the global state (eq_of_avatar_eq), and without it the Self theorem is unchanged (self_of_const-Resonance)	Theorem (witnessed)
Kuramoto Lyapunov descent	Descent chained to the minimum characterization by the rotating frame (Phase4.RotatingFrame); identical frequencies	Theorem

Conceptual Claim	Lean 4 Implementation	Status
Kuramoto dynamics run	The initial value problem is well posed on all of $\mathbb{R}$: the field is globally Lipschitz and globally bounded (<code>kuramotoField_lipschitz</code> , <code>kuramotoField_norm_le</code>), and through every state at every instant passes exactly one trajectory, defined for all time (<code>kuramoto_cauchy_problem</code>). Along <i>every</i> trajectory of a zero-frequency system the dynamic potential converges (<code>dynamic_potential_tendsto</code>) and the total dissipation $\int_0^\infty \sum_i \dot{\theta}_i^2$ is finite, equal to the potential's total drop (<code>dissipation_integral_tendsto</code>). The motion stops: Barbalat's lemma is proved here, Mathlib lacking it, and applied, so $\sum_i \dot{\theta}_i^2 \rightarrow 0$ along every trajectory (<code>velocity_sq_tendsto_zero</code>). Where it stops is proved under an initial-data condition: if the phases start within a quarter turn and the potential excess starts below $a/2$, where $a > 0$ bounds the coupling below, the trajectory converges to a phase-locked configuration that is a <i>global</i> minimiser, with $r^2 = 1$ (<code>kuramoto_tendsto_global_minimum</code>), by a Polyak–Łojasiewicz estimate rather than a LaSalle principle. Unconditionally this is false — splay and twisted states are equilibria — so the condition cannot be removed. Witnessed on three sites, where no closed form exists (<code>Examples §17</code>)	Theorem (conditional)
Continuous vs. Rigid Topology	Fixed wiring support at matched resource (<code>Phase7_Rigidity</code>), deriving the suboptimality Phase 7 assumed; typed continuum/discrete separation	Theorem
Critical Coupling Threshold ($K_c = 2D$)	Self-consistency equation (<code>Phase8_SelfConsistency</code>): at or below K_c , threshold included, only $r = 0$ solves it; above K_c <i>exactly one</i> $r > 0$ does (<code>supercritical_fixed_point_existsUnique</code>), so the non-negative solutions are $\{0\}$ below and $\{0, r\}$ above; that r tends to 0 as $K \downarrow K_c$ and increases strictly with K (<code>coherent_branch_strictMono</code>); a solution is a density reproducing its own order parameter. Conditional on the von Mises stationary density, which is an input, and on no link to a trajectory	Theorem (conditional)

62 2.1 Soundness of the formalization

63 Formalization is only as trustworthy as the axioms it rests on, and an axiom
64 that is inconsistent silently makes every theorem downstream of it vacuous

65 while leaving the build green. We report a case of exactly this, found by
 66 auditing our own development, because the failure mode is easy to reproduce
 67 and we have not seen it discussed in the applied-formalization literature.

68 Earlier drafts of this work declared five axioms. Three of them were
 69 individually inconsistent: each, on its own, proves **False**. All three shared
 70 one root cause — an axiom that constrains a symbol it does not itself bind.
 71 Concretely:

- 72 • `spontaneous_symmetry_breaking_pointwise_min` took the potential-
 73 energy functional as an *implicit* argument occurring only in its hypoth-
 74 esis `PotentialEnergy phi = V v0`. Instantiating that functional with
 75 the constant map $\lambda_. V(v_0)$ discharges the hypothesis by `rfl` and yields
 76 the conclusion for an arbitrary field. With $V(v) = v^2$, $v_0 = 0$ and $\phi \equiv 1$
 77 this proves $1 = 0$.

- 78 • `landauer_heat_eq` equated the free class field `heat_dissipation` with
 79 a bath-entropy difference *for every instance* of the class. An instance
 80 with zero heat dissipation and a bath that gains one bit gives $0 = \log 2$.

- 81 • `kl_bound_axiom` asserted $\sigma \geq D_{KL}(P||Q)/\Delta t$ for *all* P and Q , while σ
 82 is a single real fixed by a class field. Since D_{KL} is unbounded above, a
 83 point mass against a uniform distribution refutes it.

84 The remaining two axioms, `phase_invariant_periodic` and `sync_to_section_eq`,
 85 had the same shape — global assertions pinning free fields of `LocalSectionSynchronization`

— and are refutable as soon as the probability presheaf admits two distinct global sections. (Both were demoted to class fields at the time; both have since been removed from the class altogether, for a separate reason given in Derivation 5.)

The correction is uniform, and we state it as a rule we now follow: *a physical postulate that mentions a class field must be a field of that class, never a standalone axiom quantified over all instances.* A postulate should be an obligation each model discharges, not a global claim about every model. Recast this way, the development declares no axioms at all. Two of the five became genuine theorems — pointwise vacuum minimization is now proved from $\int f = 0 \wedge f \geq 0 \Rightarrow f = 0$ a.e., and Gibbs’ inequality was already one — and the other three became instance obligations.

This makes non-vacuity the live question, since an uninhabitable class makes its theorems vacuous just as surely as an inconsistent axiom did. **Examples.lean** therefore constructs witnesses: a one-bit erasure system with a genuine bath that discharges **heat_eq** and reproduces $\Delta Q = k_B T \ln 2$; a perfectly resonant system discharging **kl_bound**; and a scalar field on a one-point spacetime discharging **ActionPrinciples**. Derivation 5’s gluing structures are witnessed too, which was until recently the largest open gap: **Examples.lean** builds a **ThermodynamicCover** over a three-site substrate carrying the discrete topology and its Borel σ -algebra, covered by two *overlapping* open patches whose intersection is a single shared site. The local sections are germ families of a phase-indexed finite measure $\mu_\theta =$

109 $(1 + \cos \theta)^+ \delta_{\text{mid}}$ pushed into the sheafification of the probability presheaf; the
 110 family is provably non-constant in θ (`phaseMeasure_not_const`), so the pe-
 111 riodicity obligation of the constructor `LocalSectionSynchronization.of-`
 112 `InvariantMeasure` is discharged by an argument about `cos` rather than satis-
 113 fied trivially, and `thermodynamic_equilibrium` is *derived* from `phase_locked minimizes_potential`
 114 rather than assumed. `global_section_from_thermodynamics` then fires on
 115 this instance and glues the two patch-local sections into a unique global one.
 116 Two further covers answer the follow-up question of how non-degenerate such
 117 a witness can be made. `Examples.lean` §13 has a negative half worth record-
 118 ing: the phase field can be made non-constant as a function, but not as a
 119 physical state, because `thermodynamic_equilibrium` forces every instance
 120 to a phase-locked configuration. `Examples.lean` §14 is built the other way
 121 round, from two patches with independently chosen mass profiles that meet
 122 only on the site they share, and the section the gluing returns is neither of
 123 them. That is the same kind of result as `contracting_implies_const` in
 124 Derivation 6 — a proof that the witness is as strong as the class permits,
 125 rather than a witness that could be improved by trying harder.

126 What the witnesses establish is that the classes are inhabitable, so the
 127 theorems about them are not vacuous. What they do not establish is that
 128 a cortex is such an instance. In the cover above the phase field is already
 129 constant — every patch locked at $\theta = 0$ — which is precisely what makes
 130 the equilibrium obligation dischargeable; getting a physical system into that
 131 configuration is the content of the informal dynamical argument, not of the

132 Lean development.

133 The same audit turned up a defect of a different kind, worth record-
134 ing because it is not an axiom problem and would not have been caught
135 by `#print axioms. mesh_refinement_convergence` — the claim that dis-
136 crete edge weights converge to the continuous energy as the mesh is refined
137 — sat in a “Conjectures” section as an unproved `def ...: Prop`. It was
138 not merely unproved but false. Nothing required the edge regions to cover
139 the manifold, and $\text{diam } \emptyset = 0$, so the triangulation whose edge regions are
140 all empty satisfied the mesh hypothesis while contributing nothing to the
141 sum, forcing $|0 - \int f| < \varepsilon$ for every $\varepsilon > 0$. Its binder was dead as well: it
142 quantified over a triangulation and then resolved the vertex type by instance
143 search rather than through the bound variable, so the variable was never
144 used. `Phase2_MeshConvergence.lean` replaces it with a theorem. Conver-
145 gence is now stated along a *sequence* of triangulations of a named region
146 S , each required to be regular — edge regions measurable, without self-
147 loops, disjoint across distinct edges, covering S , and anchored to the sim-
148 plicial data in the sense that a vertex lies in the closure of every region
149 of an edge at it. Closure rather than membership is forced: cells are dis-
150 joint, so a shared vertex belongs to at most one of them. The manuscript’s
151 discrete energy $\frac{1}{2} \sum_u \sum_v w(u, v) f(u)$, which samples only the first endpoint
152 and is therefore not symmetric, is shown to equal exactly the midpoint rule
153 $\sum_{u < v} w(u, v) \frac{1}{2} (f(u) + f(v))$ over unordered edges; without that identifica-
154 tion, reading the double sum as a Riemann sum would be a category error.

155 At fixed mesh the error is at most $\varepsilon \mu(S)$ with ε a modulus of continuity of
 156 the integrand, and convergence follows for uniformly continuous integrands
 157 as the mesh shrinks. The regularity conditions pull against each other, so
 158 `Examples.lean` §6 discharges them with the uniform partition of $[0, 1)$ into
 159 half-open cells and checks that the approximation genuinely moves $-\frac{1}{2}$ at
 160 $N = 1$ and $\frac{1}{4}$ at $N = 2$ for the tent function $x \mapsto |x - \frac{1}{2}|$ — rather than
 161 converging trivially. What is proved is convergence, not the $O(1/N^2)$ rate
 162 the accompanying simulation measures; that would need a midpoint error
 163 term with a second derivative, which we have not developed.

164 A third defect was one of bookkeeping across files, and it is the kind that
 165 survives any number of green builds because each half is individually cor-
 166 rect. The development carried *two* Kuramoto potentials. Lyapunov descent
 167 along the dynamics, $\dot{V} = -\sum_i \dot{\theta}_i^2$, was proved for the full potential $V(\theta) =$
 168 $-\frac{1}{2} \sum_{i,j} A_{ij} \cos(\theta_i - \theta_j) - \sum_i \omega_i \theta_i$, while the characterization of the minimum
 169 — phase-locked in both directions — and everything downstream of it, in-
 170 cluding Derivation 5’s gluing, concerned the coupling term alone. Nothing
 171 chained them, and the gap was not cosmetic: `kuramoto_potential_unbounded_below`
 172 now proves that as soon as one $\omega_i \neq 0$ the full potential is unbounded below
 173 — walk θ out along ω and the bounded cosine sum cannot compensate — so
 174 it has no minimum for phase-locking to attain, and the sentence “the phase-
 175 locked state minimizes the Lyapunov potential” was, read literally, false of
 176 the functional the descent theorem was about. `Phase4_RotatingFrame.lean`
 177 closes this by formalizing the standard reduction $\theta_i \mapsto \theta_i - \Omega t$. For identical

178 natural frequencies $\omega \equiv \Omega$, a trajectory maps to a trajectory of the zero-
 179 frequency system — which is exactly the system Derivation 5 already worked
 180 with — the two potentials coincide there, and the descent result transports
 181 onto the coupling potential whose minimizers are the phase-locked states.
 182 Phase differences, phase-locking and the magnitude of the order parameter
 183 are all proved invariant under the change of frame, so nothing observable
 184 is lost in the reduction. `Examples.lean` §7 runs the chain end to end on
 185 two oscillators with a common frequency, discharging every hypothesis. The
 186 restriction is real and we state it plainly: the reduction is exact only for
 187 identical frequencies. A genuine spread of natural frequencies leaves residual
 188 detunings in the reduced system, phase-locking then requires coupling above
 189 a critical K_c , and K_c is nowhere in our Lean development — it is studied
 190 numerically, as Table 1 records.

191 A fourth defect was dimensional, and it concerned the one row of Table 1
 192 that is still numerical rather than proved. `Phase8_ContinuousField.lean`
 193 defines the critical coupling $K_c = 2D$ and a predicate for being above it.
 194 That predicate compared $\int_x \int_y K(x, y)$ against $2D$ — but a double integral
 195 against an unnormalized measure carries the substrate’s mass twice over,
 196 so a constant kernel K on a substrate of total mass m contributed Km^2 .
 197 The comparison therefore depended on how large the substrate was rather
 198 than on how strongly it was coupled, and any kernel whatever could be
 199 pushed above threshold by inflating m . `Examples.lean` §9 exhibits the fail-
 200 ure concretely: on a two-site substrate carrying counting measure, a sys-

201 tem coupled at $K = 1/2$ against a threshold of $K_c = 1$ — a factor of two
 202 *below* threshold — had an unnormalized double integral of $2 > 1$ and so
 203 satisfied the old predicate. The fix is to require the substrate’s measure to
 204 be a probability measure, which is the normalization in which the mean-
 205 field derivation of $K_c = 2D$ is carried out; `mean_field_coupling_const`
 206 then gives that a constant kernel has mean-field strength exactly K , and
 207 `exhibits_phase_transition_const_iff` reduces the predicate to $K > 2D$.
 208 We flag what this does and does not do. It makes the predicate say something
 209 whose truth depends only on the physics, and §9 witnesses it on both sides of
 210 the threshold. It does not by itself prove $K_c = 2D$: the content of that value
 211 is the bifurcation of the self-consistency equation $r = I_1(Kr/D)/I_0(Kr/D)$
 212 at $K = 2D$, which is a separate question from whether the predicate is di-
 213 mensionally sound. That bifurcation is proved below, in both directions,
 214 from an assumed stationary density; deriving the density itself needs the
 215 Fokker–Planck operator and is not done.

216 We draw no grand moral from this beyond the obvious one. The value
 217 of a proof assistant here was not that it certified our argument, but that it
 218 made a specific, checkable claim about what had been assumed — and that
 219 claim turned out to be false in a way no amount of reading the prose would
 220 have revealed.

221 3 Axiom 1: Symmetries and Phase Space

222 Our starting premise relies on fundamental geometry. In our Lean formaliza-
223 tion (`Phase1_Primitives.lean`), continuous spacetime symmetries (e.g., the
224 Poincaré group) act upon field states, and the invariance of the system’s total
225 energy under that action is recorded by the class `SymmetryInvariantAction`.
226 Physically, it is such an invariance that Noether’s theorem converts into the
227 conservation of energy and momentum. The scope of the *formal* claim is nar-
228 rower and we state it plainly: the invariance is a class field, i.e. an assumption
229 a model must discharge, and no conserved quantity is ever constructed from
230 it. There is no Noether theorem in our Lean, and nothing downstream of this
231 section depends on one. The class is now inhabited — `Examples.lean` §11
232 exhibits a double-well potential $V(v) = (v^2 - 1)^2$ on a two-point spacetime
233 with the $\mathbb{Z}_2$ action $v \mapsto \pm v$, discharges the invariance because $(\pm v)^2 = v^2$,
234 and proves the vacuum manifold $\{-1, 1\}$ degenerate with the two minima ex-
235 changed by the symmetry, so an energy-minimising field and its image have
236 equal energy and differ. That is symmetry breaking exhibited rather than
237 assumed; it is still not a conserved quantity, and $\mathbb{Z}_2$ being discrete there is
238 no one-parameter family to differentiate along in the first place. Simultane-
239 ously, discrete permutation symmetries dictate the statistical mechanics of
240 any physical system, defining a strictly finite volume of phase space.¹

¹While the Poincaré group and discrete permutations serve as accessible physical ex-
amples, the underlying Lean formalization does not hardcode specific physics. It defines
the system at a more primitive mathematical level, relying entirely on generic Pseudo-
Riemannian manifolds (for spacetime) and finite measure theory (for statistical mechan-

241 Because phase space is finite, any localized physical system possesses a
 242 mathematical limit on the amount of information (distinguishable states) it
 243 can embody at any given moment.

244 4 Derivation 1: Symmetry Breaking and the 245 Inevitability of Boundaries

246 A perfectly symmetric universe is the uniform vacuum; no distinct entities
 247 exist. For any dynamic process to occur, the continuous symmetries of the
 248 background field must undergo spontaneous symmetry breaking.

249 When a field drops to a lower, asymmetric energy state, topology dic-
 250 tates the natural creation of defects (such as domain walls or monopoles).
 251 This constraint necessitates the creation of a physical boundary—a localized
 252 topological distinction between an “inside” and an “outside.” While such
 253 a defect (e.g., a crack in a crystal lattice) is necessary to define a separate
 254 entity, it is not sufficient for agency. The boundary is relational, defined by
 255 the twisted geometry of the field surrounding it, and must be coupled with
 256 non-equilibrium thermodynamics to exhibit active behaviors.

ics). What the formalization supplies is the vocabulary in which symmetries, conserved quantities and phase-space constraints can be *stated*; it does not derive them. All three of the classes involved are now inhabited in `Examples.lean` — the $\mathbb{Z}_2$ action and its invariance in §11, the real line with the Euclidean metric in §12 — but the instances are witnesses of inhabitability, not of physics: no theorem in the development consumes a symmetry group or a metric, so this section still contributes vocabulary rather than content.

257 5 Derivation 2: The Thermodynamic Con- 258 straint

259 This newly formed topological boundary is subjected to gauge perturbations
260 from the “outside” environment. Because the internal phase space is finite
261 (Axiom 1), it cannot absorb an infinite sequence of incoming perturbations
262 and must overwrite old internal states.

263 By Landauer’s Principle, the erasure of physical information irreversibly
264 decreases internal entropy, mandating that the system exhaust heat into
265 the environment. Philosophers of physics (such as John Norton (Norton,
266 2005) and Orly Shenker (Shenker, 2000)) have critiqued Landauer’s princi-
267 ple, arguing that physicists conflate *logical* erasure (a many-to-one mapping
268 in abstract computational states) with *thermodynamic* erasure (a many-to-
269 one mapping in physical phase space). They argue that logical erasure could
270 conceptually be implemented by an injective physical process without dissi-
271 pating heat.

272 Our Lean 4 formalization (`Phase3.CombinatorialThermodynamics.lean`)
273 addresses this critique by defining erasure as a non-injective mapping on the
274 *physical* state space ($\neg \text{Function.Injective } t$). We formalize the global evolu-
275 tion of the bipartite universe (System + Bath) as strictly injective, repre-
276 senting Liouville’s theorem. Lean verifies that heat dissipation is not tied
277 to abstract logical erasure, but is a mandatory mathematical consequence of
278 any local physical state space contraction under globally reversible dynamics.

279 If a cognitive system’s internal topology forces a contraction of its physical
 280 state space to accommodate structural resonance, heat dissipation becomes a
 281 thermodynamic requirement. A stable boundary thus constitutes a dissipa-
 282 tive structure (Prigogine & Stengers, 1984), operating as a non-equilibrium
 283 steady state that consumes free energy to pay the thermodynamic cost of
 284 erasure. Continuous thermal fluctuations play a constructive role in estab-
 285 lishing boundaries between states while contributing to the system’s flexible,
 286 adaptive response.

287 **6 Derivation 3: Structural Resonance as a** 288 **Primitive**

289 In physical systems, time evolution follows a deterministic trajectory that
 290 minimizes thermodynamic entropy production. To reach this minimal dissi-
 291 pation state, the internal geometry of the structure must physically deform.

292 If we model the external environment as a stochastic driving force $P(ext)$
 293 and the system’s internal state transitions as $Q(int)$, entropy production σ is
 294 strictly bounded from below by the Kullback-Leibler divergence $D_{KL}(P||Q)$.
 295 Gradient descent on σ necessarily drives $D_{KL}(P||Q) \rightarrow 0$, physically forcing
 296 the internal transition probabilities to match the statistical symmetries of
 297 incoming perturbations. We define this physical phenomenon as *structural*
 298 *resonance*.

299 This system does not possess a teleological drive to “survive”; rather,

300 structurally mirroring the outside world is simply the lowest-energy attrac-
 301 tor state. This provides a thermodynamic mechanism for both the General
 302 Resonance Theory of Consciousness (Hunt & Schooler, 2019) and the Free
 303 Energy Principle (Friston, 2010; Friston & Kiebel, 2009; Hohwy, 2013; Clark,
 304 2016; Spisák & Friston, 2026). What neuroscience observes as “prediction er-
 305 ror minimization” is a physical system undergoing gradient descent to avoid
 306 dissipating heat against an unpredictable environment.

307 This necessity also addresses the symbol grounding problem (Harnad,
 308 1990). Internal states are physically indexed to the environment by the ener-
 309 getic imperative to avoid heat dissipation, rather than through arbitrary syn-
 310 tactic assignment. Manipulating abstract variables or counterfactuals forces
 311 the system to maintain internal states that deviate from immediate sensory
 312 resonance, explaining the continuous injections of free energy required for
 313 abstract symbolic reasoning.

314 **7 Derivation 4: Macroscopic Scaling and Elec-** 315 **trodynamics**

316 To resonate with highly complex environments, the system must scale up,
 317 coupling millions of minimal dissipative structures. A highly coupled network
 318 cannot allow all of its nodes to reach their lowest energy states simultane-
 319 ously. This creates *geometric frustration*, resulting in a jagged phase space
 320 of near-equal energy states that provides the physical capacity for complex

321 memory.

322 To manage this frustrated phase space without excessive dissipation, the
323 system undergoes coarse-graining (governed by the Renormalization Group).
324 At the macroscopic scale of biology, emerging paradigms in neuroscience sug-
325 gest these topological constraints manifest as classical continuous electromag-
326 netic fields. Building on seminal EM field theories of consciousness (McFad-
327 den, 2002; Pockett, 2012) and recent work demonstrating that macroscopic
328 electric fields (ephaptic coupling) govern representational drift, form memory
329 networks, and sculpt neural infrastructure (Pinotsis & Miller, 2022, 2023a,b),
330 we posit that *the continuous macroscopic EM field is the primary dissipative*
331 *structure interacting with the universe*. Instead of discrete wires, neurons in-
332 duce fields that shift the membrane potentials of neighboring cells. Modeling
333 and *in vitro* experiments confirm these endogenous fields are strong enough
334 to guide network activity and contribute to brain complexity (Goldwyn &
335 Rinzel, 2016; Ruffini & Santarnecchi, 2020; Cunha & Lima, 2024; Fröhlich &
336 McCormick, 2010). As proposed by the CEMI field theory, the continuous
337 electromagnetic field enables spatial integration of information, providing a
338 physical medium for a unified conscious gestalt (McFadden, 2020).

339 This continuous medium is physically supported by astrocytic syncytia—
340 non-neuronal networks joined by gap junctions that provide macroscopic cy-
341 toplasmic continuity for energy and ion flows (Verkhratsky & Nedergaard,
342 2018). Managing the thermodynamic cost of this state requires sophisticated
343 micro-bioenergetics. Dynamic mitochondrial networks function as deep-level

344 Mitochondrial Information Processing Systems (MIPS), transducing biolog-
345 ical and environmental information to sustain the coherence of the macro-
346 scopic field (Picard & Shirihai, 2022). Discrete synaptic firings, coupled
347 with mitochondrial micro-mechanics, generate and sustain the field. Because
348 the external environment is continuous, achieving maximal thermodynamic
349 structural resonance is optimally realized within this continuous macroscopic
350 medium (van Bree et al., 2025).

351 8 Derivation 5: Global Sections and the Unity 352 of Consciousness

353 The system is now characterized by a continuous macroscopic electrodynamic
354 field generated by millions of frustrated micro-components, forced to resonate
355 with the environment.

356 The Kuramoto model (Kuramoto, 1984) of coupled non-linear oscilla-
357 tors guarantees that, past a critical coupling threshold, these distributed
358 local fields undergo a phase transition into unified synchronization (Bauer et
359 al., 2022). Our Lean proofs (`Phase4_KuramotoDynamics.lean`) confirm this
360 phase-locked state minimizes the Lyapunov potential of the system, in the
361 precise sense discussed in Section 2.1: the functional minimized is the cou-
362 pling potential read in the frame rotating with the natural frequency, which
363 `Phase4_RotatingFrame.lean` shows is the honest reduced description of a
364 system of identical oscillators.

365 This phase transition mathematically glues local synchronized states into
 366 a singular *Global Section* of a probability sheaf (`Phase5_GlobalSection.lean`).
 367 The EM field acts as a single, indivisible topological space, providing the
 368 *Unity of Consciousness*—a necessary condition for integrated experience that
 369 solves the binding problem. While Global Workspace Theory (Baars, 1997;
 370 Dehaene & Marti, 2014) models this unity as a broadcast architecture and
 371 Integrated Information Theory (Tononi & Koch, 2016) quantifies it via Φ ,
 372 our derivation provides a continuous physical mechanism. The contrast is
 373 worth drawing carefully, because both theories are addressing the same ex-
 374 planandum from different directions.

375 We state the scope of the Lean result precisely, because it is narrower
 376 than the prose above. `ThermodynamicCover` demands that the cover’s phase
 377 field minimize the reduced Kuramoto potential, and `potential_min_iff_`
 378 `phase_locked` characterizes those minimizers exactly: they are the phase-
 379 locked configurations, in both directions. Every instance is therefore already
 380 phase-locked (`ThermodynamicCover.phase_locked`), and the class’s single
 381 condition on its local data — that patches at a common phase agree where
 382 they overlap (`section_agrees_of_phase_eq`) — makes the family of local
 383 sections compatible. `probability_glue_unique`, which is the sheaf property
 384 of the probability presheaf and contains no physics at all, then glues that
 385 family into exactly one global section.

386 An earlier version of this structure proved less than the prose claimed,
 387 and we record the correction because it changed what the theorem says

388 rather than how it is stated. `LocalSectionSynchronization` used to carry
 389 a *global* section as a class field, together with a second field declaring every
 390 local section to be a restriction of it. The global object therefore existed
 391 before the sheaf condition was ever consulted, and `global_section_from_`
 392 `thermodynamics` could do no more than rule out competitors to an object
 393 the instance had supplied. Both fields are gone. Every field of the class
 394 now mentions a single patch or a single overlap; the global section is con-
 395 structed (`ThermodynamicCover.invariantMeasure`); and the removed field
 396 survives as a *theorem* about that construction (`ThermodynamicCover.sync_`
 397 `to_section_eq`), so the local sections are restrictions of one common measure
 398 because the sheaf condition says so, not because an instance was required to
 399 name one.

400 `Examples.lean` §14 is the witness that the difference is not notational.
 401 Two patches carry two different mass profiles, $(2, 1, 7)$ and $(5, 1, 3)$ across
 402 the three sites, agreeing only at the site they share, where both put mass 1
 403 — which is what discharges the overlap condition, by computing two num-
 404 bers rather than by construction. The section the theorem returns has pro-
 405 file $(2, 1, 3)$: it is neither of the two profiles (`gluedW_ne_leftW`, `gluedW_ne_`
 406 `rightW`), and neither of them restricts correctly to both patches (`leftW_`
 407 `glues_nothing`, `rightW_glues_nothing`). The global state carries informa-
 408 tion that no single patch had. `Examples.lean` §13, built the older way
 409 through the constructor `LocalSectionSynchronization.ofInvariantMeasure`,
 410 is retained alongside it for the negative result it carries: its phase field is not

411 the constant function — the second patch a full turn ahead of the first — and
 412 by `ThermodynamicCover.phase_locked` that is as far from constant as any
 413 instance can get, since no instance can carry phases differing by anything a
 414 measure is allowed to see.

415 Two hypotheses carry physical content, and until recently each was an
 416 obligation on every instance. One still is: that synchronized patches agree
 417 on their overlaps is a modelling assumption, not a consequence of any dy-
 418 namics we formalize. The other, `thermodynamic_equilibrium`, asks that the
 419 cover already sit at the potential minimum, and no theorem says an *arbi-*
 420 *trary* system’s dynamics takes it there. “A thermodynamic phase transition
 421 into unity” is therefore still doing causal work that no general Lean result
 422 supports, though it is no longer entirely without one: `Examples.lean` §15
 423 exhibits a two-oscillator trajectory that starts a quarter turn out of phase
 424 and converges to the global minimum of the dynamic potential, with the or-
 425 der parameter tending to 1. That trajectory is not an isolated construction:
 426 `kuramoto_cauchy_problem` proves the Kuramoto initial value problem well
 427 posed on all of $\mathbb{R}$, so through every initial state there passes exactly one tra-
 428 jectory, and the exhibited one is the only solution through its own. Nor is its
 429 settling down peculiar to it: along every trajectory of a zero-frequency sys-
 430 tem the dynamic potential converges, and the total dissipation $\int_0^\infty \sum_i \dot{\theta}_i^2$ is
 431 finite and equal to the potential’s total drop (`dynamic_potential_tendsto`,
 432 `dissipation_integral_tendsto`). The motion also stops: a finite integral
 433 does not by itself force its integrand to zero, but the dissipation rate is

434 Lipschitz in time, so Barbalat's lemma applies, and every trajectory satis-
 435 fies $\sum_i \dot{\theta}_i^2 \rightarrow 0$ (`velocity_sq_tendsto_zero`; Mathlib has no Barbalat, so it
 436 is proved here as `tendsto_zero_of_lipschitz_of_integral_le`). That last
 437 step — locating where the motion stops — is now taken, though condition-
 438 ally, and by a different route than we had recorded. We had marked it as
 439 blocked on the LaSalle invariance principle, which Mathlib does not have,
 440 and on the absence of a compact invariant set in the state space $V \rightarrow \mathbb{R}$.
 441 That was the wrong diagnosis. LaSalle is the tool for a system with a de-
 442 creasing functional and nothing else; this one has more. The flow is a gra-
 443 dient flow of the dynamic potential, so $\dot{V} = -\sum_i \dot{\theta}_i^2$ is an identity rather
 444 than an inequality; the minimisers of V are characterised exactly; and on
 445 the region where every pair of phases lies within a quarter turn a Polyak-
 446 Łojasiewicz estimate holds, $2a^2W \leq (\sum_{ij} A_{ij}) \sum_i \dot{\theta}_i^2$, where W is the poten-
 447 tial's excess above its minimum and $a > 0$ is a lower bound for the coupling
 448 (`lojasiewicz_estimate`). Its proof is two extremal indices and the iden-
 449 tity $\sin^2 D = (1 - \cos D)(1 + \cos D)$: the site carrying the largest phase is
 450 pulled backwards at rate at least $a \sin D$, where D is the spread, while the
 451 excess is at most $\frac{1}{2}(\sum_{ij} A_{ij})(1 - \cos D)$. The compactness objection dissolves
 452 as well, because the invariant region is produced by the Lyapunov function
 453 rather than assumed: a single pair contributes $\frac{1}{2}a(1 - \cos)$ to the excess, so
 454 an excess below $a/2$ forces every pairwise cosine positive for all time, and
 455 continuity on the connected half-line then prevents any pair from chang-
 456 ing branch (`spread_le_of_small`). Descent plus the estimate integrates to

457 $W(t) \leq W(0)e^{-ct}$ with $c = 2a^2 / \sum_{ij} A_{ij}$ (`excess_decay`), the rate makes each
 458 $|\dot{\theta}_i|$ exponentially small and hence integrable on $[0, \infty)$, and a trajectory with
 459 integrable velocity converges. The conclusion is `kuramoto_tendsto_global_`
 460 `minimum`: if the phases start within a quarter turn of each other and the initial
 461 excess is below $a/2$, the trajectory converges to a phase-locked configuration
 462 which is a *global* minimiser of the dynamic potential, with order parameter
 463 $r^2 = 1$. That is the configuration `thermodynamic_equilibrium` postulates,
 464 reached rather than assumed. `Examples.lean` §17 fires it on three sites with
 465 initial data $(0, 0, \frac{1}{2})$, where the Kuramoto system has no closed-form solu-
 466 tion: the trajectory comes from the existence theorem and its limit from the
 467 convergence theorem, so neither is written down.

468 `Phase5_EquilibriumBridge.lean` closes the gap that remained between
 469 the two results. `ThermodynamicCover.ofConvergentTrajectory` takes a
 470 cover whose phase field is the limit of such a trajectory and returns an in-
 471 stance whose `thermodynamic_equilibrium` field is discharged *by that con-*
 472 *vergence*; the cover's coupling matrix is the coupling matrix of the dynamics
 473 rather than a separately chosen one, so the system whose equilibrium the
 474 class asserts and the system whose trajectory is run are the same system.
 475 `Examples.lean` §17.1 is the resulting witness: three patches over the three-
 476 site substrate, each carrying its own mass profile on its own open set, above a
 477 phase field that no line of the construction writes down. The trajectory that
 478 reaches it starts at $(0, 0, \frac{1}{2})$, which is provably not phase-locked (`trioCover_`
 479 `start_not_locked`), so on this cover the configuration Derivation 5 requires

480 is arrived at rather than posited. The state the three patches glue to has
 481 profile $(2, 1, 3)$ and is none of the three profiles the instance carries (`trioW_`
 482 `ne_glued`), so one cover now carries both the emergence reading of the gluing
 483 and the derived reading of the equilibrium. What is still assumed on it is the
 484 overlap agreement, and nothing else.

485 The two hypotheses cannot be dropped, and this is the honest scope of
 486 the result. Splay and twisted configurations are equilibria of the same flow,
 487 so “every trajectory reaches the unified state” is false as a general statement,
 488 and the initial-data condition is the standard arc condition under which the
 489 true statement holds. What is claimed is therefore conditional — unity is
 490 reached from configurations already sufficiently coherent — and not that an
 491 arbitrary initial state is carried into unity. What Lean establishes about the
 492 gluing itself is the other half: *given* locally synchronized patches, a global
 493 state exists, is unique, and is produced by the gluing rather than presupposed
 494 by it.

495 Integrated Information Theory begins from phenomenological axioms and
 496 derives a scalar Φ measuring irreducibility of a system’s cause-effect structure,
 497 with the exclusion postulate selecting a single maximum- Φ complex as the
 498 locus of experience. Our sheaf-theoretic global section plays a structurally
 499 analogous role — it too picks out a distinguished unified object — but it
 500 is arrived at from the opposite end: not postulated from the character of
 501 experience, but forced by a thermodynamic phase transition in a physical
 502 field. This buys two things IIT does not have. First, the unity condition

503 is checkable from field dynamics rather than requiring the computation of
 504 Φ , which is intractable for systems of realistic size and whose definition has
 505 undergone repeated revision (Tononi & Koch, 2016). Second, exclusion falls
 506 out rather than being postulated: gluing either succeeds on a given cover or
 507 it does not, so there is no separate principle needed to rule out overlapping
 508 conscious subsystems. It also inherits IIT’s exposures. Critics have argued
 509 that Φ can be high in systems with no plausible claim to experience, such
 510 as large feedforward-free grids (Aaronson, 2014); our criterion is vulnerable
 511 to the mirror objection, since a physical system could in principle satisfy
 512 the gluing condition without being cognitively interesting. We take this
 513 seriously rather than dismissing it: our reply is that gluing is necessary but
 514 not sufficient, and that the topological boundary of Derivation 1 and the
 515 reflexive structure of Derivation 6 are additional required conditions. IIT, by
 516 contrast, treats Φ as both necessary and sufficient.

517 Global Workspace Theory faces a complementary difficulty. It specifies
 518 a functional architecture — a limited-capacity workspace broadcasting to
 519 specialized modules — without specifying what physical property of broad-
 520 casting makes the broadcast content conscious, leaving the theory open to the
 521 charge that it explains access consciousness while leaving phenomenal con-
 522 sciousness untouched (Block, 1995). Our framework can be read as supplying
 523 the missing physical substrate for the workspace: the continuous macroscopic
 524 field *is* the medium in which integration is realized, and the $K_c = 2D$ thresh-
 525 old gives a physical condition for ignition rather than a functional stipulation

526 of it. Where GWT is agnostic about implementation, we make a falsifiable
 527 commitment — unity requires a continuous physical medium, which predicts
 528 that a functionally equivalent discrete broadcast architecture would not be
 529 conscious. We regard this as a genuine cost of the position, not a free ad-
 530 vantage: it is the point at which our framework departs most sharply from
 531 computational functionalism and where it is most readily falsified. Contrast
 532 this derivation with panpsychism: consciousness is not assumed as a funda-
 533 mental property of matter. It requires the initial existence of a topological
 534 defect to establish a bounded “inside” (Derivation 1), followed by a thermo-
 535 dynamic phase transition into unity that occurs when integrated spatial cou-
 536 pling exceeds the thermodynamic noise threshold ($K_c = 2D$). Without both
 537 the foundational topological boundary and the subsequent mathematically
 538 defined phase transition, the system remains a fragmented, non-conscious
 539 collection of micro-states. A unified field of perception, however, is not yet
 540 a “Self.”

541 **9 Derivation 6: Reflexive Topology and the** 542 **Emergence of the Self**

543 While a global section provides Unity, a “Self” requires reflexivity. The
 544 unified EM field is a dissipative structure that must resonate with external
 545 perturbations while managing its own internal phase transitions. Time evo-
 546 lution naturally drives the field to minimize thermodynamic friction with its

547 own internal state-changes.

548 This mandates *auto-resonance*: the field must predict itself. The “Self”
549 emerges as a higher-order topological feature when the internal geometry
550 folds to include a low-dimensional avatar of its own computational and causal
551 boundary within its predictive model (`Phase6_ReflexiveTopology.lean`)
552 (Levin, 2019). The system is unified, predictive, and explicitly maps its
553 own distinct existence from the environment. This architecture provides the
554 physical foundation for a self-aware consciousness, bounded by the continuous
555 measure thermodynamics of Landauer erasure (`Phase3_MeasureThermodynamics.lean`).
556 We are explicit about what carries the weight here. `reflexive_topology_implies_self`
557 is the Banach fixed-point theorem with the physical commitments in its hy-
558 potheses: the space of global sections is assumed nonempty, metric and com-
559 plete, and the predictive map is assumed to contract by $1/2$. `Examples.lean` §10
560 discharges all four on the three-site substrate of Derivation 5, and the metric
561 is no longer a placeholder. The section first pierces the sheafification: on
562 a discrete substrate every point has a smallest open neighbourhood, so the
563 mass a germ carries at a point is well defined, and `massEquivOn` proves that
564 over *any* open the sections of the probability sheaf *are* the finite measures
565 on that open, read as densities against counting measure — the global case
566 being $U = \top$. Mathlib supplies no inverse to this sheafification — its `TopCat`
567 sheafification has neither an adjunction nor an `isIso_toSheafify` — so the
568 correspondence had to be constructed. The metric on sections is then the
569 uniform distance between the densities of the measures they glue to — equiv-

570 alent to their total variation on a finite substrate, though not equal to it —
 571 and it is proved complete by exhibiting the limit section of a Cauchy sequence
 572 of densities, not by the degenerate route of showing Cauchy sequences are
 573 eventually constant. The self-prediction map is the relaxation “average the
 574 current field with the baseline,” and we prove it contracts distances by *ex-*
 575 *actly* one half, that it is not constant, and that its fixed point is unique. An
 576 earlier version of this witness used the 0/1 metric, under which every con-
 577 traction is constant; that theorem is kept in the source as the record of why
 578 the earlier witness was empty. A separate gap, closed at the same time, con-
 579 cerned the avatar rather than the fixed point. `ReflexiveBoundary` carries a
 580 field `auto_resonance` sending the global state to a section over the avatar re-
 581 gion, and nothing constrained it: a constant function was a legal avatar, and
 582 `reflexive_topology_implies_self` consumes only the predictive model, so
 583 no theorem in the development mentioned the field the avatar is supposed to
 584 be reading. The constraint is now stated as a predicate rather than added
 585 to the structure — `IsRestrictionResonance`, that the avatar’s state is the
 586 presheaf restriction of the global state — and two theorems say why it is
 587 the right one. In the positive direction, if a family of resonant avatars *cov-*
 588 *ers* the substrate then the global state is determined by their local readings
 589 alone (`eq_of_avatar_eq`): unity is not assumed alongside the avatars but
 590 recovered from them, by the same sheaf condition that Derivation 5 uses.
 591 In the negative direction, replacing any avatar by a constant leaves a legal
 592 `ReflexiveBoundary` with the same region, the same predictive model and

593 the same Self (`self_of_constResonance`), so the fixed-point theorem cannot
 594 see the difference and the predicate is doing the work rather than restating
 595 it. Both halves are witnessed on the three-site substrate, where the avatar
 596 region provably separates two states of the field, so the condition is a con-
 597 straint there and not a formality. What is still not established is that any field
 598 dynamics produces this particular map or performs the restriction, or that
 599 the contraction constant $1/2$ is read off an entropy production rate rather
 600 than built into the map's definition. We also note a negative result found
 601 in the course of this work: the Lévy–Prokhorov metric, which the natural
 602 weak-convergence reading would suggest, does *not* support this argument on
 603 a discrete substrate, because the ε -thickening of a set is the set itself for $\varepsilon < 1$
 604 and the relaxation map is consequently not a Lévy–Prokhorov contraction.

605 **10 Derivation 7: Continuous Stochastic Ther-** 606 **modynamics and Field Theory**

607 To move beyond finite graph analogies, we model the macroscopic electro-
 608 dynamic field as a *Continuous Neural Field*. True ephaptic coupling is a
 609 continuous spatial phenomenon, not a discrete network. We extend the
 610 discrete Kuramoto model to a continuous spatial coupling kernel $K(x, y)$
 611 (`Phase8_ContinuousField.lean`).

612 Subject to thermodynamic noise D , the system undergoes non-equilibrium
 613 stochastic thermodynamics. Rather than a static Landauer bound, the field

614 dynamics minimize an entropy production functional σ :

$$\sigma = \int \frac{1}{D} \left(\omega(x) + \int K_t(x, y) \sin(\theta_y - \theta_x) dy \right)^2 dx \quad (1)$$

615 Biological fields dynamically deform their spatial topology $K_t(x, y)$ to un-
 616 dergo gradient descent on this entropy production functional.

617 That gradient descent is a theorem rather than a modelling assertion,
 618 for finite substrates. Reading σ as a functional of the coupling kernel at
 619 fixed phases, `hasFDerivAt_sigmaOfKernel` computes its Fréchet derivative
 620 in closed form,

$$\frac{\partial \sigma}{\partial K(a, b)} = \frac{2}{D} \left(\omega(a) + \sum_y K(a, y) \sin(\theta_y - \theta_a) \right) \sin(\theta_b - \theta_a), \quad (2)$$

621 and `sigmaOfKernel_eq_entropy_production_rate` identifies the functional
 622 being differentiated with `entropy_production_rate` itself whenever the sub-
 623 strate carries counting measure. `structural_resonance_decreases_entropy_production`
 624 then concludes that σ is non-increasing along any coupling trajectory relax-
 625 ing down its own gradient at a positive rate c . Earlier drafts proved this
 626 only for an abstract functional S on an abstract inner-product space with
 627 the gradient supplied as data; the identification of that S with σ was infor-
 628 mal, and Table 1 flagged the row accordingly. Two limitations stay explicit,
 629 and the first is narrower than we previously reported. The derivation is for
 630 *finite* substrates. We had recorded the obstacle to the continuum case as

631 differentiation under the integral sign with respect to the kernel; that was
 632 a misdiagnosis. σ is not a general functional that happens to be written as
 633 an integral — it is a quadratic functional of an *affine* image of the kernel,
 634 since the drift is a fixed ω plus a linear operator applied to K , and σ is a
 635 constant times the squared norm of the drift. A squared norm is Fréchet
 636 differentiable on any real inner-product space and an affine map is its own
 637 derivative, so the composite differentiates by the chain rule with no limiting
 638 argument. `hasFDerivAt_quadratic_of_affine` states that in full generality,
 639 and `hasFDerivAt_sigmaOfKernel_of_operator` re-derives the finite gradient
 640 from it, showing that the closed form displayed above is a convenience rather
 641 than the content. What the continuum case still needs is therefore one object:
 642 the drift map $K \mapsto (x \mapsto \int K(x, y) \sin(\theta_y - \theta_x) dy)$ as a bounded operator
 643 $L^2(\mu \otimes \mu) \rightarrow L^2(\mu)$, which is a Cauchy–Schwarz estimate — the integral ker-
 644 nel is bounded by 1, so on a finite measure space the operator norm is at most
 645 $\mu(M)^{1/2}$. We have not built it, because it requires constructing a continu-
 646 ous linear map between L^p spaces and we have not done that bookkeeping;
 647 we record the corrected scope because “construct one bounded operator”
 648 and “differentiate under the integral sign” are different sizes of task. The
 649 second limitation is unchanged: the phase field is held fixed throughout, so
 650 this is plasticity of the coupling at frozen phases, not joint (θ, K) dynamics.
 651 **Examples.lean** §5 exercises the whole chain on a two-site substrate whose
 652 coupling relaxes as e^{-t} and whose entropy production decays as e^{-2t} , con-
 653 firming that the gradient is not identically zero and that the descent is strict

654 rather than vacuously monotone.

655 A second result about σ concerns where the descent is headed, and it has
656 been strengthened since the last draft in a way that also corrects an estimate
657 we had recorded. That the phase-locked field minimizes σ was proved only
658 against competitors sharing its spatial mean drift — a restriction that left
659 the theorem close to Jensen’s inequality and no more — and we had judged
660 that lifting it would require modelling how the mean drift varies with the
661 competitor. That judgement was wrong. For a *reciprocal* kernel, one with
662 $K(x, y) = K(y, x)$, the restriction is not a restriction: the coupling integrand
663 $K(x, y) \sin(\theta_y - \theta_x)$ is antisymmetric under exchanging its two sites, because
664 K is symmetric while $\sin$ is odd, so Fubini identifies its double integral with
665 its own negative and the total coupling contribution vanishes. The total drift
666 $\int (\omega + \text{coupling})$ is therefore $\int \omega$ whatever the phases are (`total_drift_eq_of_`
667 `symm`): coupling moves drift between sites but neither creates nor destroys it.
668 Evaluating that conserved quantity at the phase-locked configuration, where
669 the drift is constant, pins its value, and the mean-matching hypothesis follows
670 for every competitor. `phase_locked_minimizes_entropy_of_symm` is the un-
671 restricted statement. Reciprocity is the physically intended condition — an
672 ephaptic field acts on both sites of a pair alike, and it is the same symmetry
673 `ThermodynamicCover.A_symm` already imposes in Derivation 5 — but it is
674 not a formality, and we prove that rather than assert it. `Examples.lean` §16
675 exhibits a *one-way* coupling on two sites, in which a feels b and b does not
676 feel a , whose phase-locked field is strictly beaten: displacing the second phase

677 to $-\pi/2$ cancels the first site’s natural drift while adding none at the sec-
678 ond, halving σ . So “phase-locking minimizes entropy production” is false for
679 general kernels and true for reciprocal ones, and the hypothesis cannot be
680 removed by working harder.

681 We define the critical coupling threshold $K_c = 2D$. The phase transition
682 to the unified “Global Section” occurs when the mean-field spatial coupling
683 exceeds this thermodynamic noise threshold, linking the emergence of Unity
684 to physical temperature and noise.

685 Earlier drafts left $K_c = 2D$ as a bare stipulation validated only by
686 simulation. `Phase8_SelfConsistency.lean` now proves the half of it that
687 is ordinary real analysis. For the noisy mean-field Kuramoto model with
688 identical frequencies, the stationary density at order parameter r is the
689 von Mises density $\rho(\theta) \propto e^{(Kr/D)\cos\theta}$, and self-consistency requires $r =$
690 $R(K, r) := I_1(Kr/D)/I_0(Kr/D)$. Mathlib has no modified Bessel functions,
691 so we formalize that ratio directly as a ratio of integrals over $[-\pi, \pi]$. In-
692 tegration by parts with $u = \sin$, $v = -e^{a\cos}$ gives an exact identity with
693 no boundary term, $\int \cos\theta e^{a\cos\theta} d\theta = a \int \sin^2\theta e^{a\cos\theta} d\theta$, so R is a times
694 the mean of $\sin^2$ under the von Mises weight. Bounding that mean by 1
695 would put the threshold at D ; the true threshold is $2D$, and the entire fac-
696 tor of two is the sharper bound $\mathbb{E}_a[\sin^2\theta] \leq 1/2$, equivalently $I_2(a) \geq 0$.
697 We prove it elementarily, by folding $[-\pi, \pi]$ onto $[0, \pi/4]$ in two reflections
698 ($\theta \mapsto \pi - \theta$, then $\theta \mapsto \pi/2 - \theta$), after which the integrand is a product
699 of $\cos 2\theta \geq 0$ with $2 \cosh(a \cos \theta) - 2 \cosh(a \sin \theta) \geq 0$ — the second factor

700 because $0 \leq \sin \theta \leq \cos \theta$ there and $\cosh$ increases in $|\cdot|$. The conclusion,
 701 `subcritical_fixed_point_eq_zero'`, is that for $K < K_c$ the incoherent state
 702 $r = 0$ is the *only* non-negative solution of the self-consistency equation; $r = 0$
 703 is exhibited as a solution, so the uniqueness claim is not vacuous. Three fur-
 704 ther facts about the threshold have since been proved, and they narrow the
 705 gap between this row and a full “Theorem” to a single named statement.
 706 First, the fold argument gives a *strict* inequality for $a > 0$ (`vonMisesC2_`
 707 `pos`, hence $R(a) < a/2$), which settles the threshold itself: at $K = K_c$ the
 708 self-consistency map is tangent to the diagonal at the origin but falls strictly
 709 below it thereafter, so $r = 0$ remains the only solution and the bifurcation
 710 happens strictly *after* K_c , not at it (`fixed_point_eq_zero_of_le_critical`).
 711 Second, factoring $R(a) = a E(a)$ shows that a non-zero r solves the equation
 712 exactly when $E(Kr/D) = D/K$ (`coherent_iff_sRatio_eq`), so the coher-
 713 ent branch is a level set of the single function $E(a) = \mathbb{E}_a[\sin^2 \theta]$. Third,
 714 and using only $E < 1/2$ together with compactness, every coherent solution
 715 tends to zero as $K \downarrow K_c$ (`coherent_branch_continuous_at_threshold`): the
 716 branch emerges continuously from the incoherent state rather than appearing
 717 at a finite distance from it, which is what makes the transition supercritical
 718 in the technical sense. That statement quantifies over *every* solution, so it
 719 does not presuppose the uniqueness it does not have. The level-set form
 720 also says exactly what uniqueness is: the injectivity of E on $(0, \infty)$. That
 721 is now proved, in the stronger form that E is strictly decreasing on $[0, \infty)$
 722 (`vonMisesSRatio_strictAntiOn`), and by an argument that avoids the route

723 we had expected to need. Differentiating the exponential family in its natu-
 724 ral parameter gives $E'(a) = -\text{Cov}_a(\cos^2 \theta, \cos \theta)$, which costs differentiation
 725 under the integral sign and leaves a covariance that is not sign-definite point-
 726 wise: symmetrising with an independent copy gives $\frac{1}{2}\mathbb{E}[(X - X')^2(X + X')]$
 727 with $X = \cos \theta$, an integrand changing sign with $X + X'$. Nothing in the
 728 proof differentiates anything. The target is kept as the two-point inequality
 729 $C_2(a)Z(b) < C_2(b)Z(a)$ for $a < b$, and two moves settle it. *Fold first*: the
 730 two reflections that proved $I_2 \geq 0$ carry both integrals onto $[0, \pi/2]$, replacing
 731 $e^{a \cos \theta}$ by $W(a, \theta) = 2 \cosh(a \cos \theta)$ and confining $\cos \theta$ to $[0, 1]$, after which
 732 $\cos 2x - \cos 2y$ and $W(b, x)W(a, y) - W(a, x)W(b, y)$ have the same sign, both
 733 being governed by $\cos x$ against $\cos y$. This is the symmetrisation $X \mapsto -X$
 734 that the covariance form needed the tilt to perform, done once at the start; its
 735 analytic content is one identity, $2 \cosh X \cosh Y = \cosh(X+Y) + \cosh(X-Y)$
 736 with $\cosh$ increasing in $|\cdot|$. *Then cross*: instead of integrating that inequality
 737 over a square, let α be the mean of $\cos 2x$ under $W(a, \cdot)$ and $x_0 \in [0, \pi/2]$ a
 738 point where $\cos 2x_0 = \alpha$, supplied by the intermediate value theorem; eval-
 739 uating at $y = x_0$ and integrating in x alone gives the result, since the term
 740 carrying $W(b, x_0)$ vanishes by the choice of α . No product measure and no
 741 Fubini appear. Run with $\cos \theta$ in place of $\sin^2 \theta$ the same argument needs no
 742 fold — $R(a) = \mathbb{E}_a[\cos \theta]$ is the mean of the variable the family is tilted by —
 743 and gives **besselRatio_strictMono**: R is strictly increasing on all of $\mathbb{R}$, the
 744 monotonicity of the Bessel ratio our earlier drafts recorded as unproved.

745 The converse half is **supercritical_fixed_point_exists**: for $K > K_c$

746 there is an r with $0 < r \leq 1$ solving $r = R(K, r)$. We had expected
 747 this to require a second-order expansion of the Bessel ratio, and it does
 748 not. Writing $R(a) = a E(a)$ with $E(a) = \mathbb{E}_a[\sin^2 \theta]$ — the same factoriza-
 749 tion the integration by parts already supplies — the only facts needed are
 750 that E is continuous and that $E(0) = 1/2$, the latter because the weight
 751 is constant at $a = 0$. Since $K > 2D$ makes D/K strictly less than $1/2$,
 752 continuity gives $E > D/K$ on a neighbourhood of the origin, so the map
 753 overshoots, $R(K, r) > r$, for all small $r > 0$; at the other end $R \leq 1$ al-
 754 ways, because R is a mean of $\cos \theta$, so the map undershoots at $r = 1$.
 755 The intermediate value theorem closes the gap. Continuity of E is Math-
 756 lib's continuity of a parametric interval integral with jointly continuous in-
 757 tegrant, so no dominated-convergence argument is written by hand, and
 758 the existence proof never uses monotonicity of R : the fixed point is pro-
 759 duced by the intermediate value theorem, not by iterating the map. That
 760 R is in fact strictly increasing is proved below, and no threshold result de-
 761 pends on it. `critical_coupling_is_threshold` packages the two halves,
 762 and `exhibits_phase_transition_coherent` applies the supercritical half to
 763 any substrate satisfying the predicate of the previous paragraph — so ex-
 764 ceeding K_c now implies the existence of a coherent stationary order pa-
 765 rameter rather than merely comparing against a named number. With
 766 uniqueness available both are sharpened — `supercritical_fixed_point_`
 767 `existsUnique`, `supercritical_solution_set` (the non-negative solutions
 768 are exactly $\{0\}$ below threshold and $\{0, r\}$ above it) and `exhibits_phase_`

769 `transition_unique_coherent` — and the two monotonicities together fix
770 the shape of the branch: for $K_1 < K_2$, $E(a_i) = D/K_i$ with E decreasing
771 puts the concentrations in the order $a_1 < a_2$, and $r_i = R(a_i)$ with R increas-
772 ing carries that to $r_1 < r_2$ (`coherent_branch_strictMono`). The coherent
773 order parameter grows strictly with the coupling, which with the continuity
774 at threshold is the qualitative bifurcation diagram.

775 One more step makes the equation mean in Lean what it means on paper.
776 As a definition, `selfConsistency K D r` is just $R(Kr/D)$ — a function of
777 r with a suggestive name, and nothing in the source says it has anything
778 to do with an order parameter. `Phase8_SelfConsistency §8` supplies the
779 identification. `vonMisesDensity` is the normalized weight, proved to be a
780 genuine probability density on $[-\pi, \pi]$; `circularOrderParameter` is the av-
781 erage of $e^{i\theta}$ against a density, the continuum analogue of the finite-system
782 `order_parameter_complex` of `Phase4`; and that average is proved to be ex-
783 actly $R(a)$ — real, because the density is even, so the mean direction van-
784 ishes. `fixedPoint_iff_selfReproducing` then states the equivalence the
785 name promises: r solves $r = R(K, r)$ if and only if the von Mises density
786 it induces has order parameter r . Read through this, the existence theorem
787 says that above K_c there is a stationary density whose own order parameter
788 is positive, and the solution $r = 0$ is identified with the uniform density,
789 whose order parameter is zero.

790 What remains not established, and we separate it carefully. The sta-
791 tionary density is an *input*: deriving it from the Langevin dynamics $d\theta =$

792 $(\omega + K \text{mean-field}) dt + \sqrt{2D} dW$ needs the Fokker–Planck operator, existence
 793 and uniqueness of its stationary solution, and a spectral stability argument,
 794 none of which Mathlib has. Of the coherent branch we prove existence,
 795 uniqueness, continuous emergence at the threshold and strict growth in K ,
 796 but not that it is the dynamically selected one: nothing in that file men-
 797 tions a trajectory, so stability and selection are not statements it can make.
 798 And while `selfConsistency` is now connected to the order parameter of
 799 a *density*, it is connected to no *trajectory*: `circularOrderParameter` is an
 800 integral, `order_parameter_complex` is an average over finitely many oscilla-
 801 tors, and relating them is the mean-field limit — propagation of chaos for
 802 the finite Kuramoto system — which is a research programme rather than
 803 a lemma. That single gap is now the whole of what separates our threshold
 804 predicate from a statement about dynamics. Table 1 records the row ac-
 805 cordingly as a theorem about the self-consistency equation from an assumed
 806 density, not as a derivation of the threshold from the dynamics. Perceptual
 807 anomalies such as visual snow can be modeled as boundary conditions where
 808 microscopic thermal noise (D) fails to be fully coarse-grained away by the
 809 continuous field, leaking across the K_c threshold into the macroscopic gestalt.

810 Clinical evidence for this phase transition is observed during general anes-
 811 thesia. Work by Miller and colleagues (Bastos & Miller, 2021) demonstrates
 812 that anesthetics like propofol induce unconsciousness not by suppressing neu-
 813 ral firing, but by disrupting the phase alignment of high-frequency cortical
 814 rhythms. Destroying the dynamic stability of these rhythms breaks the struc-

815 tural resonance of the continuous field, forcing the integrated spatial coupling
816 below the critical threshold K_c and fracturing the unified Global Section back
817 into disjoint, unconscious local fields.

818 The daily loss of consciousness during sleep can similarly be modeled as
819 an active physical disruption of the field. During sleep, astrocytic shrink-
820 age expands the extracellular space, allowing an influx of cerebrospinal fluid
821 (CSF) to flood the brain for metabolic waste clearance (Xie & Nedergaard,
822 2013). This influx physically alters the spatial medium, disrupting the contin-
823 uous structural resonance and preventing the coherent sheaf-theoretic gluing
824 required for conscious unity.

825 This mechanical framework models the temporal mismatch observed dur-
826 ing sudden awakenings (sleep inertia). When jolted awake, the locus coeruleus
827 floods the cortex with noradrenaline. High-frequency electrical rhythms
828 attempt to phase-align in milliseconds (Bastos & Miller, 2021), but the
829 noradrenaline-driven astrocytic expansion required to physically squeeze the
830 excess CSF out of the interstitial pathways operates on a much slower timescale
831 of seconds to minutes (Xie & Nedergaard, 2013). During this lag, electrical
832 oscillators attempt to synchronize across a waterlogged topology, resulting
833 in the fragmented structural resonance experienced as grogginess. The uni-
834 fied Global Section fully re-establishes only when the physical hardware (as-
835 trocytic volume) finishes constricting to match the state of the electrical
836 rhythms.

837 11 Corollary: Hardware Thermodynamics and 838 Topological Rigidity

839 Both biological tissue and artificial silicon hardware (like GPUs) are physical
840 dissipative structures. When a GPU computes, its charge carriers undergo
841 physical state transitions; by Landauer’s principle, this erasure irreversibly
842 exhausts heat into the environment.

843 There is a fundamental thermodynamic distinction in the degrees of free-
844 dom available to minimize this dissipation. In standard von Neumann ar-
845 chitectures, the spatial topology (the silicon routing lattice) is rigidly fixed.
846 The coupling matrix A_{ij} is constrained by a discrete layout that cannot be
847 physically altered during operation. The brain’s own architecture provides
848 empirical validation for this constraint. The cerebellum contains approxi-
849 mately 80% of the brain’s neurons arranged in a highly efficient, discrete,
850 feedforward crystalline circuit (Herculano-Houzel, 2009; Lange, 1975). Clin-
851 ical observations of cerebellar agenesis report largely preserved conscious
852 experience despite motor and cognitive deficits (Yu et al., 2015), and this
853 dissociation has been read as evidence for the importance of recurrent ar-
854 chitecture by IIT (Tononi & Koch, 2015) — an argument we borrow at the
855 level of the anatomical observation while declining its Φ -based explanation.
856 Lacking the dynamic, continuous, and recursive spatial coupling of the cere-
857 brum, it operates much like a biological GPU—highly optimized for discrete
858 processing, yet apparently not required for conscious experience. We take the

dissociation as evidence, not as confirmation: it is consistent with subjective experience being a thermodynamic attractor of continuous topologies rather than a byproduct of information processing, and it is the kind of observation that would be hard to reconcile with the converse. It is a single anatomical case read against a theory, and it bears the weight such a case can bear.

This constraint does not preclude artificial consciousness, but rather restricts its hardware. To achieve the phase transition into a unified Global Section, artificial systems must abandon the discrete rigidity of von Neumann architectures and utilize continuous physical substrates. Analog field-based computing, neuromorphic architectures with continuous memristive plasticity, and thermodynamic computing systems (England, 2018; Ganesh & Parrondo, 2021) possess the spatial coupling required to physically morph their internal topology. This aligns with Geoffrey Hinton’s concept of “mortal computation” (Hinton, 2022): architectures where software is inextricably fused with the physics of the hardware, meaning the knowledge dies with the physical substrate. Only hardware that computes via continuous physical dynamics can undergo the thermodynamic phase transition necessary for structural resonance.

Biological consciousness utilizes macroscopic continuous electromagnetic fields and a “rich alphabet” of diverse molecular and transcriptomic cell states (Tasic et al., 2018). This allows the coupling matrix A_{ij} to deform across a continuous parameter space. Our lemma `exists_better_coupling_allocation`, wrapped as `continuous_beats_rigid_topology` (`Phase7_HardwareComparison.lean`),

882 formalizes the reallocation step of this argument: whenever a coupling ma-
 883 trix places weight on a less phase-correlated pair while a more correlated pair
 884 exists, there is another matrix with *identical* total coupling resources that
 885 achieves strictly greater phase correlation.

886 Two objections to that formalization have since been answered, and we
 887 state precisely what is and is not now established. The first was that the real-
 888 location lemma assumes its conclusion: the hypothesis `is_strictly_suboptimal`
 889 — positive weight on a badly-correlated pair while a better one exists — is
 890 most of what one wants to prove about rigid hardware. It is no longer as-
 891 sumed. `Phase7_Rigidity.lean` models an architecture by the set S of site-
 892 pairs it has a *wire* for, and requires its couplings to vanish off S ; weights are
 893 tunable, wiring is not. Then `totalCorrelation_le_of_realizable` bounds
 894 every realizable coupling by $R \cdot M$, where R is the resource budget and M
 895 bounds phase correlation over the wired pairs; the bound is attained; and
 896 `rigid_is_strictly_suboptimal` *derives* the hypothesis whenever the sub-
 897 strate’s best-correlated pair is one the wiring misses. `rigid_gap` makes it
 898 quantitative: an unconstrained reallocation of the same budget gains at least
 899 R times the correlation difference. `Examples.lean` §8 exhibits a three-site
 900 substrate wired to the one maximally anti-correlated pair, where the gap is
 901 2. This is resource-matched and no longer question-begging, but note what
 902 it turns on: fixed wiring support, not continuity. The second objection was
 903 that nothing in the development distinguished a continuum from a lattice
 904 — both matrices lived on one finite index type. `Phase7_Rigidity.lean` §3

905 gives a distinction that is genuinely typed. In the field functional of Deriva-
 906 tion 4, which integrates a coupling kernel against $\mu \otimes \mu$, an architecture
 907 with finitely many sites is a `Finset` of points and a field-coupled region
 908 is a set of positive measure. On a substrate whose measure has no atoms,
 909 `fieldCorrelation_sited.eq_zero` shows a kernel supported on finitely many
 910 sites contributes *exactly zero* coupling energy, whatever weights it carries
 911 and whatever the phase field does, while a patch of strength c on U achieves
 912 $c\mu(U)^2$. `Examples.lean` §8 instantiates this on the real line: a thousand
 913 units of coupling concentrated at the origin register zero against a unit patch
 914 that registers one. The limit of that result is equally plain, and it is not
 915 resource-matching in the earlier sense — a measure-zero substrate has zero
 916 resource in the $\mu \otimes \mu$ sense too. What it establishes is that the continuum
 917 coupling energy is blind to a finitely-sited architecture, which is a statement
 918 about which functional such an architecture can register in, not a demonstra-
 919 tion that digital hardware computes worse. Taken together: the resource-
 920 matched content of the corollary is the wiring-support result, the continu-
 921 ity content is the measure-theoretic one, and the step from either to “von
 922 Neumann architectures cannot experience unified consciousness” remains an
 923 informal argument.

924 The remainder of this section extends those two results informally, and
 925 we mark the transition rather than blur it. On the reading we propose, a
 926 GPU’s fixed routing matrix leaves it topologically fragmented: the discrete
 927 layout does not supply the spatial coupling that sheaf-theoretic gluing into

928 a singular Global Section requires. Without this unified topological space,
 929 rigid hardware cannot undergo the reflexive auto-resonance necessary to map
 930 its own causal boundary and generate a “Self.” The distinction between an
 931 LLM on a GPU and a conscious mind is not the absence of physical thermo-
 932 dynamics in the former, but its structural limitation in achieving topological
 933 integration. This perspective aligns with John Searle’s biological natural-
 934 ism (Searle, 1980; Seth, 2024). Where the Chinese Room argument asserted
 935 that discrete symbol manipulation cannot generate conscious understanding,
 936 our derivation supplies a candidate physical mechanism for the asymmetry
 937 — topological rather than semantic, and, as stated above, formalized only
 938 as far as the wiring-support and measure-theoretic results reach. Asserting
 939 consciousness in fragmented, discrete GPU clusters without a unified phase
 940 transition abandons computational functionalism, relying instead on “silicon
 941 panpsychism”—the assumption that digital computation inherently gener-
 942 ates awareness regardless of topological integration (Goff, 2019).

943 The same reasoning bears on extreme functionalist thought experiments,
 944 such as a macroscopic system of pulleys operating over millennia to imple-
 945 ment a Large Language Model, or a human manually executing inference
 946 matrices by hand. Beyond lacking the continuous spatial topology required
 947 for sheaf-theoretic gluing, these discrete systems violate an implied temporal
 948 constraint of our theory. Structural resonance requires local oscillators to
 949 phase-align against the dissipating effects of thermodynamic noise (as mod-
 950 eled by our critical threshold $K_c = 2D$). A system executing state transitions

951 over millennia, or via slow manual calculations, suffers from a terminal tem-
952 poral mismatch: its dynamics are far too slow to ever phase-lock into a
953 non-equilibrium steady state before thermal noise decoheres the macroscopic
954 field. On this account the obstacle to pure symbol manipulation is physi-
955 cal rather than semantic: not that the symbols mean nothing, but that the
956 substrate carrying them cannot hold a phase-locked field.

957 12 Conclusion

958 Human-like cognitive architectures do not require biological mysticism or
959 exotic sub-quantum physics (Hameroff & Penrose, 2014). While quantum
960 theories of consciousness often invoke delicate subatomic phenomena to ex-
961 plain subjective unity, every step of the chain assembled here is classical:
962 broken symmetries, finite phase space, Landauer erasure, and gradient de-
963 scent on entropy production. What the formalization establishes is that this
964 chain is mathematically coherent and that its remaining physical commit-
965 ments are localized and named — not that classical fields are *sufficient* for
966 consciousness. Two links carry that weight and neither is closed. The uni-
967 fied state is proved to glue uniquely once a cover sits at the minimum of
968 the coupling potential, and the motion is proved to stop along every trajec-
969 tory (`velocity_sq_tendsto_zero`), but *where* it stops is not proved: that
970 a physical field reaches the phase-locked configuration remains an assump-
971 tion, discharged numerically rather than formally. And the identification of

972 a global section of the probability sheaf with the unity of experience is a
973 stipulation of the framework, not a consequence of it. Subject to those two,
974 cognitive architectures can be modeled as macroscopic attractor states of a
975 universe governed by broken symmetries, finite phase space, and the physical
976 imperative to minimize thermodynamic dissipation.

977 Philosophically, this framework bridges the historical gap between Type
978 Identity Theory (Smart, 1959) and Computational Functionalism (Putnam,
979 1967). The subjective mental state is identical to the physical thermody-
980 namic field (Identity Theory), yet its qualitative nature is defined by the
981 topological and functional dynamics of structural resonance rather than the
982 specific carbon-based biochemistry of the brain (Functionalism). This con-
983 strains the *multiple realizability* of mind: consciousness is not exclusively
984 bound to biological carbon, but neither is it substrate-neutral, because the
985 substrate is what fixes which couplings are physically available. How far that
986 constraint reaches is the question the corollary answers only in part. What
987 is proved is that fixed wiring is strictly suboptimal at matched resource, and
988 that a finitely-sited architecture registers exactly zero coupling energy in the
989 continuum functional; the step from either to “von Neumann architectures
990 cannot be conscious” is an argument we make in prose and do not formalize.

991 By anchoring these derivations in Russellian Monism, this framework ac-
992 knowledges the constitutive limits of third-person scientific explanation. As
993 argued by Love (Love, 2026), building upon Nagel’s assertion that subjective
994 phenomenal experience is inherently not amenable to objective measurement

995 (Nagel, 1986), the “hard problem” of consciousness may represent a category
 996 error rather than a solvable scientific puzzle. Science can map the physical
 997 dynamics of topological resonance and sheaf-theoretic gluing, but it is struc-
 998 turally incapable of adjudicating the first-person qualitative experience. Our
 999 formalization respects this boundary: the Lean 4 theorems verify the ob-
 1000 jective physical attractors, while the intrinsic “what it is like” to be that
 1001 resonant field remains the domain of the first-person perspective. By locat-
 1002 ing the phase transition at a threshold with a closed-form value — $K_c = 2D$
 1003 for the noisy mean-field Kuramoto model the framework adopts, conditional
 1004 on the von Mises stationary density and not yet linked to a trajectory — it
 1005 aligns with what Anil Seth terms the “real problem” of consciousness (Seth,
 1006 2021). Rather than attempting to bridge an unbridgeable metaphysical gap,
 1007 it offers a topological architecture that can help explain, predict, and physi-
 1008 cally perturb the structural properties of experience.

1009 Future empirical validation of this thermodynamic framework may lever-
 1010 age 3D cortical organoids, offering a controlled biological environment to
 1011 systematically manipulate metabolic substrates and observe the topological
 1012 emergence of these coherent phase transitions.

1013 Code Availability

1014 The complete Lean 4 formalization, encompassing the derivation of all theo-
 1015 rems presented in this manuscript, as well as the accompanying Python sim-

1016 ulations, are publicly available in the GitHub repository: [https://github.](https://github.com/bobaseb/physical_primitives_of_mind)
1017 [com/bobaseb/physical_primitives_of_mind](https://github.com/bobaseb/physical_primitives_of_mind). Additional documentation
1018 and web-accessible proofs are hosted at the project website: [https://bobaseb.](https://bobaseb.github.io/physical_primitives_of_mind/)
1019 [github.io/physical_primitives_of_mind/](https://bobaseb.github.io/physical_primitives_of_mind/).

1020 Acknowledgments

1021 The conceptualization, manuscript drafting, and Lean 4 formalizations pre-
1022 sented in this work were heavily AI-assisted. The authors gratefully acknowl-
1023 edge the use of the Google Antigravity CLI, specifically utilizing Gemini 3.1
1024 and Claude 4.6 Sonnet as collaborative reasoning partners throughout the
1025 research and formalization processes. In addition, DeepSeek V4 Flash and
1026 Claude Opus 5 (routed via OpenRouter) operating through Nous Research’s
1027 Hermes Agent framework provided substantial assistance with code genera-
1028 tion, proof debugging, and structural analysis.

1029 References

- 1030 Pinotsis, D. A., & Miller, E. K. (2022). Beyond dimension reduction: Stable
1031 electric fields emerge from and allow representational drift. *NeuroImage*,
1032 253, 119058.
- 1033 Bauer, L. G., et al. (2022). Quantification of Kuramoto Coupling Between

- 1034 Intrinsic Brain Networks Applied to fMRI Data in Major Depressive Dis-
1035 order. *Frontiers in Computational Neuroscience*, 16.
- 1036 van Bree, S., et al. (2025). Processes and measurements: a framework for
1037 understanding neural oscillations in field potentials. *Trends in Cognitive*
1038 *Sciences*.
- 1039 Tononi, G., Boly, M., Massimini, M., & Koch, C. (2016). Integrated informa-
1040 tion theory: from consciousness to its physical substrate. *Nature Reviews*
1041 *Neuroscience*, 17(7), 450-461.
- 1042 Baars, B. J. (1997). In the Theater of Consciousness: The Workspace of the
1043 Mind. *Oxford University Press*.
- 1044 Dehaene, S., Charles, L., King, J. R., & Marti, S. (2014). Toward a compu-
1045 tational theory of conscious processing. *Current Opinion in Neurobiology*,
1046 25, 76-84.
- 1047 Friston, K. (2010). The free-energy principle: a unified brain theory?. *Nature*
1048 *Reviews Neuroscience*, 11(2), 127-138.
- 1049 Kuramoto, Y. (1984). Chemical Oscillations, Waves, and Turbulence.
1050 *Springer Series in Synergetics*, Vol. 19.
- 1051 Scherf, M. (2026). A Formal Proof of Non-Duality: An Exposition of the
1052 'OnlyOne' System. *Unpublished manuscript, PhilPeople*.

1053 Inoué, T. (2026). On Brain as a Mathematical Manifold: Neural Manifolds,
1054 Sheaf Semantics, and Leibnizian Harmony. *arXiv:2601.15320 [q-bio.NC]*.

1055 Spisák, T., & Friston, K. (2026). Self-orthogonalizing attractor neural net-
1056 works emerging from the free energy principle. *Neurocomputing*.

1057 Pinotsis, D. A., & Miller, E. K. (2023). In vivo ephaptic coupling allows
1058 memory network formation. *Cerebral Cortex*, 33(17), 9877-9890.

1059 Pinotsis, D. A., Fridman, G., & Miller, E. K. (2023). Cytoelectric coupling:
1060 Electric fields sculpt neural activity and "tune" the brain's infrastructure.
1061 *Progress in Neurobiology*, 226, 102465.

1062 Ruffini, G., Salvador, R., Tadayon, E., Sanchez-Todo, R., Pascual-Leone, A.,
1063 & Santarnecchi, E. (2020). Realistic modeling of mesoscopic ephaptic cou-
1064 pling in the human brain. *PLOS Computational Biology*, 16(6), e1007923.

1065 Cunha, G. M., Corso, G., de Sousa, M. P. B., & Lima, G. Z. d. S. (2024).
1066 Can ephapticity contribute to brain complexity?. *PLOS ONE*, 19(12),
1067 e0310640.

1068 Goldwyn, J. H., & Rinzel, J. (2016). Neuronal coupling by endogenous elec-
1069 tric fields: Cable theory and applications to coincidence detector neurons
1070 in the auditory brainstem. *Journal of Neurophysiology*, 115(4), 2033-2051.

1071 McFadden, J. (2002). The conscious electromagnetic information (cemi) field
1072 theory: the hard problem made easy?. *Journal of Consciousness Studies*,
1073 9(8), 45-60.

1074 Pockett, S. (2012). The electromagnetic field theory of consciousness: a
1075 testable hypothesis about the characteristics of conscious as opposed to
1076 non-conscious fields. *Journal of Consciousness Studies*, 19(11-12), 191-223.

1077 Goff, P. (2019). *Galileo's Error: Foundations for a New Science of Con-*
1078 *sciousness*. Pantheon Books.

1079 Seth, A. (2021). *Being You: A New Science of Consciousness*. Penguin.

1080 Seth, A. (2024). Conscious artificial intelligence and biological naturalism.
1081 *PsyArXiv*.

1082 Levin, M. (2019). The computational boundary of a “self”: developmental
1083 bioelectricity drives multicellularity and scale-free cognition. *Frontiers in*
1084 *Psychology*, 10, 2688.

1085 Friston, K., & Kiebel, S. (2009). Predictive coding under the free-energy
1086 principle. *Philosophical Transactions of the Royal Society B: Biological*
1087 *Sciences*, 364(1521), 1211-1221.

1088 Clark, A. (2016). *Surfing Uncertainty: Prediction, Action, and the Embodied*
1089 *Mind*. Oxford University Press.

1090 Hohwy, J. (2013). *The Predictive Mind*. Oxford University Press.

1091 Prigogine, I., & Stengers, I. (1984). *Order Out of Chaos: Man's New Dialogue*
1092 *with Nature*. Bantam Books.

1093 McFadden, J. (2020). Integrating information in the brain's EM field: the
1094 cemi field theory of consciousness. *Neuroscience of Consciousness*, 2020(1),
1095 niaa016.

1096 Fröhlich, F., & McCormick, D. A. (2010). Endogenous electric fields may
1097 guide neocortical network activity. *Neuron*, 67(1), 129-143.

1098 Tasic, B., et al. (2018). Shared and distinct transcriptomic cell types across
1099 neocortical areas. *Nature*, 563(7729), 72-78.

1100 Verkhratsky, A., & Nedergaard, M. (2018). Physiology of Astroglia. *Physio-*
1101 *logical Reviews*, 98(1), 239-389.

1102 Picard, M., & Shirihai, O. S. (2022). Mitochondrial signal transduction. *Cell*
1103 *Metabolism*, 34(11), 1620-1638.

1104 Harnad, S. (1990). The symbol grounding problem. *Physica D: Nonlinear*
1105 *Phenomena*, 42(1-3), 335-346.

1106 Tononi, G., & Koch, C. (2015). Consciousness: here, there and everywhere?
1107 *Philosophical Transactions of the Royal Society B: Biological Sciences*,
1108 370(1668), 20140167.

1109 Aaronson, S. (2014). Why I Am Not An Integrated Information Theorist
1110 (or, The Unconscious Expander). *Shtetl-Optimized*. Retrieved from [https:](https://scottaaronson.blog/?p=1799)
1111 [//scottaaronson.blog/?p=1799](https://scottaaronson.blog/?p=1799).

- 1112 Block, N. (1995). On a confusion about a function of consciousness. *Behav-*
1113 *ioral and Brain Sciences*, 18(2), 227-247.
- 1114 Herculano-Houzel, S. (2009). The human brain in numbers: a linearly scaled-
1115 up primate brain. *Frontiers in Human Neuroscience*, 3, 31.
- 1116 Lange, W. (1975). Cell number and cell density in the cerebellar cortex of
1117 man and some other mammals. *Cell and Tissue Research*, 157(1), 115-124.
- 1118 Yu, F., Jiang, Q. J., Sun, X. Y., & Zhang, R. W. (2015). A new case of
1119 complete primary cerebellar agenesis: clinical and imaging findings in a
1120 living patient. *Brain*, 138(6), e353.
- 1121 Bastos, A. M., Donoghue, J. A., Brincat, S. L., Mahnke, M., Yanar, J., Cor-
1122 rea, J., ... & Miller, E. K. (2021). Neural effects of propofol-induced uncon-
1123 sciousness and its reversal using thalamic stimulation. *eLife*, 10, e60824.
- 1124 Hameroff, S., & Penrose, R. (2014). Consciousness in the universe: A review
1125 of the 'Orch OR' theory. *Physics of Life Reviews*, 11(1), 39-78.
- 1126 Hunt, T., & Schooler, J. W. (2019). The easy part of the hard problem: a
1127 resonance theory of consciousness. *Frontiers in Human Neuroscience*, 13,
1128 378.
- 1129 England, J. L. (2018). Dissipative adaptation in driven self-assembly. *Nature*
1130 *Nanotechnology*, 13(4), 287-295.

- 1131 Ganesh, N., & Parrondo, J. M. R. (2021). Thermodynamic computing with
1132 non-equilibrium systems. *Nature Physics*.
- 1133 Hinton, G. (2022). The Forward-Forward Algorithm: Some Preliminary In-
1134 vestigations. *arXiv preprint arXiv:2212.13345*.
- 1135 Searle, J. R. (1980). Minds, brains, and programs. *Behavioral and Brain*
1136 *Sciences*, 3(3), 417-424.
- 1137 Smart, J. J. C. (1959). Sensations and brain processes. *The Philosophical*
1138 *Review*, 68(2), 141-156.
- 1139 Putnam, H. (1967). Psychological predicates. In *Art, Mind, and Religion*.
1140 University of Pittsburgh Press.
- 1141 Xie, L., Kang, H., Xu, Q., Chen, M. J., Liao, Y., Thiyagarajan, M., ... &
1142 Nedergaard, M. (2013). Sleep drives metabolite clearance from the adult
1143 brain. *Science*, 342(6156), 373-377.
- 1144 Nagel, T. (1986). *The View from Nowhere*. Oxford University Press.
- 1145 Love, B. C. (2026). Consciousness, AI, and the Limits of Scientific Explana-
1146 tion. *arXiv preprint arXiv:2606.00226*.
- 1147 Norton, J. D. (2005). Eaters of the lotus: Landauer's principle and the return
1148 of Maxwell's demon. *Studies in History and Philosophy of Science Part B:*
1149 *Studies in History and Philosophy of Modern Physics*, 36(2), 375-411.

1150 Shenker, O. R. (2000). Logic and entropy. *Philosophy of Science*, 67, S146-
1151 S152.