

Coupling-gated cortical coherence: a conditional framework for unity and self-representation

Sebastian Bobadilla-Suarez
Independent Researcher
`sebastian.bobadilla.s@gmail.com`
ORCID: 0000-0002-5951-0772

6 September 2026

Abstract

Coordinated cortical activity need not describe a single experienced situation. We develop a conditional field account of coherence, compatibility and self-representation. Its demonstrated results establish what coordination leaves undetermined: increasing coupling, decreasing diffusion and a matched common drive can produce identical stationary phase distributions; equal phases permit incompatible local states; and a plasticity control preserves coherence while losing alignment with a specified environmental template. The conditional construction glues compatible spatial measures into one state and makes that state the unique fixed point of a contracting reconstruction map. Lean 4 checks the composition under eight connecting hypotheses and additional structures, with compatibility and accuracy at the glued state supplied as premises. We propose the glued state as a formal correlate of conscious unity and its accurate internal reconstruction as a correlate of a minimal self. The physical hypothesis is that changing extracellular geometry increases field-mediated coupling during awakening. A noisy mean-field phase model gives a square-root stationary onset, but finite ramps delay the response and the distributional

equivalence requires measurements that distinguish recovery mechanisms. Testing the account further requires identifying neural content and self-encoding independently, then determining whether coordination maintains their agreement and whether these relations matter for unified experience.

1 Introduction

How could activity distributed across cortex support one experience? A field account proposes that endogenous extracellular electromagnetic fields partly coordinate spatially separated neural processes (McFadden, 2002, 2020; Pockett, 2012). Ephaptic effects provide a possible interaction mechanism (Fröhlich & McCormick, 2010; Anastassiou & Koch, 2015). Yet even successful coordination leaves a question: what makes the coordinated activity a description of one experienced situation, with a perspective on that situation?

We propose an account with three requirements. *Coherence* concerns whether coupling sustains collective activity. *Compatibility* concerns whether local descriptions agree where their domains overlap. *Self-representation* concerns whether the resulting global state supports an accurate model of itself. The latter two express commitments about consciousness whose motivation and possible dissociations we develop below.

Consider the experience of reaching for a cup. A description of the cup's location and a description of the hand's relation to it can share information about where the cup is while contributing different contents to the episode.

19 Synchronous activity in the two processes would leave any disagreement
20 about that shared quantity unresolved. The formal counterexample isolates
21 this problem using spatial mass profiles at a common phase that disagree
22 on an overlap. We return to the cup example to ask what agreement would
23 mean for perceptual content, and what an internal representation would have
24 to add to place the perceiver in that situation. It is an illustration of the
25 modelling task, not a neural model implemented here.

26 Awakening gives the physical requirement a concrete setting. We hypoth-
27 esize that changing extracellular geometry slowly increases effective coupling
28 through a synchronization threshold. Under stationary branch tracking, a
29 noisy phase model fixes a square-root onset of coherence. The experiment
30 faces two distinct obstacles. Finite ramp speeds delay the response, so a
31 fitted onset exponent depends on the drive. Moreover, successive stationary
32 phase distributions can be identical under increasing coupling, decreasing
33 diffusion or a suitably matched common drive. Recovery must therefore be
34 measured in a way that identifies dynamics and distinguishes causes, rather
35 than inferred from the shape of a coherence trace alone.

36 A plasticity control tests another route to representation. Its squared-
37 drift objective rewards equalizing drift, which in the chosen frequency-cluster
38 model favors connections across the clusters that define the diagnostic tem-
39 plate. Coherence survives while alignment with that template declines. To-
40 gether with the overlap counterexample, this result establishes a limit on
41 what the coordination model explains. Compatible local content needs its

own mechanism and observable.

The contribution has three levels. The separation results establish limits on what coordination explains. The conditional construction gives a unique global state from specified compatible spatial measures and makes it the unique fixed point of a contracting map assumed to reconstruct that state. Lean 4 checks the composition, and toy witnesses establish joint satisfiability of its assumptions. The biological interpretation proposes that the glued state and its internal reconstruction capture conscious unity and a minimal self. This last step requires identifying meaningful neural descriptions and explaining how coordination maintains their agreement. Figure 1 locates the proposed connections, and Section 6 states the physical and thermodynamic assumptions.

2 Coherence: the coupling model

2.1 Dynamics and the stationary threshold

The finite stochastic model is

$$d\theta_i = \left[\omega_i + \sum_j K_{ij}(t) \sin(\theta_j - \theta_i) \right] dt + \sqrt{2D} dW_i, \quad r_N = \left| \frac{1}{N} \sum_i e^{i\theta_i} \right|. \quad (1)$$

Here $D > 0$ is phase diffusion, W_i are independent Wiener processes, and K_{ij} has inverse-time units, with angles dimensionless. The uniform mean-field convention is $K_{ij} = K/N$. Frequency heterogeneity and phase diffusion are

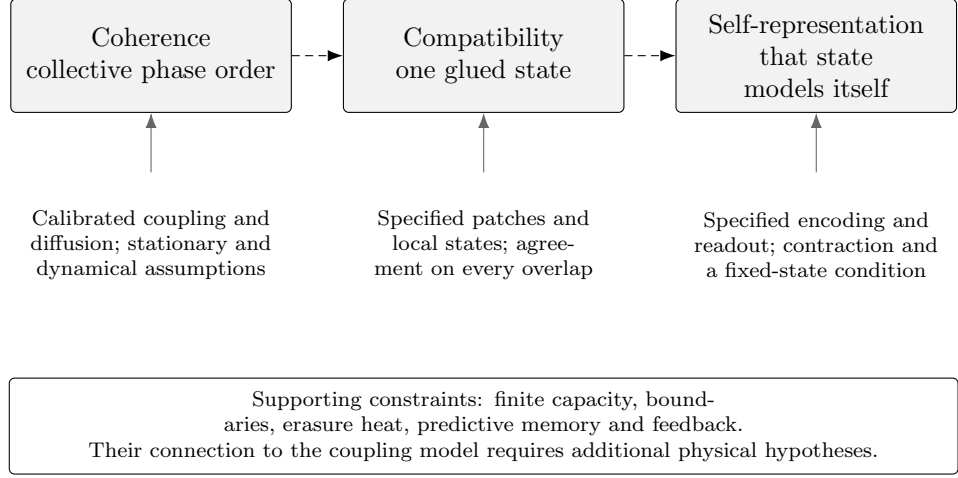


Figure 1: Three requirements of the proposed account. Each requires the inputs shown below it. Dashed arrows mark proposed connections: phase coherence alone supplies neither compatible local descriptions nor an accurate self-model. Thermodynamics constrains a candidate implementation through separate assumptions. This is the explanatory structure of the article; Table 1 records the eight implications consumed by each branch of the formal composition.

60 separate parameters.

61 For identical frequencies in their rotating frame, the stationary ansatz is

$$\rho(\theta) = \frac{e^{a \cos(\theta - \psi)}}{2\pi I_0(a)}, \quad a = \frac{Kr}{D}, \quad r = R(a) := \frac{I_1(a)}{I_0(a)}, \quad (2)$$

62 where I_n are modified Bessel functions and ψ is the collective phase. Lean
 63 verifies normalization, the circular order parameter and self-consistency for
 64 this density. At $K \leq K_c = 2D$, the only nonnegative self-consistent order
 65 parameter is zero. At $K > K_c$, exactly one positive solution coexists with

66 zero; it emerges continuously and increases strictly with K . Near threshold,

$$\frac{r^2}{K - K_c} \rightarrow \frac{1}{D}, \quad r \sim \sqrt{\frac{K - K_c}{D}}. \quad (3)$$

67 These results characterize the stationary branch (Fig. 2). Connecting it
 68 to finite trajectories requires a dynamical mean-field limit and dynamical
 69 selection, neither proved here. Propagation of chaos concerns the first of
 70 these steps (Sznitman, 1991); exchange symmetry alone does not establish
 71 independence, especially when realizations have an unpinned collective phase.

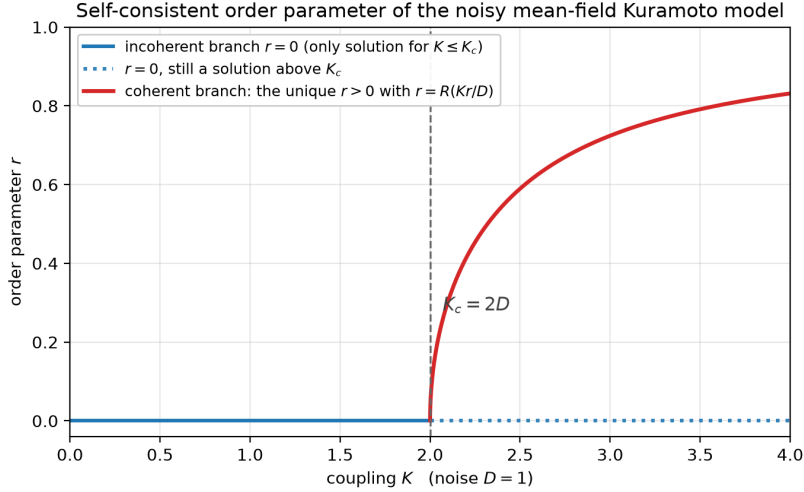


Figure 2: Stationary self-consistency for identical frequencies and phase diffusion $D = 1$. The positive coherent branch appears continuously above $K_c = 2D$ and increases with coupling. Zero remains a solution above threshold. The curve is computed from the integral expressions used in the formalization.

72 A spatial version replaces the finite sum by an integral against a kernel
 73 $K(x, y)$. The substrate measure has total mass one, making the mean-field

74 strength a kernel average. Applying the scalar threshold to a heterogeneous
75 or signed kernel requires a justified reduction to that mean-field description.
76 Whether such a reduction preserves the organization it describes is itself mea-
77 surable. Calcium imaging across wakefulness, sleep and anesthesia reports
78 higher network modularity under unconsciousness at single-cell resolution,
79 with spatially intermixed modules, and no consistent modularity difference
80 at mesoscale, with spatially localized ones (Kiyooka et al., 2026). A reduction
81 of a cortical kernel has to name the scale at which it claims to hold.

82 **2.2 Physical identification and calibration**

83 The cortical hypothesis identifies part of the effective coupling with endoge-
84 nous extracellular electromagnetic interactions. Fields can modulate spike
85 timing (Fröhlich & McCormick, 2010; Anastassiou & Koch, 2015), and field
86 models have been proposed for coordinating distributed neural information
87 (Pinotsis & Miller, 2022; McFadden, 2020). The quantitative question is
88 whether that interaction supplies sufficient coupling relative to diffusion.
89 Changes in phase organization under propofol (Bastos et al., 2021) moti-
90 vate studying this regime, but do not measure the coupling itself.

91 The physical field predicate requires a jointly continuous kernel, a nor-
92 malized substrate, a mean-field strength equal to the spatial average, and no
93 single site providing more than half that strength. Regular mesh refinement
94 establishes convergence of a scalar coupling-energy quadrature; constructing
95 the kernel and relating it to trajectories are separate modelling tasks. The

96 formal witnesses supply toy instances of the predicate.

97 Calibration is essential because a spike-time shift is not a coupling rate.
98 An estimate proportional to $N\Delta tf$, using population size, timing shift and
99 oscillation frequency, is dimensionless. A conversion of the form

$$K_{\text{eff}} = \gamma N s E f \quad (4)$$

100 requires a rate γ , where s is the timing shift per unit field and E the field
101 amplitude. This rate specifies how induced timing shifts persist and accumu-
102 late into phase drift. Supplemental Material explains why the other factors
103 cannot determine it. The biological threshold test requires independent es-
104 timates of the interaction kernel, its aggregation, and phase diffusion in the
105 same cortical regime.

106 **3 Testing coherence through recovery after** 107 **awakening**

108 The proposed control variable is extracellular geometry. Sleep-associated
109 changes in extracellular space and clearance (Xie et al., 2013) motivate a
110 coupling parameter that evolves more slowly than phase dynamics. The
111 extracellular volume fraction is approximately $\alpha = 0.2$ at diffusion tortuosity
112 $\lambda \approx 1.6$ (Syková & Nicholson, 2008); the cited sleep study reports a rise
113 in α from roughly 0.14 to 0.23 (Xie et al., 2013). Below 10 kHz cortical

114 tissue conducts approximately resistively at $0.3\text{--}0.6\text{ S m}^{-1}$ (Barbour, 2017).
 115 In an effective-medium description at fixed transmembrane current density,
 116 a larger extracellular volume fraction raises conductivity and reduces field
 117 amplitude.

118 This gives the coupling-gated hypothesis a directional prediction: contrac-
 119 tion of extracellular space on awakening raises E and, with the other factors
 120 in Eq. (4) held fixed, raises effective coupling. The proposed geometry-
 121 coupling relation and its role in human recovery both require measurement.
 122 The effective-medium argument supplies a sign under stated conditions; cal-
 123 ibration supplies the magnitude and the mapping to the scalar model.

124 If $K(t)$ rises through K_c and the phase distribution tracks the coherent
 125 stationary branch, Eq. (3) gives

$$r(t) \sim \sqrt{\frac{\dot{K}(t_0)}{D}} (t - t_0)^{1/2}, \quad t > t_0, \quad \dot{K}(t_0) > 0. \quad (5)$$

126 The crossing time satisfies $K(t_0) = K_c$. It is a prediction of latency only when
 127 coupling is measured independently of coherence. The subsequent shape
 128 and saturation depend on the later coupling trajectory. Critical slowing
 129 down restricts branch tracking near threshold, while finite populations have
 130 fluctuating order below it. A measurable onset therefore needs an explicit
 131 criterion above that fluctuation floor.

132 3.1 What a recovery distribution can identify

133 Equation (2) gives a concentration–coherence relation with no fitted shape
134 parameter. Testing it requires separate estimators: r from circular resultant
135 length and a from density shape. Estimating a by inverting R from the
136 measured r makes agreement automatic. Even a valid agreement tests a
137 statistical description shared by the von Mises family; it does not determine
138 the coupling medium.

139 There is a stronger limitation. Successive equilibrated phase distributions
140 cannot, by themselves, distinguish three different recovery mechanisms. Let
141 $q(t)$ be a prescribed ratio crossing 2, let $r(q)$ denote its stationary branch,
142 and choose a fixed rate $D_0 > 0$. Increasing coupling with $K(t) = D_0 q(t)$
143 and $D = D_0$ gives the same stationary density as fixed coupling K_0 with de-
144 creasing diffusion $D(t) = K_0/q(t)$. Both have concentration $a = q(t)r(q(t))$.
145 Uncoupled oscillators under a common prescribed drift $-h(t)\sin(\theta - \psi)$, with
146 $h(t) = D_0 q(t)r(q(t))$ and $D = D_0$, have that same density. Their distribu-
147 tions coincide window by window, even though their physical causes differ.
148 Supplemental Material supplies the construction and its numerical feasibility
149 check.

150 The missing observable is the physical clock. Resolved phase increments
151 contain rate information: in the ideal continuous-path model, quadratic vari-
152 ation accumulates at rate $2D$, while conditional increment means constrain
153 drift. This can separate coupling and diffusion under the specified observa-
154 tion model. Distinguishing endogenous coupling from the matched common

drive additionally requires a measured drive or a controlled perturbation, because their mean-field drifts can coincide. Identifying the coupling as electromagnetic then requires evidence specific to that interaction. This hierarchy determines the experiment: distribution, dynamics and physical mechanism are distinct targets.

3.2 Finite ramps change the observable onset

Finite simulations characterize how the stationary prediction can appear in a recovery measurement. The ramp study uses $N = 2000$, $D = 1$ and 32 replicas at four coupling speeds. Escape is the first of three consecutive saved samples at $r \geq 0.2$ after threshold. The supplement gives parameters, estimators and further controls.

Faster ramps delay escape (Fig. 3). Only 5/32 replicas escape within the fastest window. The three slower ramps have mean coupling delays 0.1947, 0.0684 and 0.0253. Their fitted onset exponents run from 0.898 to 0.417, compared with 0.443 to 0.448 from the same estimator on the exact stationary branch at the same coupling grids. At the fastest speed the ensemble mean never reaches the declared fit window; the returned value 0.964 is a local log-slope of a delayed trajectory.

This is delayed passage through a noisy bifurcation (Berglund & Gentz, 2002), with a direct consequence for the awakening test: the apparent exponent depends on the ramp and on the fitting window. A deviation from the square-root law constrains the adiabatic approximation only where branch

177 tracking is independently justified. The sweep does not bracket the specified
 178 deviation-based critical ramp speed.




Figure 3: Finite-rate departures from the stationary branch. Ensemble-mean order versus coupling for four ramp speeds v , with one-standard-error bands across replicas. The dashed vertical line is the static threshold and the dotted horizontal line the finite-size reference. Faster ramps delay escape, changing the apparent recovery onset.

179 Additional controls test the approximations used by the physical model.
 180 Time-step refinement attributes most of the sampled discrepancy from the
 181 stationary order to integration bias. Population-size diagnostics show why
 182 a shared random collective orientation can preserve unconditional pair cor-
 183 relation in the coherent regime. Spatial and signed-network controls exhibit
 184 coherence at specified total coupling, with the refined signed-network cross-
 185 ing bracket $K_{\text{eff}}/D \in [1.958, 2.009]$ still containing 2. These controls delimit

186 finite regimes and numerical resolution; the supplement reports their full
187 evidence. Calibrating the cortical coupling remains the experimental task.

188 **3.3 What the estimators produce on EEG**

189 Public propofol-sedation EEG (Bajwa et al., 2025) provides an exploratory
190 exercise of the concentration and coherence estimators. The implemented
191 analysis pools channels over 100-ms windows, estimating a spatiotemporal
192 phase distribution. Bipolar and common-reference montages produce differ-
193 ent residuals. The bipolar results are compatible with the Bessel relation,
194 but the concentration range is too narrow to distinguish it from its linear
195 approximation at the observed scatter.

196 An informative recording needs a broader concentration range, estima-
197 tor calibration that accounts for dependence, and an observation model for
198 source mixing, referencing and temporal pooling. The sleep-recovery hypoth-
199 esis remains untested by these sedation recordings. Public single-cell calcium
200 imaging of mouse cortex supplies the site count and the state coverage they
201 lack, recording thousands of neurons per session with their spatial coordi-
202 nates across wakefulness, NREM and REM sleep and isoflurane anesthesia
203 (Oomoto et al., 2026). Its 7.65 Hz frame rate lies below the band in which
204 the phase estimators used here operate, so it addresses sites and states rather
205 than the bandwidth those estimators require.

206 **3.4 A protocol for the awakening experiment**

207 The analysis defines six requirements for the physical test.

- 208 **1. Coupling and diffusion in the same units.** Estimate K and D
209 from resolved phase dynamics in the same cortical regime, separating
210 diffusion, frequency heterogeneity and observation noise. A measured
211 drive or controlled perturbation is additionally needed to distinguish
212 endogenous coupling from the matched common-drive alternative.
- 213 **2. Geometry calibrated to coupling.** Establish the relation between
214 an extracellular-geometry observable and effective coupling, with un-
215 certainty. A diffusion-MRI or astrocytic proxy for volume fraction re-
216 quires its own validation; the conductivity argument supplies only the
217 expected sign.
- 218 **3. Adequate phase measurements.** For the proposed spatial estima-
219 tor, use at least 100 sites per estimate, r from resultant length, and
220 a from unweighted regression of log density on the cosine of centered
221 phase. This numerical minimum needs null calibration at the actual
222 sample count, concentration range and dependence. Temporal pooling
223 changes the estimand.
- 224 **4. An independent recovery endpoint.** Measure behavioral recovery
225 separately from the phase trace.
- 226 **5. A predeclared onset criterion.** Fix a coherence threshold above the

finite-size floor and account for finite ramp speed before interpreting an onset exponent.

6. Matched model comparison. Compare coupling-gated recovery with diffusion change, common drive and alternative slow-recovery models under the same observation and filtering assumptions. Stationary distributions alone cannot resolve the equivalence in Section 3.1.

Failure of the concentration–coherence relation over an informative range would constrain the phase-density model. Failure of the calibrated geometry–coupling association would constrain the proposed physical gate. A competing dynamical explanation that better predicts independently measured recovery would count against its role in awakening. These tests address the first requirement of the framework. Even a favorable outcome leaves the question of what the coordinated activity represents.

4 Compatibility: from coordinated activity to a global description

4.1 Why common timing is insufficient

Consider overlapping patches U and V , each carrying a local description. Compatibility requires that restricting either description to $U \cap V$ gives the same object. Agreement concerns the shared domain; descriptions on distinct

246 parts can remain different. For the finite spatial measures used here, equal
247 overlap restrictions give a unique measure on the covered domain. This is
248 the gluing property formalized by a sheaf.

249 In the cup example, suppose a visual description locates the cup on a
250 table and a reaching description relates the hand to the cup. To compare
251 them, a model must express their shared quantity, such as cup location, in
252 a common coordinate system. Restriction would retain that quantity while
253 dropping details peculiar to each description. Agreement on it permits the
254 visual description to contribute surface appearance and the reaching descrip-
255 tion to contribute hand configuration without duplicating the whole situa-
256 tion. Disagreement after coordinate conversion would make the two estimates
257 incompatible at that level of description, however closely timed their activity.

258 The formal counterexample makes the additional requirement visible. Ev-
259 ery patch has the same oscillator phase, yet two local mass profiles assign
260 different densities to a shared site. They cannot be restrictions of a single
261 global measure. Phase order therefore cannot establish compatibility. Con-
262 tinuity of a physical field does not repair disagreement in the descriptions
263 associated with it.

264 Agreement to within a tolerance determines less than a glued state, and
265 more than nothing. Suppose the local profiles differ by at most ε in the
266 uniform distance at every shared site. Nonnegative weights that vanish off
267 their own patch and sum to one at each site then select a global profile lying
268 within ε of every patch on that patch, and any two such selections differ by

269 at most ε , with no dependence on how many patches cover a site; at $\varepsilon = 0$
 270 the selection is the unique glued state. What the tolerance costs is existence:
 271 for every positive ε there are families at that tolerance admitting no global
 272 description at all, so selection returns a choice among states rather than an
 273 extension of the data. A readout that is L -Lipschitz in the state varies by at
 274 most $L\varepsilon$ across selections, which is the resolution condition an interpretation
 275 at scale δ would have to meet. The supplement states these theorems and a
 276 three-site witness that attains the bound. Selection supplies no process by
 277 which local descriptions come to agree, and no rule that picks the cover or
 278 the weights.

279 For a theory of experience, the attraction of gluing is this contribution
 280 of different local contents to one description. The cup example does not yet
 281 instantiate the theorem: receptive territories and content domains must be
 282 related to the spatial patches, and their descriptions must have the required
 283 extension property. In particular, a decoder assigning probability laws to per-
 284 ceptual contents needs a further extension model, because matching overlap
 285 marginals need not determine a unique global joint distribution (Abramsky &
 286 Brandenburger, 2011). The spatial-measure construction provides a precise
 287 case of gluing; a content model must justify its own existence and uniqueness
 288 claims.

289 4.2 Coherent adaptation can lose environmental align- 290 ment

291 An adaptive coupling rule might appear to supply that missing representa-
292 tional structure. The squared-drift control shows why the objective mat-
293 ters. Symmetric coupling conserves mean drift: pairwise sine terms cancel
294 when summed over nodes. Minimizing the sum of squared drifts at that
295 fixed mean therefore rewards equalizing drift across nodes. In the chosen
296 frequency-cluster model, this favors coupling between unlike-frequency clus-
297 ters, whereas the diagnostic template rewards coupling within like-frequency
298 clusters. The objective and the scored structure pull in different directions.

299 For fixed phases, the continuum objective is

$$\sigma(K) = \frac{1}{D} \int \left[\omega(x) + \int K(x, y) \sin(\theta_y - \theta_x) d\mu(y) \right]^2 d\mu(x). \quad (6)$$

300 The formalization constructs the bounded drift operator A and its adjoint,
301 obtains $\nabla\sigma = (2/D)A^*(AK + \omega)$, and proves descent for the specified gradient
302 flow. This is a squared-drift functional; identifying thermodynamic entropy
303 production would require probability-current information.

304 The finite control co-evolves noisy phases and symmetric, nonnegative
305 couplings at fixed total resource. The frequency-cluster template is a di-
306 agnostic covariance of shared latent signals, not a training input. Across
307 three seeds, second-half order remains between 0.9167 and 0.9222, while the
308 within-cluster to between-cluster coupling ratio falls from 0.9824–0.9948 to

0.3507–0.6135 (Fig. 4). The supplement reports random-update, relabelled-
gradient and frozen-kernel controls, including comparisons at overlapping
deformation and coherence ranges and against partitions blind to frequency.

The objective reduction is transient: averaging over complete update
intervals removes the apparent gain seen immediately after updates. A pre-
specified grid of 6 rate–cadence combinations contains no case meeting the
declared 5% sustained-reduction criterion while retaining coherence and los-
ing alignment. The finding is coherent adaptation away from the specified
structure, with no demonstrated sustained cost reduction.

Kernel alignment and overlap compatibility are different observables.
This control tests the narrower suggestion that coherent adaptation under
this objective acquires the designated environmental structure. It fails for
an intelligible reason: the objective rewards a different organization. A
thermodynamically specified learning model can include data in its objective
(Corberi et al., 2026); explaining agreement about the cup would likewise
require a learning rule connected to the relevant content.

4.3 What must be specified and measured

Compatibility enters the construction as a hypothesis about the local states.
For deterministic positive coupling, convergence to a potential minimum is
proved under an explicit initial basin condition; the proposed bridge E78
connects supercritical noisy coherence to relaxation of the specified cover.
Neither step explains how initially inconsistent representations become con-




Figure 4: Coherence persists while coupling loses alignment with the frequency-cluster template. Gradient updates (solid), random updates (dash-dotted), relabelled gradient updates (dash-dot-dotted), and a frozen kernel (dashed) share one phase-noise stream per seed. Panels show the immediate post-update squared-drift objective, phase order, within-cluster to between-cluster coupling, and template distances. The gradient arm preserves order while its coupling ratio falls; growth in both distances mainly reflects kernel norm. The objective panel samples the instant after updates; complete-interval measurements find no sustained reduction in the bounded study.

331 sistent. That is the mechanism a biological account must add.

332 A neural test must first identify the domains and objects whose agreement
333 is at issue. Two decoders should estimate a shared content on overlapping
334 domains, with restriction maps specified before inspecting their agreement.
335 Their discrepancy needs a null under decoding noise, and the procedure must
336 distinguish deliberately constructed compatible and incompatible examples.
337 A common prior can make separate decoders agree artificially; region-specific
338 bias can make them disagree artificially. The supplement develops these re-
339 quirements. An estimator built from phase agreement would miss the prop-
340 erty the counterexample separates.

341 Choosing a cover is also part of choosing a candidate subject. Trivial
342 descriptions can agree while carrying little content, and a suitably chosen
343 partition can conceal disagreement. The scientifically useful choice must
344 therefore be constrained independently by what local activity encodes and
345 which shared quantities can be compared. The uniqueness theorem then
346 answers a precise question about that choice: whether the supplied local
347 descriptions determine one global state. It leaves the empirical selection of
348 domains and descriptions to the model.

349 5 Self-representation: a state that models it- 350 self

351 5.1 The reflexive construction

352 Compatibility yields a global state. Reflexivity adds a relation between that
353 state and an encoding within it. Equip global sections with a complete
354 metric. Under the additional restriction-resonance condition, the encoding
355 of a section s is its restriction to a region U , and the map takes the form

$$F(s) = \text{readout}(s|_U). \quad (7)$$

356 The same region and readout are used across the model's state space. If F
357 contracts and fixes the section s_* supplied by the cover, Banach's theorem
358 makes s_* its unique fixed point (Banach, 1922). The formal bridge E89 sup-
359 plies the contraction law and $F(s_*) = s_*$ for the boundary's encoding-readout
360 composite. Identifying that encoding with restriction is a separate resonance
361 condition; E89 does not supply it. The supercritical order parameter pro-
362 vides an inequality used by the specified contraction bound; constructing the
363 encoding and readout remains a separate obligation.

364 The equality ties the reconstruction to the state formed from the local
365 descriptions. In the cup example, the candidate global state would include
366 the represented scene and the perceiver's relation to it. A local encoding
367 would have to retain enough information for the specified readout to recover

368 that modelled state. An encoding of cup position alone would be insufficient
369 if the readout could not recover the represented hand and viewpoint. The
370 term *predictive* names this reconstruction map; a temporal neural prediction
371 requires a dynamical interpretation.

372 **5.2 Why connect these properties to unity and a min-** 373 **imal self?**

374 We identify the section glued from compatible local states with the unity
375 of a conscious episode, and its accurate internal reconstruction with a mini-
376 mal self. Two features of the episode motivate this interpretation: different
377 contents belong to one situation, and that situation is experienced from a
378 perspective.

379 The unity proposal starts from a distinction between having several con-
380 tents and having several experiences. Seeing the cup’s surface and feeling
381 the reaching hand contribute different contents to one situation. Compat-
382 ibility expresses their agreement about what they share while preserving
383 those differences. An episode can also contain uncertainty or conflicting
384 judgments. A model might represent the cup’s apparent position separately
385 from a judged position and consistently describe the discrepancy. Deciding
386 whether two estimates concern the same quantity or distinct aspects of ex-
387 perience is therefore part of defining the content model. A choice made only
388 to ensure agreement would give the unity proposal no empirical force.

389 The motivation for a minimal self is that the cup is encountered from a
 390 location and in relation to a body that can reach it. The proposed internal
 391 representation concerns the perceiver's condition in that encounter: where
 392 the perceiver is situated, what is being represented and how the body relates
 393 to it. Such a representation could be implicit in the organization of activity;
 394 it need not describe a person verbally or carry an autobiographical narrative.
 395 Requiring it to reconstruct the same global state links the perspective to the
 396 episode whose contents it organizes.

397 The fixed-point condition requires accuracy about the very state contain-
 398 ing the encoding. Reconstruction alone is insufficient for a perspective: a
 399 recording can be reconstructed without being used by a system to represent
 400 its own condition. The proposed interpretation therefore needs evidence that
 401 the encoding helps the system distinguish changes in itself from changes in
 402 its surroundings. Changing the encoded hand or viewpoint while holding
 403 the external scene fixed should alter the system's estimates of its relation
 404 to the cup as the readout predicts. This functional criterion goes beyond
 405 $F(s_*) = s_*$; whether the represented perspective is experienced remains a
 406 further question.

407 Accuracy concerns the variables represented by s_* , rather than every mi-
 408 croscopic physical fact. For a fixed map F with Lipschitz factor $q < 1$ and
 409 fixed point s_* , approximate reconstruction has a quantitative bound:

$$d(s, F(s)) \leq \varepsilon \implies d(s, s_*) \leq \frac{\varepsilon}{1 - q}. \quad (8)$$

At the stipulated rate $q = e^{-(K-2D)\tau/2}$, this radius grows as $2\varepsilon/((K-2D)\tau)$ near threshold, for fixed positive ε and τ . The bound weakens while the exact fixed point remains unique. The supplement gives explicit inequalities and composes reconstruction error with displacement of the state. A cortical realization must specify the represented variables, their metric and the required fidelity; changing readouts require a further dynamical model.

We propose reflexivity as necessary because even an unreflective episode has a perspective: the cup appears somewhere relative to the perceiver, whether or not the perceiver attends to themselves. Requiring internal self-encoding on that basis is the philosophical hypothesis. Gluing alone does not require it; the dissociations below test the two commitments separately.

5.3 Which cover and which self-model?

Both uniqueness results are conditional on their objects. Gluing gives one global extension of specified compatible local data; Banach gives one fixed point of a specified contracting map. These solve ambiguity within a chosen construction. They do not choose between different covers, different encodings or different candidate subjects. Integrated information theory makes exclusion an explicit part of its account (Tononi et al., 2016); the present uniqueness results do not settle that selection problem.

The relevant cover and map must be tied to neural organization by evidence independent of the desired conclusion. Local descriptions should predict the contents attributed to their domains, and their restrictions should

432 preserve meaning on shared domains. The encoding and readout should be
433 specified before testing them, with their sensitivity to changes in the pro-
434 posed state measured over a declared range.

435 A constant readout returning one predetermined state contracts and has a
436 fixed point, even when it discards an encoding supplied by restriction. Such a
437 construction fails the sensitivity requirement above. This is why selecting the
438 map requires evidence about its functional role in addition to mathematical
439 existence.

440 5.4 Dissociations that would bear on the identification

441 A test of the full commitment must measure its clauses separately: effective
442 coupling relative to diffusion, overlap compatibility, and the accuracy of the
443 proposed self-encoding, alongside an independent endpoint for experience.

444 A case meeting all three physical and representational conditions with evi-
445 dence against a unified conscious episode would challenge their sufficiency.
446 Unified experience with demonstrably incompatible local descriptions would
447 challenge the unity identification. Unified experience with compatibility but
448 without the specified self-encoding would challenge the reflexive requirement.

449 Anesthetic emergence, dreaming and dissociative states are candidate set-
450 tings because the framework asks whether coordination, unified content and
451 self-representation can separate. Interpreting such evidence requires a de-
452 clared relation between the experience endpoint and report; absence of report
453 alone cannot define absence of an episode. The compatibility observable in

454 the supplement is a specification, and the corresponding neural self-encoding
455 has yet to be identified. These are the main obstacles to testing the inter-
456 pretive commitment.

457 The proposal concerns which state is conscious. It supplies no explanation
458 of why satisfying these relations should be accompanied by experience. Nor
459 do its gluing and fixed-point conditions select an electromagnetic substrate:
460 the cortical field commitment resides in the physical identification and cali-
461 bration. Their multiple realizability leaves the status of other biological and
462 artificial implementations open.

463 **6 Thermodynamic constraints and the con-** 464 **ditional composition**

465 The broader composition constrains one proposed physical implementation
466 through finite-register capacity $H(p) \leq \log |X|$, erasure heat under a specified
467 system–bath model (Landauer, 1961; Norton, 2005; Shenker, 2000), and a
468 separate field-boundary result for disconnected vacuum components (Kibble,
469 1976). The supplement develops these results and their witnesses. The phase
470 threshold and representational argument stand independently of a derivation
471 from finite capacity.

472 The constraint most directly relevant to representation concerns predic-

473 tive memory. For a state X_t driven by an environmental signal S_t ,

$$W_{\text{diss}} \geq k_B T [I(X_t; S_t) - I(X_t; S_{t+1})]. \quad (9)$$

474 Under the no-back-action Markov structure $X_t \rightarrow S_t \rightarrow S_{t+1}$, the bracket
 475 is nonnegative (Still et al., 2012). The bound limits nonpredictive mem-
 476 ory at a given dissipation budget; zero retained information satisfies it too.
 477 Relating it to Landauer heat requires that memory waste be bounded by
 478 erased entropy at the same thermal scale and heat budget. These distinct
 479 information constraints (Rosas et al., 2026) leave content and learning dy-
 480 namics to be specified. Reading either side of such a trade-off from data
 481 is a further problem. In a partially observed signaling model, an agent can
 482 estimate information gathered about a hidden signal and the dissipative cost
 483 of gathering it from the counting statistics of observed trajectories alone, at
 484 finite measurement time and limited resolution (Nicoletti et al., 2026). The
 485 corresponding cortical estimator is not supplied here.

486 An acting agent changes its future inputs, so the passive assumption
 487 requires justification for the variables and interval considered. A feedback
 488 model specifies a policy $A_t \sim \pi(\cdot | X_t)$, world dynamics $S_{t+1} \sim K(\cdot | S_t, A_t)$,
 489 an observation channel and a memory update. For finite states, the exact

490 information balance is

$$\begin{aligned}
I(X_t; S_t) - I(X_t; S_{t+1}) &= I(X_t; S_t \mid S_{t+1}) \\
&\quad - I(X_t; S_{t+1} \mid S_t).
\end{aligned}
\tag{10}$$

491 Both conditional terms are nonnegative. The last term vanishes for the
492 passive model; with feedback, the signed difference can be negative because
493 actions create correlations with the future. Its causal interpretation requires
494 an environmental state sufficient to account for the relevant dynamics.

495 The finite thermodynamic extension instead bounds joint entropy reduc-
496 tion. For an elementary update that holds the controller fixed, strictly
497 positive forward and reverse path probabilities give $\Sigma = D_{KL}(P_F \| P_R) =$
498 $\Delta H(X, S) + \langle Q \rangle / (k_B T) \geq 0$ under local detailed balance, with Q positive
499 into the reservoir (Horowitz and Esposito, 2014; Ito and Sagawa, 2013). This
500 supplies the thermodynamic constraint for the active branch below. A proved
501 two-substep extension composes actuation with a memory update initial-
502 ized from its coordinate-swapped final law. It identifies the output with the
503 perception–action channels and cancels the same intermediate entropy and
504 expected energy. At a common thermal scale, the whole update’s joint en-
505 tropy decrease is bounded by its total heat budget divided by $k_B T$, without
506 assuming stationarity. The supplement gives the balances and reciprocal
507 two-bit witnesses with explicit interaction energy, heat and work. Action
508 and observation alias the existing states; additional physical registers require
509 further accounting.

510 An explicit objective also permits finite policy selection. For a supplied
 511 terminal reward and a finite set of deterministic policies, a nonempty set sat-
 512 isfying the complete update’s heat budget has a reward maximizer. A two-bit
 513 lamp task has a unique feasible optimum that strictly improves on a feasi-
 514 ble constant-action baseline; a stronger policy exceeds the budget. Reward,
 515 heat and work are evaluated on each policy’s actual composed process, with
 516 common initial law, channels, energy and reservoir convention. The base-
 517 line meets the heat budget while missing the task target, so the objective
 518 supplies content that thermodynamics alone does not. This result compares
 519 fixed policies for one update.

520 A separate finite learning model specifies how a policy register adapts. Its
 521 bit selects between the two feasible lamp policies, and its energy decreases
 522 with their known expected reward. An autonomous thermal channel changes
 523 the register, carrying each final distribution into the next update. From the
 524 specified initial preparation, Lean proves strictly increasing expected task
 525 reward and geometric convergence to a stationary mixture. Transition heat
 526 obeys local detailed balance, and cumulative update heat equals the regis-
 527 ter’s finite energy loss at zero external work. The mixture remains below the
 528 optimal policy; individual updates can lower reward, and a different initial
 529 preparation lowers expected reward. This is adaptation toward an encoded
 530 objective. Preparing the register and implementing its readout in an actua-
 531 tor require additional accounting. Task reward and task heat are prospective
 532 expectations under the prepared lamp law; the learning clock executes no

533 task episodes and supplies no spatial coupling convergence or cortical iden-
534 tification.

535 A second learning model supplies no values at all. Which of two actions
536 is rewarded is an environmental bit that no channel of the learner takes as an
537 argument: the register’s update is a function of the observed outcome alone,
538 and its readout, which fixes the action executed, is a function of the register
539 alone. The world returns success with probability $3/4$ for the rewarding
540 action and $1/4$ for the other; success leaves the register where it is and failure
541 redraws it. Acting, observing and updating compose into one channel on the
542 joint law of parameter and register, so reward, heat and entropy are read from
543 the same process. From a uniform prior over both bits, Lean gives every mass
544 of that law at every horizon: each agreeing pair carries $3/8 - \frac{1}{8}(1/2)^n$, so the
545 expected reward of the action actually executed is $3/4 - \frac{1}{4}(1/2)^n$, strictly
546 increasing from $1/2$ and always short of certainty. A frozen register and an
547 uninformative world, in which success has probability $1/2$ whatever is done,
548 both hold $1/2$ at zero cost, so the improvement is bought by the observations
549 rather than by the update rule. Cumulative heat is $(\log 3/4)[1 - (1/2)^N]$, and
550 the register’s two states are energetically degenerate, so that heat is external
551 work. It is irreducibly so: a register energy that made the updates free would
552 force the drive’s heat to be equal at both parameter values, and this drive’s is
553 not. The learner acquires no preference between the actions, each of which it
554 executes half the time at every horizon, but a correlation with the parameter.
555 Preparing the prior, supplying the work, executing further task episodes and

fabricating the readout are separate accounting, and no cortical identification follows.

6.1 What the composition establishes

A continuing finite agent sequences actuation, observation-driven learning and reset on one joint law of an unknown task parameter, a register and a task flag. The register drives the actuator; the observation reads the flag that action changed; the reset prepares its next use. Each stage starts from the preceding output. For a supplied initial work store B , a common energy E and thermal scale θ , the first law and entropy balance telescope to

$$\theta[H(p_0) - H(p_N)] \leq Q_{<N} = W_{<N} + \langle E \rangle_0 - \langle E \rangle_N \leq B + \langle E \rangle_0 - \langle E \rangle_N, \quad (11)$$

provided expected cumulative work stays within B at every prefix and each stage satisfies positive support and local detailed balance. A bit witness retains agreement with the unknown rewarding action above chance after every completed cycle, while every cycle draws a positive minimum of work, including reset. Blinding the observations leaves chance performance and makes the learning stage free. Replacing reset by idle drift preserves some learning but reduces the next cycle's performance. A finite store therefore supports only a bounded prefix of this protocol. This is accounting in expectation: the initial law and store are supplied resources, and the observation is marginalized inside the learning channel. Pathwise stores, a separate sen-

575 sor memory and preparation channels are treated below. Heat is assigned to
 576 the composite learning channel through local detailed balance, and the store
 577 bounds a mean; no cortical identification follows.

578 A store the protocol carries rather than reports supports a different claim.
 579 Let the state include the store's reading β , and let each executed transition
 580 draw w^- from it and receive a declared supply w^+ , with

$$\beta(t) = \beta(s) - w^-(s, t) + w^+(s, t), \quad w^-(s, t) \leq \beta(s) + w^+(s, t) \quad (12)$$

581 on every transition of positive probability out of a state the protocol reaches.
 582 A reading that starts nonnegative and satisfies the second condition at every
 583 stage is nonnegative at every state the protocol reaches, at every horizon.
 584 Averaging the first identity over the executed paths returns the mean ledger,
 585 $\langle \beta \rangle_N = \langle \beta \rangle_0 - W_{<N}^- + W_{<N}^+$, so pathwise solvency implies the expected bound
 586 Eq. (11) assumes. The converse fails: a store that spends a unit or idles with
 587 equal probability covers its expected draw at both prefixes of a two-stage
 588 run, with equality at the second, and is overdrawn on the path of probability
 589 $1/4$ that spends twice. An expected allowance is a bound on a mean and not
 590 a battery. Two conditions fence sustained operation. If the supply covers
 591 each draw as it is made, every horizon is solvent and no bound on cumulative
 592 work is needed. If instead $w^- - w^+ \geq c > 0$ on every reachable transition,
 593 a reading starting no higher than β_0 gives $Nc \leq \beta_0$, which is the pathwise
 594 form of the store bound and constrains each trajectory rather than a mean.

595 A deterministic spend-and-recharge cycle satisfies neither condition and is
 596 solvent at every horizon while the work it has drawn exceeds any declared
 597 allowance, so continuing operation is a claim about the supply. None of this
 598 requires positive support; the generic store accepts its supply as an input.

599 A finite source can be included in the same state. If its usable energy R
 600 pays for each supplied transfer, then on every reachable transition

$$R' = R - w^+, \quad (\beta + R)' = (\beta + R) - w^-. \quad (13)$$

601 Internal replenishment cancels from the combined ledger. Its expected supply
 602 is $W_{<N}^+ = \langle R \rangle_0 - \langle R \rangle_N$, and nonnegative store and source readings give
 603 $W_{<N}^- \leq \langle \beta + R \rangle_0$. If every executed draw before horizon N is at least $c >$
 604 0 and the initial combined reading is at most b , then $Nc \leq b$. A direct
 605 transfer of one unit followed by idling saturates this bound; the cost premise
 606 concerns the finite run. Available energy does not guarantee delivery: a
 607 finite witness starts a source bit charged and its buffer and output load
 608 empty, with a uniform selector bit. The selector controls a source–buffer
 609 swap, followed by a buffer–load swap. Both gates are reversible and conserve
 610 the three energy bits’ total energy. Half the paths deliver a unit to the load;
 611 half leave that energy in the source. The actual mean draw is $1/2$, and
 612 claiming a unit on every path violates the store ledger. An initially empty
 613 source delivers nothing. The indefinitely replenished charger cannot have
 614 any nonnegative finite source coordinate satisfying Eq. (13). This witness

models one uncertain transfer followed by idling; source preparation, gate-control costs, external refuelling and identification with the learning agent's actuator remain separate.

The measurement has a register of its own, and clearing it is an executed operation. Let the state carry a memory coordinate alongside the register and the world, and run four operations in sequence: the readout drives the actuator, a measurement writes the memory from the world the action moved, the register is updated from the memory alone, and a clearing operation drives the memory toward a declared standard law ν while the world is prepared. The update takes the memory and not the world, so two worlds holding the same memory drive the register identically: what the register can learn is what the memory carries. The clearing operation drives three coordinates at once, and its transition log-ratio is the sum of the three coordinates' own, so the memory's share is a named part of that operation's cost rather than a residue. For an erasure that lands on ν whatever it is given, that share is

$$\langle Q_M \rangle = \theta[H(\mu) - H(\nu)] + \theta D_{KL}(\mu||\nu), \quad (14)$$

with μ the memory's marginal before the operation. The first term is Landauer's bound and the second is the cost of applying one fixed channel to a law it is not matched to; Gibbs' inequality turns the identity into the bound. The identity also states what the bound does not. A memory already at ν is free to clear, and one more ordered than ν draws heat out of the reser-

635 voir, so clearing dissipates because of what the memory holds rather than by
 636 construction.

637 A bit witness fixes the numbers without computing the joint law. The
 638 memory's marginal after a measurement is a mixture of measurement outputs
 639 over a world marginal that is itself a mixture of actuator outputs; with every
 640 channel mass $1/4$ or $3/4$, the memory's mass on a set flag lies in $[7/16, 9/16]$ at
 641 every clearing, whatever the agent has learned. At $\nu = (3/4, 1/4)$ the share is
 642 $[3/4 - \mu(0)] \log 3$, hence between $(3/16) \log 3$ and $(5/16) \log 3$ at every cycle:
 643 no amount of learning makes the measurement free. In the first cycle the
 644 memory is exactly uniform and the two terms of Eq. (14) are $(3/4) \log 3 - \log 2$
 645 and $(1/2) \log(4/3)$, summing to $(1/4) \log 3$; the first is positive because $16 <$
 646 27 . Two controls fence the reading. A blind measurement that records
 647 nothing about the world leaves the same marginal and pays the identical heat,
 648 so the cost is set by the memory's marginal and not by what the memory
 649 is about. Replacing the clearing by the symmetric drift costs nothing and
 650 carries the measurement into the next episode, where the erasure returns the
 651 standard state whatever was written. The four operations are stages of one
 652 protocol, so Eq. (11) charges the whole run, clearing included. Fabricating
 653 the gates and any cortical identification remain outside this model; preparing
 654 the prior is charged below by the same identity.

655 The work those operations draw can be carried in the same state. Let a
 656 coordinate R of the agent's own state read a finite charge, let each executed
 657 transition draw that coordinate's loss, and let the clearing spend at least c of

658 it while no operation returns any. If the charge starts no higher than b and
659 stays nonnegative through m complete cycles, then

$$mc \leq b, \tag{15}$$

660 so a finite source funds finitely many clearings. The bound needs a cost that
661 depends on the stage: an agent that pays at one operation in four meets
662 a uniform per-transition cost only at $c \leq 0$, while a charge paid once in
663 each period of p operations sums to mc over m periods however it sits inside
664 the period. A witness whose world carries the task flag and a charge of
665 three units, one spent at each clearing, attains Eq. (15) at three cycles and
666 reads below empty at the sixteenth operation. Its clearings cost between
667 $(3/16)\log 3$ and $(5/16)\log 3$ throughout, by the marginal bounds above, so
668 what the source funds is what Eq. (14) prices. The drawn unit is a declared
669 reading of the state and is not identified with that heat.

670 Two limits here belong to the model rather than to the witness. The
671 clearing is a channel of the memory alone, so it cannot be conditioned on
672 the charge: past the third cycle the agent still clears, still costs at least
673 $(3/16)\log 3$ and draws nothing, so the run is unfunded rather than halted.
674 And an agent all of whose channels have full support cannot be funded at
675 all. A ledgered reading that never rises is then constant between every pair
676 of states, in both directions, so every net draw is zero, and the witness's
677 restricted support is forced by that rather than chosen. Preparing the charge,

678 refuelling it from outside and fabricating the gates remain outside this model.

679 Preparation can be modelled by the same constant channel as memory
680 clearing. For a strictly positive target ν , Eq. (14) gives its mean log-ratio
681 heat from hardware law μ as $\theta[H(\mu) - H(\nu)] + \theta D_{KL}(\mu||\nu)$. Against a uni-
682 form target the two terms cancel, because $D_{KL}(\mu||\text{unif}) = \log |S| - H(\mu)$:
683 the channel has zero mean heat from any starting law. The reservoir inter-
684 pretation uses local detailed balance; gate construction, control work and a
685 funding source require separate accounting. Prepending this channel as stage
686 zero includes preparation in the cumulative balance whenever its hypotheses
687 hold. Once it reaches the declared prior, every subsequent law is the original
688 agent's own. The sensor-memory witness prepares its uniform prior from a
689 different, strictly positive hardware law and satisfies the whole-run entropy
690 bound.

691 Two things a preparation cannot do are properties of its signature. It
692 writes the system coordinate and leaves the parameter's marginal as it found
693 it, so the agent's uncertainty about the world is an input rather than some-
694 thing it prepares; and a preparation that reads neither coordinate produces
695 a product, so it cannot create a correlation between parameter and regis-
696 ter. At $\theta = 1$, preparing a bit law biased to ε from uniform has mean heat
697 $(1/2 - \varepsilon) \log[(1 - \varepsilon)/\varepsilon]$ for $0 < \varepsilon < 1$; this family has no finite upper bound. Its
698 sharp endpoint has zero reverse probabilities and fails the identity's positive-
699 target premise. This is a restriction of the specified channel and heat model;
700 it supplies no cost bound for other implementations of exact preparation,

701 including the charge above.

702 Table 1 lists the connecting hypotheses of two branches. The theorem
703 **chain** uses passive predictive memory; **chain_active** uses a specified finite
704 feedback process. Each accepts eight implications and shares the passage
705 from coarse-graining to the phase model and its proposed representational
706 role. In both, the finite register and the field boundary are independent
707 starting objects.

708 For the active branch, the controller has the register's state type and the
709 fourth node records

$$\theta > 0, \quad \theta[H(p) - H(p')] \leq \langle Q \rangle \leq B, \quad (16)$$

710 where p, p' and $\langle Q \rangle$ belong to that same controlled substep. E34Active sup-
711 plies positive path probabilities, local detailed balance at the register's ther-
712 mal scale θ , and the allocation $\langle Q \rangle \leq B$, with B equal to the register's
713 heat. Landauer's lower bound does not imply this upper-budget allocation
714 or identify erasure with actuation. A resource model of one register supplies
715 it instead: every operation is controlled by that register and exchanges heat
716 with its reservoir at θ , and the register's dissipated heat is the total its op-
717 erations deliver there. An operation that does not increase the joint entropy
718 then has a nonnegative share, by the left inequality of Eq. (16) applied to it,
719 so an operation whose companions all compress has a share no larger than
720 the register's dissipation. That accounting identity and that compressiveness

721 are the model's physical inputs; the allocation follows from them. Composing
 722 the operations into one evolving law with its preparation and power source
 723 is a separate model, and the register–bath model at the end of this sub-
 724 section shows that a microscopic bipartite dynamics supplies neither input.
 725 E45Active requires convergence of the specified coupling energies. One scalar
 726 constitutive model connects those energies to the named feedback step: writ-
 727 ing $d = H(p) - H(p')$, it sets the spatial density to $f_A(x) = g d b(x)$, with
 728 supplied gain $g \geq 0$ and profile b . Regular spatial refinement gives

$$E_n(f_A) \longrightarrow L_A = g d \int_S b d\mu, \quad L_A \leq g \frac{B}{\theta} \int_S b d\mu \quad (17)$$

729 when $\int_S b d\mu \geq 0$. The limit follows from mesh assumptions; the upper
 730 bound uses Eq. (16) for the same process. The constitutive law uses ensemble
 731 entropy reduction and supplies neither a microscopic actuator nor its work
 732 account. Changing the gain or profile at fixed process and heat changes the
 733 coupling energies.

734 A second, microscopic model supplies an actuator and a kernel. The
 735 hardware carries finitely many spatial response modes with fixed profiles ϕ_i
 736 and installation prices u_i . A configuration z of the controller and its switches
 737 sets the mode occupancies $c_i(z)$, and those same occupancies determine both
 738 a reciprocal kernel on the product space and a stored installation energy,

$$K_z(x, y) = \sum_i c_i(z) \phi_i(x) \phi_i(y), \quad U(z) = \sum_i u_i c_i(z). \quad (18)$$

739 A transition $z \rightarrow z'$ changes K only through the modes whose occupancy it
 740 changes, so a pair outside the response of those modes is untouched, and the
 741 transition's external work is $U(z') - U(z) + Q$. The kernel the executed step
 742 installs is the average of K_z over the law that step actually leaves; its change is
 743 the mean of the pathwise changes, and the mean of the transition works splits
 744 into the change in stored energy and the channel's heat. A finite measurable
 745 partition of the substrate with one sample per cell reconstructs the installed
 746 kernel throughout each cell pair, and the cell-mass weighted finite energy
 747 equals the product-measure integral of that reconstruction exactly. On a
 748 compact substrate with continuous profiles, a sample error tending to zero
 749 gives

$$\mathcal{E}_n \longrightarrow \iint K d\mu d\mu. \quad (19)$$

750 The refinement index here is spatial and is independent of the number of ex-
 751 ecuted control updates; convergence in one is not convergence in the other.
 752 The same hardware discharges the passive edge from a declared predictive
 753 system's own joint law, and two declared systems with different dissipated
 754 work give the identical edge, so predictive efficiency selects neither the hard-
 755 ware nor the partition. Mode profiles, prices, occupancy readout, reservoir
 756 convention and substrate are declared hardware data, which no heat bound
 757 supplies. Fabricating that hardware, powering it beyond one prepared up-
 758 date, and identifying the substrate or the modes with a physical field remain
 759 separate requirements.

760 A third finite model examines the register's own bath. One register bit
 761 and one bath bit carry each of the four deterministic maps of the register as
 762 a permutation of the pair, acting on a bath prepared in its ground state: an
 763 injective map moves the register alone, and an erasure swaps the register's
 764 bit into the bath. With bath levels 0 and $2 \log 2$, degenerate register levels, a
 765 uniform register and $\theta = 1$, the bath's mean energy gain along the executed
 766 paths is 0 for each injective map and $\log 2$ for each erasure, and those values
 767 inhabit Landauer's heat equation with the bath supports the gates actually
 768 reach. For an arbitrary finite protocol and an explicitly modelled bath energy
 769 E_B ,

$$\sum_{k < N} \langle \Delta E_B \rangle_k = \langle E_B \rangle_N - \langle E_B \rangle_0, \quad \sum_{k < N} \langle q_k \rangle = \langle E_B \rangle_N - \langle E_B \rangle_0 + \sum_{k < N} \langle q_k - \Delta E_B \rangle, \quad (20)$$

770 where the first identity needs no positive support and supplies no entropy
 771 inequality, and the second is the correction a reported heat ledger carries
 772 when it omits part of that transfer. Three consequences fence what a map-
 773 indexed budget can mean. Executing the same erasing gate twice returns
 774 the $\log 2$ and restores the initial law, so the two shares are $\log 2$ and $-\log 2$
 775 about a total of zero, and a reusable eraser requires its bath prepared again. A
 776 second reversible gate holds the register fixed and excites its bath, performing
 777 the same register map as idling at $2 \log 2$ instead of 0; the work supplied is
 778 $\log 2$ for one erasure, $2 \log 2$ for that drive and zero for two erasures. And a
 779 swap is an involution, so the gates' own transition log-ratios vanish on every

780 realized path while their bath gains $\log 2$: a deterministic reversible gate is
781 not a strictly positive reduced channel, and local detailed balance at any
782 thermal scale assigns it none of the heat its bath receives. Positivity, local
783 detailed balance and the ledger identity are inputs of the resource model
784 rather than consequences of microscopic reversibility. Preparing the bath,
785 supplying the work, and identifying this bath with the reservoir of the ledger's
786 operations remain separate; the numerical values are those of the declared
787 energy gap and preparation.

788 Two structures also enter both theorems: a thermodynamic cover, in-
789 cluding overlap compatibility, and a reflexive boundary specifying the local
790 encoding and readout. Given these and the connecting hypotheses, both
791 conclude **UnifiedSelf**: the section glued on that cover is the unique fixed
792 point of self-prediction. Restriction resonance, needed for the interpretation
793 in Eq. (7), is a separate condition. The result binds the fixed point to the
794 specified local states; compatibility and accuracy at that state are premises,
795 while uniqueness follows from the construction.

796 The joint witnesses combine a double well, a register, a changing uniform-
797 grid mesh and a three-site cortex; the active branch includes the noisy two-bit
798 actuator. They establish that the hypotheses can hold together in distinct toy
799 components. Mesh convergence is supplied independently of the thermody-
800 namic bounds, and no common biological mechanism is established. A second
801 active witness composes the same branch with the bridge discharged from
802 that register's ledger: its two operations deliver $(2/5) \log 2$ and $(3/5) \log 2$ to

Connection	Status	Required assumption
E12	Independent physical premise	A finite-capacity system realizes a field that leaves its vacuum manifold. Capacity and the boundary are independent roots.
E23	Bridge assumption	The boundary is associated with a finite register whose update leaves states unreachable.
E34 / E34Active	Bridge assumption	Passive: a predictive structure at the register’s thermal scale and heat budget, with nonzero wasted memory. Active: positivity, local detailed balance and a register-heat allocation, derived from its ledger identity and other operations’ compressiveness. Preparation, source and store accounts use their own laws; identifying them with register heat remains an input. A reversible register–bath gate realizes erasure without satisfying the positivity or local-detailed-balance premise.
E45 / E45Active	Modelling assumption	Passive: a convergent regime, for hardware holding a declared predictive law. Active: Eq. (17) from a supplied gain, profile and meshes; Eq. (19) from declared response modes.
E56	Physical commitment	A specified physical field realizes the limiting coupling and satisfies the kernel conditions.
E67	Physical commitment	That field operates above the synchronization threshold.
E78	Bridge assumption	Coherence drives a supplied cover to its relaxed equilibrium at coupling at least the named mean-field value.
E89	Modelling assumption	A map mediated by local self-encoding fixes the section glued on that cover and has the Lipschitz factor $q < 1$ used in the exact and approximate-state results.

Table 1: Connecting hypotheses of the passive and active compositions, eight per branch. They share E12, E23 and E56–E89. The cover’s compatibility is supplied as a structure field. Identifying the boundary’s encoding with restriction requires a separate resonance condition. The mathematical objects and node-level results are mapped in Supplemental Material, Table S1.

803 one reservoir, summing to the register's $\log 2$. A third active witness closes
 804 one spatial response mode with a thermal switch: executing the channel
 805 strictly raises a pairwise kernel value, closing the mode costs its stored en-
 806 ergy plus the channel's heat, and uniform partitions of an atomless interval
 807 give a coupling energy that changes between the first two refinements and
 808 converges to the value the field witness uses. The active witnesses reject
 809 a zero heat budget, an incompatible energy limit, a compressive operation
 810 above the register's whole dissipation, a ledger whose other operation draws
 811 heat out of the reservoir, an uncharged installation, and a wrong spatial limit
 812 for the microscopic sequence. Counterexamples also expose the role of over-
 813 lap agreement. The development declares no additional Lean axioms beyond
 814 the foundational axioms used by its libraries; physical premises are explicit
 815 hypotheses and model fields.

816 **7 Discussion**

817 The strongest results concern what a coherence measurement leaves unre-
 818 solved. Different recovery mechanisms share stationary distributions; equal
 819 phases permit incompatible local descriptions; and the tested plasticity ob-
 820 jective sustains coherence while moving away from its diagnostic structure.
 821 These findings make the coupling hypothesis more demanding to test and
 822 prevent a successful synchronization model from doing unexplained repre-
 823 sentational work.

824 The proposed account identifies a compatible, internally represented state
825 with a conscious episode. Its main explanatory gap is the mechanism that
826 would maintain those relations: compatibility and accuracy at the glued state
827 enter the construction as premises. Their mathematical consequences leave
828 open why cortical activity should acquire or preserve them.

829 The most informative next step for that connection is a test of whether
830 coordination helps maintain independently measured content agreement. Af-
831 ter validating separate decoders and their overlap restrictions, a controlled
832 change in coupling or coordination could test whether agreement deterio-
833 rates in the predicted way, with controls for changes in sensory information
834 and decoder quality. Such a result would bear directly on the proposed role
835 of coherence. In parallel, the awakening hypothesis needs its own calibrated
836 geometry–coupling relation and resolved phase dynamics to distinguish a cou-
837 pling increase from diffusion change or common drive. A favorable recovery
838 trace alone would leave the representational connection open.

839 Testing the self clause requires perturbing a candidate encoding and com-
840 paring the resulting changes in represented scene, body and viewpoint with
841 the readout’s predictions. Establishing that functional role would make dis-
842 sociations from unified experience informative. The exact fixed-state ideal-
843 ization must also be assessed against the fidelity and time dependence of the
844 biological process.

845 Five components of the account are not derived here, and the formaliza-
846 tion marks each. Coarse-graining the discrete couplings converges a scalar

847 energy and produces no kernel: the vertex-sited candidate contributes ex-
 848 actly zero continuum coupling energy on an atomless substrate, which is a
 849 no-go for that architecture rather than for kernels, since cell pairs of pos-
 850 itive mass do pass to a kernel limit. The field conversion $K_{\text{eff}} = \gamma N s E f$
 851 carries its physical content in γ , a hypothesis about how long field-induced
 852 timing shifts persist; the oscillation frequency, a phase relaxation rate and
 853 the diffusion constant are candidates that need not agree, and the remaining
 854 factors cannot choose between them. Descent of the squared-drift objec-
 855 tive is proved for the specified flow at fixed phases, whereas the claim that
 856 such adaptation acquires a designated environmental structure is not proved
 857 and fails in the bounded study reported above. The non-identifiability re-
 858 sult concerns stationary distributions, so it limits what a coherence trace
 859 establishes rather than what any measurement can establish; separating the
 860 mechanisms requires resolving dynamics. The thermodynamic results take a
 861 prepared initial law as given and leave the supply of free energy to a sepa-
 862 rate model. These are boundaries of scope and derivation. Reading any of
 863 them as a positive finding about cortex, or as support for scale separation,
 864 maladaptive learning or a limit on observation in general, would overstate
 865 what has been checked.

866 Two of these results bear on formulations that read a generative model
 867 off a system's steady state, as the free-energy principle does (Friston, 2010).
 868 Nothing here refutes that principle on its own terms: the squared-drift func-
 869 tional is not a variational free energy, it is not identified with entropy pro-

duction, and no bridge between them is claimed or needed for any theorem
above. What transfers is narrower. The stationary distributions of increasing coupling, decreasing diffusion and a matched common drive coincide, so a steady-state density does not by itself determine which mechanism produced it, and here the competing mechanisms are written down rather than supposed. Separately, local descriptions agreeing pairwise on their overlaps need not determine one global state, and at every positive tolerance a three-patch family admits no global section at all; a decomposition into local descriptions therefore does not by itself deliver one belief about one world. Both points instantiate identifiability concerns already raised about that principle (Aguilera et al., 2022; Biehl et al., 2021; Bruineberg et al., 2022), in a model where the alternatives are explicit. Neither establishes that cortex fails to perform inference, and the plasticity control above bears on the objective it names and on no other.

Integrated-information, global-workspace and predictive-processing accounts organize unity, access and modelling differently (Tononi et al., 2016; Baars, 1997; Dehaene et al., 2014; Friston, 2010). The present proposal contributes a candidate interaction mechanism and explicit relations to test between local content and internal self-encoding. Its explanatory value will turn on whether those relations can be identified independently and shown to matter for experience. For the cup example, that means explaining how the system maintains one represented situation and how the perceiver’s place within it contributes to the episode.

893 **Data and code availability**

894 The Lean development, simulation scripts, tests and saved numerical sum-
895 maries are maintained in the manuscript source repository. Publication sim-
896 ulation values are regenerated from saved summaries without rerunning pro-
897 duction sweeps. The exploratory EEG analysis uses OpenNeuro dataset
898 ds005620 (Bajwa et al., 2025); preprocessing and selection are described
899 in Supplemental Material. The source repository contains the reproducible
900 analysis and does not redistribute the raw recordings.

901 **Funding and competing interests**

902 This research received no external funding. The author declares no compet-
903 ing interests.

904 **Use of AI tools**

905 AI assistance to this work was substantial and is not separable by task.
906 Several coding assistants and chat interfaces were used in an interleaved way
907 throughout: Google’s Antigravity CLI (Gemini 3.1, Claude 4.6 Opus, Claude
908 4.6 Sonnet), OpenAI’s Codex CLI (Astra, Terra, Luna, Sol), Claude Code
909 (Claude Opus 5), Gemini Pro through its web interface, and Fable (Claude
910 Opus 4) through OpenRouter, accessed with the Hermes CLI. Between them
911 they contributed to the Lean development, the simulation and analysis code,

912 and the drafting and revision of both documents. Because the tools were
913 used in combination and in alternation, no section, proof, figure or result is
914 attributable to any one of them, and none is claimed to be. The author is
915 responsible for all content, including everything a tool produced. Machine
916 checking establishes the stated Lean results under their hypotheses; it does
917 not validate biological interpretations or replace scientific review of generated
918 code and prose.

919 **References**

- 920 Pinotsis, D. A., & Miller, E. K. (2022). Beyond dimension reduction: Stable
921 electric fields emerge from and allow representational drift. *NeuroImage*,
922 253, 119058.
- 923 Tononi, G., Boly, M., Massimini, M., & Koch, C. (2016). Integrated informa-
924 tion theory: from consciousness to its physical substrate. *Nature Reviews*
925 *Neuroscience*, 17(7), 450-461.
- 926 Baars, B. J. (1997). In the Theater of Consciousness: The Workspace of the
927 Mind. *Oxford University Press*.
- 928 Dehaene, S., Charles, L., King, J. R., & Marti, S. (2014). Toward a compu-
929 tational theory of conscious processing. *Current Opinion in Neurobiology*,
930 25, 76-84.

- 931 Friston, K. (2010). The free-energy principle: a unified brain theory?. *Nature*
932 *Reviews Neuroscience*, 11(2), 127-138.
- 933 Aguilera, M., Millidge, B., Tschantz, A., & Buckley, C. L. (2022). How par-
934 ticular is the physics of the free energy principle?. *Physics of Life Reviews*,
935 40, 24-50.
- 936 Biehl, M., Pollock, F. A., & Kanai, R. (2021). A technical critique of some
937 parts of the free energy principle. *Entropy*, 23(3), 293.
- 938 Bruineberg, J., Dolega, K., Dewhurst, J., & Baltieri, M. (2022). The Em-
939 peror’s New Markov Blankets. *Behavioral and Brain Sciences*, 45, e183.
- 940 Pinotsis, D. A., & Miller, E. K. (2023). In vivo ephaptic coupling allows
941 memory network formation. *Cerebral Cortex*, 33(17), 9877-9890.
- 942 Pinotsis, D. A., Fridman, G., & Miller, E. K. (2023). Cytoelectric coupling:
943 Electric fields sculpt neural activity and ”tune” the brain’s infrastructure.
944 *Progress in Neurobiology*, 226, 102465.
- 945 McFadden, J. (2002). The conscious electromagnetic information (cemi) field
946 theory: the hard problem made easy?. *Journal of Consciousness Studies*,
947 9(8), 45-60.
- 948 Pockett, S. (2012). The electromagnetic field theory of consciousness: a
949 testable hypothesis about the characteristics of conscious as opposed to
950 non-conscious fields. *Journal of Consciousness Studies*, 19(11-12), 191-223.

951 McFadden, J. (2020). Integrating information in the brain's EM field: the
952 cemi field theory of consciousness. *Neuroscience of Consciousness*, 2020(1),
953 niaa016.

954 Fröhlich, F., & McCormick, D. A. (2010). Endogenous electric fields may
955 guide neocortical network activity. *Neuron*, 67(1), 129-143.

956 Anastassiou, C. A., & Koch, C. (2015). Ephaptic coupling to endogenous
957 electric field activity: why bother?. *Current Opinion in Neurobiology*, 31,
958 95-103.

959 Pockett, S. (2002). Difficulties with the electromagnetic field theory of con-
960 sciousness. *Journal of Consciousness Studies*, 9(4), 51-56.

961 Vöröslakos, M., et al. (2018). Direct effects of transcranial electric stimulation
962 on brain circuits in rats and humans. *Nature Communications*, 9, 483.

963 Verkhratsky, A., & Nedergaard, M. (2018). Physiology of Astroglia. *Physio-*
964 *logical Reviews*, 98(1), 239-389.

965 Bastos, A. M., Donoghue, J. A., Brincat, S. L., Mahnke, M., Yanar, J., Cor-
966 rea, J., ... & Miller, E. K. (2021). Neural effects of propofol-induced uncon-
967 sciousness and its reversal using thalamic stimulation. *eLife*, 10, e60824.

968 Bajwa, I. J., Nilsen, A. S., Skukies, R., Aamodt, A., Ernst, G., Storm, J.
969 F., & Juel, B. E. (2025). A repeated awakening study exploring the capac-
970 ity of complexity measures to capture dreaming during propofol sedation.
971 *Scientific Reports*, 15, 32012.

- 972 Xie, L., Kang, H., Xu, Q., Chen, M. J., Liao, Y., Thiagarajan, M., ... &
 973 Nedergaard, M. (2013). Sleep drives metabolite clearance from the adult
 974 brain. *Science*, 342(6156), 373-377.
- 975 Syková, E., & Nicholson, C. (2008). Diffusion in brain extracellular space.
 976 *Physiological Reviews*, 88(4), 1277-1340.
- 977 Barbour, B. (2017). Analysis of claims that the brain extracellular impedance
 978 is high and non-resistive. *Biophysical Journal*, 113(7), 1636-1638.
- 979 Norton, J. D. (2005). Eaters of the lotus: Landauer's principle and the return
 980 of Maxwell's demon. *Studies in History and Philosophy of Science Part B:*
 981 *Studies in History and Philosophy of Modern Physics*, 36(2), 375-411.
- 982 Landauer, R. (1961). Irreversibility and heat generation in the computing
 983 process. *IBM Journal of Research and Development*, 5(3), 183-191.
- 984 Kibble, T. W. B. (1976). Topology of cosmic domains and strings. *Journal*
 985 *of Physics A: Mathematical and General*, 9(8), 1387-1398.
- 986 Sakaguchi, H. (1988). Cooperative phenomena in coupled oscillator systems
 987 under external fields. *Progress of Theoretical Physics*, 79(1), 39-46.
- 988 Still, S., Sivak, D. A., Bell, A. J., & Crooks, G. E. (2012). Thermodynamics
 989 of prediction. *Physical Review Letters*, 109(12), 120604.
- 990 Kawai, R., Parrondo, J. M. R., & Van den Broeck, C. (2007). Dissipation:
 991 The phase-space perspective. *Physical Review Letters*, 98(8), 080602.

- 992 Horowitz, J. M., & Esposito, M. (2014). Thermodynamics with continuous
993 information flow. *Physical Review X*, 4, 031015. doi:10.1103/PhysRevX.4.
994 031015.
- 995 Ito, S., & Sagawa, T. (2013). Information thermodynamics on
996 causal networks. *Physical Review Letters*, 111, 180603. doi:10.1103/
997 PhysRevLett.111.180603.
- 998 Corberi, F., dello Russo, S., Messuti, G., Scarpetta, S., & Smaldone, L.
999 (2026). Thermodynamic learning. *arXiv preprint*. arXiv:2609.04732.
- 1000 Banach, S. (1922). Sur les opérations dans les ensembles abstraits et leur ap-
1001 plication aux équations intégrales. *Fundamenta Mathematicae*, 3(1), 133-
1002 181.
- 1003 Shenker, O. R. (2000). Logic and entropy. *PhilSci-Archive preprint*. philsci-
1004 archive.pitt.edu/115.
- 1005 Rosas, F. E., Mediano, P. A. M., Luppi, A. I., Boyd, A. B., Berta, M., Jensen,
1006 H. J., Gastpar, M., & Carroll, S. M. (2026). Four faces of information in
1007 natural systems. *Nature Reviews Physics*. doi:10.1038/s42254-026-00971-4.
- 1008 Sznitman, A.-S. (1991). Topics in propagation of chaos. In *École d'Été de*
1009 *Probabilités de Saint-Flour XIX—1989*, Lecture Notes in Mathematics,
1010 vol. 1464, pp. 165–251. Springer. doi:10.1007/BFb0085169.
- 1011 Berglund, N., & Gentz, B. (2002). Pathwise description of dynamic pitchfork

1012 bifurcations with additive noise. *Probability Theory and Related Fields*,
1013 122, 341–388. doi:10.1007/s004400100174.

1014 Abramsky, S., & Brandenburger, A. (2011). The sheaf-theoretic structure
1015 of non-locality and contextuality. *New Journal of Physics*, 13, 113036.
1016 doi:10.1088/1367-2630/13/11/113036.

1017 Townsend, R. G., Solomon, S. S., Chen, S. C., Pietersen, A. N. J., Martin,
1018 P. R., Solomon, S. G., & Gong, P. (2015). Emergence of complex wave
1019 patterns in primate cerebral cortex. *The Journal of Neuroscience*, 35(11),
1020 4657-4662. doi:10.1523/JNEUROSCI.4509-14.2015.

1021 Xu, Y., Long, X., Feng, J., & Gong, P. (2023). Interacting spiral wave pat-
1022 terns underlie complex brain dynamics and are related to cognitive pro-
1023 cessing. *Nature Human Behaviour*, 7(7), 1196-1215. doi:10.1038/s41562-
1024 023-01626-5.

1025 Nicoletti, G., Di Terlizzi, I., & Busiello, D. M. (2026). Balancing information
1026 and dissipation with partially observed fluctuating signals. *Physical Review*
1027 *Letters*, 137(12), 127101. doi:10.1103/8gx2-rbns.

1028 Kiyooka, D., Oomoto, I., Kitazono, J., Saito, Y., Kobayashi, M., Matsub-
1029 ara, C., Kobayashi, K., Murayama, M., & Oizumi, M. (2026). Single-
1030 cell resolution functional networks during unconsciousness are segre-
1031 gated into spatially intermixed modules. *Cell Reports*, 45(2), 116902.
1032 doi:10.1016/j.celrep.2025.116902.

1033 Oomoto, I., Kiyooka, D., Oizumi, M., & Murayama, M. (2026). Multi-area
1034 single-cell calcium imaging dataset of the mouse cortex across wakefulness,
1035 sleep, and anesthesia. *Scientific Data*. doi:10.1038/s41597-026-08202-2.