

1 Physical Primitives of Mind: Modeling
2 Consciousness via Symmetries and Landauer
3 Erasure

4 Sebastian Bobadilla-Suarez
 sebastian.bobadilla.s@gmail.com

5 31 August 2026

6 **Abstract**

7 We present a mathematical synthesis and formal Lean 4 verifica-
8 tion of the dependencies in a candidate route from established phys-
9 ical primitives to cognitive attractor states, without introducing new
10 fundamental physics. Starting from the finiteness of a localized sys-
11 tem’s phase space, we construct a framework for predictive processing;
12 the spacetime symmetries are the setting rather than a premise, and
13 no theorem consumes them. Spontaneous symmetry breaking gener-
14 ates necessary topological boundaries, and finite phase space sub-
15 jects these boundaries to Landauer’s thermodynamic erasure. Agency

16 arises when systems evolve along paths of least action, minimizing en-
 17 tropy production via internal structural resonance with external per-
 18 turbations. Scaling this dynamic to macroscopic levels aligns with
 19 emerging paradigms in neuroscience—specifically classical electrody-
 20 namic fields (ephaptic coupling)—and Kuramoto phase synchroniza-
 21 tion, culminating in a topological fixed point: a unified “Self.” We
 22 state the claim narrowly, because the subject matter invites a wider
 23 reading than the work supports: no new physics is derived here. Every
 24 physical result the chain invokes is established—Picard–Lindelöf, Kib-
 25 ble’s classification of topological defects, Landauer’s erasure bound,
 26 Gibbs’ inequality, Still’s bound on dissipation by nonpredictive in-
 27 formation, Barbalat’s lemma, Sakaguchi’s threshold, Banach’s fixed
 28 point. What is new is that they are stated in one formal language,
 29 their consistency mechanically checked, and the points where the chain
 30 relies on stipulation rather than derivation located exactly. One condi-
 31 tional theorem composes eight named hypotheses to its endpoint; the
 32 eight unproved steps are its arguments rather than sentences about
 33 it, and they are jointly satisfiable rather than jointly empty. The de-
 34 velopment declares no axioms and every structure carrying physical
 35 content is inhabited by a witness; the synchronization step is proved
 36 from a dynamics for initial data inside an explicit basin and assumed
 37 outside it, and the identification of a global section of the probability
 38 sheaf with the unity of experience is a stipulation, marked as such
 39 throughout rather than absorbed into the derivation. The step from

a coarse-grained coupling to the brain’s endogenous electromagnetic field is likewise an empirical commitment and not a deduction; we separate it from the coarse-graining theorem that precedes it, weaken it to a modulatory-but-integrative role, and give four conditions that would falsify it. The framework is nonetheless falsifiable in one place that is its own and no rival’s: recovery from sudden awakening should follow a delayed sigmoid rather than an exponential, because the variable gating it is astrocytic volume on a timescale of seconds to minutes rather than cortical phase realignment on a timescale of milliseconds; and, more sharply, the concentration and the coherence of the phase distribution must collapse onto a single parameter-free curve, the Bessel ratio, throughout that recovery. As a corollary, we show why rigid von Neumann architectures (e.g., GPUs) are thermodynamically disadvantaged or structurally constrained from reaching these specific continuous thermodynamic attractors.

1 Introduction

The pursuit of a physical theory of consciousness typically involves searching for the biological correlates of mind or treating subjective experience as an abstract computational algorithm. This paper asks a more fundamental question: If we assume only basic physical primitives—continuous symmetries (the Poincaré group), geometry, and thermodynamics—is the emergence of a unified, perceiving entity a natural physical attractor?

62 We propose a framework in which the structural properties of conscious-
 63 ness are organized as candidate physical attractor states. Its narrative runs
 64 from continuous symmetries and boundaries, through dissipation and phys-
 65 ical resonance, to the macroscopic electromagnetic unity proposed by field
 66 theories of mind. This is not a derivation from basic principles: the for-
 67 mal result is a conditional composition schema whose eight unproved arrows
 68 are explicit assumptions. To avoid critiques of “math theatre,” we delin-
 69 eate between literal physical mechanisms (e.g., thermodynamic erasure via
 70 Landauer’s principle, boundaries from spontaneous symmetry breaking) and
 71 theoretical mappings, utilizing Lean 4 formalization as our mathematical
 72 anchor. We adopt a perspective aligned with Russellian Monism: objec-
 73 tive measurements of macroscopic electromagnetic fields constitute the third-
 74 person “outside,” while the system’s internal physical deformation to reach
 75 structural resonance constitutes the first-person intrinsic experience.

76 **What is new here, and what is not**

77 A framework paper about consciousness invites a particular suspicion, and it
 78 is better answered at the start than at the end. **No new physics is derived**
 79 **here.** Every physical result the chain invokes is established, and most of
 80 it is old: Picard–Lindelöf for the well-posedness of the coupled dynamics
 81 (Coddington & Levinson, 1955); Kibble’s classification of topological defects
 82 by the homotopy of the vacuum manifold (Kibble, 1976); Landauer’s bound
 83 on the heat cost of erasure (Landauer, 1961); Gibbs’ inequality for the relative

84 entropy (Cover & Thomas, 2006); Barbalat’s lemma for the asymptotics of a
85 dissipating trajectory (Barbalat, 1959); the threshold $K_c = 2D$ for the noisy
86 mean-field Kuramoto model (Sakaguchi, 1988); the bound of Still, Sivak,
87 Bell and Crooks relating dissipation to nonpredictive information (Still et
88 al., 2012); the identity of Kawai, Parrondo and Van den Broeck equating
89 dissipated work with the relative entropy between forward and time-reversed
90 ensembles (Kawai et al., 2007); and Banach’s fixed-point theorem (Banach,
91 1922). A physicist reading this paper will recognize all of it and would
92 establish the point in three paragraphs if we did not. Conceding it costs
93 nothing.

94 What is new is of a different kind, and it is a smaller claim than a theory
95 of consciousness usually makes. These results have been stated in one formal
96 language, their mutual consistency has been mechanically checked, and the
97 places where the chain relies on stipulation rather than derivation have been
98 located exactly. The value of a machine-checked chain is not that any link
99 is novel; it is that the claim being made *about* the chain is itself checkable,
100 and we state it in the form a reader can check rather than in the form that
101 asks to be believed. Four things. The development declares **no axioms**:
102 **#print axioms** on any result reports only **propext**, **Classical.choice** and
103 **Quot.sound**. Every physical postulate is a field of a class that each model
104 discharges for itself rather than a global assertion, and every such class is
105 inhabited by an explicit witness, so no theorem is vacuous for want of a
106 model. The links form a single *conditional composition theorem* — **chain**,

107 in `Chain.lean` — with **eight named hypotheses**: four formalization gaps,
 108 two modelling assumptions and two physical commitments. Its conclusion is
 109 `UnifiedSelf`: the particular section produced by sheaf gluing is the unique
 110 fixed point of self-prediction, rather than `Unity` and `Self` being witnessed by
 111 unrelated sections. The $n8 \rightarrow n9$ hypothesis explicitly carries both the non-
 112 linear Lipschitz law and this state identification; neither is derived from field
 113 dynamics. The eight hypotheses are jointly satisfiable on one toy substrate
 114 (`chain_nonvacuous`), which rules out inconsistency but is not evidence that
 115 one cortical mechanism realizes the composition. The seams — where physics
 116 stops and interpretation begins — are thereby fixed in the source rather than
 117 negotiable in the prose.

118 The schema exposes two missing arrows. The capacity bound of the first
 119 link has no hypotheses at all, so an implication from it to the existence of a
 120 boundary is equivalent to assuming the boundary outright: those two links
 121 are independent starting points and not consecutive steps. What finiteness
 122 does imply directly is the heat cost, skipping the boundary entirely. Neither
 123 of the two arrows that are theorems is one the framework’s usual narration
 124 draws. Section 2.1 reports the corresponding soundness result: three tempt-
 125 ing global axiom shapes each prove **False**, from one shared root cause.

126 The framework is not thereby unfalsifiable, and one of its predictions is
 127 its own rather than a rival’s. Section 11 argues that recovery from sudden
 128 awakening should follow a *delayed sigmoid* rather than an exponential, be-
 129 cause the variable gating recovery is astrocytic volume and cerebrospinal-fluid

clearance, on a timescale of seconds to minutes, while the cortical rhythms whose realignment is usually taken to constitute the recovery move in milliseconds. That is a claim about functional form; it is fittable against data of a kind already collected; and a purely electrical account of the same recovery predicts a different form. Because a delayed sigmoid is what any slow saturating gate would give, the section also states the sharper claim the same mathematics supports: at every instant of the recovery the concentration and the coherence of the phase distribution must lie on the fixed curve $r = I_1(a)/I_0(a)$, the same curve for every subject, with no free parameters and with the threshold read off the ratio a/r rather than fitted.

2 Formalization in Lean 4

A core contribution of this work is the machine-checked dependency accounting of the framework’s conditional composition schema using the Lean 4 interactive theorem prover. This approach builds on recent efforts to formalize consciousness and brain semantics, such as the ‘OnlyOne’ system (Scherf, 2026) and sheaf-theoretic models of neural manifolds (Inoué, 2026). It is necessary to explicitly distinguish which claims have been mathematically derived as theorems within the proof assistant, and which remain irreducible physical postulates. Figure 1 shows the schema with the status of every node *and of every arrow* marked, and Table 1 states in seven rows what each link assumes. The complete claim-by-claim map to Lean identifiers, which runs to

several pages, is Table S1 in the supplement; the main text does not depend on reading it.

2.1 Soundness of the formalization

Formalization is only as trustworthy as the axioms it rests on, and an axiom that is inconsistent silently makes every theorem downstream of it vacuous while leaving the build green. We report a case of exactly this, found by auditing our own development, because the failure mode is easy to reproduce and we have not seen it discussed in the applied-formalization literature.

The failure has one shape, and naming it is more useful than cataloguing its instances: *an axiom that constrains a symbol it does not itself bind*. Such a declaration is not a postulate about one system. It is a constraint on every model of the class the symbol belongs to, including models built to violate it, and a class that has such a model is thereby inconsistent. Three concrete cases, each of which proves `False` on its own:

- `spontaneous_symmetry_breaking_pointwise_min` took the potential-energy functional as an *implicit* argument occurring only in its hypothesis `PotentialEnergy phi = V v0`. Instantiating that functional with the constant map $\lambda_. V(v_0)$ discharges the hypothesis by `rfl` and yields the conclusion for an arbitrary field. With $V(v) = v^2$, $v_0 = 0$ and $\phi \equiv 1$ this proves $1 = 0$.
- `landauer_heat_eq` equated the free class field `heat_dissipation` with




Figure 1: The conditional composition schema, with eight assumed vertical arrows and the status of each node and of each arrow. *Nodes.* ■ theorem; ■ theorem resting on a class field that each model must discharge; ■ empirical commitment, neither formalized nor derived; ■ stipulation; ■ setting, consumed by no theorem. *Arrows.* A **solid** arrow is a Lean theorem. A **dashed** arrow is a named hypothesis of chain (Chain.lean): an implication between the two node statements which the development does not prove

Claim	Status	What is assumed
A finite phase space has bounded information capacity	Theorem	Finiteness of the state space, and nothing else — no symmetry group, no metric. $H \leq \log X $, with equality only at the uniform law. On an infinite state space the step from “unreachable state” to “erasure” fails, and the failure is exhibited rather than assumed away.
Broken symmetry forces a boundary	Theorem	Substrate connected; vacuum manifold disconnected. No hypothesis mentions energy, a potential, or a symmetry.
Erasure costs heat	Theorem + class field	Global system–bath dynamics injective (Liouville). The heat identity <code>heat.eq</code> is discharged per instance, not asserted for all models.
Dissipation is bounded below by nonpredictive information	Theorem + class field	Still et al.’s bound is discharged per instance; the environment evolves by a Markov kernel the system does not act on.
The discrete coupling energy converges to its continuum integral; the field carrying that coupling is the cortical EM field	Theorem; then empirical	Convergence holds along regular refining triangulations, for the energy functional and not the dynamics. The limit is a scalar: no theorem produces a kernel of two continuum points, and a kernel sited at the triangulation’s vertices contributes zero coupling energy on an atomless substrate. The physical identification is a commitment with stated failure conditions (Section 7).
Above $K_c = 2D$ a unique coherent state exists, and synchronized patches glue into a unique global section	Theorem 10	Von Mises stationary density assumed rather than derived from the Fokker–Planck operator; couplings positive and symmetric, so unfrustrated. Gluing needs a cover at the potential minimum whose synchronized patches agree on overlaps; the first is derived from a convergent trajectory, the second is an instance obligation throughout.
The coherent order parameter carries a unique reflexive fixed point:	Theorem	The self-prediction map is the avatar’s read-out of its own encoding, and its Lipschitz rate

172 a bath-entropy difference *for every instance* of the class. An instance
 173 with zero heat dissipation and a bath that gains one bit gives $0 = \log 2$.

174 • `kl_bound_axiom` asserted $\sigma \geq D_{KL}(P\|Q)/\Delta t$ for *all* P and Q , while σ
 175 is a single real fixed by a class field. Since D_{KL} is unbounded above, a
 176 point mass against a uniform distribution refutes it.

177 Two further declarations of the same shape are worth recording be-
 178 cause their refutation is different in kind. `phase_invariant_periodic` and
 179 `sync_to_section_eq`, as global assertions pinning free fields of `LocalSectionSynchronization`,
 180 are refutable as soon as the probability presheaf admits two distinct global
 181 sections — so the refutation needs no arithmetic at all, only that the sheaf is
 182 not trivial. Neither is a postulate in the development: the periodicity condi-
 183 tion is an obligation of the constructor that builds a cover from an invariant
 184 measure, and `sync_to_section_eq` is a theorem about the global section the
 185 sheaf condition produces (Derivation 5).

186 The rule that prevents it is uniform, and we state it as the rule we follow:
 187 *a physical postulate that mentions a class field must be a field of that class,*
 188 *never a standalone axiom quantified over all instances.* A postulate is then
 189 an obligation each model discharges, not a global claim about every model.
 190 Under that rule the development declares no axioms at all. Two of the
 191 statements above are not postulates in the first place — pointwise vacuum
 192 minimization is proved from $\int f = 0 \wedge f \geq 0 \Rightarrow f = 0$ a.e., and Gibbs’
 193 inequality is Mathlib’s — and the genuinely physical ones, Landauer’s heat

194 identity and the two information bounds, are instance obligations.
 195 This makes non-vacuity the live question, since an uninhabitable class
 196 makes its theorems vacuous just as surely as an inconsistent axiom did.
 197 `Examples.lean` therefore constructs witnesses: a one-bit erasure system with
 198 a genuine bath that discharges `heat_eq` and reproduces $\Delta Q = k_B T \ln 2$; a
 199 perfectly resonant system discharging `kl_bound`; and a scalar field on a one-
 200 point spacetime discharging `ActionPrinciples`. Derivation 5's gluing struc-
 201 tures are the most demanding case and are witnessed too: `Examples.lean`
 202 builds a `ThermodynamicCover` over a three-site substrate carrying the dis-
 203 crete topology and its Borel σ -algebra, covered by two *overlapping* open
 204 patches whose intersection is a single shared site. The local sections are
 205 germ families of a phase-indexed finite measure $\mu_\theta = (1 + \cos \theta)^+ \delta_{\text{mid}}$ pushed
 206 into the sheafification of the probability presheaf; the family is provably
 207 non-constant in θ (`phaseMeasure_not_const`), so the periodicity obligation
 208 of the constructor `LocalSectionSynchronization.ofInvariantMeasure` is
 209 discharged by an argument about `cos` rather than satisfied trivially, and
 210 `thermodynamic_equilibrium` is *derived* from `phase_locked_minimizes_potential`
 211 rather than assumed. `global_section_from_thermodynamics` then fires on
 212 this instance and glues the two patch-local sections into a unique global one.
 213 Two further covers answer the follow-up question of how non-degenerate such
 214 a witness can be made. `Examples.lean` §13 has a negative half worth record-
 215 ing: the phase field can be made non-constant as a function, but not as a
 216 physical state, because `thermodynamic_equilibrium` forces every instance

217 to a phase-locked configuration. `Examples.lean` §14 is built the other way
 218 round, from two patches with independently chosen mass profiles that meet
 219 only on the site they share, and the section the gluing returns is neither of
 220 them. That is the same kind of result as `contracting_implies_const` in
 221 Derivation 6 — a proof that the witness is as strong as the class permits,
 222 rather than a witness that could be improved by trying harder.

223 What the witnesses establish is that the classes are inhabitable, so the
 224 theorems about them are not vacuous. What they do not establish is that
 225 a cortex is such an instance. In the cover above the phase field is already
 226 constant — every patch locked at $\theta = 0$ — which is precisely what makes
 227 the equilibrium obligation dischargeable; getting a physical system into that
 228 configuration is the content of the informal dynamical argument, not of the
 229 Lean development.

230 The same audit turned up a defect of a different kind, worth record-
 231 ing because it is not an axiom problem and would not have been caught
 232 by `#print axioms.mesh_refinement_convergence` — the claim that dis-
 233 crete edge weights converge to the continuous energy as the mesh is refined
 234 — sat in a “Conjectures” section as an unproved `def ...: Prop`. It was
 235 not merely unproved but false. Nothing required the edge regions to cover
 236 the manifold, and $\text{diam } \emptyset = 0$, so the triangulation whose edge regions are
 237 all empty satisfied the mesh hypothesis while contributing nothing to the
 238 sum, forcing $|0 - \int f| < \varepsilon$ for every $\varepsilon > 0$. Its binder was dead as well: it
 239 quantified over a triangulation and then resolved the vertex type by instance

240 search rather than through the bound variable, so the variable was never
 241 used. `Phase2_MeshConvergence.lean` replaces it with a theorem. Conver-
 242 gence is now stated along a *sequence* of triangulations of a named region
 243 S , each required to be regular — edge regions measurable, without self-
 244 loops, disjoint across distinct edges, covering S , and anchored to the sim-
 245 plicial data in the sense that a vertex lies in the closure of every region
 246 of an edge at it. Closure rather than membership is forced: cells are dis-
 247 joint, so a shared vertex belongs to at most one of them. The manuscript’s
 248 discrete energy $\frac{1}{2} \sum_u \sum_v w(u, v) f(u)$, which samples only the first endpoint
 249 and is therefore not symmetric, is shown to equal exactly the midpoint rule
 250 $\sum_{u < v} w(u, v) \frac{1}{2} (f(u) + f(v))$ over unordered edges; without that identifica-
 251 tion, reading the double sum as a Riemann sum would be a category error.
 252 At fixed mesh the error is at most $\varepsilon \mu(S)$ with ε a modulus of continuity of
 253 the integrand, and convergence follows for uniformly continuous integrands
 254 as the mesh shrinks. The regularity conditions pull against each other, so
 255 `Examples.lean` §6 discharges them with the uniform partition of $[0, 1)$ into
 256 half-open cells and checks that the approximation genuinely moves — $\frac{1}{2}$ at
 257 $N = 1$ and $\frac{1}{4}$ at $N = 2$ for the tent function $x \mapsto |x - \frac{1}{2}|$ — rather than
 258 converging trivially. What is proved is convergence, not the $O(1/N^2)$ rate
 259 the accompanying simulation measures; that would need a midpoint error
 260 term with a second derivative, which we have not developed.

261 A third defect was one of bookkeeping across files, and it is the kind that
 262 survives any number of green builds because each half is individually cor-

263 rect. The development carried *two* Kuramoto potentials. Lyapunov descent
 264 along the dynamics, $\dot{V} = -\sum_i \dot{\theta}_i^2$, was proved for the full potential $V(\theta) =$
 265 $-\frac{1}{2} \sum_{i,j} A_{ij} \cos(\theta_i - \theta_j) - \sum_i \omega_i \theta_i$, while the characterization of the minimum
 266 — phase-locked in both directions — and everything downstream of it, in-
 267 cluding Derivation 5’s gluing, concerned the coupling term alone. Nothing
 268 chained them, and the gap was not cosmetic: `kuramoto_potential_unbounded_below`
 269 now proves that as soon as one $\omega_i \neq 0$ the full potential is unbounded below
 270 — walk θ out along ω and the bounded cosine sum cannot compensate — so
 271 it has no minimum for phase-locking to attain, and the sentence “the phase-
 272 locked state minimizes the Lyapunov potential” was, read literally, false of
 273 the functional the descent theorem was about. `Phase4_RotatingFrame.lean`
 274 closes this by formalizing the standard reduction $\theta_i \mapsto \theta_i - \Omega t$. For identical
 275 natural frequencies $\omega \equiv \Omega$, a trajectory maps to a trajectory of the zero-
 276 frequency system — which is exactly the system Derivation 5 already worked
 277 with — the two potentials coincide there, and the descent result transports
 278 onto the coupling potential whose minimizers are the phase-locked states.
 279 Phase differences, phase-locking and the magnitude of the order parameter
 280 are all proved invariant under the change of frame, so nothing observable
 281 is lost in the reduction. `Examples.lean` §7 runs the chain end to end on
 282 two oscillators with a common frequency, discharging every hypothesis. The
 283 restriction is real and we state it plainly: the reduction is exact only for
 284 identical frequencies. A genuine spread of natural frequencies leaves residual
 285 detunings in the reduced system, phase-locking then requires coupling above

286 a critical K_c , and that K_c — Kuramoto’s threshold for a spread of frequen-
 287 cies, not the noise threshold $2D$ of Section 10 — is nowhere in our Lean
 288 development. It is studied numerically, in Figure 5, as Table S1 records.

289 A fourth defect was dimensional, and it concerned the one row of Table S1
 290 that is still numerical rather than proved. `Phase8_ContinuousField.lean`
 291 defines the critical coupling $K_c = 2D$ and a predicate for being above it.
 292 That predicate compared $\int_x \int_y K(x, y)$ against $2D$ — but a double integral
 293 against an unnormalized measure carries the substrate’s mass twice over, so
 294 a constant kernel K on a substrate of total mass m contributed Km^2 . The
 295 comparison therefore depended on how large the substrate was rather than
 296 on how strongly it was coupled, and any kernel whatever could be pushed
 297 above threshold by inflating m . `Examples.lean` §9 exhibits the failure con-
 298 cretely: on a two-site substrate carrying counting measure, a system coupled
 299 at $K = 1/2$ against a threshold of $K_c = 1$ — a factor of two *below* threshold
 300 — has an unnormalized double integral of $2 > 1$, and would satisfy any pred-
 301 icate stated with that integral. The fix is to require the substrate’s measure
 302 to be a probability measure, which is the normalization in which the mean-
 303 field derivation of $K_c = 2D$ is carried out; `mean_field_coupling_const`
 304 then gives that a constant kernel has mean-field strength exactly K , and
 305 `exhibits_phase_transition_const_iff` reduces the predicate to $K > 2D$.
 306 We flag what this does and does not do. It makes the predicate say something
 307 whose truth depends only on the physics, and §9 witnesses it on both sides of
 308 the threshold. It does not by itself prove $K_c = 2D$: the content of that value

309 is the bifurcation of the self-consistency equation $r = I_1(Kr/D)/I_0(Kr/D)$
 310 at $K = 2D$, which is a separate question from whether the predicate is di-
 311 mensionally sound. That bifurcation is proved below, in both directions,
 312 from an assumed stationary density; deriving the density itself needs the
 313 Fokker–Planck operator and is not done.

314 We draw no grand moral from this beyond the obvious one. The value
 315 of a proof assistant here was not that it certified our argument, but that it
 316 made a specific, checkable claim about what had been assumed — and that
 317 claim turned out to be false in a way no amount of reading the prose would
 318 have revealed.

319 **3 Axiom 1: Symmetries as Setting, and the** 320 **Capacity of a Finite Phase Space**

321 Our starting premise relies on fundamental geometry. In our Lean formaliza-
 322 tion (`Phase1Primitives.lean`), continuous spacetime symmetries (e.g., the
 323 Poincaré group) act upon field states, and the invariance of the system’s total
 324 energy under that action is recorded by the class `SymmetryInvariantAction`.
 325 Physically, it is such an invariance that Noether’s theorem converts into the
 326 conservation of energy and momentum. The scope of the *formal* claim is nar-
 327 rower and we state it plainly: the invariance is a class field, i.e. an assumption
 328 a model must discharge, and no conserved quantity is ever constructed from
 329 it. There is no Noether theorem in our Lean, and nothing downstream of this

330 section depends on one. The class is now inhabited — `Examples.lean` §11
 331 exhibits a double-well potential $V(v) = (v^2 - 1)^2$ on a two-point spacetime
 332 with the $\mathbb{Z}_2$ action $v \mapsto \pm v$, discharges the invariance because $(\pm v)^2 = v^2$,
 333 and proves the vacuum manifold $\{-1, 1\}$ degenerate with the two minima ex-
 334 changed by the symmetry, so an energy-minimising field and its image have
 335 equal energy and differ. That is symmetry breaking exhibited rather than
 336 assumed; it is still not a conserved quantity, and $\mathbb{Z}_2$ being discrete there is
 337 no one-parameter family to differentiate along in the first place. Simultane-
 338 ously, discrete permutation symmetries dictate the statistical mechanics of
 339 any physical system, defining a strictly finite volume of phase space.¹

340 Because phase space is finite, any localized physical system possesses a
 341 mathematical limit on the amount of information (distinguishable states) it
 342 can embody at any given moment. That sentence is the section’s contribution
 343 to the chain, and three theorems carry it, in `Phase1_PhaseSpaceCapacity.lean`,
 344 whose only hypothesis about the system is that its state space is finite: no
 345 symmetry group, no metric, no action functional appears anywhere in them.

346 The first is the capacity bound itself. A probability distribution p on a

¹While the Poincaré group and discrete permutations serve as accessible physical ex-
 amples, the underlying Lean formalization does not hardcode specific physics. It defines
 the system at a more primitive mathematical level, relying entirely on generic Pseudo-
 Riemannian manifolds (for spacetime) and finite measure theory (for statistical mechan-
 ics). What the formalization supplies is the vocabulary in which symmetries, conserved
 quantities and phase-space constraints can be *stated*; it does not derive them. All three
 of the classes involved are now inhabited in `Examples.lean` — the $\mathbb{Z}_2$ action and its in-
 variance in §11, the real line with the Euclidean metric in §12 — but the instances are
 witnesses of inhabitability, not of physics: no theorem in the development consumes a
 symmetry group or a metric, and that half of the premise is setting rather than content.
 The other half of the premise is not setting, and the paragraphs below are what it does.

347 finite state space X satisfies $H(p) = -\sum_x p(x) \log p(x) \leq \log |X|$ (`shannon_`
 348 `entropy_le_log_card`), with equality if and only if p is uniform (`shannon_`
 349 `entropy_eq_log_card_iff`). The “only if” half is the part worth stating: it
 350 says $\log |X|$ is the exact capacity rather than a convenient over-estimate, so
 351 a system holding any distribution other than the flat one over its states holds
 352 strictly less than its size suggests. The proof is Gibbs’ inequality against the
 353 uniform law — $\log x \leq x - 1$ applied at $x = 1/(|X|p(x))$, with the $p(x) = 0$
 354 case carrying the slack that makes the equality condition sharp — and it is
 355 elementary. What was missing here was the statement, not the difficulty.

356 The second says the limit binds. Let the boundary absorb a stream of
 357 perturbations drawn from an alphabet P ; a stream of length k has $|P|^k$
 358 possible histories. Once $|X| < |P|^k$, the map from histories to final states
 359 cannot be injective (`absorb_not_injective`), so two distinct histories leave
 360 the system in the same state, and the entropy of the source strictly exceeds
 361 anything the system can hold (`source_entropy_exceeds_capacity`). The
 362 update map is arbitrary and is allowed to be reversible at every individual
 363 step; the collision is forced by counting alone, which is the sense in which the
 364 premise, and not the model of the dynamics, is doing the work.

365 The third is what Derivation 2 consumes, and it is the one place where
 366 finiteness does something no other hypothesis could. On a finite state space,
 367 a map that fails to reach every state is not injective (`is_erasure_of_not_`
 368 `surjective`) — which is exactly the erasure condition Landauer’s bound re-
 369 quires — so any update of a finite phase space leaving some state unreachable

370 dissipates strictly positive heat (`finite_phase_space_dissipates`). On an
 371 infinite phase space the step simply fails: the successor map on $\mathbb{N}$ misses 0 and
 372 is injective. `Examples.lean` §19 records that failure next to the witnesses, so
 373 the finiteness hypothesis is fenced rather than assumed to be load-bearing,
 374 and runs all three results on a single bit: the fair coin attains $\log 2$, a $3/4$ –
 375 $1/4$ bit falls strictly short, two binary perturbations already outrun a one-bit
 376 memory, and §1’s one-bit thermodynamic system pays heat for a state it
 377 cannot reach.

378 None of this derives the finiteness of phase space. Finiteness is assumed,
 379 as it is in any statistical-mechanical treatment, and the Poincaré group re-
 380 mains an illustration in prose. What has been removed is the gap between
 381 assuming the premise and using it.

382 **4 Derivation 1: Symmetry Breaking and the** 383 **Inevitability of Boundaries**

384 A perfectly symmetric universe is the uniform vacuum; no distinct entities
 385 exist. For any dynamic process to occur, the continuous symmetries of the
 386 background field must undergo spontaneous symmetry breaking.

387 When a field drops to a lower, asymmetric energy state, topology dic-
 388 tates the natural creation of defects (such as domain walls or monopoles).
 389 This constraint necessitates the creation of a physical boundary—a localized
 390 topological distinction between an “inside” and an “outside.” While such

391 a defect (e.g., a crack in a crystal lattice) is necessary to define a separate
 392 entity, it is not sufficient for agency. The boundary is relational, defined by
 393 the twisted geometry of the field surrounding it, and must be coupled with
 394 non-equilibrium thermodynamics to exhibit active behaviors.

395 That first sentence is a claim about topology, and it needs a different
 396 argument from the energy claim beside it. The energy half — an energy-
 397 minimizing field sits in the vacuum manifold almost everywhere (`pointwise_`
 398 `vacuum_of_global_min`) — together with the observation that a boundary
 399 defect *obstructs* extension over a contractible interior, produces no defect:
 400 both are statements about a field already given. What produces one is
 401 `exists_notMem_of_connectedComponentIn`: if the substrate is connected
 402 and a continuous field takes values in two *different connected components*
 403 of the vacuum manifold M , then somewhere the field is not in M at all.
 404 That excluded region is the wall. This is the π_0 case of Kibble’s classifica-
 405 tion of topological defects by the homotopy of the vacuum manifold (Kibble,
 406 1976), and the proof is elementary in exactly the way it is elementary in the
 407 physics literature: the continuous image of a connected space is connected, a
 408 connected subset of M lies inside one component of M , and the field’s range
 409 would be such a subset if it never left M . No hypothesis mentions energy, a
 410 potential, or a symmetry; connectedness of the substrate and disconnected-
 411 ness of the vacuum are the whole of it.

412 We formalize the π_0 case and stop there. The higher cases — $\pi_1 \neq 0$
 413 forcing vortex lines, $\pi_2 \neq 0$ forcing monopoles — are the rest of Kibble’s

414 classification and Mathlib’s homotopy-group library would support at least
 415 the next one; they are available and not taken, because the argument this
 416 paper runs uses domain walls and nothing else. The witness is the double
 417 well of `Examples.lean` §11.1, where the vacuum manifold is *computed* to be
 418 exactly $\{-1, 1\}$ (`wellVacuum_eq`) rather than merely known to contain both
 419 minima, and its disconnection is proved in the form the theorem consumes:
 420 no preconnected subset of it contains both, since a preconnected subset of $\mathbb{R}$
 421 is an interval and an interval spanning -1 and 1 contains 0 , which is the top
 422 of the barrier. `wellV_domain_wall` then quantifies over *every* continuous field
 423 on *any* connected substrate that reaches both minima, with no formula for
 424 the field anywhere in the statement. That the class is non-empty is checked
 425 separately, because a theorem quantified over all fields is worthless if none
 426 satisfies its hypotheses: the clipped identity on $\mathbb{R}$ does, and the wall it is
 427 forced to have sits at the origin (`kink_zero_notMem`).

428 **5 Derivation 2: The Thermodynamic Con-** 429 **straint**

430 This newly formed topological boundary is subjected to gauge perturbations
 431 from the “outside” environment. Because the internal phase space is finite
 432 (Axiom 1), it cannot absorb an infinite sequence of incoming perturbations
 433 and must overwrite old internal states.

434 By Landauer’s Principle (Landauer, 1961), the erasure of physical in-

435 formation irreversibly decreases internal entropy, mandating that the sys-
 436 tem exhaust heat into the environment. Philosophers of physics (such as
 437 John Norton (Norton, 2005) and Orly Shenker (Shenker, 2000)) have cri-
 438 tiqued Landauer’s principle, arguing that physicists conflate *logical* erasure
 439 (a many-to-one mapping in abstract computational states) with *thermody-*
 440 *namic* erasure (a many-to-one mapping in physical phase space). They argue
 441 that logical erasure could conceptually be implemented by an injective phys-
 442 ical process without dissipating heat.

443 Our Lean 4 formalization (`Phase3.CombinatorialThermodynamics.lean`)
 444 addresses this critique by defining erasure as a non-injective mapping on the
 445 *physical* state space ($\neg \text{Function.Injective } t$). We formalize the global evolu-
 446 tion of the bipartite universe (System + Bath) as strictly injective, repre-
 447 senting Liouville’s theorem. Lean verifies that heat dissipation is not tied
 448 to abstract logical erasure, but is a mandatory mathematical consequence
 449 of any local physical state space contraction under globally reversible dy-
 450 namics. If a cognitive system’s internal topology forces a contraction of its
 451 physical state space to accommodate structural resonance, heat dissipation
 452 becomes a thermodynamic requirement. The finiteness premise enters at ex-
 453 actly this point. What a modeller can usually check about a physical update
 454 is not that it is many-to-one but that some state has become unreachable;
 455 on a finite phase space those are the same condition, and `finite_phase_`
 456 `space_dissipates` is the resulting statement. That identification preserves
 457 the Norton–Shenker distinction rather than blurring it: an update that keeps

458 the incoming perturbation on the record is injective and dissipates nothing
 459 (`keepRecord_injective`), while one that refreshes the register holding it is
 460 not (`absorbStep_is_erasure`). Losing *which* history occurred is forced by
 461 counting alone and is free; overwriting the register that held it is what is
 462 charged $k_B T \ln 2$. A stable boundary thus constitutes a dissipative structure
 463 (Prigogine & Stengers, 1984), operating as a non-equilibrium steady state
 464 that consumes free energy to pay the thermodynamic cost of erasure. Contin-
 465 uous thermal fluctuations play a constructive role in establishing boundaries
 466 between states while contributing to the system’s flexible, adaptive response.

467 **6 Derivation 3: Prediction as the Thermody-** 468 **namic Cost of Memory**

469 A dissipative structure is driven by its environment. Write X_t for the state
 470 of the structure and S_t for the environmental signal driving it, and let the
 471 signal evolve by a dynamics of its own, $S_t \rightarrow S_{t+1}$, on which the structure
 472 does not act. Two quantities are then defined at each step: the *memory*
 473 $I(X_t; S_t)$, the mutual information between the structure’s present state and
 474 the signal’s present value, and the *predictive information* $I(X_t; S_{t+1})$, which
 475 is what that same state says about the signal’s next value.

476 **What the entropy-production bound does not give.** There is a nat-
 477 ural route from entropy production to predictive processing, and it does

478 not work. Model the environment as a stochastic driving force $P(\text{ext})$ and
 479 the structure's internal transitions as $Q(\text{int})$. Entropy production is then
 480 bounded below by $D_{KL}(P||Q)/\Delta t$ — this is Gibbs' inequality, it is carried
 481 in Lean as an instance obligation, and §3 of the supplement states it. Re-
 482 arranged it reads $D_{KL}(P||Q) \leq \Delta t \sigma$, so a system descending on σ lowers
 483 a *bound* on the divergence, and it is tempting to conclude that the internal
 484 transition probabilities come to mirror the environment's statistical symme-
 485 tries. They do not, because a driven system does not relax to $\sigma = 0$. It
 486 relaxes to a non-equilibrium steady state at strictly positive entropy produc-
 487 tion, and the bound then yields only $D_{KL} \leq \Delta t \sigma_{\text{NESS}}$, a positive number.
 488 The inequality is true and does real work — combined with Gibbs it gives
 489 $\sigma \geq 0$ — but no limit may be read off it, and a framework that wants predic-
 490 tive processing has to get it from somewhere else. It gets it from the bound
 491 below.

492 **Where predictive processing comes from: a published bound.** Still,
 493 Sivak, Bell and Crooks (Still et al., 2012) prove that for a system driven by
 494 an environmental signal the work dissipated over one step of the drive obeys

$$W_{\text{diss}} \geq k_B T [I(X_t; S_t) - I(X_t; S_{t+1})]. \quad (1)$$

495 The bracket is the *nonpredictive* information: what the structure retains
 496 about the signal's present that says nothing about the signal's future. Dis-
 497 sipation is bounded below by it. Memory that fails to predict is not merely

498 useless to the structure; it is thermodynamic waste, and the waste is quan-
499 tified in energy.

500 **Three consequences, and they are theorems rather than readings.**

501 The bound is carried in Lean as a class field, `PredictiveDissipation`.
502 `still_bound (Phase3_PredictiveThermodynamics)`, for the same reason ev-
503 ery other physical postulate here is: its `axiom` form, quantified over all joint
504 laws with the dissipated work fixed, is refutable, and `Examples §18.4` refutes
505 it. What is *derived* is everything downstream.

506 First, *the bracket is non-negative*, and this is proved rather than assumed.
507 If the structure does not act on its environment, then $X_t \rightarrow S_t \rightarrow S_{t+1}$
508 is a Markov chain and the data processing inequality gives $I(X_t; S_{t+1}) \leq$
509 $I(X_t; S_t)$: a system cannot know more about the signal's future than about its
510 present (`predictiveInfo_le_mutualInfo`). Mathlib carries the Kullback-
511 Leibler divergence but no mutual information, so the latter is defined here
512 as the divergence of the joint law from the product of its marginals, which is
513 the standard definition and makes the library's data processing lemma apply
514 directly.

515 Second, *the second law follows* instead of being assumed. Dissipation is
516 bounded below by a quantity now known to be non-negative, so $W_{\text{diss}} \geq 0$
517 (`dissipatedWork_nonneg`). This mirrors the role Gibbs' inequality plays
518 for the entropy-production bound, with one difference that matters: Gibbs'
519 inequality compares a system to an arbitrary reference distribution, whereas

520 the data processing inequality compares the system to its own future, which
 521 is a quantity selection can act on.

522 Third, and this is the point of the replacement, *the inference now runs*
 523 *in the valid direction.* The quantity of interest sits on the small side of
 524 the inequality: $I(X_t; S_t) - I(X_t; S_{t+1}) \leq W_{\text{diss}}/k_B T$. No limit is required.
 525 Any bound on dissipation is immediately a bound on wasted memory, so a
 526 structure held to an energy budget is thereby held to a prediction standard:
 527 this is `nonpredictive_le_dissipation`. In the idealised quasi-static case
 528 $W_{\text{diss}} = 0$, the theorem `predictive_eq_of_no_dissipation` gives that every
 529 retained bit is predictive. This is a weaker statement than the one it replaces,
 530 and it is the one that is true.

531 **The hypothesis is load-bearing, and we exhibit what breaks with-**
 532 **out it.** The data processing step requires that the structure not write to
 533 its environment. In the formalization this is not an assumption stated beside
 534 the theorem but the shape of the evolution itself: the signal's step is applied
 535 to the pair as the parallel kernel $\text{id} \parallel \kappa$, which acts as the identity on the
 536 system coordinate and through κ on the signal coordinate, so no channel
 537 from X_t to S_{t+1} exists. **Examples** §18.3 shows the restriction is not decora-
 538 tion: a Markov kernel on the *pair* which copies the system's state into the
 539 signal carries an independent two-bit law to a perfectly correlated one, rais-
 540 ing mutual information from zero to a positive value. An agent that steers
 541 its environment can indeed know more about that environment's future than

542 about its present, and Still's bound is not about such an agent.

543 **Both regimes are exhibited, not just described.** Examples §18 builds
544 two systems over one correlated two-bit joint law, differing only in the signal's
545 dynamics. When the signal is frozen, every remembered bit is a predictive
546 bit, the nonpredictive information is zero, and the instance is permitted to
547 dissipate nothing — with the memory itself proved strictly positive, so the
548 conclusion is about a system that does remember something. When the signal
549 is redrawn from its stationary law at each step, the predictive information is
550 exactly zero, the *entire* memory is nonpredictive, and the instance is *forced*
551 to dissipate $k_B T$ times its whole mutual information. Establishing that the
552 mutual information of the witness is neither zero nor infinite is where the
553 work of that section lies.

554 **An equality, where the entropy-production route had only a bound.**
555 Kawai, Parrondo and Van den Broeck (Kawai et al., 2007) show that the
556 average dissipated work is exactly $k_B T$ times the relative entropy between the
557 forward and time-reversed phase-space densities. Carried here as a predicate
558 on an instance rather than as a further postulate, it replaces the abstract
559 dissipation field with a quantity an experiment can estimate — the statistical
560 distinguishability of a recording of the system from the same recording run
561 backwards — and yields the statement that the irreversibility of the dynamics
562 is at least the amount of memory that fails to earn its keep (`nonpredictive_`
563 `le_arrow`).

564 **What this does not establish.** It does not deliver the Free Energy Prin-
565 ciple. Still’s bound says that prediction is thermodynamically favoured; it
566 says nothing about hierarchical generative models, variational free energy,
567 or Bayesian inference, and we do not claim to have derived any of them.
568 The Free Energy Principle (Friston, 2010; Friston & Kiebel, 2009; Hohwy,
569 2013; Clark, 2016; Spisák & Friston, 2026) and the General Resonance The-
570 ory of Consciousness (Hunt & Schooler, 2019) are consonant with what is
571 proved here, which is a weaker relation than derivation and the only one
572 we assert. Nor does it solve the symbol grounding problem (Harnad, 1990).
573 Maximal predictive information is compatible with the structure storing a
574 lossy, compressed or otherwise unrecognisable function of the signal; a state
575 that predicts is not thereby a state that means. The narrower claim is the
576 one worth making: whatever internal states a physical structure holds, those
577 that do not help predict its environment are paid for in heat, so the pressure
578 toward states indexed to the environment is energetic rather than semantic.
579 We continue to call the phenomenon *structural resonance*, and there is no
580 teleological drive to “survive” in it: predicting the environment is simply the
581 cheaper of the two ways to hold a memory.

582 **7 Macroscopic Scaling: A Theorem and a** 583 **Commitment**

584 This is the chain’s fourth step, and it is deliberately not numbered as a
585 derivation. Two steps of different kinds meet here and are easy to run to-
586 gether: a coarse-graining argument, which is a piece of mathematics, and
587 the identification of the coarse-grained medium with the brain’s endogenous
588 electromagnetic field, which is an empirical claim about cortex. Presenting
589 the second as though it inherited the licence of the first would make it look
590 like a deduction, and it is not one. Kept apart, the empirical claim is the
591 framework’s central commitment and the place where it is most easily wrong
592 — which is where a claim of this kind belongs.

593 **The formal half: coarse-graining is a theorem, and a** 594 **narrow one**

595 A network of discretely coupled units has a coupling energy that is a sym-
596 metric double sum over pairs. `Phase2.MeshConvergence.lean` proves that
597 this sum converges to a continuous integral: along a sequence of regular
598 triangulations of a region S whose edge regions shrink to points, the dis-
599 crete energy $\frac{1}{2} \sum_u \sum_v \mu(R_{uv}) f(x_u)$ converges to $\int_S f d\mu$ (`mesh_refinement_`
600 `convergence`), witnessed on the uniform partition of $[0, 1)$ in **Examples §6**.
601 The identification that makes this a coarse-graining rather than an arbitrary

602 Riemann sum is `sum_sum_mul_of_symm`: because the summand is symmetric,
603 the double sum is exactly the midpoint rule over unordered edges, so the
604 discrete energy is a quadrature of the continuous one rather than a different
605 quantity that happens to converge to it.

606 What this establishes is worth stating exactly, because it is less than the
607 phrase “coarse-graining, governed by the Renormalization Group” suggests.
608 It is a statement about an *energy functional*, not about a dynamics: nothing
609 here shows that trajectories of the discrete Kuramoto system converge
610 to solutions of a field equation, and the earlier appeal to the Renormaliza-
611 tion Group was decorative — no renormalization group flow is constructed
612 anywhere in this development. The limit object is a scalar density on the sub-
613 strate, in the intended reading the magnitude of the stress-energy tensor, and
614 it is not an electromagnetic field. The continuous object the later sections
615 use, `ContinuousNeuralField`, is *posited* as a structure with a coupling kernel
616 K ; no theorem produces it from the discrete system. The gap between the
617 finite Kuramoto model and its continuum limit is the propagation-of-chaos
618 problem, which we have not addressed and record as open.

619 **The empirical half: which physical field realizes the ker-**
620 **nel**

621 The formal half says that a large network of symmetrically coupled units
622 admits a continuum description with a symmetric kernel. It says nothing




Figure 2: The coarse-graining theorem, exercised numerically. The discrete coupling energy of N units evenly spaced on the ring, with a Gaussian kernel of periodic distance and phase field $\theta(x) = \sin x$, is compared against the continuous integral it is proved to converge to. The error falls as $\mathcal{O}(N^{-2})$, the rate of the midpoint rule, which is the quadrature the Lean statement uses. The theorem asserts convergence, not a rate; the rate is what the simulation adds, and it is the ordinary rate for this quadrature rather than anything the framework predicts. The convergence is of an *energy functional* — nothing here says the discrete trajectories converge to the continuous ones.

623 about which physical field in cortex that kernel describes. Our answer is an
 624 empirical commitment, and we state it as one:

625 *In the mammalian cortex, the field realizing the coarse-grained*
 626 *coupling kernel is the endogenous electromagnetic field.*

627 This is a claim about biology, not a consequence of anything above it.
 628 It is what makes the framework falsifiable and it is what we expect to be
 629 attacked.

630 **What supports it.** Neurons do not only signal through synapses: trans-
631 membrane currents generate extracellular fields that shift the membrane po-
632 tentials of neighbouring cells, a mechanism known as ephaptic coupling. En-
633 dogenous fields of 1–5 mV/mm have been shown *in vitro* to entrain spike
634 timing (Fröhlich & McCormick, 2010; Anastassiou & Koch, 2015); realistic
635 head models put mesoscopic ephaptic effects within the range where they
636 can bias network dynamics (Ruffini & Santarnecchi, 2020); cable-theoretic
637 treatments show field coupling is sufficient for coincidence detection in the
638 auditory brainstem (Goldwyn & Rinzel, 2016); and stable electric fields have
639 been argued to carry information that survives representational drift in the
640 underlying spiking population (Pinotsis & Miller, 2022, 2023a,b; Cunha &
641 Lima, 2024). The field is spatially continuous and superposes contributions
642 from large populations, which is the property the framework needs: it is a
643 medium in which spatially separated activity is combined without a further
644 mechanism for combining it. This is the integrative role the CEMI field
645 theory assigns it (McFadden, 2002, 2020; Pockett, 2012).

646 **What the evidence supports, and what it does not.** The endoge-
647 nous field’s effect on any individual neuron is *modulatory*. Anastassiou and
648 Koch report that endogenous fields of 1–5 mV/mm shift spike timing by one
649 to three milliseconds (Anastassiou & Koch, 2015), an effect that is real, re-
650 producible, and subthreshold: it biases when a cell fires; it does not decide
651 whether it fires. Nothing in the work cited above supports calling the field

652 the *primary* dissipative structure of the system, and we do not claim it. The
653 claim is that the field is *modulatory in its effect on individual neurons and*
654 *integrative in its spatial reach* — weak per synapse, and the only variable in
655 the system that is defined everywhere at once. What the framework needs
656 is the reach, not the strength, and the two are easy to conflate.

657 **Why reach is the property that matters.** A theory of consciousness
658 needs an account of why many local processes make one experience, and most
659 supply it by stipulation or by a measure. The account here is structural, and
660 it is the framework’s strongest argument. The field is the only variable in the
661 system that is *defined everywhere at once*. Unity is not asserted of it; unity
662 is what a field is. That is what makes Derivation 5’s sheaf-theoretic step
663 available at all — local states glue to one global section because there is one
664 object over which gluing can happen — so the answer to the combination
665 problem is a piece of geometry rather than a postulate. Compare “integra-
666 tion” in Integrated Information Theory, which is a number, or “broadcast” in
667 Global Workspace Theory, which is a metaphor about a workspace. It is also
668 what the identification above is carrying: if the coarse-grained coupling is
669 realized by something that is *not* defined everywhere at once — by synaptic
670 connectivity, say — then this argument is unavailable, and the framework
671 loses the property that distinguishes it from its rivals rather than merely
672 losing a mechanism.

673 **Its critics.** The objection with the longest history is that the fields are
674 too weak to matter, and it has been pressed from within the field-theoretic
675 camp: Pockett, who has herself defended an electromagnetic theory of con-
676 sciousness (Pockett, 2012), catalogues its difficulties (Pockett, 2002) — the
677 measurement problem of relating gross field structure to reported experience,
678 the absence of a constant relationship between the two, and the question of
679 whether the field does causal work or merely reflects it. The quantitative case
680 for scepticism has strengthened since. Vöröslakos and colleagues measured
681 intracranial fields during transcranial stimulation in rats and human cadav-
682 ers and found that the great majority of applied current is shunted through
683 the scalp, so that conventional protocols deliver fields at or below the 1 V/m
684 range in cortex (Vöröslakos et al., 2018). That is a bound on *applied* fields
685 rather than endogenous ones, but it constrains how much causal work a field
686 of that magnitude can be doing, and it is the reason we do not claim the
687 field drives spiking.

688 **What would falsify it.** We state the failure conditions rather than leaving
689 them to a referee.

690 1. *Epiphenomenality.* If the endogenous field can be cancelled or shunted
691 — by a conductive medium that nulls the extracellular potential with-
692 out altering transmembrane currents — and neither behaviour nor re-
693 port changes, the identification fails. This is the decisive experiment
694 and it is, in principle, available.

695 2. *A magnitude gap.* If the spatial integration the framework requires
696 demands field effects an order of magnitude above the 1–5 mV/mm
697 measured in cortex, the commitment fails on arithmetic. The relevant
698 number is not whether ephaptic coupling exists but whether its coupling
699 strength suffices to place the system above the critical coupling K_c of
700 Derivation 7.

701 3. *A timescale mismatch in the wrong direction.* The prediction of Sec-
702 tion 11 turns on the field’s dynamics being fast relative to the astrocytic
703 and clearance variables. If recovery from sleep inertia tracks an expo-
704 nential set by the electrical timescale rather than the delayed sigmoid
705 the slow variable predicts, the framework loses the one prediction that
706 is its own.

707 4. *Dissociation of field coherence from reported unity.* If large-scale field
708 coherence is preserved across states in which experiential unity is lost
709 — or absent in states where it is retained — the identification of the
710 global section with unified experience has no physical carrier.

711 A framework that survives all four is not thereby true. It is merely still
712 standing, which is more than most theories of consciousness arrange to be
713 able to lose.

714 **Geometric frustration, and the regime our theorems ac-**
715 **tually cover**

716 The system’s memory capacity is standardly grounded in *geometric frustra-*
717 *tion*: a highly coupled network cannot bring all its nodes to their lowest en-
718 ergy simultaneously, producing a jagged landscape of near-degenerate states
719 with the capacity to store many patterns. That argument is standard and it
720 is not what this development formalizes. Confronting the discrepancy is the
721 honest course, so we do it here rather than in a footnote.

722 Frustration requires competing couplings — in the Kuramoto setting,
723 a sign structure that no phase assignment can satisfy on every edge at
724 once. The formal setting admits this: `KuramotoSystem` carries $A : V \rightarrow$
725 $V \rightarrow \mathbb{R}$, symmetric and otherwise unconstrained, so signed couplings are
726 representable. The question is therefore not whether the development can
727 speak about a frustrated network but *where along the chain* the sign of A
728 is constrained, and that question has an exact answer. Positivity is not a
729 hypothesis distributed through the development; it is a *field of one class*.
730 `ThermodynamicCover` (`Phase5_GlobalSection`) demands `A_pos` of every in-
731 stance, so it is demanded at the node where local sections are glued and
732 at the edge that builds a cover, and it is demanded nowhere upstream of
733 that node. The theorems that feed the class carry it with them — `phase_`
734 `locked_minimizes_potential` assumes $A_{ij} > 0$, and the characterization of
735 the minimizers as the phase-locked configurations inherits it, while the bridge

of `Phase5_EquilibriumBridge` and the convergence theorem of `Phase4_RotatingFrame` §7 require the strictly stronger uniform bound $0 < a \leq A_{ij}$, since a enters the Łojasiewicz constant. A network with all-positive couplings is an unfrustrated ferromagnet whose potential has one minimum up to a global phase shift. It is the opposite of a jagged landscape, and it has no capacity to store multiple patterns at all.

Locating the constraint that precisely narrows the objection, and the narrower version is the one we accept. It is not that the framework describes an unfrustrated brain throughout; the mean-field reduction of Derivation 7 has no sign structure to be frustrated at all, since its coupling is a scalar K rather than a matrix. It is that **one step** of the framework — the gluing of Derivation 5 — describes an unfrustrated system, and every step before it does not care.

The two halves are worth listing by name, because the surviving half is larger than the concession usually made. What holds with no assumption whatever on the sign of the coupling is the entire dissipative account: well-posedness on all of $\mathbb{R}$ (`is_kuramoto_trajectory_exists`, `_unique`); the Lyapunov identity itself, that the potential’s time derivative is minus the summed squared velocity (`dV_dt_le_zero`), which is the theorem the whole descent argument rests on and which mentions no positivity hypothesis; covariance under passage to a rotating frame; monotone descent of the potential (`dynamic_potential_antitone`) and its convergence to a finite limit (`dynamic_potential_tendsto`); finiteness of the dissipation integral; and —

759 through Barbalat’s lemma — that the motion stops along every trajectory
 760 (`velocity_sq_tendsto_zero`). Every one of these is stated for an arbitrary
 761 symmetric A , so a frustrated network runs, dissipates, and comes to rest
 762 inside this development exactly as an unfrustrated one does. What requires
 763 positivity is the identification of the resting place: the three minimization
 764 results of `Phase4_KuramotoDynamics`, which between them characterize the
 765 potential’s minimizers as exactly the phase-locked configurations; the class
 766 field `ThermodynamicCover.A_pos`; and the convergence theorem `kuramoto_`
 767 `tendsto_global_minimum`. *That* the motion stops is frustration-agnostic;
 768 *where* it stops is not, and the claims this framework makes about the des-
 769 tination all live on the positive side of that line. On the continuum side
 770 the accounting differs slightly: `phase_locked_minimizes_entropy_of_symm`
 771 assumes only that the kernel is symmetric and so tolerates signed kernels,
 772 but it is conditional on a dynamically phase-locked field existing — which is
 773 precisely what frustration may prevent.

774 We therefore scope the claim rather than repeat it. The memory capacity
 775 of a frustrated cortical network is not supported by any theorem here, and we
 776 do not assert it as part of the chain. `Phase5_TwistedGluing` tests the natural
 777 repair: let local states agree on overlaps up to a phase and ask whether
 778 their offset cocycle is a coboundary. The result is negative in the physically
 779 immediate case. Locked, twisted and splay configurations all assign a phase
 780 ϕ_i to every patch, hence their offsets are $\phi_i - \phi_j$, automatically a coboundary
 781 (`isCoboundary_of_phaseField`); the family therefore glues after patchwise

782 rotation (`gluesUpToPhase_of_phaseField`). Frustration alone supplies no
783 non-trivial obstruction and hence no content label. A genuine Čech class
784 would require additional transition or gauge data not determined by the
785 oscillator phases, a proper nerve-indexed cochain complex, and non-trivial
786 nerve topology. Those are new modelling ingredients, not consequences of
787 the frustrated Kuramoto dynamics, so they remain outside the chain.

788 **The supporting physiology**

789 The continuous medium is physically supported by astrocytic syncytia—non-
790 neuronal networks joined by gap junctions that provide macroscopic cytoplas-
791 mic continuity for energy and ion flows (Verkhatsky & Nedergaard, 2018).
792 Managing the thermodynamic cost of this state requires sophisticated micro-
793 bioenergetics. Dynamic mitochondrial networks function as deep-level Mi-
794 tochondrial Information Processing Systems (MIPS), transducing biological
795 and environmental information to sustain the coherence of the macroscopic
796 field (Picard & Shirihaï, 2022). Discrete synaptic firings, coupled with mito-
797 chondrial micro-mechanics, generate and sustain the field. Because the exter-
798 nal environment is continuous, achieving maximal thermodynamic structural
799 resonance is optimally realized within this continuous macroscopic medium
800 (van Bree et al., 2025).

801 8 Derivation 5: Global Sections and the Unity 802 of Consciousness

803 The system is now characterized by a continuous macroscopic electrodynamic
804 field generated by millions of coupled micro-components, forced to resonate
805 with the environment. Following Section 7, the coupling regime assumed
806 from here on is the unfrustrated one — all couplings positive — which is
807 what every theorem below requires and what the frustrated network of the
808 memory argument is not.

809 The Kuramoto model (Kuramoto, 1984) of coupled non-linear oscilla-
810 tors guarantees that, past a critical coupling threshold, these distributed
811 local fields undergo a phase transition into unified synchronization (Bauer et
812 al., 2022). Our Lean proofs (`Phase4_KuramotoDynamics.lean`) confirm this
813 phase-locked state minimizes the Lyapunov potential of the system, in the
814 precise sense discussed in Section 2.1: the functional minimized is the cou-
815 pling potential read in the frame rotating with the natural frequency, which
816 `Phase4_RotatingFrame.lean` shows is the honest reduced description of a
817 system of identical oscillators.

818 This phase transition mathematically glues local synchronized states into
819 a singular *Global Section* of a probability sheaf (`Phase5_GlobalSection.lean`).
820 The EM field acts as a single, indivisible topological space, providing the
821 *Unity of Consciousness*—a necessary condition for integrated experience that
822 solves the binding problem. While Global Workspace Theory (Baars, 1997;

823 Dehaene & Marti, 2014) models this unity as a broadcast architecture and
 824 Integrated Information Theory (Tononi & Koch, 2016) quantifies it via Φ ,
 825 our derivation provides a continuous physical mechanism. The contrast is
 826 worth drawing carefully, because both theories are addressing the same ex-
 827 planandum from different directions.

828 We state the scope of the Lean result precisely, because it is narrower
 829 than the prose above. `ThermodynamicCover` demands that the cover’s phase
 830 field minimize the reduced Kuramoto potential, and `potential_min_iff_`
 831 `phase_locked` characterizes those minimizers exactly: they are the phase-
 832 locked configurations, in both directions. Every instance is therefore already
 833 phase-locked (`ThermodynamicCover.phase_locked`), and the class’s single
 834 condition on its local data — that patches at a common phase agree where
 835 they overlap (`section_agrees_of_phase_eq`) — makes the family of local
 836 sections compatible. `probability_glue_unique`, which is the sheaf property
 837 of the probability presheaf and contains no physics at all, then glues that
 838 family into exactly one global section.

839 The shape of the class is what makes this an emergence theorem rather
 840 than a uniqueness theorem, and the difference is easy to lose. Were `LocalSectionSynchronization`
 841 to carry a *global* section as a class field, together with a second field declar-
 842 ing every local section to be a restriction of it, the global object would exist
 843 before the sheaf condition was ever consulted, and `global_section_from_`
 844 `thermodynamics` could do no more than rule out competitors to an object
 845 the instance had supplied. It carries neither. Every field of the class men-

846 tions a single patch or a single overlap; the global section is constructed
 847 (`ThermodynamicCover.invariantMeasure`); and the statement that the lo-
 848 cal sections are restrictions of one common measure is a *theorem* about that
 849 construction (`ThermodynamicCover.sync_to_section.eq`), true because the
 850 sheaf condition says so, not because an instance was required to name a
 851 measure.

852 `Examples.lean` §14 is the witness that the difference is not notational.
 853 Two patches carry two different mass profiles, $(2, 1, 7)$ and $(5, 1, 3)$ across
 854 the three sites, agreeing only at the site they share, where both put mass 1
 855 — which is what discharges the overlap condition, by computing two num-
 856 bers rather than by construction. The section the theorem returns has pro-
 857 file $(2, 1, 3)$: it is neither of the two profiles (`gluedW_ne_leftW`, `gluedW_ne_`
 858 `rightW`), and neither of them restricts correctly to both patches (`leftW_`
 859 `glues_nothing`, `rightW_glues_nothing`). The global state carries informa-
 860 tion that no single patch had. `Examples.lean` §13, built instead through
 861 the constructor `LocalSectionSynchronization.ofInvariantMeasure`, sits
 862 alongside it for the negative result it carries: its phase field is not the con-
 863 stant function — the second patch a full turn ahead of the first — and by
 864 `ThermodynamicCover.phase_locked` that is as far from constant as any in-
 865 stance can get, since no instance can carry phases differing by anything a
 866 measure is allowed to see.

867 Two hypotheses carry physical content. One of them is an obligation on
 868 every instance: that synchronized patches agree on their overlaps is a mod-

869 elling assumption, not a consequence of any dynamics we formalize. The
 870 other, `thermodynamic_equilibrium`, asks that the cover already sit at the
 871 potential minimum, and no theorem says an *arbitrary* system's dynamics
 872 takes it there. “A thermodynamic phase transition into unity” is therefore
 873 doing causal work that no general Lean result supports, though not work
 874 without any support at all: `Examples.lean` §15 exhibits a two-oscillator tra-
 875 jectory that starts a quarter turn out of phase and converges to the global
 876 minimum of the dynamic potential, with the order parameter tending to 1.
 877 That trajectory is not an isolated construction: `kuramoto_cauchy_problem`
 878 proves the Kuramoto initial value problem well posed on all of $\mathbb{R}$, so through
 879 every initial state there passes exactly one trajectory, and the exhibited one
 880 is the only solution through its own. Nor is its settling down peculiar to
 881 it: along every trajectory of a zero-frequency system the dynamic potential
 882 converges, and the total dissipation $\int_0^\infty \sum_i \dot{\theta}_i^2$ is finite and equal to the po-
 883 tential's total drop (`dynamic_potential_tendsto`, `dissipation_integral_`
 884 `tendsto`). The motion also stops: a finite integral does not by itself force
 885 its integrand to zero, but the dissipation rate is Lipschitz in time, so Bar-
 886 balat's lemma applies, and every trajectory satisfies $\sum_i \dot{\theta}_i^2 \rightarrow 0$ (`velocity_`
 887 `sq_tendsto_zero`; Mathlib has no Barbalat, so it is proved here as `tendsto_`
 888 `zero_of_lipschitz_of_integral_le`). That last step — locating where the
 889 motion stops — is taken conditionally, and by a route that avoids the LaSalle
 890 invariance principle, which Mathlib does not carry. LaSalle is the tool for
 891 a system with a decreasing functional and nothing else; this one has more,

892 and the extra structure is what makes the conclusion reachable without it.
 893 The flow is a gradient flow of the dynamic potential, so $\dot{V} = -\sum_i \dot{\theta}_i^2$ is an
 894 identity rather than an inequality; the minimisers of V are characterised ex-
 895 actly; and on the region where every pair of phases lies within a quarter turn
 896 a Polyak–Łojasiewicz estimate holds, $2a^2W \leq (\sum_{ij} A_{ij}) \sum_i \dot{\theta}_i^2$, where W is
 897 the potential’s excess above its minimum and $a > 0$ is a lower bound for the
 898 coupling (`lojasiewicz_estimate`). Its proof is two extremal indices and the
 899 identity $\sin^2 D = (1 - \cos D)(1 + \cos D)$: the site carrying the largest phase
 900 is pulled backwards at rate at least $a \sin D$, where D is the spread, while the
 901 excess is at most $\frac{1}{2}(\sum_{ij} A_{ij})(1 - \cos D)$. The compactness objection dissolves
 902 as well, because the invariant region is produced by the Lyapunov function
 903 rather than assumed: a single pair contributes $\frac{1}{2}a(1 - \cos)$ to the excess, so
 904 an excess below $a/2$ forces every pairwise cosine positive for all time, and
 905 continuity on the connected half-line then prevents any pair from chang-
 906 ing branch (`spread_le_of_small`). Descent plus the estimate integrates to
 907 $W(t) \leq W(0)e^{-ct}$ with $c = 2a^2 / \sum_{ij} A_{ij}$ (`excess_decay`), the rate makes each
 908 $|\dot{\theta}_i|$ exponentially small and hence integrable on $[0, \infty)$, and a trajectory with
 909 integrable velocity converges. The conclusion is `kuramoto_tendsto_global_`
 910 `minimum`: if the phases start within a quarter turn of each other and the initial
 911 excess is below $a/2$, the trajectory converges to a phase-locked configuration
 912 which is a *global* minimiser of the dynamic potential, with order parameter
 913 $r^2 = 1$. That is the configuration `thermodynamic_equilibrium` postulates,
 914 reached rather than assumed. `Examples.lean` §17 fires it on three sites with

915 initial data $(0, 0, \frac{1}{2})$, where the Kuramoto system has no closed-form solu-
 916 tion: the trajectory comes from the existence theorem and its limit from the
 917 convergence theorem, so neither is written down.

918 `Phase5_EquilibriumBridge.lean` closes the gap that remained between
 919 the two results. `ThermodynamicCover.ofConvergentTrajectory` takes a
 920 cover whose phase field is the limit of such a trajectory and returns an in-
 921 stance whose `thermodynamic_equilibrium` field is discharged *by that con-*
 922 *vergence*; the cover’s coupling matrix is the coupling matrix of the dynamics
 923 rather than a separately chosen one, so the system whose equilibrium the
 924 class asserts and the system whose trajectory is run are the same system.
 925 `Examples.lean` §17.1 is the resulting witness: three patches over the three-
 926 site substrate, each carrying its own mass profile on its own open set, above a
 927 phase field that no line of the construction writes down. The trajectory that
 928 reaches it starts at $(0, 0, \frac{1}{2})$, which is provably not phase-locked (`trioCover_`
 929 `start_not_locked`), so on this cover the configuration Derivation 5 requires
 930 is arrived at rather than posited. The state the three patches glue to has
 931 profile $(2, 1, 3)$ and is none of the three profiles the instance carries (`trioW_`
 932 `ne_glued`), so one cover now carries both the emergence reading of the gluing
 933 and the derived reading of the equilibrium. What is still assumed on it is the
 934 overlap agreement, and nothing else.

935 **The step that carries the positivity.** `ThermodynamicCover` demands
 936 one thing more, and it is worth naming here rather than only in the limita-

937 tions section, because this is the step at which it enters the chain. The class
 938 field `A_pos` requires every coupling to be strictly positive. Nothing upstream
 939 of this node requires it: the trajectory exists, the potential descends, the
 940 dissipation integral is finite and the motion stops for an arbitrary symmetric
 941 coupling matrix, signed entries included. The unfrustrated regime is there-
 942 fore not an assumption the framework makes about the brain throughout; it
 943 is an assumption Derivation 5 makes about the configuration it glues, and it
 944 is the one place a frustrated network would part company with the develop-
 945 ment. The consequences are set out under *Geometric frustration* above.

946 The two hypotheses cannot be dropped, and this is the honest scope of
 947 the result. Splay and twisted configurations are equilibria of the same flow,
 948 so “every trajectory reaches the unified state” is false as a general statement,
 949 and the initial-data condition is the standard arc condition under which the
 950 true statement holds. What is claimed is therefore conditional — unity is
 951 reached from configurations already sufficiently coherent — and not that an
 952 arbitrary initial state is carried into unity. What Lean establishes about the
 953 gluing itself is the other half: *given* locally synchronized patches, a global
 954 state exists, is unique, and is produced by the gluing rather than presupposed
 955 by it.

956 Integrated Information Theory begins from phenomenological axioms and
 957 derives a scalar Φ measuring irreducibility of a system’s cause-effect structure,
 958 with the exclusion postulate selecting a single maximum- Φ complex as the
 959 locus of experience. Our sheaf-theoretic global section plays a structurally

960 analogous role — it too picks out a distinguished unified object — but it
961 is arrived at from the opposite end: not postulated from the character of
962 experience, but forced by a thermodynamic phase transition in a physical
963 field. This buys two things IIT does not have. First, the unity condition
964 is checkable from field dynamics rather than requiring the computation of
965 Φ , which is intractable for systems of realistic size and whose definition has
966 undergone repeated revision (Tononi & Koch, 2016). Second, exclusion falls
967 out rather than being postulated: gluing either succeeds on a given cover or
968 it does not, so there is no separate principle needed to rule out overlapping
969 conscious subsystems. It also inherits IIT’s exposures. Critics have argued
970 that Φ can be high in systems with no plausible claim to experience, such
971 as large feedforward-free grids (Aaronson, 2014); our criterion is vulnerable
972 to the mirror objection, since a physical system could in principle satisfy
973 the gluing condition without being cognitively interesting. We take this
974 seriously rather than dismissing it: our reply is that gluing is necessary but
975 not sufficient, and that the topological boundary of Derivation 1 and the
976 reflexive structure of Derivation 6 are additional required conditions. IIT, by
977 contrast, treats Φ as both necessary and sufficient.

978 Global Workspace Theory faces a complementary difficulty. It specifies
979 a functional architecture — a limited-capacity workspace broadcasting to
980 specialized modules — without specifying what physical property of broad-
981 casting makes the broadcast content conscious, leaving the theory open to the
982 charge that it explains access consciousness while leaving phenomenal con-

983 sciousness untouched (Block, 1995). Our framework can be read as supplying
 984 the missing physical substrate for the workspace: the continuous macroscopic
 985 field *is* the medium in which integration is realized, and the $K_c = 2D$ thresh-
 986 old gives a physical condition for ignition rather than a functional stipulation
 987 of it. Where GWT is agnostic about implementation, we make a falsifiable
 988 commitment — unity requires a continuous physical medium, which predicts
 989 that a functionally equivalent discrete broadcast architecture would not be
 990 conscious. We regard this as a genuine cost of the position, not a free ad-
 991 vantage: it is the point at which our framework departs most sharply from
 992 computational functionalism and where it is most readily falsified. Contrast
 993 this derivation with panpsychism: consciousness is not assumed as a funda-
 994 mental property of matter. It requires the initial existence of a topological
 995 defect to establish a bounded “inside” (Derivation 1), followed by a thermo-
 996 dynamic phase transition into unity that occurs when integrated spatial cou-
 997 pling exceeds the thermodynamic noise threshold ($K_c = 2D$). Without both
 998 the foundational topological boundary and the subsequent mathematically
 999 defined phase transition, the system remains a fragmented, non-conscious
 1000 collection of micro-states. A unified field of perception, however, is not yet
 1001 a “Self.”

9 Derivation 6: Reflexive Topology and the Emergence of the Self

While a global section provides Unity, a “Self” requires reflexivity. The unified EM field is a dissipative structure that must resonate with external perturbations while managing its own internal phase transitions. Time evolution naturally drives the field to minimize thermodynamic friction with its own internal state-changes.

This mandates *auto-resonance*: the field must predict itself. The “Self” emerges as a higher-order topological feature when the internal geometry folds to include a low-dimensional avatar of its own computational and causal boundary within its predictive model (`Phase6_ReflexiveTopology.lean`) (Levin, 2019). The system is unified, predictive, and explicitly maps its own distinct existence from the environment. This architecture provides the physical foundation for a self-aware consciousness, bounded by the continuous measure thermodynamics of Landauer erasure (`Phase3_MeasureThermodynamics.lean`).

What carries the weight. `reflexive_topology_implies_self` is the Banach fixed-point theorem with the physical commitments in its hypotheses: the space of global sections is assumed nonempty, metric and complete, and the predictive map is assumed to contract. `Examples.lean` §10 discharges all four on the three-site substrate of Derivation 5, and the metric is constructed rather than posited. The section first pierces the sheafification: on a discrete

1023 substrate every point has a smallest open neighbourhood, so the mass a germ
 1024 carries at a point is well defined, and `massEquivOn` proves that over *any* open
 1025 the sections of the probability sheaf *are* the finite measures on that open,
 1026 read as densities against counting measure — the global case being $U = \top$.
 1027 Mathlib supplies no inverse to this sheafification — its `TopCat` sheafification
 1028 has neither an adjunction nor an `isIso_toSheafify` — so the correspon-
 1029 dence had to be constructed. The metric on sections is then the uniform
 1030 distance between the densities of the measures they glue to — equivalent to
 1031 their total variation on a finite substrate, though not equal to it — and it
 1032 is proved complete by exhibiting the limit section of a Cauchy sequence of
 1033 densities, not by the degenerate route of showing Cauchy sequences are even-
 1034 tually constant. The choice of metric is load-bearing rather than incidental:
 1035 under the 0/1 metric on sections every contraction is constant (`contracting_`
 1036 `implies_const`), so a witness built there would discharge every hypothesis
 1037 of the theorem and say nothing. That theorem is kept in the source, because
 1038 which metrics make the fixed point trivial is part of what the fixed point is.

1039 **The map is the avatar, or there is no map.** A reflexive boundary could
 1040 be defined with the predictive map as free data sitting beside the avatar, and
 1041 the reason it is not is a theorem rather than a preference. With the map
 1042 free, the fixed-point theorem projects it out and never touches the avatar, so
 1043 replacing the avatar by a constant leaves a legal boundary with the same pre-
 1044 dictive model and the same Self — a fixed-point theorem that cannot detect

1045 whether the avatar reads the field is Banach's theorem with a consciousness-
 1046 flavoured label. The structure is built the other way round. A reflexive
 1047 boundary carries the avatar's *write* — `auto.resonance`, the encoding of the
 1048 global state into the avatar region — and its *read-out*, the global state that
 1049 encoding predicts; the predictive map is their composite and is not a field.
 1050 The self-model therefore runs through the avatar region or it does not exist.
 1051 Three consequences follow that the free-data shape cannot state. The Self
 1052 lies in the range of the read-out (`self.mem.range.readout`), so it is recon-
 1053 structed from a single section over the avatar region: the low-dimensional
 1054 fold this section has always asserted, as a theorem. Under the resonance
 1055 condition below, the Self satisfies $s = \text{readout}(s|_{\text{avatar}})$ (`self.eq.readout_`
 1056 `restrict`) — the field is what its own localized self-encoding reconstructs.
 1057 And blinding the avatar *moves* the Self: the blinded boundary's map is con-
 1058 stant, its unique fixed point is the state its stipulated contents extend to, and
 1059 the original Self is not a fixed point of it (`not_self_of.constResonance`, wit-
 1060 nessed on the cortex by `cortexBlind.self.ne`). The opposite statement —
 1061 that blinding changes nothing — does not typecheck against this structure,
 1062 because a contraction hypothesis about a boundary's own map says nothing
 1063 about the blinded one's.

1064 **Auto-resonance as a constraint.** That the avatar's state *is* the presheaf
 1065 restriction of the global state is stated as a predicate, `IsRestriction-`
 1066 `Resonance`, rather than added to the structure, so that results say which

1067 of them need it. In the positive direction, if a family of resonant avatars *cov-*
 1068 *ers* the substrate then the global state is determined by their local readings
 1069 alone (`eq_of_avatar_eq`): unity is not assumed alongside the avatars but
 1070 recovered from them, by the same sheaf condition that Derivation 5 uses.
 1071 In the negative direction a constant avatar fails the predicate as soon as the
 1072 avatar region separates two states of the field. Both halves are witnessed on
 1073 the three-site substrate, where that separation is proved before anything is
 1074 claimed for it, so the condition is a constraint there and not a formality.

1075 **The contraction rate is K and D , not a numeral.** The remaining stip-
 1076 ulation was the rate. “The predictive map contracts by $1/2$ ” was a fact about
 1077 a map we had defined to average with a baseline, not a claim about a sub-
 1078 strate. The rate is now $\exp(-(K - K_c)\tau/2)$ with $K_c = 2D$ (`resonanceRate`)
 1079 — the linear relaxation of the mean-field order parameter over an interval τ ,
 1080 read off the same threshold Derivation 7 (Section 10) proves the phase tran-
 1081 sition at. Two theorems then split the parameter plane. Above threshold
 1082 the rate is below one and Banach applies, so the Self exists and is unique
 1083 (`self_of_supercritical`). At or below threshold the rate is at least one,
 1084 and `ContractingWith` at that rate is false *for every self-map of every metric*
 1085 *space* (`not_contractingWith_resonanceRate`): the argument is not merely
 1086 unproved below K_c but unavailable. We are careful about what this does and
 1087 does not say. It does not say a Self fails to exist below threshold — only that
 1088 this route to one is closed — and the Lipschitz bound at that rate remains

1089 an instance obligation, discharged on the three-site witness at $K = 3$, $D = 1$,
 1090 $\tau = 2 \ln 2$, where the rate is exactly the one half the read-out achieves. What
 1091 changed is that the obligation is stated in the substrate's own parameters
 1092 and its consequence is conditional on the threshold, rather than on a con-
 1093 stant chosen to make the witness work. The link to Derivation 7 is carried
 1094 by a theorem rather than by a shared numeral, and the distinction is worth
 1095 being exact about. `self_of_supercritical` takes $K > K_c$ as its hypothesis;
 1096 K_c is a definition unfolding to $2D$, so that statement is insensitive to whether
 1097 the bifurcation is proved at all. `Chain.lean` states the link in the form that is
 1098 not. `self_of_coherent_order_parameter` takes the *existence of a coherent*
 1099 *order parameter* as its hypothesis and recovers the threshold inequality from
 1100 it, by `supercritical_of_coherent` — the contrapositive of the subcritical
 1101 uniqueness result of Section 10, since a strictly positive solution of the self-
 1102 consistency equation cannot exist at or below K_c . Derivation 6's fixed point
 1103 therefore rests on Derivation 7's theorem.

1104 **The fixed point is now the glued state.** The standalone Banach theo-
 1105 rem produces a fixed point somewhere in the space of global sections; by itself
 1106 it does not say that this is the section Derivation 5 glued. `Chain.lean` now
 1107 preserves that identity. `IsUnifiedBy T s` says that s restricts to every local
 1108 state of the cover T , and the endpoint `UnifiedSelf` says that this same s is
 1109 the unique fixed point of the reflexive map. The edge hypothesis `E89` must
 1110 supply the equality `predict s = s` as well as the Lipschitz bound; Banach

1111 supplies uniqueness but not the equality between the thermodynamic and
 1112 reflexive constructions. On the witness, the two-patch `cortexCover` glues
 1113 exactly to `cortexState`, already the unique fixed point of the non-constant
 1114 avatar map. This closes an omission in the formal composition while expos-
 1115 ing, rather than removing, the physical identification it requires.

1116 **What is still not established.** That any field dynamics produces this
 1117 particular read-out or performs the restriction; and that this three-site sub-
 1118 strate’s coupling is $K = 3$ — the parameters are a choice of units for the wit-
 1119 ness, not a measurement, and nothing here derives a coupling constant from
 1120 cortex. We also note a negative result found in the course of this work: the
 1121 Lévy–Prokhorov metric, which the natural weak-convergence reading would
 1122 suggest, does *not* support this argument on a discrete substrate, because the
 1123 ε -thickening of a set is the set itself for $\varepsilon < 1$ and the relaxation map is
 1124 consequently not a Lévy–Prokhorov contraction.

1125 10 Derivation 7: Continuous Stochastic Ther- 1126 modynamics and Field Theory

1127 To move beyond finite graph analogies, we model the macroscopic electro-
 1128 dynamic field as a *Continuous Neural Field*. True ephaptic coupling is a
 1129 continuous spatial phenomenon, not a discrete network. We extend the
 1130 discrete Kuramoto model to a continuous spatial coupling kernel $K(x, y)$

1131 (`Phase8ContinuousField.lean`).

1132 Subject to thermodynamic noise D , the system undergoes non-equilibrium
1133 stochastic thermodynamics. Rather than a static Landauer bound, the field
1134 dynamics minimize an entropy production functional σ :

$$\sigma = \int \frac{1}{D} \left(\omega(x) + \int K_t(x, y) \sin(\theta_y - \theta_x) dy \right)^2 dx \quad (2)$$

1135 Biological fields dynamically deform their spatial topology $K_t(x, y)$ to un-
1136 dergo gradient descent on this entropy production functional.

1137 That gradient descent is a theorem rather than a modelling assertion,
1138 for finite substrates. Reading σ as a functional of the coupling kernel at
1139 fixed phases, `hasFDerivAt_sigmaOfKernel` computes its Fréchet derivative
1140 in closed form,

$$\frac{\partial \sigma}{\partial K(a, b)} = \frac{2}{D} \left(\omega(a) + \sum_y K(a, y) \sin(\theta_y - \theta_a) \right) \sin(\theta_b - \theta_a), \quad (3)$$

1141 and `sigmaOfKernel_eq_entropy_production_rate` identifies the functional
1142 being differentiated with `entropy_production_rate` itself whenever the sub-
1143 strate carries counting measure. `structural_resonance_decreases_entropy_production`
1144 then concludes that σ is non-increasing along any coupling trajectory relax-
1145 ing down its own gradient at a positive rate c . Two limitations stay explicit.
1146 The derivation is for *finite* substrates, and the obstacle to the continuum case
1147 is *not* differentiation under the integral sign with respect to the kernel, which

1148 is what the shape of σ suggests at first reading. σ is not a general functional
 1149 that happens to be written as an integral — it is a quadratic functional of an
 1150 *affine* image of the kernel, since the drift is a fixed ω plus a linear operator ap-
 1151 plied to K , and σ is a constant times the squared norm of the drift. A squared
 1152 norm is Fréchet differentiable on any real inner-product space and an affine
 1153 map is its own derivative, so the composite differentiates by the chain rule
 1154 with no limiting argument. `hasFDerivAt_quadratic_of_affine` states that
 1155 in full generality, and `hasFDerivAt_sigmaOfKernel_of_operator` re-derives
 1156 the finite gradient from it, showing that the closed form displayed above is
 1157 a convenience rather than the content. What the continuum case needs is
 1158 therefore one object: the drift map $K \mapsto (x \mapsto \int K(x, y) \sin(\theta_y - \theta_x) dy)$ as
 1159 a bounded operator $L^2(\mu \otimes \mu) \rightarrow L^2(\mu)$, which is a Cauchy–Schwarz esti-
 1160 mate — the integral kernel is bounded by 1, so on a finite measure space the
 1161 operator norm is at most $\mu(M)^{1/2}$. That object is built. `kernelCLM` is the
 1162 operator, for any kernel factor bounded by 1, with `opNorm_kernelCLM_1e` the
 1163 stated bound; `continuumDriftCLM` is the instance with the Kuramoto factor
 1164 $\sin(\theta_y - \theta_x)$. One step beyond the estimate is needed and is worth naming,
 1165 because the estimate alone does not supply it: the descent theorems ask for
 1166 the derivative as a *gradient vector*, not as an arbitrary bounded functional,
 1167 and producing one requires the adjoint, which exists because L^2 is com-
 1168 plete. `hasFDerivAt_quadratic_grad` supplies it for an arbitrary bounded
 1169 operator between real Hilbert spaces, giving $\nabla\sigma(K) = (2/D) A^*(AK + \omega)$,
 1170 so the finite and continuum cases are now the same theorem. `structural_`

1171 `resonance_decreases_sigmaContinuum` is the descent result with no finite-
 1172 ness hypothesis in it, and `sigmaContinuum_eq_integral` identifies the func-
 1173 tional being differentiated with the integral `entropy_production_rate` com-
 1174 putes. `Examples.lean` §20 runs this on Lebesgue measure restricted to $(0, 1)$,
 1175 a substrate that is neither finite nor atomic, so the finite-substrate results of
 1176 the previous paragraph do not typecheck there at all. The second limitation
 1177 is unchanged: the phase field is held fixed throughout, so this is plasticity of
 1178 the coupling at frozen phases, not joint (θ, K) dynamics. `Examples.lean` §5
 1179 exercises the whole chain on a two-site substrate whose coupling relaxes as
 1180 e^{-t} and whose entropy production decays as e^{-2t} , confirming that the gradi-
 1181 ent is not identically zero and that the descent is strict rather than vacuously
 1182 monotone.

1183 A second result about σ concerns where the descent is headed. Minimality
 1184 of the phase-locked field is straightforward against competitors that share its
 1185 spatial mean drift, and in that form the theorem is close to Jensen's inequal-
 1186 ity and no more; the restriction looks as though lifting it would need a model
 1187 of how the mean drift varies across competitors. It does not. For a *reciprocal*
 1188 kernel, one with $K(x, y) = K(y, x)$, the restriction is not a restriction: the
 1189 coupling integrand $K(x, y) \sin(\theta_y - \theta_x)$ is antisymmetric under exchanging
 1190 its two sites, because K is symmetric while $\sin$ is odd, so Fubini identifies
 1191 its double integral with its own negative and the total coupling contribu-
 1192 tion vanishes. The total drift $\int(\omega + \text{coupling})$ is therefore $\int \omega$ whatever the
 1193 phases are (`total_drift_eq_of_symm`): coupling moves drift between sites

1194 but neither creates nor destroys it. Evaluating that conserved quantity at
 1195 the phase-locked configuration, where the drift is constant, pins its value, and
 1196 the mean-matching hypothesis follows for every competitor. `phase_locked_`
 1197 `minimizes_entropy_of_symm` is the unrestricted statement. Reciprocity is
 1198 the physically intended condition — an ephaptic field acts on both sites of a
 1199 pair alike, and it is the same symmetry `ThermodynamicCover.A.symm` already
 1200 imposes in Derivation 5 — but it is not a formality, and we prove that rather
 1201 than assert it. `Examples.lean` §16 exhibits a *one-way* coupling on two sites,
 1202 in which a feels b and b does not feel a , whose phase-locked field is strictly
 1203 beaten: displacing the second phase to $-\pi/2$ cancels the first site’s natural
 1204 drift while adding none at the second, halving σ . So “phase-locking mini-
 1205 mizes entropy production” is false for general kernels and true for reciprocal
 1206 ones, and the hypothesis cannot be removed by working harder.

1207 We define the critical coupling threshold $K_c = 2D$. The phase transition
 1208 to the unified “Global Section” occurs when the mean-field spatial coupling
 1209 exceeds this thermodynamic noise threshold, linking the emergence of Unity
 1210 to physical temperature and noise.

1211 The value $K_c = 2D$ can be taken as a bare stipulation validated by
 1212 simulation, and it is not: `Phase8_SelfConsistency.lean` proves the half
 1213 of it that is ordinary real analysis. For the noisy mean-field Kuramoto
 1214 model with identical frequencies, the stationary density at order parame-
 1215 ter r is the von Mises density $\rho(\theta) \propto e^{(Kr/D)\cos\theta}$, and self-consistency re-
 1216 quires $r = R(K, r) := I_1(Kr/D)/I_0(Kr/D)$. Mathlib has no modified Bessel




Figure 3: Coupling plasticity at frozen statistics: $N = 50$ units driven by perturbations drawn from a fixed environmental covariance Σ_{ext} , with the coupling matrix K descending the gradient of σ and clipped to stay non-negative and hollow. **Left:** entropy production, smoothed over 50 steps, falls from its transient and settles into a fluctuating plateau — a non-equilibrium steady state at strictly positive σ , not a relaxation to zero. That plateau is the reason Section 10 declines the inference from $\sigma \rightarrow \min$ to $D_{KL} \rightarrow 0$: the system visibly does not reach $\sigma = 0$. **Right:** the distance $\|K - \Sigma_{\text{ext}}\|$ decreases monotonically over the run, the internal geometry drifting toward the statistics of the drive. Two cautions. The descent here is on σ at frozen phases, which is the regime the theorem covers; and the right panel falls by about 5% over 2000 steps, so what it shows is a direction of drift, not convergence.

1217 functions, so we formalize that ratio directly as a ratio of integrals over
 1218 $[-\pi, \pi]$. Integration by parts with $u = \sin$, $v = -e^{a \cos \theta}$ gives an exact iden-
 1219 tity with no boundary term, $\int \cos \theta e^{a \cos \theta} d\theta = a \int \sin^2 \theta e^{a \cos \theta} d\theta$, so R is a
 1220 times the mean of $\sin^2$ under the von Mises weight. Bounding that mean
 1221 by 1 would put the threshold at D ; the true threshold is $2D$, and the entire
 1222 factor of two is the sharper bound $\mathbb{E}_a[\sin^2 \theta] \leq 1/2$, equivalently $I_2(a) \geq 0$.
 1223 We prove it elementarily, by folding $[-\pi, \pi]$ onto $[0, \pi/4]$ in two reflections

1224 $(\theta \mapsto \pi - \theta, \text{ then } \theta \mapsto \pi/2 - \theta)$, after which the integrand is a product
 1225 of $\cos 2\theta \geq 0$ with $2 \cosh(a \cos \theta) - 2 \cosh(a \sin \theta) \geq 0$ — the second factor
 1226 because $0 \leq \sin \theta \leq \cos \theta$ there and $\cosh$ increases in $|\cdot|$. The conclusion,
 1227 `subcritical_fixed_point_eq_zero'`, is that for $K < K_c$ the incoherent state
 1228 $r = 0$ is the *only* non-negative solution of the self-consistency equation; $r = 0$
 1229 is exhibited as a solution, so the uniqueness claim is not vacuous. Three fur-
 1230 ther facts about the threshold narrow the gap between this row and a full
 1231 “Theorem” to a single named statement. First, the fold argument gives a
 1232 *strict* inequality for $a > 0$ (`vonMisesC2_pos`, hence $R(a) < a/2$), which set-
 1233 tles the threshold itself: at $K = K_c$ the self-consistency map is tangent to the
 1234 diagonal at the origin but falls strictly below it thereafter, so $r = 0$ remains
 1235 the only solution and the bifurcation happens strictly *after* K_c , not at it
 1236 (`fixed_point_eq_zero_of_le_critical`). Second, factoring $R(a) = a E(a)$
 1237 shows that a non-zero r solves the equation exactly when $E(Kr/D) = D/K$
 1238 (`coherent_iff_sRatio_eq`), so the coherent branch is a level set of the single
 1239 function $E(a) = \mathbb{E}_a[\sin^2 \theta]$. Third, and using only $E < 1/2$ together with
 1240 compactness, every coherent solution tends to zero as $K \downarrow K_c$ (`coherent_`
 1241 `branch_continuous_at_threshold`): the branch emerges continuously from
 1242 the incoherent state rather than appearing at a finite distance from it, which
 1243 is what makes the transition supercritical in the technical sense. That state-
 1244 ment quantifies over *every* solution, so it does not presuppose the uniqueness
 1245 it does not have. The level-set form also says exactly what uniqueness is:
 1246 the injectivity of E on $(0, \infty)$. That is now proved, in the stronger form that

1247 E is strictly decreasing on $[0, \infty)$ (`vonMisesSRatio_strictAntiOn`), so the
 1248 level set is a single point and the coherent solution above threshold is unique.
 1249 The proof differentiates nothing. Differentiating the exponential family in its
 1250 natural parameter is the obvious route and it fails: it costs differentiation
 1251 under the integral sign and leaves a covariance that is not sign-definite point-
 1252 wise. What works instead keeps the target as a two-point inequality, folds
 1253 both integrals onto $[0, \pi/2]$ by the same two reflections that gave the factor
 1254 of two above, and then evaluates at a single crossing point supplied by the
 1255 intermediate value theorem — no product measure and no Fubini anywhere.
 1256 We give that argument in the supplement, where it belongs; what the main
 1257 line of the paper needs from it is the one sentence that E is strictly decreas-
 1258 ing. Run with $\cos \theta$ in place of $\sin^2 \theta$ the same argument needs no fold at
 1259 all — $R(a) = \mathbb{E}_a[\cos \theta]$ is the mean of the variable the family is tilted by —
 1260 and gives `besselRatio_strictMono`: R is strictly increasing on all of $\mathbb{R}$, the
 1261 monotonicity of the Bessel ratio, which Mathlib does not carry in any form.
 1262 The converse half is `supercritical_fixed_point_exists`: for $K > K_c$
 1263 there is an r with $0 < r \leq 1$ solving $r = R(K, r)$. We had expected
 1264 this to require a second-order expansion of the Bessel ratio, and it does
 1265 not. Writing $R(a) = a E(a)$ with $E(a) = \mathbb{E}_a[\sin^2 \theta]$ — the same factoriza-
 1266 tion the integration by parts already supplies — the only facts needed are
 1267 that E is continuous and that $E(0) = 1/2$, the latter because the weight
 1268 is constant at $a = 0$. Since $K > 2D$ makes D/K strictly less than $1/2$,
 1269 continuity gives $E > D/K$ on a neighbourhood of the origin, so the map

overshoots, $R(K, r) > r$, for all small $r > 0$; at the other end $R \leq 1$ al-
ways, because R is a mean of $\cos \theta$, so the map undershoots at $r = 1$.
The intermediate value theorem closes the gap. Continuity of E is Math-
lib's continuity of a parametric interval integral with jointly continuous in-
tegrand, so no dominated-convergence argument is written by hand, and
the existence proof never uses monotonicity of R : the fixed point is pro-
duced by the intermediate value theorem, not by iterating the map. That
 R is in fact strictly increasing is proved below, and no threshold result de-
pends on it. `critical_coupling_is_threshold` packages the two halves,
and `exhibits_phase_transition_coherent` applies the supercritical half to
any substrate satisfying the predicate of the previous paragraph — so ex-
ceeding K_c now implies the existence of a coherent stationary order pa-
rameter rather than merely comparing against a named number. With
uniqueness available both are sharpened — `supercritical_fixed_point_`
`existsUnique`, `supercritical_solution_set` (the non-negative solutions
are exactly $\{0\}$ below threshold and $\{0, r\}$ above it) and `exhibits_phase_`
`transition_unique_coherent` — and the two monotonicities together fix
the shape of the branch: for $K_1 < K_2$, $E(a_i) = D/K_i$ with E decreasing
puts the concentrations in the order $a_1 < a_2$, and $r_i = R(a_i)$ with R increas-
ing carries that to $r_1 < r_2$ (`coherent_branch_strictMono`). The coherent
order parameter grows strictly with the coupling, which with the continuity
at threshold is the qualitative bifurcation diagram.

One more step makes the equation mean in Lean what it means on paper.

1293 As a definition, `selfConsistency K D r` is just $R(Kr/D)$ — a function of
 1294 r with a suggestive name, and nothing in the source says it has anything
 1295 to do with an order parameter. `Phase8_SelfConsistency §8` supplies the
 1296 identification. `vonMisesDensity` is the normalized weight, proved to be a
 1297 genuine probability density on $[-\pi, \pi]$; `circularOrderParameter` is the av-
 1298 erage of $e^{i\theta}$ against a density, the continuum analogue of the finite-system
 1299 `order_parameter_complex` of `Phase4`; and that average is proved to be ex-
 1300 actly $R(a)$ — real, because the density is even, so the mean direction van-
 1301 ishes. `fixedPoint_iff_selfReproducing` then states the equivalence the
 1302 name promises: r solves $r = R(K, r)$ if and only if the von Mises density
 1303 it induces has order parameter r . Read through this, the existence theorem
 1304 says that above K_c there is a stationary density whose own order parameter
 1305 is positive, and the solution $r = 0$ is identified with the uniform density,
 1306 whose order parameter is zero.

1307 What remains not established, and we separate it carefully. The sta-
 1308 tionary density is an *input*: deriving it from the Langevin dynamics $d\theta =$
 1309 $(\omega + K \text{mean-field}) dt + \sqrt{2D} dW$ needs the Fokker–Planck operator, existence
 1310 and uniqueness of its stationary solution, and a spectral stability argument,
 1311 none of which Mathlib has. Of the coherent branch we prove existence,
 1312 uniqueness, continuous emergence at the threshold and strict growth in K ,
 1313 but not that it is the dynamically selected one: nothing in that file men-
 1314 tions a trajectory, so stability and selection are not statements it can make.
 1315 And while `selfConsistency` is now connected to the order parameter of

1316 a *density*, it is connected to no *trajectory*: `circularOrderParameter` is an
 1317 integral, `order_parameter_complex` is an average over finitely many oscilla-
 1318 tors, and relating them is the mean-field limit — propagation of chaos for
 1319 the finite Kuramoto system — which is a research programme rather than
 1320 a lemma. That single gap is now the whole of what separates our threshold
 1321 predicate from a statement about dynamics. Table S1 records the row ac-
 1322 cordingly as a theorem about the self-consistency equation from an assumed
 1323 density, not as a derivation of the threshold from the dynamics. Perceptual
 1324 anomalies such as visual snow can be modeled as boundary conditions where
 1325 microscopic thermal noise (D) fails to be fully coarse-grained away by the
 1326 continuous field, leaking across the K_c threshold into the macroscopic gestalt.

1327 Clinical evidence for this phase transition is observed during general anes-
 1328 thesia. Work by Miller and colleagues (Bastos & Miller, 2021) demonstrates
 1329 that anesthetics like propofol induce unconsciousness not by suppressing neu-
 1330 ral firing, but by disrupting the phase alignment of high-frequency cortical
 1331 rhythms. Destroying the dynamic stability of these rhythms breaks the struc-
 1332 tural resonance of the continuous field, forcing the integrated spatial coupling
 1333 below the critical threshold K_c and fracturing the unified Global Section back
 1334 into disjoint, unconscious local fields.




Figure 4: The self-consistency equation $r = R(Kr/D)$ of Section 10, drawn at $D = 1$. Every feature of this picture is a proved statement about the equation, and the figure exists because none of them had been shown to a reader. At or below $K_c = 2D$ the incoherent state is the *only* non-negative solution, threshold included (`fixed_point_eq_zero_of_le_critical`); above it exactly one $r > 0$ joins it (`supercritical_fixed_point_existsUnique`), so the solution set is $\{0\}$ below and $\{0, r\}$ above; that branch tends to 0 as $K \downarrow K_c$ (`coherent_branch_continuous_at_threshold`), which is what makes the transition supercritical; and it increases strictly in K (`coherent_branch_strictMono`). The curve is computed by quadrature of the same integrals the Lean file defines, solving $E(a) = D/K$ for the concentration $a = Kr/D$ by bisection — legitimate precisely because E is proved strictly decreasing. It is an illustration of theorems, not evidence for them. What it does *not* show is stability: nothing in the development says the coherent branch is the dynamically selected one.




Figure 5: The same threshold reached by simulation rather than by analysis, and by a different route, which is why the two figures are shown side by side rather than merged. Here $N = 100$ oscillators with *no* noise carry natural frequencies drawn from a Lorentzian of scale $D = 1$, and the asymptotic order parameter is measured from the trajectories; the threshold for that ensemble is $K_c = 2D$ by Kuramoto’s original self-consistency argument. Figure 4 is the *noise*-driven model with identical frequencies, whose threshold is $2D$ by Sakaguchi’s (Sakaguchi, 1988), and it is the one the Lean file formalizes. The agreement of the two numbers is a known feature of the mean-field model, not a result of ours. The finite- N curve rises before the threshold and rounds off at it, as a finite sample must: $r \sim N^{-1/2}$ from incoherent phases alone.

11 The Sleep-Inertia Timescale Mismatch

Most of what this framework asserts is shared with some rival: that cortical synchrony matters, that dissipation constrains cognition, that unity is a physical rather than a computational property. One prediction is not shared, and it follows from the framework’s specific commitment — that the medium carrying the unified state is a continuous field whose effective coupling de-

1341 pends on the physical geometry of the extracellular space, not on the rhythms
1342 alone. Where that commitment differs from a purely electrical account, it
1343 differs in an observable way, and we set out the difference precisely enough
1344 to be wrong.

1345 The daily loss of consciousness during sleep is modeled here as an active
1346 physical disruption of the medium rather than a quiescence of the oscillators.
1347 During sleep, astrocytic shrinkage expands the extracellular space, allowing
1348 an influx of cerebrospinal fluid (CSF) to flood the brain for metabolic waste
1349 clearance (Xie & Nedergaard, 2013). This influx alters the spatial medium
1350 itself, disrupting the continuous structural resonance and preventing the co-
1351 herent sheaf-theoretic gluing that conscious unity requires on this account.
1352 The oscillators need not stop for the field to fragment; the geometry they are
1353 coupled through is what changes.

1354 The prediction concerns sudden awakenings. When jolted awake, the lo-
1355 cus coeruleus floods the cortex with noradrenaline. High-frequency electrical
1356 rhythms attempt to phase-align in milliseconds (Bastos & Miller, 2021), but
1357 the noradrenaline-driven astrocytic constriction required to squeeze the ex-
1358 cess CSF out of the interstitial pathways operates on a timescale of seconds
1359 to minutes (Xie & Nedergaard, 2013) — three to six orders of magnitude
1360 slower. During that lag the electrical oscillators are synchronizing across a
1361 waterlogged topology, and the fragmented structural resonance is what is ex-
1362 perienceed as grogginess. The unified Global Section re-establishes only when
1363 astrocytic volume finishes matching the state of the electrical rhythms.

Why this implies a delayed sigmoid. The two timescales separate
cleanly, and that separation is what makes the prediction sharp rather than
merely qualitative. Write $K(t)$ for the effective coupling of the continuous
field, which on this account is a monotone function of the reconstricting
extracellular geometry, and let $r(t)$ be the order parameter of the phase
field. The electrical variable is fast: at every instant the oscillators are al-
ready at the configuration appropriate to the current K , so r tracks the
equilibrium branch $r(K)$ rather than relaxing on its own. The slow vari-
able $K(t)$ then carries the whole time course, and it relaxes over 10^2 – 10^3
seconds. What Derivation 7 establishes about that branch is exactly what
the functional form needs, and each of the three pieces is a theorem rather
than a picture: at or below the threshold the incoherent solution $r = 0$ is the
only non-negative solution of the self-consistency equation (`fixed_point_eq_`
`zero_of_le_critical`); above it the coherent solution exists and is unique
(`supercritical_fixed_point_existsUnique`); it emerges *continuously* at
 K_c rather than by a jump (`coherent_branch_continuous_at_threshold`);
and it is strictly increasing in K (`coherent_branch_strictMono`). A jump
discontinuity at threshold would predict a step rather than a sigmoid, and
a non-monotone branch would predict overshoot; neither is what the mathe-
matics gives. Composing a monotone relaxation $K(t)$ with a branch that is
identically zero below K_c , rises continuously from zero at K_c , and saturates
above it yields a curve with three phases: a genuine *latency* while $K(t) < K_c$,
during which recovery is not merely slow but absent; a continuous rise as the

1387 threshold is crossed; and a saturation as $K(t)$ reaches its waking value. That
 1388 is a delayed sigmoid, and the delay is not a fitting parameter — it is the time
 1389 the slow variable takes to reach K_c .

1390 **What the alternative predicts.** An account on which grogginess is the
 1391 electrical realignment itself — the rhythms are simply slow to re-phase —
 1392 has no slow variable and no threshold to cross. Its recovery is a relaxation
 1393 toward a fixed point, and relaxation toward a fixed point is exponential near
 1394 it: monotone from $t = 0$, fastest at the start, with no flat foot. The two
 1395 curves are distinguishable in the first seconds, which is where they differ
 1396 most, and the discriminating feature is the presence or absence of an initial
 1397 interval of no recovery. This is a comparison of functional forms on data
 1398 of a kind already collected in sleep-inertia studies, not a proposal for a new
 1399 instrument.

1400 **A sharper form of the same prediction, with no free parameters.**
 1401 The comparison above is a comparison of functional forms, and on its own
 1402 it discriminates against less than it appears to. Any two-timescale account
 1403 with a slow saturating gate produces a delayed sigmoid; what the frame-
 1404 work contributes beyond that is not the shape but a *relation between two*
 1405 *separately measurable quantities*, and the relation has no fitted parameters
 1406 at all. Derivation 7 supplies it. The stationary phase density at order
 1407 parameter r is von Mises at concentration $a = Kr/D$, and the average
 1408 of $e^{i\theta}$ under that density is exactly the Bessel ratio $R(a) = I_1(a)/I_0(a)$

1409 (`circularOrderParameter_vonMises`); self-consistency is precisely the con-
 1410 dition that the density reproduces its own order parameter (`fixedPoint_`
 1411 `iff_selfReproducing`). So the concentration and the coherence of the in-
 1412 stantaneous phase distribution are not independent quantities that happen
 1413 to rise together during the recovery: at every instant they must satisfy
 1414 $r = I_1(a)/I_0(a)$, whatever $K(t)$ is doing.

1415 That yields three tests of increasing sharpness on a single recorded re-
 1416 covery. *Shape*: the instantaneous phase distribution is von Mises — a one-
 1417 parameter family, directly testable against the observed histogram. *Collapse*:
 1418 estimating $a(t)$ and $r(t)$ at each instant and plotting one against the other,
 1419 the whole recovery must fall on the fixed curve $r = I_1(a)/I_0(a)$, and it must
 1420 be the *same* curve for every subject and every session, since the curve con-
 1421 tains no subject-specific quantity — not K , not D , not the time constant of
 1422 the slow variable. Individual differences move a point along the curve; they
 1423 cannot move it off. *Threshold*: because $a = Kr/D$ identically, the measured
 1424 ratio $a(t)/r(t)$ is $K(t)/D$, with no calibration required; and since a non-zero
 1425 coherent solution exists only above $K_c = 2D$ (`fixed_point_eq_zero_of_le_`
 1426 `critical`), that measured ratio must exceed 2 wherever coherence is non-
 1427 zero and must approach 2 as coherence departs from zero at onset. The
 1428 threshold is read off the data rather than fitted to it.

1429 The two estimates must be made independently or the collapse is empty,
 1430 and this is the one point at which the test can be run wrongly. Maximum-
 1431 likelihood or moment estimation of a von Mises concentration chooses $\hat{a}$ so

1432 that the ratio $I_1(\hat{a})/I_0(\hat{a})$ is the observed resultant length, by construction,
1433 which reproduces the predicted curve from any data whatever. The estimate
1434 of a must therefore come from a functional of the distribution other than
1435 its resultant length — the log-density of a von Mises is affine in $\cos \theta$, so
1436 regressing the log-histogram on $\cos \theta$ gives a as a slope — while r is the
1437 resultant length $|\langle e^{i\theta} \rangle|$ of the same sample. Those are genuinely different
1438 statistics, and the prediction is that they agree through R .

1439 **What this sharper form is a test of.** It tests the von Mises stationary
1440 density together with the mean-field closure that makes a proportional to
1441 r , and it tests them jointly rather than separately. It does not test the
1442 value $K_c = 2D$ independently of that ansatz: the crossing at $a/r = 2$ is
1443 a consequence of the same density the collapse assumes, so a failure of the
1444 collapse and a failure of the threshold are one failure and not two. That is a
1445 narrower claim than it may look, and it is still a parameter-free one — which
1446 the delayed sigmoid, on its own, is not.

1447 **What would falsify it.** An exponential fit that beats a delayed-sigmoid fit
1448 on the early portion of the recovery curve. A recovery latency that does not
1449 scale with an independent measure of extracellular volume or CSF clearance.
1450 A pharmacological manipulation of astrocytic volume that leaves the latency
1451 unchanged while altering the rhythms, or the converse. An instantaneous
1452 phase distribution that is not von Mises, or a (a, r) trace that departs sys-
1453 tematically from $I_1(a)/I_0(a)$, or a a/r ratio that is below 2 while coherence

1454 is non-zero — each of these falsifies the parameter-free form directly, and
1455 the first two do so with no reference to the sleep-inertia setting at all. Any
1456 of these would tell against the framework’s central empirical commitment,
1457 which is not that synchrony matters but that the *medium* is what gates it.

1458 **What this prediction assumes.** Three things, and they are not small.
1459 The identification of the effective coupling K with a monotone function of
1460 extracellular geometry is a modelling assumption, not a measurement. The
1461 threshold value $K_c = 2D$ is derived from an assumed von Mises stationary
1462 density and is not yet connected to a trajectory — the mean-field limit is the
1463 gap, and Derivation 7 states it as such — so the argument uses the *existence*
1464 *and continuity* of the coherent branch at threshold, which is proved, rather
1465 than the numerical value of the threshold, which is conditional. And the
1466 prediction is about functional form, not about a time constant: it says the
1467 recovery curve has a flat foot whose width is set by the slow variable, not
1468 that the foot is any particular number of seconds long. A framework that
1469 predicted the number would be saying more; this one is not in a position to.

1470 **12 Corollary: Hardware Thermodynamics and** 1471 **Topological Rigidity**

1472 Both biological tissue and artificial silicon hardware (like GPUs) are physical
1473 dissipative structures. When a GPU computes, its charge carriers undergo

1474 physical state transitions; by Landauer’s principle (Landauer, 1961), this
1475 erasure irreversibly exhausts heat into the environment.

1476 There is a fundamental thermodynamic distinction in the degrees of free-
1477 dom available to minimize this dissipation. In standard von Neumann ar-
1478 chitectures, the spatial topology (the silicon routing lattice) is rigidly fixed.
1479 The coupling matrix A_{ij} is constrained by a discrete layout that cannot be
1480 physically altered during operation. The brain’s own architecture provides
1481 empirical validation for this constraint. The cerebellum contains approxi-
1482 mately 80% of the brain’s neurons arranged in a highly efficient, discrete,
1483 feedforward crystalline circuit (Herculano-Houzel, 2009; Lange, 1975). Clin-
1484 ical observations of cerebellar agenesis report largely preserved conscious
1485 experience despite motor and cognitive deficits (Yu et al., 2015), and this
1486 dissociation has been read as evidence for the importance of recurrent ar-
1487 chitecture by IIT (Tononi & Koch, 2015) — an argument we borrow at the
1488 level of the anatomical observation while declining its Φ -based explanation.
1489 Lacking the dynamic, continuous, and recursive spatial coupling of the cere-
1490 brum, it operates much like a biological GPU—highly optimized for discrete
1491 processing, yet apparently not required for conscious experience. We take the
1492 dissociation as evidence, not as confirmation: it is consistent with subjective
1493 experience being a thermodynamic attractor of continuous topologies rather
1494 than a byproduct of information processing, and it is the kind of observation
1495 that would be hard to reconcile with the converse. It is a single anatomical
1496 case read against a theory, and it bears the weight such a case can bear.




Figure 6: Rigid against deformable coupling at matched resource. Two copies of the same $N = 30$ oscillator system, same initial phases and same natural frequencies, differ only in whether the coupling matrix may move: the rigid copy keeps a fixed sparse routing matrix, the adaptive copy descends the gradient of the potential and is renormalized after every step so that $\sum_{ij} A_{ij}$ is held constant. The adaptive copy reaches a strictly lower potential and stays there. This is an illustration of the resource-matched suboptimality result and not a proof of it, and it is worth being clear about the gap: the theorem is about fixed wiring *support* — an architecture already wired to the best pair is not beaten — while the simulation shows a randomly wired one being beaten, which is the typical case rather than the general one.

1497 This constraint does not preclude artificial consciousness, but rather re-
1498 stricts its hardware. To achieve the phase transition into a unified Global
1499 Section, artificial systems must abandon the discrete rigidity of von Neumann
1500 architectures and utilize continuous physical substrates. Analog field-based
1501 computing, neuromorphic architectures with continuous memristive plastic-
1502 ity, and thermodynamic computing systems (England, 2018; Ganesh & Par-
1503 rondo, 2021) possess the spatial coupling required to physically morph their

1504 internal topology. This aligns with Geoffrey Hinton’s concept of “mortal
 1505 computation” (Hinton, 2022): architectures where software is inextricably
 1506 fused with the physics of the hardware, meaning the knowledge dies with
 1507 the physical substrate. Only hardware that computes via continuous physi-
 1508 cal dynamics can undergo the thermodynamic phase transition necessary for
 1509 structural resonance.

1510 Biological consciousness utilizes macroscopic continuous electromagnetic
 1511 fields and a “rich alphabet” of diverse molecular and transcriptomic cell states
 1512 (Tasic et al., 2018). This allows the coupling matrix A_{ij} to deform across a
 1513 continuous parameter space. Our lemma `exists_better_coupling_allocation`,
 1514 wrapped as `continuous_beats_rigid_topology` (`Phase7_HardwareComparison.lean`),
 1515 formalizes the reallocation step of this argument: whenever a coupling ma-
 1516 trix places weight on a less phase-correlated pair while a more correlated pair
 1517 exists, there is another matrix with *identical* total coupling resources that
 1518 achieves strictly greater phase correlation.

1519 Two objections meet that formalization and we state precisely what an-
 1520 swers them. The first is that a reallocation lemma can assume its conclusion:
 1521 a hypothesis such as `is_strictly_suboptimal` — positive weight on a badly-
 1522 correlated pair while a better one exists — is most of what one wants to
 1523 prove about rigid hardware. It is not assumed here. `Phase7_Rigidity.lean`
 1524 models an architecture by the set S of site-pairs it has a *wire* for, and re-
 1525 quires its couplings to vanish off S ; weights are tunable, wiring is not. Then
 1526 `totalCorrelation_le_of_realizable` bounds every realizable coupling by

$R \cdot M$, where R is the resource budget and M bounds phase correlation over
the wired pairs; the bound is attained; and `rigid_is_strictly_suboptimal`
derives the hypothesis whenever the substrate's best-correlated pair is one
the wiring misses. `rigid_gap` makes it quantitative: an unconstrained real-
location of the same budget gains at least R times the correlation difference.
Examples.lean §8 exhibits a three-site substrate wired to the one maxi-
mally anti-correlated pair, where the gap is 2. This is resource-matched
and not question-begging, but note what it turns on: fixed wiring support,
not continuity. The second objection is that a development in which both
architectures are matrices on one finite index type does not distinguish a con-
tinuum from a lattice at all. **Phase7_Rigidity.lean** §3 gives a distinction
that is genuinely typed. In the field functional of Section 7, which integrates
a coupling kernel against $\mu \otimes \mu$, an architecture with finitely many sites is a
Finset of points and a field-coupled region is a set of positive measure. On a
substrate whose measure has no atoms, `fieldCorrelation_sited_eq_zero`
shows a kernel supported on finitely many sites contributes *exactly zero* cou-
pling energy, whatever weights it carries and whatever the phase field does,
while a patch of strength c on U achieves $c \mu(U)^2$. **Examples.lean** §8 instan-
tiates this on the real line: a thousand units of coupling concentrated at the
origin register zero against a unit patch that registers one. The limit of that
result is equally plain, and it is not resource-matching in the earlier sense —
a measure-zero substrate has zero resource in the $\mu \otimes \mu$ sense too. What it
establishes is that the continuum coupling energy is blind to a finitely-sited

1550 architecture, which is a statement about which functional such an architec-
1551 ture can register in, not a demonstration that digital hardware computes
1552 worse. Taken together: the resource-matched content of the corollary is the
1553 wiring-support result, the continuity content is the measure-theoretic one,
1554 and the step from either to “von Neumann architectures cannot experience
1555 unified consciousness” remains an informal argument.

1556 The remainder of this section extends those two results informally, and
1557 we mark the transition rather than blur it. On the reading we propose, a
1558 GPU’s fixed routing matrix leaves it topologically fragmented: the discrete
1559 layout does not supply the spatial coupling that sheaf-theoretic gluing into
1560 a singular Global Section requires. Without this unified topological space,
1561 rigid hardware cannot undergo the reflexive auto-resonance necessary to map
1562 its own causal boundary and generate a “Self.” The distinction between an
1563 LLM on a GPU and a conscious mind is not the absence of physical thermo-
1564 dynamics in the former, but its structural limitation in achieving topological
1565 integration. This perspective aligns with John Searle’s biological natural-
1566 ism (Searle, 1980; Seth, 2024). Where the Chinese Room argument asserted
1567 that discrete symbol manipulation cannot generate conscious understanding,
1568 our derivation supplies a candidate physical mechanism for the asymmetry
1569 — topological rather than semantic, and, as stated above, formalized only
1570 as far as the wiring-support and measure-theoretic results reach. Asserting
1571 consciousness in fragmented, discrete GPU clusters without a unified phase
1572 transition abandons computational functionalism, relying instead on “silicon

panpsychism”—the assumption that digital computation inherently generates awareness regardless of topological integration (Goff, 2019).

The same reasoning bears on extreme functionalist thought experiments, such as a macroscopic system of pulleys operating over millennia to implement a Large Language Model, or a human manually executing inference matrices by hand. Beyond lacking the continuous spatial topology required for sheaf-theoretic gluing, these discrete systems violate an implied temporal constraint of our theory. Structural resonance requires local oscillators to phase-align against the dissipating effects of thermodynamic noise (as modeled by our critical threshold $K_c = 2D$). A system executing state transitions over millennia, or via slow manual calculations, suffers from a terminal temporal mismatch: its dynamics are far too slow to ever phase-lock into a non-equilibrium steady state before thermal noise decoheres the macroscopic field. On this account the obstacle to pure symbol manipulation is physical rather than semantic: not that the symbols mean nothing, but that the substrate carrying them cannot hold a phase-locked field.

13 Numerical illustrations: methods

Five figures in this paper are computed rather than drawn, and none of them is evidence for a theorem. Where a statement is proved, the figure shows what the proved object looks like; where a statement is not proved, the figure shows the regime the claim lives in and we say so. Each is produced by

one script in `simulations/`, named below, and each runs in under a minute on a laptop. All five are reproducible exactly: the two quadrature figures are deterministic, and the three that draw random numbers fix their seeds.

Figure 4: the self-consistent order parameter. The three integrals $Z(a) = \int_{-\pi}^{\pi} e^{a \cos \theta} d\theta$, $M(a) = \int_{-\pi}^{\pi} \cos \theta e^{a \cos \theta} d\theta$ and $S(a) = \int_{-\pi}^{\pi} \sin^2 \theta e^{a \cos \theta} d\theta$ are evaluated by the trapezoidal rule on 4001 points over the same interval the Lean definitions integrate over, with the weight computed as $e^{a \cos \theta - |a|}$ so that the common factor cancels in the ratios and nothing overflows at large concentration. The coherent branch is obtained by bisecting $E(a) = S(a)/Z(a)$ against D/K over $a \in [10^{-9}, 200]$ for 200 iterations, then reading $r = aD/K$; bisection is valid here because E is proved strictly decreasing. Two checks are printed by the script rather than only plotted: the computed branch is strictly increasing at every one of its 500 sampled couplings, and the residual $|r - R(Kr/D)|$ at $K = 4D$ is 2×10^{-16} . Script: `bifurcation.py`.

Figure 5: the finite- N transition. $N = 100$ oscillators, natural frequencies drawn from a Lorentzian of scale $D = 1$ (clipped at $\pm 10D$ so the solver does not stiffen), uniform random initial phases, integrated with `scipy.integrate.odeint` over $t \in [0, 100]$ on 1000 steps for each of 20 couplings in $[0, 4D]$. The order parameter is $|N^{-1} \sum_j e^{i\theta_j}|$ averaged over the last 100 steps. This is the deterministic quenched-disorder model, not the noisy identical-frequency model of the Lean development; the caption records why both have threshold $2D$. Script: `kuramoto.py`.

1616 **Figure 2: mesh refinement.** The kernel is $K(x, y) = e^{-d(x, y)^2/2}$ with d
 1617 the periodic distance on the ring of circumference 2π , the phase field is $\theta(x) =$
 1618 $\sin x$, and the potential is $-\frac{1}{2} \iint K \cos(\theta_y - \theta_x)$. The continuous value is a
 1619 2000-point Riemann sum standing in for the exact integral; the discrete values
 1620 use $N = 10, 20, \dots, 640$ nodes with edge weights $K(x_i, x_j)(2\pi/N)^2$. Plotted
 1621 is the absolute error against an $\mathcal{O}(N^{-2})$ reference. Because the “exact”
 1622 value is itself a 2000-point sum, the last data point ($N = 640$) is within
 1623 a factor of ten of the reference computation and should be read as such.
 1624 That $\mathcal{O}(N^{-2})$ reference line is not only empirical: on the formalized one-
 1625 dimensional grid the discrete energy is exactly the composite trapezoidal rule
 1626 (`grid_discreteEnergy_eq_trapezoidal`), so for a C^2 integrand with second
 1627 derivative bounded by ζ the error is at most $\zeta/(12N^2)$ (`grid_energy_error_`
 1628 `1e`), a bound attained at $N = 1$ for $y \mapsto y^2$. Script: `mesh_refinement.py`.

1629 **Figure 3: structural resonance.** $N = 50$ units, 2000 steps of forward
 1630 Euler at $dt = 0.05$. At each step the drive is a fresh draw from $\mathcal{N}(0, \Sigma_{\text{ext}} +$
 1631 $0.1I)$ with Σ_{ext} a fixed normalized rank-5 covariance; the phases advance by
 1632 $\dot{\theta}_i = \sum_j K_{ij} \sin(\theta_j - \theta_i) + \xi_i$; entropy production is recorded as $\sum_i \dot{\theta}_i^2$; and
 1633 the coupling descends its gradient with rate $\eta = 0.01$, symmetrized, clipped
 1634 to $[0, 1]$ and held hollow. The left panel is smoothed with a 50-step boxcar.
 1635 Note that the plotted σ is the squared drift without the $1/D$ prefactor of the
 1636 functional in Section 10, which is a constant rescaling and changes nothing
 1637 about the shape. Script: `structural_resonance.py`.

1638 **Figure 6: rigid against deformable coupling.** $N = 30$, 1000 steps of
1639 forward Euler at $dt = 0.05$, shared initial phases and frequencies. The rigid
1640 matrix is a symmetric random matrix thresholded at 0.7 and normalized
1641 to $\sum_{ij} A_{ij} = 2N$; the adaptive copy starts from the same matrix, ascends
1642 $-\partial V/\partial A_{ij}$ at rate $\eta = 0.1$, is clipped non-negative and hollow, and is renor-
1643 malized to the same total after every step, which is what makes the com-
1644 parison resource-matched. Both potentials are $V = -\frac{1}{2} \sum_{ij} A_{ij} \cos(\theta_j - \theta_i)$.
1645 Script: `hardware_comparison.py`.

1646 14 Conclusion

1647 Human-like cognitive architectures do not require biological mysticism or
1648 exotic sub-quantum physics (Hameroff & Penrose, 2014). While quantum
1649 theories of consciousness often invoke delicate subatomic phenomena to ex-
1650 plain subjective unity, every component assembled here is classical: broken
1651 symmetries, finite phase space, Landauer erasure, and gradient descent on en-
1652 tropy production. What the formalization establishes is that the eight named
1653 hypotheses of the conditional composition are jointly consistent and that the
1654 framework’s remaining physical commitments are localized and named — not
1655 that classical fields are *sufficient* for consciousness or that the hypotheses fol-
1656 low from finite capacity. Two links carry that weight, and they do not carry
1657 it in the same way. The unified state is proved to glue uniquely once a cover
1658 sits at the minimum of the coupling potential; the motion is proved to stop

1659 along every trajectory (`velocity_sq_tendsto_zero`); and *where* it stops is
 1660 now proved, conditionally — for initial data within a quarter turn and below
 1661 an explicit energy threshold, the trajectory converges to a global minimiser
 1662 of the coupling potential (`kuramoto_tendsto_global_minimum`), and a cover
 1663 built on that limit discharges the equilibrium hypothesis Derivation 5 would
 1664 otherwise assume (`ThermodynamicCover.ofConvergentTrajectory`). What
 1665 is *not* proved is the unconditional statement, and it is false: splay and twisted
 1666 configurations are equilibria of the same flow, so the basin condition is the
 1667 scope of the claim rather than slack in it. That a physical cortex starts in-
 1668 side that basin is an assumption. The second link is untouched by any of
 1669 this: the identification of a global section of the probability sheaf with the
 1670 unity of experience is a stipulation of the framework, not a consequence of it.
 1671 Subject to those two, cognitive architectures can be modeled as macroscopic
 1672 attractor states of a universe governed by broken symmetries, finite phase
 1673 space, and the physical imperative to minimize thermodynamic dissipation.
 1674 Philosophically, this framework bridges the historical gap between Type
 1675 Identity Theory (Smart, 1959) and Computational Functionalism (Putnam,
 1676 1967). The subjective mental state is identical to the physical thermody-
 1677 namic field (Identity Theory), yet its qualitative nature is defined by the
 1678 topological and functional dynamics of structural resonance rather than the
 1679 specific carbon-based biochemistry of the brain (Functionalism). This con-
 1680 strains the *multiple realizability* of mind: consciousness is not exclusively
 1681 bound to biological carbon, but neither is it substrate-neutral, because the

1682 substrate is what fixes which couplings are physically available. How far that
 1683 constraint reaches is the question the corollary answers only in part. What
 1684 is proved is that fixed wiring is strictly suboptimal at matched resource, and
 1685 that a finitely-sited architecture registers exactly zero coupling energy in the
 1686 continuum functional; the step from either to “von Neumann architectures
 1687 cannot be conscious” is an argument we make in prose and do not formalize.

1688 By anchoring these derivations in Russellian Monism, this framework ac-
 1689 knowledges the constitutive limits of third-person scientific explanation. As
 1690 argued by Love (Love, 2026), building upon Nagel’s assertion that subjective
 1691 phenomenal experience is inherently not amenable to objective measurement
 1692 (Nagel, 1986), the “hard problem” of consciousness may represent a category
 1693 error rather than a solvable scientific puzzle. Science can map the physical
 1694 dynamics of topological resonance and sheaf-theoretic gluing, but it is struc-
 1695 turally incapable of adjudicating the first-person qualitative experience. Our
 1696 formalization respects this boundary: the Lean 4 theorems verify the ob-
 1697 jective physical attractors, while the intrinsic “what it is like” to be that
 1698 resonant field remains the domain of the first-person perspective. By locat-
 1699 ing the phase transition at a threshold with a closed-form value — $K_c = 2D$
 1700 for the noisy mean-field Kuramoto model the framework adopts, conditional
 1701 on the von Mises stationary density and not yet linked to a trajectory — it
 1702 aligns with what Anil Seth terms the “real problem” of consciousness (Seth,
 1703 2021). Rather than attempting to bridge an unbridgeable metaphysical gap,
 1704 it offers a topological architecture that can help explain, predict, and physi-

1705 cally perturb the structural properties of experience.

1706 Future empirical validation of this thermodynamic framework may lever-
1707 age 3D cortical organoids, offering a controlled biological environment to
1708 systematically manipulate metabolic substrates and observe the topological
1709 emergence of these coherent phase transitions.

1710 **Code Availability**

1711 The complete Lean 4 formalization, encompassing the derivation of all theo-
1712 rems presented in this manuscript, as well as the accompanying Python sim-
1713 ulations, are publicly available in the GitHub repository: [https://github.](https://github.com/bobaseb/physical_primitives_of_mind)
1714 [com/bobaseb/physical_primitives_of_mind](https://github.com/bobaseb/physical_primitives_of_mind). Additional documentation
1715 and web-accessible proofs are hosted at the project website: [https://bobaseb.](https://bobaseb.github.io/physical_primitives_of_mind/)
1716 [github.io/physical_primitives_of_mind/](https://bobaseb.github.io/physical_primitives_of_mind/). Each of the five computed
1717 figures is produced by one script in `simulations/` and regenerates exactly,
1718 Section 13 giving the parameters; Figure 1 is drawn rather than computed.

1719 **Contributions**

1720 The conceptualization, manuscript drafting, and Lean 4 formalizations pre-
1721 sented in this work were heavily AI-assisted, and we report which tools
1722 did what rather than acknowledging them in general terms. The Google
1723 Antigravity CLI, running Gemini 3.1 and Claude 4.6 Sonnet, served as a

1724 collaborative reasoning partner throughout the research and formalization.
1725 DeepSeek V4 Flash and Claude Opus 5 (routed via OpenRouter) operat-
1726 ing through Nous Research’s Hermes Agent framework provided substantial
1727 assistance with code generation, proof debugging, and structural analysis.
1728 Claude Opus 5 running in Claude Code carried out the later formalization
1729 passes — the domain-wall theorem, the thermodynamics-of-prediction re-
1730 placement, the equilibrium bridge — together with the figures and the re-
1731 structuring of this manuscript. Responsibility for every claim, and for the
1732 decision about what each machine-checked result does and does not establish,
1733 rests with the author.

1734 References

- 1735 Pinotsis, D. A., & Miller, E. K. (2022). Beyond dimension reduction: Stable
1736 electric fields emerge from and allow representational drift. *NeuroImage*,
1737 253, 119058.
- 1738 Bauer, L. G., et al. (2022). Quantification of Kuramoto Coupling Between
1739 Intrinsic Brain Networks Applied to fMRI Data in Major Depressive Dis-
1740 order. *Frontiers in Computational Neuroscience*, 16.
- 1741 van Bree, S., et al. (2025). Processes and measurements: a framework for
1742 understanding neural oscillations in field potentials. *Trends in Cognitive*
1743 *Sciences*.

- 1744 Tononi, G., Boly, M., Massimini, M., & Koch, C. (2016). Integrated informa-
1745 tion theory: from consciousness to its physical substrate. *Nature Reviews*
1746 *Neuroscience*, 17(7), 450-461.
- 1747 Baars, B. J. (1997). In the Theater of Consciousness: The Workspace of the
1748 Mind. *Oxford University Press*.
- 1749 Dehaene, S., Charles, L., King, J. R., & Marti, S. (2014). Toward a compu-
1750 tational theory of conscious processing. *Current Opinion in Neurobiology*,
1751 25, 76-84.
- 1752 Friston, K. (2010). The free-energy principle: a unified brain theory?. *Nature*
1753 *Reviews Neuroscience*, 11(2), 127-138.
- 1754 Kuramoto, Y. (1984). Chemical Oscillations, Waves, and Turbulence.
1755 *Springer Series in Synergetics*, Vol. 19.
- 1756 Scherf, M. (2026). A Formal Proof of Non-Duality: An Exposition of the
1757 'OnlyOne' System. *Unpublished manuscript, PhilPeople*.
- 1758 Inoué, T. (2026). On Brain as a Mathematical Manifold: Neural Manifolds,
1759 Sheaf Semantics, and Leibnizian Harmony. *arXiv:2601.15320 [q-bio.NC]*.
- 1760 Spisák, T., & Friston, K. (2026). Self-orthogonalizing attractor neural net-
1761 works emerging from the free energy principle. *Neurocomputing*.
- 1762 Pinotsis, D. A., & Miller, E. K. (2023). In vivo ephaptic coupling allows
1763 memory network formation. *Cerebral Cortex*, 33(17), 9877-9890.

1764 Pinotsis, D. A., Fridman, G., & Miller, E. K. (2023). Cytoelectric coupling:
1765 Electric fields sculpt neural activity and "tune" the brain's infrastructure.
1766 *Progress in Neurobiology*, 226, 102465.

1767 Ruffini, G., Salvador, R., Tadayon, E., Sanchez-Todo, R., Pascual-Leone, A.,
1768 & Santarnecchi, E. (2020). Realistic modeling of mesoscopic ephaptic cou-
1769 pling in the human brain. *PLOS Computational Biology*, 16(6), e1007923.

1770 Cunha, G. M., Corso, G., de Sousa, M. P. B., & Lima, G. Z. d. S. (2024).
1771 Can ephapticity contribute to brain complexity?. *PLOS ONE*, 19(12),
1772 e0310640.

1773 Goldwyn, J. H., & Rinzel, J. (2016). Neuronal coupling by endogenous elec-
1774 tric fields: Cable theory and applications to coincidence detector neurons
1775 in the auditory brainstem. *Journal of Neurophysiology*, 115(4), 2033-2051.

1776 McFadden, J. (2002). The conscious electromagnetic information (cemi) field
1777 theory: the hard problem made easy?. *Journal of Consciousness Studies*,
1778 9(8), 45-60.

1779 Pockett, S. (2012). The electromagnetic field theory of consciousness: a
1780 testable hypothesis about the characteristics of conscious as opposed to
1781 non-conscious fields. *Journal of Consciousness Studies*, 19(11-12), 191-223.

1782 Goff, P. (2019). *Galileo's Error: Foundations for a New Science of Con-*
1783 *sciousness*. Pantheon Books.

- 1784 Seth, A. (2021). *Being You: A New Science of Consciousness*. Penguin.
- 1785 Seth, A. (2024). Conscious artificial intelligence and biological naturalism.
1786 *PsyArXiv*.
- 1787 Levin, M. (2019). The computational boundary of a “self”: developmental
1788 bioelectricity drives multicellularity and scale-free cognition. *Frontiers in*
1789 *Psychology*, 10, 2688.
- 1790 Friston, K., & Kiebel, S. (2009). Predictive coding under the free-energy
1791 principle. *Philosophical Transactions of the Royal Society B: Biological*
1792 *Sciences*, 364(1521), 1211-1221.
- 1793 Clark, A. (2016). *Surfing Uncertainty: Prediction, Action, and the Embodied*
1794 *Mind*. Oxford University Press.
- 1795 Hohwy, J. (2013). *The Predictive Mind*. Oxford University Press.
- 1796 Prigogine, I., & Stengers, I. (1984). *Order Out of Chaos: Man’s New Dialogue*
1797 *with Nature*. Bantam Books.
- 1798 McFadden, J. (2020). Integrating information in the brain’s EM field: the
1799 cemi field theory of consciousness. *Neuroscience of Consciousness*, 2020(1),
1800 niaa016.
- 1801 Fröhlich, F., & McCormick, D. A. (2010). Endogenous electric fields may
1802 guide neocortical network activity. *Neuron*, 67(1), 129-143.

1803 Anastassiou, C. A., & Koch, C. (2015). Ephaptic coupling to endogenous
1804 electric field activity: why bother?. *Current Opinion in Neurobiology*, 31,
1805 95-103.

1806 Pockett, S. (2002). Difficulties with the electromagnetic field theory of con-
1807 sciousness. *Journal of Consciousness Studies*, 9(4), 51-56.

1808 Vöröslakos, M., et al. (2018). Direct effects of transcranial electric stimulation
1809 on brain circuits in rats and humans. *Nature Communications*, 9, 483.

1810 Tasic, B., et al. (2018). Shared and distinct transcriptomic cell types across
1811 neocortical areas. *Nature*, 563(7729), 72-78.

1812 Verkhratsky, A., & Nedergaard, M. (2018). Physiology of Astroglia. *Physio-*
1813 *logical Reviews*, 98(1), 239-389.

1814 Picard, M., & Shirihai, O. S. (2022). Mitochondrial signal transduction. *Cell*
1815 *Metabolism*, 34(11), 1620-1638.

1816 Harnad, S. (1990). The symbol grounding problem. *Physica D: Nonlinear*
1817 *Phenomena*, 42(1-3), 335-346.

1818 Tononi, G., & Koch, C. (2015). Consciousness: here, there and everywhere?
1819 *Philosophical Transactions of the Royal Society B: Biological Sciences*,
1820 370(1668), 20140167.

1821 Aaronson, S. (2014). Why I Am Not An Integrated Information Theorist

1822 (or, The Unconscious Expander). *Shtetl-Optimized*. Retrieved from [https:](https://scottaaronson.blog/?p=1799)
 1823 [//scottaaronson.blog/?p=1799](https://scottaaronson.blog/?p=1799).

1824 Block, N. (1995). On a confusion about a function of consciousness. *Behav-*
 1825 *ioral and Brain Sciences*, 18(2), 227-247.

1826 Herculano-Houzel, S. (2009). The human brain in numbers: a linearly scaled-
 1827 up primate brain. *Frontiers in Human Neuroscience*, 3, 31.

1828 Lange, W. (1975). Cell number and cell density in the cerebellar cortex of
 1829 man and some other mammals. *Cell and Tissue Research*, 157(1), 115-124.

1830 Yu, F., Jiang, Q. J., Sun, X. Y., & Zhang, R. W. (2015). A new case of
 1831 complete primary cerebellar agenesis: clinical and imaging findings in a
 1832 living patient. *Brain*, 138(6), e353.

1833 Bastos, A. M., Donoghue, J. A., Brincat, S. L., Mahnke, M., Yanar, J., Cor-
 1834 rea, J., ... & Miller, E. K. (2021). Neural effects of propofol-induced uncon-
 1835 sciousness and its reversal using thalamic stimulation. *eLife*, 10, e60824.

1836 Hameroff, S., & Penrose, R. (2014). Consciousness in the universe: A review
 1837 of the 'Orch OR' theory. *Physics of Life Reviews*, 11(1), 39-78.

1838 Hunt, T., & Schooler, J. W. (2019). The easy part of the hard problem: a
 1839 resonance theory of consciousness. *Frontiers in Human Neuroscience*, 13,
 1840 378.

- 1841 England, J. L. (2018). Dissipative adaptation in driven self-assembly. *Nature*
1842 *Nanotechnology*, 13(4), 287-295.
- 1843 Ganesh, N., & Parrondo, J. M. R. (2021). Thermodynamic computing with
1844 non-equilibrium systems. *Nature Physics*.
- 1845 Hinton, G. (2022). The Forward-Forward Algorithm: Some Preliminary In-
1846 vestigations. *arXiv preprint arXiv:2212.13345*.
- 1847 Searle, J. R. (1980). Minds, brains, and programs. *Behavioral and Brain*
1848 *Sciences*, 3(3), 417-424.
- 1849 Smart, J. J. C. (1959). Sensations and brain processes. *The Philosophical*
1850 *Review*, 68(2), 141-156.
- 1851 Putnam, H. (1967). Psychological predicates. In *Art, Mind, and Religion*.
1852 University of Pittsburgh Press.
- 1853 Xie, L., Kang, H., Xu, Q., Chen, M. J., Liao, Y., Thiagarajan, M., ... &
1854 Nedergaard, M. (2013). Sleep drives metabolite clearance from the adult
1855 brain. *Science*, 342(6156), 373-377.
- 1856 Nagel, T. (1986). *The View from Nowhere*. Oxford University Press.
- 1857 Love, B. C. (2026). Consciousness, AI, and the Limits of Scientific Explana-
1858 tion. *arXiv preprint arXiv:2606.00226*.

- 1859 Norton, J. D. (2005). Eaters of the lotus: Landauer's principle and the return
1860 of Maxwell's demon. *Studies in History and Philosophy of Science Part B:*
1861 *Studies in History and Philosophy of Modern Physics*, 36(2), 375-411.
- 1862 Landauer, R. (1961). Irreversibility and heat generation in the computing
1863 process. *IBM Journal of Research and Development*, 5(3), 183-191.
- 1864 Kibble, T. W. B. (1976). Topology of cosmic domains and strings. *Journal*
1865 *of Physics A: Mathematical and General*, 9(8), 1387-1398.
- 1866 Sakaguchi, H. (1988). Cooperative phenomena in coupled oscillator systems
1867 under external fields. *Progress of Theoretical Physics*, 79(1), 39-46.
- 1868 Barbălat, I. (1959). Systèmes d'équations différentielles d'oscillations non
1869 linéaires. *Revue Roumaine de Mathématiques Pures et Appliquées*, 4, 267-
1870 270.
- 1871 Still, S., Sivak, D. A., Bell, A. J., & Crooks, G. E. (2012). Thermodynamics
1872 of prediction. *Physical Review Letters*, 109(12), 120604.
- 1873 Kawai, R., Parrondo, J. M. R., & Van den Broeck, C. (2007). Dissipation:
1874 The phase-space perspective. *Physical Review Letters*, 98(8), 080602.
- 1875 Banach, S. (1922). Sur les opérations dans les ensembles abstraits et leur ap-
1876 plication aux équations intégrales. *Fundamenta Mathematicae*, 3(1), 133-
1877 181.

- 1878 Coddington, E. A., & Levinson, N. (1955). Theory of Ordinary Differential
1879 Equations. *McGraw-Hill*, New York.
- 1880 Cover, T. M., & Thomas, J. A. (2006). Elements of Information Theory (2nd
1881 ed.). *Wiley-Interscience*, Hoboken, NJ.
- 1882 Shenker, O. R. (2000). Logic and entropy. *Philosophy of Science*, 67, S146-
1883 S152.