

Coupling-gated cortical coherence: a conditional framework for unity and self-representation

Sebastian Bobadilla-Suarez
Independent Researcher
`sebastian.bobadilla.s@gmail.com`
ORCID: 0000-0002-5951-0772

20 September 2026

Abstract

Collective phase order can coordinate distributed cortex, but does not establish shared content or self-representation. We distinguish coherence, compatibility and accurate reconstruction. Shared content requires independently identified quantities on overlapping regions; compatible local descriptions determine a unique global state. A contraction fixes that state uniquely if it reconstructs the state exactly; fidelity across other states separately requires discriminating codes. Counterexamples show why a phase resultant and fixed-point uniqueness cannot replace these measurements. Lean 4 verifies the conditional implications and a model jointly satisfies their eight connecting hypotheses. This establishes mathematical consistency, not cortical realization. A declared actuator's stored mode energy bounds mean coupling and yields a necessary condition for a coherent stationary branch; cortical mode prices remain unknown. We propose a held-out comparison of field-score and structural-connectivity onset models during anaesthetic emergence. A predictive advantage would support the specified spatial model; causal field identification requires intervention. Unity and a minimal self remain interpretations without an independent measure of experience.

1 Introduction

Endogenous extracellular fields can affect spike timing and may coordinate activity across cortical regions (McFadden, 2002, 2020; Pockett, 2012; Fröhlich & McCormick, 2010; Anastassiou & Koch, 2015). Coordination raises a further question: when do those regions describe one situation, and when does the resulting state represent its own place in it?

Consider reaching for a cup. Vision and movement can describe different features while agreeing on the cup’s location. Matching their timing leaves agreement about that shared quantity untested. An internal description of the perceiver’s relation to the cup asks for still more. We call the three requirements *coherence* (collective phase order), *compatibility* (agreement of local descriptions on overlaps), and *self-representation* (accurate reconstruction from an internal encoding). The mathematical states are spatial measures and fields; identifying them with perceptual content requires independent decoding evidence.

The familiar noisy mean-field oscillator supplies a stationary coherence threshold (Sakaguchi, 1988; Strogatz & Mirollo, 1991). We add a resource condition for a declared actuator, bounds that connect phase order to content only through specified encoders, and conditions under which compatible descriptions glue and an accurate reconstruction fixes the glued state. A Lean 4 development checks these implications and a joint witness satisfies their eight connecting hypotheses in a specified model. It proves neither that

23 cortex satisfies those hypotheses nor that the glued state is experienced. The
24 latter identification motivates the framework and needs its own endpoint.

25 The empirical question is which measurements could support the pro-
26 posed connections. A phase resultant can miss a locally ordered winding field;
27 stationary phase distributions can leave the coupling mechanism unidenti-
28 fied; and coherent adaptation can lose alignment with a designated content
29 template. Independent decoders must identify shared quantities before their
30 agreement can be tested. We therefore distinguish stationary branches from
31 finite observed onset, and compare prespecified field and connectivity onset
32 models on held-out emergence recordings. The comparison tests a specified
33 spatial prediction, not the presence of experience.

34 Sections 2–5 develop the three requirements; Section 6 applies their mea-
35 surement obligations to a digital candidate. Section 7 states the resource and
36 composition premises, and Sections 8–9 examine observables and the cortical
37 test. Table 1 lists the eight connecting hypotheses; Supplemental Material
38 gives their derivations and controls.

39 **2 Limits of phase order**

40 Several limits hold without the cortical hypothesis. Equal phases can coexist
41 with incompatible local descriptions; phase order constrains shared content
42 only through encoders related on overlaps. A winding field can have zero
43 global resultant while remaining locally ordered. A constant readout can




Figure 1: Three requirements of the proposed account. Each requires the inputs shown below it. Dashed arrows mark proposed connections: phase coherence alone supplies neither compatible local descriptions nor an accurate self-model. Thermodynamics constrains a candidate implementation through separate assumptions. This is the explanatory structure of the article; Table 1 records the eight implications consumed by each branch of the formal composition.

44 have a unique fixed point without reconstructing its input. A region that
 45 does not distinguish relevant states cannot encode them, and a site cannot
 46 report a change outside its causal past by a declared deadline. These lim-
 47 its reflect familiar phase oscillator and information constraints (Acebrón et
 48 al., 2005; Ott & Antonsen, 2008; Dörfler & Bullo, 2014; Breakspear, 2017).
 49 The sections below state each one beside the condition or measurement it
 50 motivates.

3 Coherence: physical coupling and its re- source condition

3.1 Dynamics and the stationary threshold

The finite stochastic model is

$$d\theta_i = \left[\omega_i + \sum_j K_{ij}(t) \sin(\theta_j - \theta_i) \right] dt + \sqrt{2D} dW_i, \quad r_N = \left| \frac{1}{N} \sum_i e^{i\theta_i} \right|. \quad (1)$$

Here θ_i is the phase of oscillator i , ω_i its natural frequency, $D > 0$ is phase diffusion, W_i are independent Wiener processes, and K_{ij} has inverse-time units, with angles dimensionless. Uniform mean-field coupling is $K_{ij} = K/N$; frequency heterogeneity and diffusion are distinct parameters. The stationary calculation below uses identical frequencies and this scalar coupling. Applying it to a spatial kernel requires an aggregation rule and a check of the resulting approximation.

For identical frequencies in their rotating frame, the positive normalized stationary densities of the one-body Fokker–Planck equation are exactly

$$\rho(\theta) = \frac{e^{a \cos(\theta - \psi)}}{2\pi I_0(a)}, \quad a = \frac{Kr}{D}, \quad r = R(a) := \frac{I_1(a)}{I_0(a)}, \quad (2)$$

where I_n are modified Bessel functions and ψ is the collective phase. The formal derivation uses the absence of stationary current for a periodic gradient drift to obtain this density, then verifies normalization and self-consistency.

67 At $K \leq K_c = 2D$, only $r = 0$ is self-consistent (Sakaguchi, 1988; Strogatz &
68 Mirollo, 1991). Above threshold, exactly one positive solution coexists with
69 zero and increases strictly with K . Its onset is

$$\frac{r^2}{K - K_c} \rightarrow \frac{1}{D}, \quad r \sim \sqrt{\frac{K - K_c}{D}}. \quad (3)$$

70 Figure 2 shows this stationary branch.

71 Linearization about the uniform density gives rate $K/2 - D$ for the first
72 harmonic and $-Dn^2$ for harmonics $n \geq 2$. A growing mode therefore exists
73 exactly above threshold. About the coherent density, the quadratic form
74 is strictly dissipative transverse to collective rotation, whose rate is zero.
75 These linear estimates concern the limiting equation. Relating them to finite
76 stochastic trajectories requires a dynamical mean-field limit and branch selec-
77 tion; exchange symmetry alone does not establish the needed independence
78 (Sznitman, 1991).

79 **3.2 Installed energy bounds mean coupling**

80 A spatial model replaces the finite sum by integration against a kernel. One
81 declared implementation carries response modes with continuous profiles ϕ_i ,
82 installation prices u_i and nonnegative occupancies $c_i(\Xi)$ in configuration Ξ :

$$K_\Xi(x, y) = \sum_i c_i(\Xi) \phi_i(x) \phi_i(y), \quad U(\Xi) = \sum_i u_i c_i(\Xi). \quad (4)$$




Figure 2: Stationary self-consistency for identical frequencies and phase diffusion $D = 1$. The positive coherent branch appears continuously above $K_c = 2D$ and increases with coupling. Zero remains a solution above threshold. The curve is computed from the integral expressions used in the formalization.

83 The same configuration sets the reciprocal kernel and stored energy. An
84 executed transition changes the kernel through its changed occupancies and
85 requires work $U(\Xi') - U(\Xi) + Q$, with Q positive into the reservoir. The
86 installed kernel is averaged over the resulting configuration law $p(\Xi)$. Cell-
87 pair reconstruction on refining spatial partitions converges to its product-
88 measure integral under the specified compactness and sampling assumptions;
89 the refinement index is independent of the number of control updates.

90 Normalize the substrate measure μ to one. Write $m_i = \int \phi_i d\mu$ and
91 declare $\kappa > 0$ with $m_i^2 \leq \kappa u_i$ for every mode. Factorizing each mode's

92 integral gives

$$\begin{aligned}
 U_{\text{inst}} &= \sum_{\Xi} p(\Xi) U(\Xi), \\
 K_{\text{mf}} &= \iint K(x, y) d\mu(x) d\mu(y) = \sum_{\Xi} p(\Xi) \sum_i c_i(\Xi) m_i^2 \leq \kappa U_{\text{inst}}.
 \end{aligned}
 \tag{5}$$

93 If this mean coupling is the K of the scalar model, $\kappa U_{\text{inst}} \leq 2D$ excludes a
 94 coherent stationary branch and any growing Fourier mode about the uniform
 95 density. Thus coherence on that branch requires $U_{\text{inst}} > 2D/\kappa$. The condition
 96 is necessary: energy above it need not produce sufficient coupling because
 97 the bound can be loose.

98 The theorem concerns stored energy in the declared actuator modes. The
 99 estimated cortical signalling expenditure and extracellular electrostatic en-
 100 ergy in Supplemental Material measure different quantities; neither identifies
 101 the occupancies, mode prices or their relation to metabolic supply. Conse-
 102 quently the bound supplies a necessary condition for a declared actuator,
 103 but no cortical value of κ or demonstration that cortex meets it. Applying
 104 the scalar threshold to heterogeneous or signed spatial dynamics also needs
 105 a justified reduction.

106 4 Compatibility: from coordinated activity 107 to a global description

108 4.1 Why common timing is insufficient

109 For overlapping patches U and V , compatibility means that restricting either
110 local description to $U \cap V$ gives the same object. Descriptions can agree on
111 shared quantities while contributing different information elsewhere. Com-
112 patible sections of the specified sheaf have a unique global extension. On
113 a finite cover, local finite spatial measures agreeing as measures on every
114 overlap also have a unique finite-measure extension whose sheaf image is the
115 glued section. No mass-one normalization is imposed.

116 In the cup example, visual and reaching descriptions must express a
117 shared quantity, such as cup location, in common coordinates. Restriction
118 retains that quantity while dropping details particular to each description.
119 Agreement permits surface appearance and hand configuration to contribute
120 to one situation. The mathematical counterexample assigns the same phase
121 to every patch while two local mass profiles disagree on a shared site. Those
122 profiles have no common extension. Because they are chosen independently
123 of phase, synchrony places no constraint on their disagreement.

124 This distinction limits the biological interpretation. Receptive territories
125 and content domains must be related to spatial patches, and the descriptions
126 must support the required extension. Matching overlap marginals of de-

127 coded probability laws need not determine a unique global joint distribution
 128 (Abramsky & Brandenburger, 2011).

129 The cover is itself a choice, and the construction constrains it in one way
 130 only: synchronized patches agree where they overlap. Where they do not
 131 overlap there is nothing to compare, so a cover by pairwise disjoint patches
 132 satisfies that obligation for *every* assignment of local descriptions and phases
 133 and the gluing theorem still returns a unique global section, which is the
 134 tuple it was handed. Membership of the class then reports nothing about the
 135 system beyond the patches covering it. What a gluing claim asks of a cover
 136 is therefore a non-degeneracy the theorem does not assume: that distinct
 137 patches meet, and that what they share is a quantity fixed independently of
 138 the agreement it is being used to establish.

139 4.2 When coherence does constrain content

140 A declared relation between phase and content gives a positive result. Let
 141 contents take values in a real normed vector space $\mathcal{C}$. Patch i reads an
 142 assigned oscillator through $e_{i,x}$ at site x , evaluated on the circle point
 143 $(\cos \theta, \sin \theta)$. Assume a uniform Lipschitz constant L in chord distance and
 144 require $e_{i,x} = e_{j,x}$ wherever patches i and j overlap. Encoders may vary across
 145 sites; their equality on overlaps is a supplied relation between descriptions of
 146 the shared quantity. The order-parameter identity $N^2 r^2 = \sum_{ij} \cos(\theta_i - \theta_j)$
 147 implies $1 - \cos(\theta_i - \theta_j) \leq N^2(1 - r^2)$ for each pair. Consequently every

148 overlap discrepancy is bounded by

$$\delta_r = L\sqrt{2} N\sqrt{1-r^2}. \quad (6)$$

149 Here discrepancy is the norm of the difference, and r and N are read on
150 whichever population the bound is applied to. That choice is what decides
151 whether the estimate says anything. A patch carries an order parameter of
152 its own, and reading the bound there carries that patch's resultant and that
153 patch's site count, so a large population observed through small patches is
154 constrained by the small count. Read instead on the whole array, the factor
155 N makes the estimate weak, because it draws a pointwise guarantee from
156 an average over N^2 pairs. At $r = 1$ the observations agree exactly at either
157 scale. A shared scalar encoder is a special case. Dropping the overlap rela-
158 tion permits disagreement even at perfect locking. A common nonexpansive
159 linear projection preserves agreement, but agreement of a projection does
160 not establish agreement of the full vector.

161 What Eq. (6) bounds is a chord, which never exceeds 2, so the estimate
162 constrains anything only where the right-hand side falls below that. At
163 the sheet's own patch order, the bound is informative for at most 3 sites;
164 beyond that, the coherence-to-content link requires the bridge assumption
165 (E78) of the cortical hypothesis. Supplemental Material provides the detailed
166 derivation of this saturation, and states three conditional derivations that
167 replace the uniform factor N by a count of disagreeing pairs, a chain across

168 a cover's shared sites, or a spectral gap at a frequency-locked configuration.

169 Mix retained fields with observations at gain $\eta \in [0, 1]$: $s_i \mapsto (1 - \eta)s_i +$
 170 ηo_i . Content and observation discrepancies bounded by ε and δ give a next
 171 discrepancy at most $(1 - \eta)\varepsilon + \eta\delta$. Repeated observations bounded by δ
 172 therefore give

$$\varepsilon_n \leq (1 - \eta)^n \varepsilon_0 + \delta. \quad (7)$$

173 Thus agreeing observations preserve exact compatibility and, for $\eta > 0$, re-
 174 duce an initial mismatch. Approximate observations leave a nondecaying
 175 term in the error bound. Taking a uniform bound on Eq. (6) connects this
 176 update to the phase process.

177 4.3 What approximate agreement determines

178 Exact gluing requires exact restrictions. For nonnegative profiles on a finite
 179 substrate, suppose every overlap discrepancy is at most ε in uniform distance.
 180 Weights that vanish outside their patches and sum to one at each site select
 181 a global profile by weighted averaging. The selection lies within ε of each
 182 patch there, and any two such selections differ by at most ε , independently
 183 of cover multiplicity. An L -Lipschitz readout consequently varies by at most
 184 $L\varepsilon$ across selections.

185 At zero tolerance, the selection is the unique exact extension. At every
 186 positive tolerance there are families with no exact extension; selection then
 187 gives an approximation whose dependence on weights is bounded. Supple-

188 mental Material gives a three-site witness attaining the bound. The mass-
 189 profile construction uses the scalar case of the content dynamics and addi-
 190 tionally requires nonnegativity of the initial profiles and observations. The
 191 general vector result supplies no identification with the spatial-measure sheaf.

192 These constructions concern declared objects: a specified cover, contents
 193 valued in a normed space, and encoders related on overlaps. They select
 194 neither the neural cover nor the decoder, and identifying their contents with
 195 perceptual content, or establishing that the identification bears on experi-
 196 ence, is a further modelling task; Section 8 sets out what a measurement of
 197 it would have to establish.

198 **5 Self-representation: a state that models it-** 199 **self**

200 **5.1 The reflexive construction**

201 Compatibility gives a global state. Reflexivity relates it to an encoding within
 202 that state. Equip global sections with a complete metric. If the additional
 203 restriction-resonance condition identifies encoding with restriction to a region
 204 $\mathcal{A}$, the reconstruction map is

$$F(s) = \text{readout}(s|_{\mathcal{A}}). \tag{8}$$

205 The region and readout are fixed across the state space. If F contracts and
206 fixes the section s_* formed from the local descriptions, Banach's theorem
207 makes s_* its unique fixed point (Banach, 1922). The formal bridge E89 sup-
208 plies the contraction and $F(s_*) = s_*$; restriction resonance is separate. Su-
209 percritical coupling makes the stipulated contraction factor $\Lambda = e^{-(K-2D)\tau/2}$
210 less than one for a relaxation time $\tau > 0$. Constructing a map with that
211 factor remains a modelling obligation.

212 In the cup example, the global state includes the represented scene and the
213 perceiver's relation to it. A local encoding must retain enough information
214 for the readout to recover that modelled state; cup position alone would be
215 insufficient if hand configuration and viewpoint could not be reconstructed.

216 For the same map with Lipschitz factor $\Lambda < 1$, approximate reconstruc-
217 tion gives

$$d(s, F(s)) \leq \varepsilon \implies d(s, s_*) \leq \frac{\varepsilon}{1 - \Lambda}. \quad (9)$$

218 At the stipulated rate this radius grows as $2\varepsilon/((K-2D)\tau)$ near threshold, for
219 fixed positive ε and τ . Exact uniqueness survives while the error guarantee
220 weakens.

221 5.2 Reconstruction across a family of states

222 Sensitivity can be required quantitatively. Declare relevant macrostates, an
223 encoder Φ and readout R , and require $d(s, R(\Phi(s))) \leq \varepsilon$ throughout that

224 family. If two states have the same code, the triangle inequality gives

$$\Phi(s) = \Phi(t) \implies \max\{d(s, R(\Phi(s))), d(t, R(\Phi(t)))\} \geq \frac{1}{2}d(s, t). \quad (10)$$

225 Thus M relevant states pairwise more than 2ε apart require at least M
226 distinguishable codes; an alphabet of at most 2^b codes requires $M \leq 2^b$.
227 Applied to the reflexive boundary, E is its local encoding and $R \circ E$ is the
228 existing self-prediction map.

229 This is an additional accuracy criterion on a family, with no requirement
230 for multiple exact fixed points. A blind encoder or constant readout fails it
231 for separated states despite having a unique fixed point.

232 Where the encoding is read as the state's own restriction to a region, the
233 criterion becomes a condition on the region. Restriction resonance requires
234 the encoder and that restriction to identify the same pairs of states, so a
235 boundary reconstructing two relevant states more than 2ε apart while its
236 region reads them identically is not resonant. Read forwards, a resonant
237 boundary that reconstructs a pairwise separated family restricts injectively
238 on it. Which states a region reads identically is fixed by where the region
239 sits, before an encoding is chosen, so the resolution the family demands is
240 the region's and no readout supplies it. The implication runs one way: a
241 region that does resolve the family is not thereby shown to carry a resonant
242 encoding, and E89 supplies no such resolution.

243 Those two demands are jointly satisfiable with nothing folded anywhere.

244 Take the region to be the whole substrate: restriction is then the identity, so
 245 the boundary is resonant, reconstructs every declarable family exactly, and
 246 restricts injectively on all of them. A properness condition on the region is
 247 therefore a separate requirement rather than a consequence of resonance and
 248 resolution, and what a fold into a local region is asked for. Its quantitative
 249 form replaces the alphabet by a metric. With an L -Lipschitz readout, a
 250 family reconstructed within ε whose members lie pairwise more than $d_{\min}$
 251 apart has δ -separated codes whenever $\delta L \leq d_{\min} - 2\varepsilon$, so the family is no
 252 larger than the δ -packing number of what the region carries. That is the
 253 shape a continuous family of states needs, and at the whole substrate it says
 254 nothing, because the region's range is everything. The gain L is declared and
 255 derived here for no physical readout, and at unbounded gain the bound says
 256 nothing.

257 **5.3 Agreement after a change takes communication**

258 Snapshot compatibility does not specify how quickly descriptions respond
 259 when a shared quantity changes. In a synchronous network whose updates
 260 read only the site's own state and declared neighbours' messages, two initial
 261 configurations agreeing on a site's causal past through T rounds produce the
 262 same state there at round T . The deadline consequence is the one Section 2
 263 states, and it follows from the reconstruction bound in Section 5.2: a report
 264 from that site is a readout of a code that the two interventions leave identical.

265 A three-site delay line witnesses failure after one round and exact report-

ing after two; cutting the path prevents delivery at every round. A report
 may still happen to be correct about one fixed world. The bound concerns
 responsiveness across interventions.

Throughout this section *prediction* names a reconstruction map, and the
 relevant family of states, the metric, the tolerance, the communication graph
 and the deadline are all declared inputs, as are the represented variables and
 the fidelity counted as accurate. Reading the map as temporal neural pre-
 diction would require dynamics; reading the resulting code count as physical
 memory capacity would require a hardware model, which a state-dependent
 readout would have to extend.

6 What the conditions constrain in a digital candidate

The results of Section 2 hold for any candidate, and for a large language
 model (LLM) running on graphics processing units (GPUs) four of them are
 checkable rather than merely stateable. Those four are where this framework
 bears on claims made for such systems. It does not exclude them, and no re-
 sult here could; what it does is fix what a claim of unity or self-representation
 commits its maker to measuring.

A quantitative bound for a linear bottleneck. Let a finite spatial
 coupling operator K have orthonormal eigenmodes e_i with $|\lambda_1| \geq \dots \geq |\lambda_N|$.

286 Every linear approximation A of rank at most d satisfies

$$\sum_{i=1}^N \|Ke_i - Ae_i\|^2 \geq \sum_{i=d+1}^N \lambda_i^2. \quad (11)$$

287 Retaining the first d modes attains equality. A linear encoder and readout
 288 factoring through d real coordinates therefore pay this error. Applying the
 289 bound requires a spectrum for the target and a rank constraint on the actual
 290 approximation. Query/key dimension alone supplies no such constraint on
 291 softmax attention: even a two-position rank-one logit matrix can yield rank-
 292 two attention. The result prices an explicitly linear bottleneck in a finite
 293 representation of a spatial kernel; it gives no universal gap between trans-
 294 formers and continuous fields. Supplemental Material derives the bound and
 295 checks the softmax counterexample.

296 A fixed synthetic task makes the scope visible. Two local views share
 297 scene and self coordinates; the target is the global state and a self-relative
 298 coordinate. On the parity-held-out family, a rank-two linear decoder’s input-
 299 weighted predicted training-error floor is 2.000, and its held-out error is 2.000.
 300 Equal numerical values do not make the training-law floor a held-out guaran-
 301 tee. On a separate scene-family shift, the fitted full linear decoder has train-
 302 ing error 0.000 and held-out error 2.667. The task’s declared target map
 303 reconstructs the noiseless held-out inputs with error 0.000; the restricted
 304 training family leaves one direction unidentified. A small causal attention
 305 decoder and a coupled phase code also lose held-out accuracy on that shift

under their stated CPU fits; neither is governed by the linear rank floor. The measured codes, overlap perturbations, self-state interventions and communication deadlines are described in Supplemental Material. These task results rank no architecture family and supply no hardware-energy or experiential claim.

Codes, not reports. Reconstruction across a declared family requires the encoder to distinguish states more than 2ε apart when the target error is ε (Eq. (10)). Verbal reports alone do not establish that condition. For a proposed self-model, declare the family and fidelity, identify the internal code under intervention, then test whether it separates the family. A token sequence is one possible code; an activation region is another and requires its own measurement.

The same interface has an information ceiling. For any finite joint law of an internal state X and a length- k token sequence Y over vocabulary V , $I(X; Y) \leq H(Y) \leq k \log_2 |V|$ bits. The proof uses the full joint law and allows dependent emissions and zero-probability outcomes. A binary copy attains one bit per emitted binary symbol; an independent output conveys zero. This counts the declared output sequence, so a cache, timing signal or other accessible state must be included separately when it forms part of the channel.

The region, not the readout. If the code is a restriction of the candidate’s state to a region, that region must distinguish the relevant family

328 before a readout is chosen. In one forward pass of a decoder-only transformer,
329 the causal mask restricts position t to positions at or before t ; an intervention
330 outside that causal past cannot be reconstructed from its activation alone.
331 Across passes, prior output adds a separate channel. The bounded-fan-in
332 communication result of Section 5 applies to a declared physical intercon-
333 nect and deadline; a full-support attention graph does not satisfy its sparse-
334 neighbourhood premise. Supplemental Material gives the graph bounds and
335 their hardware scope.

336 **Agreement about what.** Unity requires the candidate’s local descriptions
337 to agree on independently identified shared quantities (Section 4). Common
338 timing, a shared residual stream and fluent output supply none of that, for
339 the reason Section 2 gives: coordination bounds content disagreement only
340 through encoders already related on overlaps. Which quantities two parts
341 of a model describe in common is a decoding question, carrying the failure
342 modes of Section 8. The cover is a choice, and choosing it to secure agreement
343 empties the claim.

344 **Physical scope.** Digital representations can satisfy compatibility and re-
345 construction. The coupling-gated route additionally requires a calibrated
346 physical coupling, its noise regime and an independently identified relation
347 from phase order to shared content. An attention matrix or a shared clock
348 supplies none of these; the cortical proposal carries the same physical obliga-
349 tions. Finite-resolution capacity depends on code-space volume, resolution

350 and readout gain, so continuity alone favours no substrate. Supplemental
 351 Material gives the capacity and hardware-support comparisons. Neither the
 352 installed-energy condition nor those comparisons exclude a digital realization.
 353 The conditions impose measurements on a digital candidate but settle
 354 no question of machine experience. Supplemental Material places this com-
 355 parison beside biological naturalism, substrate-dependent computation and
 356 current proposals about LLM organization.

357 **7 Thermodynamics and composition**

358 The phase and representation results concern coupling, content disagreement
 359 and reconstruction error. A finite change in phase order has a probability-
 360 current cost bound under the conditions below. Thermodynamic heat re-
 361 quires a separate physical identification of the current and reservoir. The
 362 installed-energy condition of Section 3.2 is necessary for its declared actu-
 363 ator; whether it restricts a physical candidate depends on measured mode
 364 prices, coupling reduction and diffusion. No composition edge derives co-
 365 herence, compatibility or reconstruction from a heat bound. Finite-register
 366 capacity obeys $H(p) \leq \log |X|$; erasure heat depends on specified system-
 367 bath assumptions (Landauer, 1961; Norton, 2005; Shenker, 2000). A separate
 368 boundary result concerns a field reaching disconnected vacuum components
 369 (Kibble, 1976). Table 1 states the connecting assumptions.

370 **7.1 Changing order carries a probability-current cost**

371 For a smooth positive normalized periodic phase density ρ satisfying $\partial_t \rho =$
 372 $-\partial_\theta J$, with diffusion $D > 0$, define $\sigma_J = \int_{-\pi}^{\pi} J^2 / (D\rho) d\theta$ and $r = |\int e^{i\theta} \rho d\theta|$.
 373 Differentiation under the integral gives $\dot{r}^2 \leq D\sigma_J(1-r^2)$ wherever $0 < r < 1$.
 374 A smooth path staying in that interval therefore obeys

$$\int_0^\tau \sigma_J dt \geq \frac{[\arcsin r(\tau) - \arcsin r(0)]^2}{D\tau}. \quad (12)$$

375 Endpoint order, duration, diffusion and the probability current determine
 376 this floor. In an exploratory finite-volume onset calculation, a threshold cho-
 377 sen after inspecting pilot runs gives current cost 0.120 versus a floor of 0.046.
 378 These numbers describe the discretized model, not thermal or metabolic
 379 power.

380 The current $J = v\rho - D\partial_\theta \rho$ also satisfies

$$\left(\int_{-\pi}^{\pi} J d\theta \right)^2 \leq D\sigma_J. \quad (13)$$

381 A uniform density under constant drive ω attains $\sigma_J = \omega^2/D$. The identical-
 382 frequency gradient model has positive self-consistent stationary order above
 383 $K = 2D$ with $J = 0$ and $\sigma_J = 0$. Stationary order therefore has no positive
 384 current-cost floor. Nor does the bound price an externally chosen $\dot{K}$: in the
 385 identical-frequency zero-drive model, the uniform density remains an exact
 386 zero-current solution under any coupling schedule. A cortical power claim

387 needs the physical current, reservoir and thermal conversion to be identified.
 388 This current functional is distinct from the substrate squared-drift objective
 389 used in the plasticity controls.

390 **7.2 Feedback changes the information constraint**

391 For passive memory X_t driven by an environmental signal S_t , the predictive
 392 bound is

$$W_{\text{diss}} \geq k_B T [I(X_t; S_t) - I(X_t; S_{t+1})]. \quad (14)$$

393 The Markov structure $X_t \rightarrow S_t \rightarrow S_{t+1}$ makes the bracket nonnegative (Still
 394 et al., 2012). Zero retained information also satisfies it, so a task and a
 395 learning rule require additional specification. Relating this bound to Lan-
 396 dauer heat requires a common thermal scale and budget, and memory waste
 397 bounded by erased entropy.

398 An acting agent changes the signal it will receive. Specify policy, world,
 399 observation and memory-update channels. For their finite laws,

$$\begin{aligned} I(X_t; S_t) - I(X_t; S_{t+1}) &= I(X_t; S_t \mid S_{t+1}) \\ &\quad - I(X_t; S_{t+1} \mid S_t). \end{aligned} \quad (15)$$

400 The second conditional term vanishes in the passive model. With feedback,
 401 the signed difference can be negative as action creates correlations with the
 402 future. A causal reading requires an environmental state sufficient to describe
 403 the relevant dynamics.

404 For a controlled substep, the active thermodynamic bound instead uses
405 the joint entropy. Forward and reverse path laws give $\Sigma = D_{KL}(P_F \| P_R) =$
406 $\Delta H(X, S) + \langle Q \rangle / (k_B T) \geq 0$ under local detailed balance, with thermal scale
407 $k_B T > 0$ and heat positive into the reservoir (Horowitz and Esposito, 2014;
408 Ito and Sagawa, 2013). Every forward path of positive probability must have
409 a positive-probability reverse for this finite identity. A missing reverse gives
410 infinite divergence. The supplement develops the support conditions and
411 composes actuation with memory updates on their shared intermediate law.

412 **7.3 Learning and repeated operation need explicit re-** 413 **sources**

414 The finite models in Supplemental Material separate policy selection, learning
415 from a phase-derived sensor and repeated operation. A binary closed-loop
416 witness lets a register install the phase energy landscape, lets the phase evolve
417 under a bath-supported channel, senses the new phase and updates that same
418 register. A fixed decoder trained on one family predicts the next phase on
419 a held-out family with accuracy 0.783 under an informative sensor, versus
420 0.552 under a matched blind sensor. Their marginal phase order is the same.
421 The executed law also charges installation, sensing, reset and preparation
422 against a finite work store; phase-flip bath heat is recorded separately. The
423 symbol task and mode prices are supplied toy inputs, so this does not identify
424 a cortical mechanism or a heat cost for the full sensing and reset protocol.

425 For a complete cycle with a common energy observable $\mathcal{E}$ and thermal
 426 scale $k_B T$, positive forward and reverse path support, and local detailed
 427 balance at each stage, the conditional joint-law budget is

$$k_B T [H(p_0) - H(p_N)] \leq Q_{<N} = W_{<N} + \langle \mathcal{E} \rangle_0 - \langle \mathcal{E} \rangle_N \leq B + \langle \mathcal{E} \rangle_0 - \langle \mathcal{E} \rangle_N, \quad (16)$$

428 when expected cumulative work stays within the supplied allowance B at
 429 every prefix. This bounds expectations. The numerical witness specifies a
 430 bath for the phase transition; its sensing, reset and preparation prices do
 431 not establish reverse protocols or heat for the complete cycle. Separate con-
 432 structions track a store on each trajectory and specify when a mean allowance
 433 represents an available resource; an expected budget alone permits overdrawn
 434 paths.

435 For the declared process, a finite source with positive net cost funds
 436 finitely many cycles. A common physical ledger for the sensing and reset
 437 channels, and a cortical calibration of every mode price, remain separate
 438 obligations.

439 **7.4 What the composition establishes**

440 The passive and feedback composition theorems each accept eight connecting
 441 implications (Table 1). In the active branch, a named substep at the register's

442 thermal scale satisfies

$$k_B T > 0, \quad k_B T [H(p) - H(p')] \leq \langle Q \rangle \leq B, \quad (17)$$

443 with B the register's heat. The upper allocation follows from a ledger whose
444 operations deliver that total heat to one reservoir and whose other opera-
445 tions do not increase joint entropy. Those ledger assumptions are physical
446 inputs. Neither Landauer's lower bound nor microscopic reversibility sup-
447 plies them. Finite two-bit witnesses exhibit positive forward and reverse
448 elementary paths, local detailed balance at one thermal scale, and interac-
449 tion energies of order $k_B T$; these conditions alone do not quantify proximity
450 to equilibrium. Cortex has not been shown to satisfy these premises. The
451 illustrative cortical metabolic expenditure over one region and one phase-
452 diffusion time is 2.4×10^{14} times kT at body temperature, but identifying it
453 with the register heat budget B requires a common physical ledger. Read-
454 ing the register at the scale where its premises are physical, such as a single
455 channel, requires E45Active to carry that step's budget to the convergence
456 of a continuum kernel. The active branch is an alternative route to the same
457 node rather than a stronger one, and no result of Section 3.2 passes through
458 it: the installed-energy condition is derived from E56 and E67.

Table 1: Connecting hypotheses, eight per branch, and the installed-energy consequence. The branches share E12, E23 and E56–E89. Exact overlap compatibility and the encoding–readout structure are additional inputs. Supplemental Material, Table S1, maps the node-level results.

Connection	Status	Required assumption
E12 <i>capacity and boundary</i>	Physical premise	A finite-capacity system realizes a field leaving its vacuum manifold. Capacity and boundary are independent roots.
E23 <i>boundary to register</i>	Bridge assumption	The boundary is associated with a finite register whose update leaves states unreachable.
E34 / E34Active <i>register to dissipation</i>	Bridge assumption	Passive: predictive memory at the register’s thermal scale and heat budget, with nonzero waste. Active: positive paths, local detailed balance and allocation of register heat to the named substep; the allocation uses the ledger identity and other operations’ compressiveness. Both are exhibited at thermal-scale heats on finite state spaces with positive reverse paths, which is not the regime a cortical register occupies.
E45 / E45Active <i>couplings to continuum</i>	Modelling assumption	Convergent spatial coupling energies. Declared mode hardware supplies a kernel limit for either branch. A learning witness computes installation and work from its own executed law. Convergence follows from mesh refinement rather than from the budget, which selects no regime at macroscopic scale. Refinement and physical identifications remain inputs.
E56 <i>field identification</i>	Physical commitment	A specified physical field realizes the limiting coupling and satisfies the kernel conditions.
E67 <i>above threshold</i>	Physical commitment	That field operates above the synchronization threshold.

Table 1 (continued)

Connection	Status	Required assumption
E78 <i>coherence to cover</i>	Bridge assumption	Coherence relaxes the supplied cover at coupling at least the named mean-field value. Shared phase encoders bound vector-content mismatch; the exact cover and timely responsiveness remain separate requirements. Patches and oscillators carry one index, so one oscillator carries exactly one content patch: a modelling assumption rather than a consequence, and not an innocent one, since field-mediated coupling is spatial while two territories describing one quantity need not be neighbours. Where the cover's nerve is connected the same edge carries through the sites its patches share, at a residual linear in that nerve's diameter; where the configuration is frequency-locked it carries through the coupling graph's algebraic connectivity λ_2 instead of through any patch count, which asks the cover for nothing and is declared hardware data in the standing of κ .
E89 <i>contraction to fixed point</i>	Modelling assumption	A map mediated by local self-encoding fixes the glued section and has the stipulated Lipschitz factor $\Lambda < 1$. Reconstruction over a declared family imposes an additional code-separation criterion.
Installed energy	Conditional theorem	Under Section 3.2's occupancy and mode-price conditions, E56 and E67 require $U_{\text{inst}} > 2D/\kappa$. The learning witness clears this necessary bound; its matched control does not. Neither physical premise follows.

Both branches also take a thermodynamic cover carrying exact overlap compatibility and a reflexive boundary specifying encoding and read-out. They conclude that the glued section is the unique fixed point of self-

462 prediction. Compatibility and accuracy at that section enter as premises;
463 uniqueness follows. Identifying encoding with restriction needs the separate
464 resonance condition of Eq. (8).

465 Each of the eight edges is supplied rather than derived from the one above
466 it. E12 is equivalent to its conclusion, so capacity and boundary are inde-
467 pendent roots rather than a step from one to the other; both E45 discharges
468 take their convergence from refinement data and not from the thermody-
469 namic bound; E78 states that coherence relaxes the supplied cover, and the
470 coherence-to-content estimate of Eq. (6) bounds disagreement only for en-
471 coders already shared on overlaps. The inferential content of the composition
472 is its last step, where contraction and a fixed glued section give uniqueness.
473 What composing them establishes is that the premises are consistent, named
474 in one place and machine-checked together, and not that coherence follows
475 from installed energy or self-representation from coherence.

476 Joint witnesses combine a register, boundary, refining mesh and small
477 spatial cover; active witnesses add controlled thermal updates and a priced
478 installation. They establish satisfiability under their stated identifications.
479 Mesh convergence follows from refinement data, and the witnesses establish
480 no common cortical mechanism. Supplemental Material maps the claims to
481 their Lean declarations and records the counterexamples. The development
482 declares no additional axioms beyond its libraries' foundational axioms.

483 8 What the observations establish

484 The composition holds under its premises. Whether observation can estab-
485 lish those premises is a separate question, and the three requirements it
486 connects concern different objects: phase dynamics, local descriptions and
487 internal encodings. Their observables must retain these distinctions. The
488 following controls show which mechanisms remain indistinguishable, which
489 temporal approximations fail, which states one observable cannot separate,
490 and what independent content measurements must establish. Each is a finite
491 numerical regime on a declared model, with full parameters and controls in
492 Supplemental Material, and none of them calibrates a cortical quantity.

493 8.1 Stationary distributions leave the mechanism un- 494 determined

495 Equation (2) gives a concentration–coherence relation with no fitted shape
496 parameter. Estimate r from circular resultant length and a from density
497 shape: inverting R from the measured r would make agreement automatic.
498 Even independently estimated agreement identifies a statistical family shared
499 by different mechanisms.

500 Let $q(t)$ cross 2, let $r(q)$ be its stationary branch, and fix $D_0 > 0$. Increas-
501 ing coupling $K(t) = D_0 q(t)$ at diffusion D_0 and decreasing diffusion $D(t) =$
502 $K_0/q(t)$ at coupling K_0 both give concentration $a = q(t)r(q(t))$. Uncoupled
503 oscillators with prescribed drift $-h(t)\sin(\theta - \psi)$, where $h(t) = D_0 q(t)r(q(t))$,

504 and diffusion D_0 have that same density. Their stationary distributions co-
 505 incide window by window despite different causes.

506 Resolved dynamics supply additional information. In the ideal model with
 507 continuous paths, quadratic variation accumulates at rate $2D$ and conditional
 508 increment means constrain drift. Separating endogenous coupling from a
 509 matched common drive additionally requires a measured drive or controlled
 510 perturbation; identifying electromagnetic coupling requires evidence specific
 511 to that interaction. Supplemental Material gives the equivalence construction
 512 and its numerical check.

513 **8.2 Finite ramps change the observable onset**

514 The square-root law in Eq. (3) describes stationary solutions. A changing
 515 coupling requires time for the phase distribution to approach those solutions.
 516 In a deterministic mean-field control that crosses the threshold symmetrically,
 517 none of the scanned rates meets the declared tracking tolerance: slowing
 518 the crossing does not remove the terminal residual under that initialization.
 519 A slow control variable therefore supplies no general guarantee of branch
 520 tracking at onset. Supplemental Material gives the residual, rate range and
 521 admissible rates away from threshold.

522 The finite-population ramp study uses $N = 2000$, $D = 1$ and 32 replicas
 523 at 8 coupling speeds from $v = 0.1$ to $v = 0.0001$. Escape is the first of three
 524 consecutive saved samples at $r \geq 0.2$ after threshold. Faster ramps delay
 525 escape (Fig. 3); only 5/32 replicas escape within the fastest window. The

526 6 uncensored legs supply the delay scaling below; the decade-spaced ones
 527 among them have mean coupling delays 0.1947, 0.0684 and 0.0253. Fitted
 528 onset exponents differ from those of the exact stationary branch evaluated
 529 on the same grids. At the fastest speed the ensemble mean never reaches the
 530 fit window, so its fitted slope does not estimate stationary onset.

531 Delayed passage through a noisy bifurcation (Berglund & Gentz, 2002)
 532 makes the apparent exponent depend on ramp speed and fitting window.
 533 The specified deviation-based critical ramp speed is not bracketed by this
 534 sweep. With fluctuating coupling, the relevant stationary comparison also
 535 depends on rate: the mean coupling approximates order in the tested fast
 536 regime, while averaging stationary order over coupling works for the tested
 537 slow drives confined to one side of threshold. Slow crossing drives can satisfy
 538 neither approximation.

539 **The prediction.** Over the 6 uncensored legs a finite-range fit gives $\Delta K \propto$
 540 $v^{0.444}$. The interval from fitting those means is $[0.394, 0.494]$; resampling
 541 replicas independently within each speed gives a 95% interval $[0.398, 0.496]$,
 542 with fraction 0.019 of 10000 bootstrap fits at or above $1/2$. This fraction
 543 describes the resampling distribution under the fitted protocol; it is not a
 544 test of the ideal exponent. A deterministic mean-field run seeded at the
 545 fluctuation floor gives exponent 0.425 with the same estimator and escape
 546 criterion; tightening its criterion to $r \geq 0.05$ gives 0.490. These controls
 547 show that the criterion affects the fitted exponent, but do not prove that it

accounts for the entire finite-population shortfall. The larger-population low-threshold fit has substantial log residuals; matched-speed fits do not converge toward one-half across the tested sizes and criteria. Some trajectories also cross the absolute escape level before K_c , so the first postcritical crossing can yield a zero or grid-limited delay. This finite sweep establishes no common asymptotic exponent.

The ideal delayed-onset calculation motivates a square-root scaling hypothesis; its exponent is generic to suitable bifurcations and does not identify the coupling medium. A within-subject emergence design can vary the propofol taper and test how an operational onset changes with that manipulation, but interpreting the result as a coupling-rate law requires a measured or calibrated relation between taper and $\dot{K}$. The escape criterion must be set relative to each recording's pre-threshold fluctuation scale. Section 8.1 describes how D could be estimated from resolved phase traces.

Additional controls address integration bias, population size and spatial or signed coupling; the refined signed-network crossing bracket $K_{\text{eff}}/D \in [1.958, 2.009]$ contains 2.

8.3 A winding state leaves the order parameter ambiguous

The global order parameter r discards spatial arrangement. On a periodic sheet, a phase field can wind once around the sheet and have $r = 0$ while




Figure 3: Finite-rate departures from the stationary branch. Curves show ensemble-mean order for four ramp speeds v , with one-standard-error bands. Coloured points mark each replica’s first sustained escape, plotted at its recorded order; censored replicas have no point. The dashed vertical line is the static threshold, the dash-dotted horizontal line is the escape criterion, and the dotted horizontal line is the finite-size reference. The replica-bootstrap interval for the delay exponent is reported in the text.

neighbouring sites remain locked. At identical frequencies, every such winding is stationary for an isotropic symmetric kernel: contributions at opposite separations cancel. It violates the quarter-turn initial-spread premise of the convergence result, which is why that result can still select the uniform state.

In the spatial control, the winding persists below the interpolated decay length 0.1871 mm and unwinds above it (Fig. 4). At 0.1 mm, uniform and winding starts under the same coupling, diffusion and kernel end at $r = 0.9657$ and $r = 0.0083$; patch-local order is 0.9015. Under quenched fre-

577 quency spread 1.5 rad s^{-1} , the measured retention boundary is 0.1421 mm,
 578 and the corresponding order comparison is 0.9418 against 0.0348 over 3 frozen
 579 realizations. These are finite controls, not a cortical wave identification or a
 580 derived transition.




Figure 4: A winding phase field on the spatial sheet, at identical frequencies. Left: global order (solid) against patch-local order (dashed) across a subset of coupling decay lengths. Right: the twisted and uniform initial conditions compared across lengths, with the shaded 0.1–0.3 mm band and the interpolated retention boundary. Below that boundary the two initial conditions reach different stationary values of r at the same coupling, diffusion and kernel, so the resultant does not report the coupling. The sweep is sampled at the lengths marked; it establishes no transition and no biological identification.

581 Linearization about a winding reweights each separation’s coupling by
 582 the cosine of its phase advance. The resulting symmetric kernel has explicit
 583 mode rates: nonnegative reweighted coupling gives no growing mode, while
 584 a positive rate destabilizes that winding. The criterion concerns linear rates,
 585 not asymptotic stability, and the measured retention boundary need not
 586 equal its limit. An amplitude extension adds an existence condition: $\Gamma + \lambda >$

0, where Γ is linear gain and $\lambda = \sum_d f_d(\cos \Delta\theta_d - 1)$ comes from the same
kernel. For the nonnegative kernels used here, a broad kernel can remove a
winding from that model even though the phase-only equations retain it as
a stationary state.

Neither model turns these states into measured cortical spirals. The pe-
riodic sheet cannot carry a lone phase defect, the relaxational dynamics has
no propagation speed, and the amplitude variable has no calibrated relation
to a recorded envelope. Supplemental Material gives the spectral, amplitude
and defect results. Locally, Eq. (6) can still bound content disagreement on a
winding patch when shared encoders are supplied, although its global $r = 0$
form is uninformative. A nonzero ring winding number q also obstructs a
global real-valued phase lift; it has no identified content meaning here. Thus
an onset criterion based only on r can register a change of basin as well
as a change of coupling. The protocol of Section 9.3 needs a spatial phase
observable.

8.4 Coherent adaptation can lose environmental align- ment

An adaptive coupling rule might supply representational structure, but its
objective must reward the relevant agreement. For symmetric coupling, pair-
wise sine terms cancel in the population mean drift. Minimizing squared drift
at that fixed mean therefore rewards equalizing drift across nodes. In the

frequency-cluster control, this favors coupling between unlike-frequency clusters, whereas the diagnostic environmental template rewards coupling within clusters.

The formalization proves descent of the squared-drift objective for a specified gradient flow at fixed phases. Its functional differs from thermodynamic entropy production: at a coherent stationary density the probability current vanishes while phase-averaged squared drift is positive. Supplemental Material gives the operator, its gradient and the conditions of the descent result.

The finite control co-evolves noisy phases and symmetric, nonnegative couplings at fixed total resource. Its template is a diagnostic covariance of latent signals, never a training input. Across three illustrative seeds, second-half order remains between 0.9167 and 0.9222, while the within-cluster to between-cluster coupling ratio falls from 0.9824–0.9948 to 0.3507–0.6135 (Fig. 5). Random-update, permuted-gradient and frozen-kernel controls are reported in Supplemental Material.

The immediate post-update objective reduction disappears when averaged over complete update intervals. A prespecified grid of 6 rate–cadence combinations contains no case meeting the 5% sustained-reduction criterion while retaining coherence and losing alignment. The finding is coherent adaptation away from the specified structure, with no demonstrated sustained cost reduction.

Kernel alignment and overlap compatibility are different observables. This control addresses acquisition of its designated template under its par-

631 ticular objective. Learning models can explicitly include data in an objective
632 (Corberi et al., 2026); an account of agreement about the cup must specify
633 the content its learning rule is meant to preserve.




Figure 5: Coherence persists while coupling loses alignment with the frequency-cluster template. Gradient updates (solid), random updates (dash-dotted), permuted gradient updates (dash-dot-dotted), and a frozen kernel (dashed) share one phase-noise stream per seed. Panels show the immediate post-update squared-drift objective, phase order, within-cluster to between-cluster coupling, and template distances. The gradient arm preserves order while its coupling ratio falls; growth in both distances mainly reflects kernel norm. The objective panel samples the instant after updates; complete-interval measurements find no sustained reduction in the bounded study.

634 8.5 Content agreement needs independent decoders

635 A neural test needs independently identified domains, local descriptions and
636 restriction maps. Two decoders should estimate a shared content on over-
637 lapping domains before their agreement is examined. Discrepancy requires a
638 null under decoding noise and controls using deliberately compatible and in-
639 compatible examples. An implementation on synthetic data separates these
640 cases at $AUC = 0.986$, on a 95% interval of $[0.980, 0.992]$ resampled by trial;
641 the compatible construction scores 0.476 on an interval containing one half.
642 Shared-prior shrinkage can nevertheless manufacture agreement, and regional
643 bias can manufacture disagreement. Per-patch decoding diagnostics detect
644 these failures; the compatibility statistic alone does not.

645 Phase agreement cannot replace that measurement. On constructed con-
646 tent, phase carries a share 1.000 of the content entropy when content is built
647 from it and 0.000 when it is not. A phase-derived statistic reports com-
648 patibility on phase-locked patches whose local descriptions disagree. The
649 implemented content statistic detects their incompatibility. These controls
650 validate a measurement on known constructions; a neural content model and
651 decoder remain to be established.

652 Choosing the cover also chooses a candidate subject: trivial descriptions
653 can agree while carrying little content, and a partition can conceal disagree-
654 ment. The useful choice must therefore be constrained by what neural ac-
655 tivity encodes. The exact composition assumes compatibility of its supplied
656 cover. The content mechanisms above identify additional hypotheses under

657 which agreement can arise; they leave the empirical choice of content and
658 the connection to that cover explicit.

659 **8.6 Estimator calibration on exploratory EEG**

660 The concentration range and the dependence structure decide whether the
661 concentration-coherence relation can be tested at all. Under the specified
662 log-density estimator and bias-floor criterion, discrimination at the observed
663 concentration ceiling $a = 0.542$ requires 3000 independent sites in the scanned
664 designs, rising to 31000 when sites share a mean in groups of 10. These
665 counts describe that estimator and dependence model; the intended recording
666 requires its own calibration. Supplemental Material reports the bin-count
667 controls and the full dependence-aware design.

668 That ceiling comes from exercising the estimators on public propofol-
669 sedation EEG (Bajwa et al., 2025), pooling channels in 100-ms win-
670 dows. Bipolar and common-reference montages produce different residuals.
671 Bipolar-montage scalp EEG sits in the linear Bessel regime ($a_{\max} = 0.542$,
672 deviation from $a/2 < 1\%$). The exercise confirms the estimator is well-
673 behaved in this regime but does not test the nonlinear prediction. The
674 bipolar results are compatible with the Bessel relation and with monotone
675 alternatives that agree to this precision over the sampled range. Temporal
676 pooling changes the estimand, and sedation recordings leave the awakening
677 hypothesis untested.

678 An informative dataset therefore needs sufficient concentration range or

precision, calibrated estimators and an observation model for dependence,
source mixing and reference choice. Intracranial recordings (ECoG, sEEG or
LFP) or narrowband burst analysis may reach a more informative concentration range, but that range must be measured. At the observed ceiling, the specified design requires 3000 independent sites or 31000 sites under its clustered dependence model; these are design calculations, not sample requirements transferable to another recording. Public mouse calcium imaging provides thousands of spatially located neurons across waking, sleep and anesthesia (Oomoto et al., 2026), but its 7.65 Hz frame rate is below the phase-estimation band used here. Site count, state coverage and bandwidth are separate requirements.

9 The cortical hypothesis and an awakening test

The cortical proposal identifies a physical source of coupling and asks whether its changes help explain recovery. The tests in Section 8 determine what observations must distinguish before a phase trace can support this identification.

696 9.1 Physical identification and calibration

697 The cortical hypothesis identifies part of effective coupling with endogenous
698 extracellular electromagnetic interactions. Fields can modulate spike timing
699 (Fröhlich & McCormick, 2010; Anastassiou & Koch, 2015), and field models
700 propose coordination of distributed neural information (Pinotsis & Miller,
701 2022; McFadden, 2020). Phase organization under propofol motivates study-
702 ing this regime (Bastos et al., 2021), but does not measure coupling.

703 The formal field predicate requires a jointly continuous kernel, a normal-
704 ized substrate, coupling equal to the kernel average, and no single site pro-
705 viding more than half that strength. Toy constructions satisfy it. A cortical
706 realization must specify the scale at which the kernel and its reduction ap-
707 ply: calcium imaging reports different state-dependent modularity patterns
708 at single-cell and mesoscale resolution (Kiyooka et al., 2026).

709 Cortical columns are the obvious candidate unit, and they answer one
710 of the two demands the formalism makes. As the oscillator index they are
711 a serious proposal: they are delineated by response properties and anatomy
712 rather than by the agreement they would be used to establish (Mountcastle,
713 1997), and a phase can be estimated from a column-sized territory at the
714 electrode spacings used for cortical wave measurements (Townsend et al.,
715 2015). As the content cover they are a poor one, for three separate reasons.
716 A parcellation into non-overlapping columns is a disjoint cover, which satisfies
717 the compatibility obligation without constraining anything (Section 4). The
718 cover condition additionally requires the patches to exhaust the substrate,

719 while much of association cortex has no accepted columnar parcellation. And
720 the column's standing as a canonical cortical unit is contested (Horton &
721 Adams, 2005; Rakic, 2008). Offering columns for both roles at once therefore
722 supplies the oscillator index; the content cover remains an open empirical
723 question, which the bridge assumption E78 records.

724 The spatial reduction is also testable. A row-sum aggregation at zero
725 diagonal tracks the uniform-coupling control over the tested spatial decay
726 range, while frequency heterogeneity displaces the operational threshold. The
727 fit window biases both estimates relative to $2D$; agreement with the control
728 is therefore the supported conclusion. Supplemental Material reports the
729 reduction, its numerical error and the distinct assumptions needed for the
730 kernel-average identification in Eq. (5).

731 The field commitment requires physical evidence beyond a choice of math-
732 ematical representation, and the available support comparisons do not supply
733 it. A finite matrix spread over positive-mass cells preserves its total cou-
734 pling and phase correlation exactly, while point-supported vanishing under
735 an atomless measure and a disadvantage from excluding a better-correlated
736 pair concern specified support constraints; Section 6 reads the same compar-
737 isons against a digital candidate. Neither supplies energy prices.

738 A spike-time shift alone supplies no coupling rate. The dimensionless
739 product of population size N , timing sensitivity $\mathcal{T}$, field amplitude E and

740 frequency f needs a conversion rate γ :

$$K_{\text{eff}} = \gamma N \mathcal{T} E f. \quad (18)$$

741 This rate describes how induced timing shifts accumulate into drift. Nei-
742 ther it nor the installation factor κ follows from a heat bound. Testing the
743 physical threshold requires independent estimates of the interaction kernel,
744 its aggregation and diffusion in the same cortical regime.

745 **9.2 Extracellular geometry and recovery**

746 The proposed control variable is extracellular geometry. Extracellular vol-
747 ume fraction is approximately $\alpha = 0.2$ at tortuosity $\lambda \approx 1.6$ (Syková &
748 Nicholson, 2008); a sleep study reports a rise from roughly 0.14 to 0.23 (Xie
749 et al., 2013). Below 10 kHz, cortical tissue conducts approximately resistively
750 at $0.3\text{--}0.6 \text{ S m}^{-1}$ (Barbour, 2017). At fixed transmembrane current density,
751 an effective-medium description predicts that a larger extracellular fraction
752 raises conductivity and reduces field amplitude.

753 The coupling-gated hypothesis consequently predicts that extracellular
754 contraction on awakening raises E and, holding the other factors in Eq. (18)
755 fixed, raises effective coupling. This argument supplies a direction; the
756 geometry–coupling relation, its magnitude and its role in human recovery
757 require measurement.

758 For a smooth crossing at nonzero speed, evaluating the stationary branch

759 along the changing control gives a deterministic tracking approximation that
 760 fails at onset in the tested crossing control:

$$r_{\text{ss}}(K(t)) \sim \sqrt{\frac{\dot{K}(t_0)}{D}} (t - t_0)^{1/2}, \quad K(t_0) = K_c, \quad \dot{K}(t_0) > 0. \quad (19)$$

761 This is a deterministic tracking approximation from Eq. (3) at fixed D ; the
 762 crossing controls in Section 8.2 show that it fails as an observed-onset ap-
 763 proximation under the tested initialization. A separation of geometry and
 764 phase time scales can support tracking away from threshold. The awakening
 765 hypothesis needs a calibrated dynamical prediction for an operational onset
 766 above the finite-size fluctuation floor. Crossing K_c alone does not supply a
 767 recovery latency.

768 **9.3 A protocol for the awakening experiment**

769 Neural inertia describes a behavioural barrier and hysteresis between anaes-
 770 thetic induction and emergence (Friedman et al., 2010); recovery can pass
 771 through metastable activity states (Hudson et al., 2014; Proekt & Hudson,
 772 2018). The $v^{1/2}$ exponent of an ideal delayed bifurcation is generic (Baer
 773 et al., 1989; Berglund & Gentz, 2002) and cannot alone discriminate a field
 774 mechanism from a synaptic one. A possible discriminator is spatial onset
 775 relative to independently estimated local field coupling and network connec-
 776 tivity. The present mean-field model has no spatially varying κ and derives
 777 no nucleation order, so a field-based onset map needs a separate model spec-

778 ified before the observations are inspected.

779 For that comparison, partition the recording into prespecified regions i
780 and measure pre-onset order $r_i(t_0)$ and diffusion D_i . Define a field score
781 $h_i = N_i \mathcal{T}_i E_i f_i$ from independent pre-onset field and timing measurements,
782 and a connectivity score c_i from a separately estimated structural network.
783 For the same measured emergence drive $g(t)$, compare $K_i^F(t) = a_F g(t) h_i$ with
784 $K_i^C(t) = a_C g(t) c_i$. Each model has one fitted scale. The local first-harmonic
785 approximation predicts

$$\hat{r}_i^M(t) = r_i(t_0) \exp\left(\int_{t_0}^t [K_i^M(s)/2 - D_i] ds\right), \quad M \in \{F, C\}. \quad (20)$$

786 Predicted onset is the first crossing of a region-specific threshold fixed from
787 pre-onset fluctuations. Fit the two scales on training sessions, then com-
788 pare held-out onset times and site order under the same observation model.
789 Larger field scores predict earlier onset only when baseline order and diffu-
790 sion do not reverse the ranking. This declared approximation omits interre-
791 gional coupling and saturation; a successful ranking would support its field
792 covariate, not prove ephaptic causation. Supplemental Material specifies the
793 controls and failure conditions.

794 A within-subject protocol can manipulate emergence rate, with spatially
795 resolved recordings to compare the proposed field and connectivity models:

- 796 **1. Manipulate emergence rate.** Vary the propofol infusion taper
797 across sessions within the same subject, recording drug concentrations

798 or a validated pharmacokinetic proxy to quantify the manipulation.

799 **2. Measure behavioural recovery.** Record time-to-response to a stan-
800 dard stimulus as the independent behavioural endpoint, rather than
801 estimating ΔK directly.

802 **3. Resolve spatial onset.** Identify anatomical locations where coher-
803 ence first crosses the baseline-scaled threshold and score the two maps
804 of Eq. (20) on held-out sessions. The two covariate maps are fixed
805 before the outcome sessions are inspected.

806 **4. Match physical controls.** Confirm that the observed coherence-
807 to-amplitude relation holds under the specific local diffusion D and
808 observation noise, estimable from the quadratic variation of resolved
809 phase traces.

810 Predeclare the recording window, regional fluctuation thresholds, and a
811 shared timing-error model. Fit the two scales and that observation model
812 on training sessions only; on held-out sessions, score the observed site-onset
813 times and sites that never cross by the same regional predictive rule. Report
814 site order as a secondary outcome and behavioural time-to-response sepa-
815 rately. Sessions without a crossing are censored observations, not zero delays
816 or exclusions. Compare the two scores only where independently estimated
817 field and connectivity maps predict distinguishable onset patterns at the
818 recording resolution. A held-out field-model advantage would support that
819 specified spatial predictor; it would not identify an ephaptic cause, shared

820 content or experience. Supplemental Material gives the controls and pilot
821 quantities needed to judge whether the comparison is informative.

822 **10 Discussion**

823 **10.1 Unity and a minimal self as interpretive motiva-** 824 **tion**

825 We propose the state glued from compatible local descriptions as a formal
826 correlate of conscious unity, and its accurate internal reconstruction as a cor-
827 relate of a minimal self. In the cup example, this would mean one description
828 that preserves the cup’s relation to the perceiver as well as the scene. Com-
829 patibility permits distinct contents and uncertainty; the shared quantities
830 must be chosen from independent evidence about what each region encodes.
831 The reconstruction need not be a verbal or autobiographical account. Its
832 proposed relation to experienced perspective remains an interpretation. Nei-
833 ther identification is among the eight hypotheses in Table 1, and satisfying
834 them would not measure experience.

835 **10.2 What a test of the full proposal must establish**

836 A test needs separate estimates of effective coupling and diffusion, shared
837 content on overlapping regions, and reconstruction from a specified encoding
838 at a declared fidelity. The cover, relevant state family and decoder must

839 be fixed without selecting them for agreement. An independent experiential
840 endpoint is needed to evaluate the interpretation; absence of a report alone
841 does not define absence of experience. A unified episode with incompati-
842 ble decoded descriptions would challenge the unity identification. A unified
843 episode lacking the proposed self-encoding would challenge the minimal-self
844 claim. Conversely, satisfying the mathematical conditions without such an
845 episode would challenge their sufficiency as correlates.

846 For the physical route, cortical actuator modes and prices, the geometry-
847 coupling relation and a dynamical onset model remain to be measured. The
848 field/connectivity comparison in Section 9.3 evaluates one spatial onset pre-
849 diction. Controlled changes in coordination and independently decoded con-
850 tent agreement would address the representational connection. The synthetic
851 decoder and EEG exercises calibrate parts of these measurements; they do
852 not identify a neural content code or test experience.

853 **10.3 Explanatory scope**

854 Integrated-information, global-workspace and predictive-processing accounts
855 organize unity, access and modelling differently (Tononi et al., 2016; Baars,
856 1997; Dehaene et al., 2014; Friston, 2010). This account specifies conditional
857 relations among coupling, local content and reconstruction. Its formal field
858 condition requires a continuous kernel, a normalized substrate, a specified
859 average coupling and noise, and a site-contribution bound. Those conditions
860 do not identify an electromagnetic substrate; the extracellular interpretation

needs the separate calibration of Section 9. The joint witness shows that the formal hypotheses can hold together in a declared model. A cortical realization and an experiential interpretation require additional evidence.

The remaining explanatory question is whether these relations occur in neural activity and matter for experience. The framework identifies candidate correlates and tests of their physical and representational premises. It does not explain why satisfying them would be accompanied by experience.

Notation

Table 2: Symbols used throughout, with the section that defines each. Where one letter carries more than one meaning, the meanings are separated by the section they are read in.

Symbol	Defined	Meaning
θ_i, ω_i	Eq. (1)	Phase and natural frequency of oscillator i
N	Eq. (1)	Number of oscillators in the population the quantity is read on
D	§2, Supp.	Phase diffusion in inverse time; phase spread $D \leq \pi/2$
K_{ij}, K	Eq. (1)	Pairwise and scalar mean-field coupling, $K_{ij} = K/N$
K_c	Eq. (2)	Synchronization threshold, $K_c = 2D$
r, r_N	Eq. (1)	Resultant length (order parameter), on the population it is read on
a, ψ	Eq. (2)	Concentration $a = Kr/D$ and collective phase of the stationary density
q	Section 8.3	Winding number of a phase field read around a ring
μ	Eq. (4)	Substrate measure, normalized to one
Ξ, ϕ_i, u_i, c_i	Eq. (4)	Actuator configuration; mode profile, installation price and occupancy of mode i

Table 2 (continued)

Symbol	Defined	Meaning
U_{inst}, κ	Eq. (5)	Installed energy and the mode-data conversion factor bounding $K_{\text{mf}} \leq \kappa U_{\text{inst}}$
L	Section 4.2	Lipschitz constant of an encoder or readout
ε, δ	Section 4.2	Content and observation discrepancy tolerances
M, b	Eq. (10)	Size of a declared family of states and code length in bits
Λ, τ	Eq. (8)	Contraction factor $\Lambda = e^{-(K-2D)\tau/2}$ and its relaxation time
T	§3.2, §5	Temperature in the thermal scale $k_B T$; number of communication rounds to a deadline
ν	Section 6	Fan-in of a communication graph: the most sites any one site hears from
$k_B T, B$	Eq. (17)	Thermal scale, read as one symbol, and heat budget of a register
$\gamma, \mathcal{T}, f$	Eq. (18)	Field-to-coupling conversion rate, timing sensitivity and frequency
α	Section 9	Extracellular volume fraction
λ	§8.3, §9	Eigenvalue the coupling kernel supplies at a winding; extracellular tortuosity
Ω, λ_2	Section 4.2	Common rotation rate of a frequency-locked configuration, and the algebraic connectivity of the coupling graph
β, k	Section 4.2	Chord diameter of a patch and the diameter of a cover's nerve, in hops
$\mathcal{C}, e_{i,x}$	Section 4.2	Content space and the phase encoder patch i reads at site x
Φ, R	Eq. (10)	Encoder and readout of the reconstruction criterion
$\mathcal{A}, F$	Eq. (8)	Encoding region and the reconstruction map $F(s) = \text{readout}(s _{\mathcal{A}})$
$\mathcal{E}$	Eq. (16)	Energy observable shared by the stages of a continuing protocol
z_i, Γ, A_{ij}	Section 8.3	Complex amplitude, linear gain and coupling matrix of the amplitude layer
$\Delta\theta_d$	Section 8.3	Phase a winding advances across separation d

Table 2 (continued)

Symbol	Defined	Meaning
E	Eq. (18)	Field amplitude, in the cortical identification
U	§3.2	Stored energy $U(\Xi)$ of a configuration
s, s_*	§4–5	A global section and the one the reconstruction fixes

869 Data and code availability

870 The Lean development, simulation scripts, tests and saved numerical
871 summaries are maintained in the GitHub repository `bobaseb/coupling_
872 gated_cortical_coherence`. The formalization uses Lean 4.34.0-rc2 and
873 Mathlib commit `5ce203ecd53f8c3e47ee5530d7c86b94a55f0183`. Publi-
874 cation simulation values are generated from saved summaries without re-
875 running production sweeps. The exploratory EEG analysis uses Open-
876 Neuro dataset `ds005620` (Bajwa et al., 2025) (subjects 1010, 1016, 1022,
877 1033, 1045, 1056, resting awake and sedation runs 1–2); preprocessing
878 and selection are described in Supplemental Material and implemented in
879 `simulations/empirical_collapse.py`. The repository contains the repro-
880 ducible analysis and does not redistribute the raw recordings.

881 Funding and competing interests

882 This research received no external funding. The author declares no compet-
883 ing interests.

884 Use of AI tools

885 AI assistance to this work was substantial and is not separable by task. Sev-
886 eral coding assistants and chat interfaces were used in an interleaved way
887 throughout: Google’s Antigravity CLI (Gemini 3.1 Pro, Claude 4.6 Opus,
888 Claude 4.6 Sonnet), OpenAI’s Codex CLI (GPT-6 Astra; GPT-5.6 Terra,
889 Luna and Sol), Claude Code (Claude Opus 5), Gemini Pro through its web in-
890 terface, and Anthropic’s Fable and DeepSeek v4 Flash through OpenRouter,
891 accessed with the Hermes CLI. Between them they contributed to the Lean
892 development, the simulation and analysis code, and the drafting and revision
893 of both documents. Because the tools were used in combination and in alter-
894 nation, no section, proof, figure or result is attributable to any one of them,
895 and none is claimed to be. The author is responsible for all content, including
896 everything a tool produced. Machine checking establishes the stated Lean
897 results under their hypotheses; it does not validate biological interpretations
898 or replace scientific review of generated code and prose.

899 References

- 900 Pinotsis, D. A., & Miller, E. K. (2022). Beyond dimension reduction: Stable
901 electric fields emerge from and allow representational drift. *NeuroImage*,
902 253, 119058.
- 903 Tononi, G., Boly, M., Massimini, M., & Koch, C. (2016). Integrated informa-

tion theory: from consciousness to its physical substrate. *Nature Reviews Neuroscience*, 17(7), 450-461.

Baars, B. J. (1997). In the Theater of Consciousness: The Workspace of the Mind. *Oxford University Press*.

Dehaene, S., Charles, L., King, J. R., & Marti, S. (2014). Toward a computational theory of conscious processing. *Current Opinion in Neurobiology*, 25, 76-84.

Friston, K. (2010). The free-energy principle: a unified brain theory?. *Nature Reviews Neuroscience*, 11(2), 127-138.

Milinkovic, B., & Aru, J. (2026). On biological and artificial consciousness: A case for biological computationalism. *Neuroscience & Biobehavioral Reviews*, 181, 106524. doi:10.1016/j.neubiorev.2025.106524.

Aru, J., Larkum, M. E., & Shine, J. M. (2023). The feasibility of artificial consciousness through the lens of neuroscience. *Trends in Neurosciences*, 46(12), 1008–1017. doi:10.1016/j.tins.2023.09.009.

Seth, A. K. (2026). Conscious artificial intelligence and biological naturalism. *Behavioral and Brain Sciences*, 49, e315. doi:10.1017/S0140525X25000032.

Seth, A. K. (2026). The stuff matters: Consciousness, computation, and biology. *Behavioral and Brain Sciences*, 49, e366. doi:10.1017/S0140525X26106906.

924 Gurnee, W., et al. (2026). Verbalizable Representations Form a
925 Global Workspace in Language Models. *Transformer Circuits Thread*.
926 arXiv:2607.15495.

927 Aguilera, M., Millidge, B., Tschantz, A., & Buckley, C. L. (2022). How par-
928 ticular is the physics of the free energy principle?. *Physics of Life Reviews*,
929 40, 24-50.

930 Biehl, M., Pollock, F. A., & Kanai, R. (2021). A technical critique of some
931 parts of the free energy principle. *Entropy*, 23(3), 293.

932 Bruineberg, J., Dolega, K., Dewhurst, J., & Baltieri, M. (2022). The Em-
933 peror’s New Markov Blankets. *Behavioral and Brain Sciences*, 45, e183.

934 Pinotsis, D. A., & Miller, E. K. (2023). In vivo ephaptic coupling allows
935 memory network formation. *Cerebral Cortex*, 33(17), 9877-9895.

936 Pinotsis, D. A., Fridman, G., & Miller, E. K. (2023). Cytoelectric coupling:
937 Electric fields sculpt neural activity and ”tune” the brain’s infrastructure.
938 *Progress in Neurobiology*, 226, 102465.

939 McFadden, J. (2002). The conscious electromagnetic information (cemi) field
940 theory: the hard problem made easy?. *Journal of Consciousness Studies*,
941 9(8), 45-60.

942 Pockett, S. (2012). The electromagnetic field theory of consciousness: a
943 testable hypothesis about the characteristics of conscious as opposed to
944 non-conscious fields. *Journal of Consciousness Studies*, 19(11-12), 191-223.

- 945 McFadden, J. (2020). Integrating information in the brain's EM field: the
 946 cemi field theory of consciousness. *Neuroscience of Consciousness*, 2020(1),
 947 niaa016.
- 948 Fröhlich, F., & McCormick, D. A. (2010). Endogenous electric fields may
 949 guide neocortical network activity. *Neuron*, 67(1), 129-143.
- 950 Anastassiou, C. A., & Koch, C. (2015). Ephaptic coupling to endogenous
 951 electric field activity: why bother?. *Current Opinion in Neurobiology*, 31,
 952 95-103.
- 953 Pockett, S. (2002). Difficulties with the electromagnetic field theory of con-
 954 sciousness. *Journal of Consciousness Studies*, 9(4), 51-56.
- 955 Vöröslakos, M., et al. (2018). Direct effects of transcranial electric stimulation
 956 on brain circuits in rats and humans. *Nature Communications*, 9, 483.
- 957 Verkhratsky, A., & Nedergaard, M. (2018). Physiology of Astroglia. *Physio-*
 958 *logical Reviews*, 98(1), 239-389.
- 959 Mountcastle, V. B. (1997). The columnar organization of the neocortex.
 960 *Brain*, 120(4), 701-722. doi:10.1093/brain/120.4.701.
- 961 Horton, J. C., & Adams, D. L. (2005). The cortical column: a structure
 962 without a function. *Philosophical Transactions of the Royal Society B*,
 963 360(1456), 837-862. doi:10.1098/rstb.2005.1623.

964 Rakic, P. (2008). Confusing cortical columns. *Proceedings of the National*
965 *Academy of Sciences*, 105(34), 12099-12100. doi:10.1073/pnas.0807271105.

966 Bastos, A. M., Donoghue, J. A., Brincat, S. L., Mahnke, M., Yanar, J., Cor-
967 rea, J., ... & Miller, E. K. (2021). Neural effects of propofol-induced uncon-
968 sciousness and its reversal using thalamic stimulation. *eLife*, 10, e60824.

969 Bajwa, I. J., Nilsen, A. S., Skukies, R., Aamodt, A., Ernst, G., Storm, J.
970 F., & Juel, B. E. (2025). A repeated awakening study exploring the capac-
971 ity of complexity measures to capture dreaming during propofol sedation.
972 *Scientific Reports*, 15, 32746.

973 Xie, L., Kang, H., Xu, Q., Chen, M. J., Liao, Y., Thiagarajan, M., ... &
974 Nedergaard, M. (2013). Sleep drives metabolite clearance from the adult
975 brain. *Science*, 342(6156), 373-377.

976 Syková, E., & Nicholson, C. (2008). Diffusion in brain extracellular space.
977 *Physiological Reviews*, 88(4), 1277-1340.

978 Barbour, B. (2017). Analysis of claims that the brain extracellular impedance
979 is high and non-resistive. *Biophysical Journal*, 113(7), 1636-1638.

980 Logothetis, N. K., Kayser, C., & Oeltermann, A. (2007). In vivo measure-
981 ment of cortical impedance spectrum in monkeys: implications for signal
982 propagation. *Neuron*, 55(5), 809-823. doi:10.1016/j.neuron.2007.07.027.

983 Gabriel, S., Lau, R. W., & Gabriel, C. (1996). The dielectric properties of bio-
984 logical tissues: II. Measurements in the frequency range 10 Hz to 20 GHz.

- 985 *Physics in Medicine and Biology*, 41(11), 2251–2269. doi:10.1088/0031-
986 9155/41/11/002.
- 987 Norton, J. D. (2005). Eaters of the lotus: Landauer’s principle and the return
988 of Maxwell’s demon. *Studies in History and Philosophy of Science Part B:*
989 *Studies in History and Philosophy of Modern Physics*, 36(2), 375–411.
- 990 Landauer, R. (1961). Irreversibility and heat generation in the computing
991 process. *IBM Journal of Research and Development*, 5(3), 183–191.
- 992 Kibble, T. W. B. (1976). Topology of cosmic domains and strings. *Journal*
993 *of Physics A: Mathematical and General*, 9(8), 1387–1398.
- 994 Sakaguchi, H. (1988). Cooperative phenomena in coupled oscillator systems
995 under external fields. *Progress of Theoretical Physics*, 79(1), 39–46.
- 996 Strogatz, S. H., & Mirollo, R. E. (1991). Stability of incoherence in a popu-
997 lation of coupled oscillators. *Journal of Statistical Physics*, 63, 613–635.
- 998 Acebrón, J. A., Bonilla, L. L., Pérez Vicente, C. J., Ritort, F., &
999 Spigler, R. (2005). The Kuramoto model: A simple paradigm for
1000 synchronization phenomena. *Reviews of Modern Physics*, 77, 137–185.
1001 doi:10.1103/RevModPhys.77.137.
- 1002 Ott, E., & Antonsen, T. M. (2008). Low dimensional behavior of
1003 large systems of globally coupled oscillators. *Chaos*, 18, 037113.
1004 doi:10.1063/1.2930766.

1005 Dörfler, F., & Bullo, F. (2014). Synchronization in complex net-
1006 works of phase oscillators: A survey. *Automatica*, 50, 1539–1564.
1007 doi:10.1016/j.automatica.2014.04.012.

1008 Breakspear, M. (2017). Dynamic models of large-scale brain activity. *Nature*
1009 *Neuroscience*, 20, 340–352. doi:10.1038/nn.4497.

1010 Still, S., Sivak, D. A., Bell, A. J., & Crooks, G. E. (2012). Thermodynamics
1011 of prediction. *Physical Review Letters*, 109(12), 120604.

1012 Kawai, R., Parrondo, J. M. R., & Van den Broeck, C. (2007). Dissipation:
1013 The phase-space perspective. *Physical Review Letters*, 98(8), 080602.

1014 Horowitz, J. M., & Esposito, M. (2014). Thermodynamics with continuous
1015 information flow. *Physical Review X*, 4, 031015. doi:10.1103/PhysRevX.4.
1016 031015.

1017 Ito, S., & Sagawa, T. (2013). Information thermodynamics on
1018 causal networks. *Physical Review Letters*, 111, 180603. doi:10.1103/
1019 PhysRevLett.111.180603.

1020 Corberi, F., dello Russo, S., Messuti, G., Scarpetta, S., & Smaldone, L.
1021 (2026). Thermodynamic learning. *arXiv preprint*. arXiv:2609.04732.

1022 Banach, S. (1922). Sur les opérations dans les ensembles abstraits et leur ap-
1023 plication aux équations intégrales. *Fundamenta Mathematicae*, 3(1), 133-
1024 181.

- 1025 Shenker, O. R. (2000). Logic and entropy. *PhilSci-Archive preprint*. philsci-
1026 archive.pitt.edu/115.
- 1027 Rosas, F. E., Mediano, P. A. M., Luppi, A. I., Boyd, A. B., Berta, M., Jensen,
1028 H. J., Gastpar, M., & Carroll, S. M. (2026). Four faces of information in
1029 natural systems. *Nature Reviews Physics*. doi:10.1038/s42254-026-00971-4.
- 1030 Sznitman, A.-S. (1991). Topics in propagation of chaos. In *École d'Été de*
1031 *Probabilités de Saint-Flour XIX—1989*, Lecture Notes in Mathematics,
1032 vol. 1464, pp. 165–251. Springer. doi:10.1007/BFb0085169.
- 1033 Dai Pra, P., & den Hollander, F. (1996). McKean–Vlasov limit for interacting
1034 random processes in random media. *Journal of Statistical Physics*, 84, 735–
1035 772. doi:10.1007/BF02179656.
- 1036 Bertini, L., Giacomin, G., & Pakdaman, K. (2010). Dynamical aspects of
1037 mean field plane rotators and the Kuramoto model. *Journal of Statistical*
1038 *Physics*, 138, 270–290. doi:10.1007/s10955-009-9908-9.
- 1039 Berglund, N., & Gentz, B. (2002). Pathwise description of dynamic pitchfork
1040 bifurcations with additive noise. *Probability Theory and Related Fields*,
1041 122, 341–388. doi:10.1007/s004400100174.
- 1042 Abramsky, S., & Brandenburger, A. (2011). The sheaf-theoretic structure
1043 of non-locality and contextuality. *New Journal of Physics*, 13, 113036.
1044 doi:10.1088/1367-2630/13/11/113036.

1045 Townsend, R. G., Solomon, S. S., Chen, S. C., Pietersen, A. N. J., Martin,
1046 P. R., Solomon, S. G., & Gong, P. (2015). Emergence of complex wave
1047 patterns in primate cerebral cortex. *The Journal of Neuroscience*, 35(11),
1048 4657-4662. doi:10.1523/JNEUROSCI.4509-14.2015.

1049 Xu, Y., Long, X., Feng, J., & Gong, P. (2023). Interacting spiral wave pat-
1050 terns underlie complex brain dynamics and are related to cognitive pro-
1051 cessing. *Nature Human Behaviour*, 7(7), 1196-1215. doi:10.1038/s41562-
1052 023-01626-5.

1053 Nicoletti, G., Di Terlizzi, I., & Busiello, D. M. (2026). Balancing information
1054 and dissipation with partially observed fluctuating signals. *Physical Review*
1055 *Letters*, 137(12), 127101. doi:10.1103/8gx2-rbns.

1056 Kiyooka, D., Oomoto, I., Kitazono, J., Saito, Y., Kobayashi, M., Matsub-
1057 ara, C., Kobayashi, K., Murayama, M., & Oizumi, M. (2026). Single-
1058 cell resolution functional networks during unconsciousness are segre-
1059 gated into spatially intermixed modules. *Cell Reports*, 45(2), 116902.
1060 doi:10.1016/j.celrep.2025.116902.

1061 Oomoto, I., Kiyooka, D., Oizumi, M., & Murayama, M. (2026). Multi-area
1062 single-cell calcium imaging dataset of the mouse cortex across wakefulness,
1063 sleep, and anesthesia. *Scientific Data*. doi:10.1038/s41597-026-08202-2.

1064 Attwell, D., & Laughlin, S. B. (2001). An energy budget for signaling in the

1065 grey matter of the brain. *Journal of Cerebral Blood Flow & Metabolism*,
1066 21(10), 1133-1145. doi:10.1097/00004647-200110000-00001.

1067 Harris, J. J., Jolivet, R., & Attwell, D. (2012). Synaptic energy use and
1068 supply. *Neuron*, 75(5), 762-777. doi:10.1016/j.neuron.2012.08.019.

1069 Nicholls, D. G., & Ferguson, S. J. (2013). *Bioenergetics 4*. Academic Press.

1070 Baer, S. M., Erneux, T., & Rinzel, J. (1989). The slow passage through a
1071 Hopf bifurcation: delay, memory effects, and resonance. *SIAM Journal on*
1072 *Applied Mathematics*, 49(1), 55–71. doi:10.1137/0149003.

1073 Friedman, E. B., Sun, Y., Moore, J. T., Hung, H.-T., Meng, Q. C., Perera, P.,
1074 Joiner, W. J., Thomas, S. A., Eckenhoff, R. G., Sehgal, A., & Kelz, M. B.
1075 (2010). A conserved behavioral state barrier impedes transitions between
1076 anesthetic-induced unconsciousness and wakefulness: evidence for neural
1077 inertia. *PLoS ONE*, 5(7), e11903. doi:10.1371/journal.pone.0011903.

1078 Hudson, A. E., Calderon, D. P., Pfaff, D. W., & Proekt, A. (2014). Recovery
1079 of consciousness is mediated by a network of discrete metastable activity
1080 states. *Proceedings of the National Academy of Sciences*, 111(25), 9283–
1081 9288. doi:10.1073/pnas.1408296111.

1082 Proekt, A., & Hudson, A. E. (2018). A stochastic basis for neural inertia
1083 in emergence from general anaesthesia. *British Journal of Anaesthesia*,
1084 121(1), 86–94. doi:10.1016/j.bja.2018.02.035.