

# Supplemental Material: Coupling-gated cortical coherence

## Overview

This document supplies the formal constructions, numerical methods and exploratory EEG analysis for the conditional field model in the main article. Mathematical results are distinguished from model assumptions and physical identifications. The Lean development uses only its library's foundational axioms; physical postulates are hypotheses or fields of models with explicit witnesses. The soundness section below explains why an unrestricted postulate about model-dependent quantities can be inconsistent. Source code and saved simulation summaries are maintained at [https://github.com/bobaseb/coupling\\_gated\\_cortical\\_coherence](https://github.com/bobaseb/coupling_gated_cortical_coherence).

## 1 Symmetry Breaking and the Emergence of Boundaries

The framework's first physical motif is that a perfectly symmetric universe cannot support agency. The formal result is narrower: under the stated vacuum and connectedness hypotheses, a field reaching two vacuum components must leave the vacuum somewhere.

Let  $\mathcal{M}$  be a pseudo-Riemannian manifold representing spacetime, and let  $\phi : \mathcal{M} \rightarrow \mathcal{V}$  be a physical field mapping to a vacuum manifold  $\mathcal{V}$  (the space of energy minima). Let  $G$  be a continuous symmetry group acting on  $\mathcal{V}$ .

**Theorem 1 (Domain-wall obstruction):** A continuous field on a connected substrate whose image reaches two different connected components of the vacuum manifold must leave that manifold somewhere. This is the  $\pi_0$  case of the topological-defect argument. A separate extension obstruction states that non-null-homotopic boundary data admit no continuous extension over the specified contractible interior.

*Lean 4 Implementation (Phase1\_Primitives.lean):* We formalize this by defining the set of degenerate energy minima `DynamicalVacuum` and proving `boundary_defect_forces_interior_vacuum_break`: if a boundary field configuration  $\phi_\partial$  is not null-homotopic, then it admits no continuous extension over a contractible interior. This is an obstruction to extending the specified boundary data. Note the direction: this result establishes that a defect *obstructs extension*, not that a defect must exist. The companion lemma `contractible_interior_forces_trivial_boundary` proves the contrapositive.

A direction that produces a defect is proved, at the level of  $\pi_0$ . `exists_notMem_of_no_common_preconnected` states it in the form a witness can discharge — if no preconnected subset of the vacuum manifold  $M$  contains both  $\phi(x_1)$  and  $\phi(x_2)$ , and the substrate is connected, then  $\phi$  leaves  $M$  somewhere — and `exists_notMem_of_connectedComponentIn` restates it with  $\phi(x_2) \notin$  the connected component of  $M$  containing  $\phi(x_1)$ . The proof uses the connectedness of continuous images and `IsPreconnected.subset_connectedComponentIn`, and nothing else; no hypothesis mentions energy, a potential or a symmetry. This is the domain-wall case of Kibble's classification and we take only that case:  $\pi_1 \neq 0$  forcing vortices and  $\pi_2 \neq 0$  forcing monopoles are recorded as available and not taken, since the manuscript's argument uses domain walls alone. Separately, `spontaneous_symmetry_breaking` shows a global energy minimiser lies in the vacuum manifold almost everywhere; this is proved without additional axioms, the pointwise step being `pointwise_vacuum_of_global_min`. `Examples.lean` §11 witnesses that result on the double well  $V(v) = (v^2 - 1)^2$  over a two-point spacetime carrying a Dirac probability measure. Two features of the witness are deliberate. Its minimisation hypothesis is *proved* rather than assumed — everywhere else in the development the existence of a global minimiser is a hypothesis, since no compactness or direct-method argument is formalized. And the almost-everywhere in the conclusion is shown necessary: the minimising field sits at the vacuum on the support and at 0 off it, so the pointwise version of the statement is false of the very field the theorem applies to. The same section instantiates the  $\mathbb{Z}_2$  action  $v \mapsto \pm v$  and discharges the energy invariance, giving the first instance of `SymmetryInvariantAction`; the vacuum manifold  $\{-1, 1\}$  is proved degenerate and the two minima exchanged by the symmetry, so a minimiser and its image have equal energy and differ. `Examples.lean` §11.1 then supplies the

witness for the  $\pi_0$  theorem on the same well. The vacuum manifold is *computed* there, `wellVacuum.eq` proving `DynamicalVacuum`  $V = \{-1, 1\}$  exactly rather than merely containing both minima — which matters, because a larger vacuum set might have been connected — and `wellVacuum.separated` proves the disconnection in the form the theorem consumes, by the observation that a preconnected subset of  $\mathbb{R}$  spanning  $-1$  and  $1$  contains  $0$ , which is not a minimum. `wellV.domain.wall` is the resulting statement about *every* continuous field on *any* connected substrate reaching both minima. Non-emptiness of that class is checked separately rather than assumed: the clipped identity on  $\mathbb{R}$  satisfies the hypotheses, and `kink.zero.notMem` locates its forced wall at the origin. No conserved quantity is constructed from any of this, and there is still no Noether theorem.

## 2 The Thermodynamic Constraint: Landauer’s Principle

For a finite effective register, the number of distinguishable histories that can be retained is bounded. This does not imply that every update erases. The capacity theorem states that a distribution on a finite state space  $X$  carries at most  $\log |X|$  nats, with equality only at the uniform law (`shannon.entropy.eq.log.card.iff`), and a perturbation stream of length  $k$  over an alphabet  $P$  with  $|P|^k > |X|$  therefore cannot be recorded injectively (`absorb.not.injective`), whatever the update map is and even when every individual step of it is reversible.

Let the global phase space  $\Omega = \Omega_{sys} \times \Omega_{bath}$  undergo globally reversible Hamiltonian dynamics, modeled as an injective mapping  $f : \Omega \rightarrow \Omega$ .

**Theorem 2 (Thermodynamic Erasure):** Under the specified finite-register and reversible system–bath assumptions, the heat identity gives a lower bound proportional to the entropy erased by the register. Erasing one unbiased bit gives  $\Delta Q \geq k_B T \ln 2$ . Non-injectivity alone does not specify a universal one-bit entropy loss; the bound uses the erased entropy of the particular update.

*The same heat, in three vocabularies (Phase3.LandauerBridge.lean).* Derivation 2 charges a register for destroying information and Derivation 3 charges a structure for keeping information it cannot use. `PredictiveDissipation.ofLandauer` builds the predictive structure of a given register: its thermal scale is that register’s temperature, its dissipated work is that register’s Landauer heat, and Still’s bound is derived from `landauer.bound` rather than assumed a second time. The single hypothesis is the identification  $(I_{mem} - I_{pred}) \leq S(id) - S(t)$  — the memory wasted is no larger than the entropy erased — and it is a restriction, not bookkeeping: on an injective update the right-hand side is zero and the inequality forces the waste to vanish (`nonpredictive.eq.zero.of.injective`), which is the Norton–Shenker distinction in informational form. The measure-theoretic statement of Landauer (`Phase3.MeasureThermodynamics`) is specialised in the same file rather than left standing beside the combinatorial one: at counting measure on a finite register, `continuous.entropy` is `boltzmann.entropy`, volume contraction is non-injectivity, and `landauer.heat.bound` is  $T$  times the entropy erased (`landauer.heat.bound.count.eq`). The witness is the one-bit eraser, where a register destroying  $\log 2$  pays for a memory wasting  $\log 2$ , with equality (`Examples §18.5`).

*Lean 4 Implementation (Phase3.CombinatorialThermodynamics.lean):* We explicitly decouple logical erasure from physical erasure. Lean verifies that any contraction in the *physical* state space ( $\neg \text{Function.Injective}(t_{sys})$ ) under globally reversible dynamics mathematically necessitates heat dissipation. The combinatorial core, `landauer.from.reversibility` and `second.law`, is fully derived. Finiteness of the phase space is what converts the physically checkable hypothesis into the one the bound needs: on a finite state space a map that fails to reach every state is not injective (`is.eraser.of.not.surjective`), so an unreachable state suffices for  $\Delta Q > 0$  (`finite.phase.space.dissipates`). The hypothesis is fenced: `Nat.succ` misses  $0$  and is injective, so the implication is false on an infinite phase space, and `Examples.lean §19` records this beside the witnesses. The Norton–Shenker distinction survives the identification — retaining the incoming perturbation is injective and free (`keepRecord.injective`), refreshing the register that held it is not (`absorbStep.is.eraser`). Converting it into a statement about heat requires the postulate  $\Delta Q = T \Delta S_{bath}$ , which is carried as the class field `StatisticalMechanics.heat.eq`; `Examples.lean` exhibits a one-bit erasure system with an explicit bath that discharges it, so the postulate is demonstrably satisfiable. A register that satisfies the stated erasure and thermal hypotheses dissipates; identifying it with a biological boundary is an additional modelling step.

### 3 Structural Resonance and Least Action

A bound relating environmental and internal distributions must be a physical assumption about specified processes, not a consequence of Gibbs’ inequality alone. The class field `StructuralResonance.kl_bound` states  $D_{KL}(P||Q) \leq \Delta t \sigma$  for an instance. Even under this assumption, descent of  $\sigma$  does not imply  $D_{KL} \rightarrow 0$  unless the upper bound vanishes. A positive limiting upper bound permits a positive divergence. No convergence of internal geometry to environmental structure follows.

**Theorem 3 (Nonpredictive information bounds dissipation):** Let  $X_t$  be the system’s state and  $S_t$  an environmental signal evolving by a dynamics of its own on which the system does not act. Then

$$I(X_t; S_{t+1}) \leq I(X_t; S_t), \quad W_{\text{diss}} \geq k_B T [I(X_t; S_t) - I(X_t; S_{t+1})], \quad (1)$$

the second being the bound of Still et al. (2012). Consequently  $W_{\text{diss}} \geq 0$ , and any bound on the dissipated work is a bound on the nonpredictive information  $I(X_t; S_t) - I(X_t; S_{t+1})$ .

The first inequality is the data processing inequality for the Markov chain  $X_t \rightarrow S_t \rightarrow S_{t+1}$ , proved in `Phase3.PredictiveThermodynamics.lean` as `predictiveInfo.le_mutualInfo`; mutual information is defined there from Mathlib’s `klDiv` as the divergence of the joint law from the product of its marginals, since Mathlib carries no mutual information of its own. The second inequality is the class field `PredictiveDissipation.still_bound`, an irreducible physical postulate for the same reason as every other one here — its axiom form is refutable, and `Examples.lean` §18.4 refutes it. Everything downstream is derived: non-negativity of the dissipated work, the bound of the nonpredictive information by the dissipation budget, and the statement that a quasi-static drive leaves every retained bit predictive. `Examples.lean` §18 inhabits the class twice over one correlated two-bit law — a frozen signal, permitted to dissipate nothing, and a scrambled one, forced to dissipate its entire memory — and fences the theorem with a Markov kernel on the pair that writes the system’s state into the signal and thereby *raises* mutual information, so the no-back-action hypothesis is shown to carry weight rather than to decorate. The equality of Kawai et al. (2007),  $\langle W_{\text{diss}} \rangle = k_B T D_{KL}(\text{forward}||\text{reverse})$ , is carried as a predicate on an instance and yields `nonpredictive.le_arrow`: the irreversibility of the dynamics is at least the memory that fails to earn its keep.

The class field in `Phase3.KLBound.lean` is separate from Gibbs’ nonnegativity theorem. The witness `boolDetuned` attains the assumed bound with  $D_{KL} = \sigma = \log 2$ , showing that a positive divergence is permitted.

The fixed-phase squared-drift functional in `Phase8.ContinuousField.lean` is a quadratic functional of an affine kernel map. The bounded operators `kernelCLM` and `continuumDriftCLM`, together with the Hilbert-space adjoint, give  $\nabla \sigma = (2/D)A^*(AK + \omega)$ . The theorem `structural_resonance.decreases_sigmaContinuum` proves descent for the specified gradient flow, with an atomless witness on the unit interval. The identifier `entropy_production_rate` names this squared-drift functional in the source; the name does not establish its identification with thermodynamic entropy production. That interpretation requires probability-current information absent from this calculation.

For a symmetric kernel, antisymmetry of  $K(x, y) \sin(\theta_y - \theta_x)$  makes its double integral zero. Thus the mean drift is independent of the phase field. A dynamically phase-locked field attains the associated squared-drift lower bound under the stated hypotheses. The theorem `phase_locked_minimizes_entropy_of_symm` expresses this minimization. A one-way two-site coupling supplies a counterexample without symmetry. These are statements about the objective and its constrained mean, not convergence of joint noisy phase/plasticity dynamics.

### 4 Macroscopic Scaling and Kuramoto Phase Synchronization

We posit a continuous macroscopic coupling description. We model the local phases  $\theta_i$  as coupled nonlinear oscillators via the Kuramoto model.

*Scope of the continuum description.* `mesh_refinement_convergence` proves convergence of one scalar coupling energy. It produces neither a continuum kernel nor field dynamics: `edge_region` has no  $M \times M$  product structure from which such a kernel could be recovered, and the vertex-supported candidate is invisible to the continuum functional on an atomless substrate. The later symmetric kernel is therefore a modelling input. Its identification with the brain’s endogenous *electromagnetic* field is a further empirical commitment about cortex, argued in the main text from the ephaptic-coupling literature with explicit calibration and observational limitations. Nothing in this section mathematically establishes either step.

For a system with coupling strength  $K$  and thermodynamic noise  $D$ , the dynamic evolution is:

$$\dot{\theta}_i = \omega_i + \frac{K}{N} \sum_{j=1}^N \sin(\theta_j - \theta_i) + \eta_i(t) \quad (2)$$

**Theorem 4 (Stationary coherent branch):** Under the identical-frequency von Mises stationary ansatz, the self-consistency equation has exactly one positive solution when  $K > 2D$  and none when  $K \leq 2D$ . This does not assert dynamical selection or noiseless phase locking in the stochastic model.

*Lean 4 Implementation (Phase4.KuramotoDynamics.lean):* We formalize the Lyapunov potential  $V(\theta) = -\frac{1}{2} \sum_{i,j} A_{ij} \cos(\theta_j - \theta_i)$  of the coupled field and prove two directions: `phase_locked_minimizes_potential` (the phase-locked configuration attains the global minimum whenever all couplings are positive) and `potential_min_implies_phase_locked` (any minimizer is phase-locked). Note the scope: this is a statement about the *static* potential landscape and holds for all positive coupling matrices. The critical threshold  $K_c = 2D$  does not appear in this file. It is formalized separately, in `Phase8-SelfConsistency.lean` (see below), as a statement about the mean-field self-consistency equation with the stationary density assumed; deriving that density from the noise-driven stochastic dynamics would require the spectral theory of the associated Fokker-Planck operator, which is not available in Mathlib. The numerical controls below distinguish identical-frequency diffusion from the separate quenched-frequency model illustrated by `simulations/kuramoto.py`.

*Positivity, and the regime these theorems cover.* The hypothesis “all couplings positive” is not a technical convenience and it is not free. `KuramotoSystem` carries  $A : V \rightarrow V \rightarrow \mathbb{R}$ , symmetric and otherwise unconstrained, so signed — frustrated — couplings are representable; but every result in this development that identifies *where* the dynamics ends assumes positivity. `phase_locked_minimizes_potential` and `potential_min_iff_phase_locked` assume  $A_{ij} > 0$ ; the class field `ThermodynamicCover.A_pos` demands it of every instance, and so of every conclusion of Theorem 5; and `kuramoto_tendsto_global_minimum` requires the strictly stronger uniform bound  $0 < a \leq A_{ij}$ , since  $a$  enters the Łojasiewicz constant. Such a network is an unfrustrated ferromagnet with one minimum up to a global phase shift — exactly not the near-degenerate landscape a memory argument from geometric frustration requires. By contrast, nothing about the sign of  $A$  is used by `is_kuramoto_trajectory_exists` and its uniqueness counterpart, by the Lyapunov identity `dV_dt_le_zero` on which the whole descent argument rests, by `dynamic_potential_antitone` and `dynamic_potential_tendsto`, or by `velocity_sq_tendsto_zero`. Since positivity is a *field* of `ThermodynamicCover` rather than a hypothesis distributed across the development, it enters the chain at one node and not before it. So *that* the motion stops holds for frustrated systems as well; *where* it stops does not. On the continuum side `phase_locked_minimizes_entropy_of_symm` assumes only symmetry and so admits signed kernels, but it takes the existence of a dynamically phase-locked field as a hypothesis, which is precisely what frustration threatens. The main text scopes the frustration claim out of the chain accordingly, rather than repeating it beside theorems that exclude it.

*Twisted gluing does not recover content from phase frustration (Phase5.TwistedGluing.lean).* We also tested the natural weakening of exact overlap agreement: local states may differ by a phase on each overlap. `gluesUpToPhase_of_isCoboundary` proves that coboundary offsets glue after one rotation per patch, while `isCoboundary_of_gluesUpToPhase` proves the converse when the phase action is free on overlaps. Any Kuramoto configuration, including a twisted or splay state, supplies absolute patch phases  $\phi_i$ , so its offsets  $\phi_i - \phi_j$  are automatically a coboundary (`isCoboundary_of_phaseField`) and the twisted family glues (`gluesUpToPhase_of_phaseField`). The formal `PhaseObstruction` quotient is deliberately not called Čech cohomology because its total cochains still assign irrelevant values to patch pairs with empty overlap.

The data give a narrower positive answer. Townsend et al. (2015) identify phase singularities in multi-electrode primate LFP recordings by an integer winding number around electrode loops, and the associated complex waves modulate spiking; Xu et al. (2023) find task-dependent spiral centres and rotation directions in human fMRI phase maps. Around such a zero-amplitude point the circle-valued phase need not admit one real-valued lift. Local lifts differ on overlaps by integer multiples of  $2\pi$ , supplying measured transition data without an arbitrary gauge field. `Phase5.PhaseLifts.lean` formalizes the finite electrode loop: `loopWinding` sums its integer transitions, `loopWinding.eq_zero_of_hasGlobalLift` proves that transitions from one global lift telescope to zero, and `threeElectrodeWinding.not_hasGlobalLift` gives a winding-one witness. This does not rescue `PhaseObstruction`: that quotient uses circle-valued differences of absolute patch phases, while the measured obstruction uses integer transitions between local lifts.

The cover-level specification has cortical surface  $M$ , analytic signal  $z = Ae^{i\phi}$ , singularity set  $Z_t$ , punctured domain  $M_t^* = M \setminus Z_t$ , and measurement map  $q_t = z/|z| : M_t^* \rightarrow S^1$ . A trial-independent

electrode or source-space cover of  $M_t^*$  supplies local real lifts; their differences divided by  $2\pi$  are integer transitions on connected overlaps, and nerve loops are the inputs to `loopWinding`. In finite-resolution data the amplitude threshold defining  $Z_t$ , the loop geometry and the required cross-scale persistence are fixed using training or baseline data, never the content labels of test trials.

The proposed content-sensitive test compares nested preregistered classifiers. The control model contains stimulus, task, behavioural response, eye movement and arousal; the augmented model adds signed winding, source-localized singularity position, rotation direction and lifetime. Evaluation uses held-out participants and sessions, with label permutations within task, stimulus and arousal strata. The content hypothesis requires replicated incremental out-of-sample discrimination, including under a changed or absent motor report. This controls the known task dependence rather than mistaking it for content. A positive result would be an empirical association, not an identification of a topological sector with phenomenal content; a null increment rejects this proposed role. The construction and test remain outside the conditional chain, and a full nerve-indexed integer complex is deferred until the measurement protocol warrants it.

*Lean 4 Implementation (Phase8.ContinuousField.lean):* The continuous-field file defines  $K_c = 2D$  and a predicate for a substrate above that threshold. Two cautions. First, the predicate is not connected to `is_continuous_kuramoto_trajectory`, `entropy_production_rate`, or the order parameter of Phase4. What it is connected to is the mean-field self-consistency equation  $r = I_1(Kr/D)/I_0(Kr/D)$ , whose bifurcation at  $K = 2D$  we formalize in both directions (see the next paragraph): `exhibits_phase_transition_coherent` shows that a substrate satisfying the predicate admits a solution with  $0 < r \leq 1$ , and `exhibits_phase_transition_unique_coherent` that it admits exactly one. So the threshold is not a bare stipulation, but what is proved is a statement about a stationary density rather than about a trajectory. Second, the predicate requires the substrate's `volume` to be a probability measure, and this hypothesis is load-bearing rather than decorative. It compares `mean_field_coupling`, the kernel averaged over both arguments, against  $2D$ ; on a substrate of total mass  $m$  that average carries a factor  $m^2$ , so without normalization the comparison tracks the substrate's size rather than its coupling and can be satisfied by any kernel at all. `Examples.lean` §9 exhibits a system a factor of two below threshold whose unnormalized double integral nonetheless clears  $2D$ . Under the hypothesis, `mean_field_coupling_const` gives that a constant kernel has strength exactly  $K$  and `exhibits_phase_transition_const_iff` reduces the predicate to  $K > 2D$ , witnessed in §9 on both sides of the threshold.

*Lean 4 Implementation (Phase8.SelfConsistency.lean):* The self-consistency equation behind  $K_c = 2D$  is formalized as real analysis. With the von Mises stationary density  $\rho(\theta) \propto e^{a \cos \theta}$ ,  $a = Kr/D$ , taken as given (the algebraic step from the stationary Fokker–Planck equation to this density is verified independently at `simulations/fokker_planck_von_mises.py` in the repository), `besselRatio` is  $I_1(a)/I_0(a)$  written as a ratio of interval integrals, Mathlib having no modified Bessel functions. `vonMisesM_eq_mul_vonMisesS` is the integration-by-parts identity  $\int_{-\pi}^{\pi} \cos \theta e^{a \cos \theta} = a \int_{-\pi}^{\pi} \sin^2 \theta e^{a \cos \theta}$ , exact because  $\sin(\pm\pi) = 0$  kills the boundary term. `vonMisesC2_nonneg` proves  $I_2(a) \geq 0$  by folding  $[-\pi, \pi]$  onto  $[0, \pi/4]$  through the reflections  $\theta \mapsto \pi - \theta$  and  $\theta \mapsto \pi/2 - \theta$ , where the integrand becomes a product of two manifestly non-negative factors; `vonMisesS_le_half_vonMisesZ` restates it as  $\mathbb{E}_a[\sin^2 \theta] \leq 1/2$  and `besselRatio_le_half_self` as  $R(a) \leq a/2$ . That factor of  $1/2$  is exactly what separates a threshold at  $D$  from one at  $2D$ ; `subcritical_fixed_point_eq_zero` is the crude form obtained from  $\sin^2 \leq 1$  and is kept to make the point visible. The subcritical theorem is `subcritical_fixed_point_eq_zero'`: for  $0 \leq K < \text{critical\_coupling } D$  and  $D > 0$ , the only  $r \geq 0$  with  $r = R(K, r)$  is  $r = 0$ . The supercritical theorem is `supercritical_fixed_point_exists`: for  $K > \text{critical\_coupling } D$  there is an  $r$  with  $0 < r \leq 1$  and  $r = R(K, r)$ . Its proof needs no series expansion. Factoring  $R(a) = a E(a)$  with  $E(a) = \mathbb{E}_a[\sin^2 \theta]$  — the factorization the integration by parts supplies — it suffices that  $E$  is continuous with  $E(0) = 1/2$ : since  $D/K < 1/2$ , the map overshoots for small  $r > 0$ , while `besselRatio_le_one` ( $R$  is a mean of  $\cos \theta$ ) makes it undershoot at  $r = 1$ , and `intermediate_value_1cc'` supplies the crossing. Continuity of  $E$  is Mathlib's continuity of a parametric interval integral, applied to the jointly continuous integrand. `critical_coupling_is_threshold` states the two halves together, and §9 of the file discharges the hypotheses of each at concrete values ( $K = 3$ ,  $D = 1$  above threshold;  $K = 3$ ,  $D = 2$  below it), so neither is vacuous. Section 8 of the file identifies the equation with the physics it is named for. `vonMisesDensity` is the normalized weight, proved a probability density on  $[-\pi, \pi]$  (`vonMisesDensity_integral_eq_one`); `circularOrderParameter` is the average of  $e^{i\theta}$  against a density, the continuum analogue of `order_parameter_complex` of Phase4; `circularOrderParameter_vonMises` proves that average equals  $R(a)$ , real because `vonMises_mean_sin` gives  $\mathbb{E}_a[\sin \theta] = 0$ ; and `fixedPoint_iff_selfReproducing` states the equivalence:  $r$  solves  $r = R(K, r)$  exactly when the density it induces has order parameter  $r$ . `incoherent_density_uniform` identifies the solution  $r = 0$  with

the uniform density, and `supercritical_coherent_density` restates the existence theorem as the existence of a stationary density with positive order parameter. Three results added since sharpen the threshold. `vonMisesC2_pos` upgrades the fold argument to a strict inequality for  $a > 0$  — `foldB` is continuous, non-negative on  $[0, \pi/4]$  and strictly positive on its interior, so `intervalIntegral_pos_of_pos_on` applies — giving  $\mathbb{E}_a[\sin^2 \theta] < 1/2$  and  $R(a) < a/2$  for every  $a > 0$ . That closes the threshold case: `fixed_point_eq_zero_of_le_critical` covers all  $K \leq 2D$ , so the incoherent and coherent regimes now exhaust  $K \geq 0$  with no gap, and  $K = 2D$  falls on the incoherent side. `coherent_iff_sRatio_eq` reformulates the coherent branch: for  $r \neq 0$ , solving  $r = R(K, r)$  is equivalent to  $E(Kr/D) = D/K$ , so the branch is a level set of  $E$ . `coherent_branch_continuous_at_threshold` then proves that for every  $\varepsilon > 0$  there is a  $\delta > 0$  such that on  $(2D, 2D + \delta)$  every coherent solution satisfies  $r < \varepsilon$  — the branch emerges continuously from zero, excluding a discontinuous jump. Its proof needs no monotonicity of  $E$ : on a compact interval bounded away from the origin the extreme value theorem gives a maximum strictly below  $1/2$ , and beyond that the crude bound  $E(a) \leq 1/a$  suffices. Section 7 of the file closes uniqueness, which the paragraph above reduced to the injectivity of  $E$  on  $(0, \infty)$ , and closes it in the stronger form `vonMisesSRatio_strictAntiOn`:  $E$  is strictly decreasing on  $[0, \infty)$ . The expected route — differentiate the exponential family, obtain  $E'(a) = -\text{Cov}_a(\cos^2 \theta, \cos \theta)$ , and show that covariance positive — is not the route taken, and it was the harder one: it needs differentiation under the integral sign, and the covariance is not sign-definite pointwise, since the symmetrised form  $\frac{1}{2}\mathbb{E}[(X - X')^2(X + X')]$  with  $X = \cos \theta$  has an integrand changing sign with  $X + X'$ . Instead the target is kept as the two-point inequality  $C_2(a)Z(b) < C_2(b)Z(a)$  for  $a < b$ , and two moves settle it. *Fold first*: `vonMisesZ_eq_foldW` and `vonMisesC2_eq_foldA` carry both integrals onto  $[0, \pi/2]$  by the two reflections already used for  $I_2 \geq 0$ , replacing  $e^{a \cos \theta}$  by  $W(a, \theta) = 2 \cosh(a \cos \theta)$  and confining  $\cos \theta$  to  $[0, 1]$ ; there `crossIntegrand_nonneg` shows  $\cos 2x - \cos 2y$  and  $W(b, x)W(a, y) - W(a, x)W(b, y)$  have the same sign, both being governed by  $\cos x$  against  $\cos y$ . This is the same  $X \mapsto -X$  symmetrisation the covariance form needed, performed once at the start rather than inside a double integral, and its analytic content is one identity: `cosh_mul_cosh_cross`, i.e.  $\cosh(ap) \cosh(bq) \leq \cosh(bp) \cosh(aq)$  for  $0 \leq a \leq b$ ,  $0 \leq q \leq p$ , from  $2 \cosh X \cosh Y = \cosh(X+Y) + \cosh(X-Y)$  and  $\cosh$  increasing in  $|\cdot|$ . *Then cross*: rather than integrating that inequality over a square, `foldA_mul_foldW_lt` takes  $\alpha$  to be the mean of  $\cos 2x$  under  $W(a, \cdot)$  and  $x_0 \in [0, \pi/2]$  a point where  $\cos 2x_0 = \alpha$ , supplied by `intermediate_value_Icc'`; evaluating at  $y = x_0$  and integrating in  $x$  alone gives the conclusion, the  $W(b, x_0)$  term vanishing by the choice of  $\alpha$ . Strictness comes from `crossIntegrand_pos` on a subinterval missing  $x_0$ . The same crossing argument run with  $\cos \theta$  in place of  $\sin^2 \theta$  needs no fold, since  $R(a) = \mathbb{E}_a[\cos \theta]$  is the mean of the tilting variable itself, and gives `besselRatio_strictMono`:  $R$  is strictly increasing on all of  $\mathbb{R}$ . The consequences are `coherent_fixed_point_unique` (at most one positive solution), `supercritical_fixed_point_existsUnique` and `supercritical_solution_set` (the non-negative solutions are exactly  $\{0\}$  at or below threshold and  $\{0, r\}$  above it), and `coherent_branch_strictMono`, which uses both halves:  $E$  decreasing puts the concentrations of two couplings in the order  $a_1 < a_2$  through  $E(a_i) = D/K_i$ , and  $R$  increasing carries that order to  $r_1 < r_2$ . Scope: the density is assumed, not derived from the Fokker–Planck equation; nothing in this file differentiates anything, so  $E$  and  $R$  are reached as strictly monotone rather than as smooth, and the rate at which the branch leaves zero is the business of `Phase8_CriticalExponent.lean` below; dynamical selection is not established, and could not be here, since no theorem in the file mentions a trajectory; and the order parameter reached is that of a density, not of a trajectory — no theorem relates `circularOrderParameter` to `order_parameter_complex`, which would be the mean-field limit. Every declaration in the file depends only on `propext`, `Classical.choice` and `Quot.sound`.

*Lean 4 Implementation (Phase8\_CriticalExponent.lean)*: The rate is what separates the bifurcation from a generic saturating gate, and it costs the differentiation the file above avoids. `vmMoment n a` is the cosine moment  $\int_{-\pi}^{\pi} \cos^n \theta e^{a \cos \theta}$ , and `hasDerivAt_vmMoment` differentiates it in  $a$  under the integral sign, raising the cosine power by one; the interchange is Mathlib’s dominated-convergence form for interval integrals, dominated by the constant  $e^{|a|+1}$  on the unit ball around the parameter, which is where the compactness of  $[-\pi, \pi]$  and the boundedness of  $\cos$  do their work. Iterating gives `contDiff_two_vmMoment`, and `vonMisesSRatio_eq_moments` rewrites  $E(a) = 1 - \text{vmMoment } 2a / \text{vmMoment } 0a$ , so  $E$  inherits two derivatives from a quotient whose denominator is positive. The moments at  $a = 0$  are elementary integrals of powers of  $\cos$ :  $2\pi, 0, \pi, 0$  and  $3\pi/4$ . They give  $E(0) = 1/2$ ,  $E'(0) = 0$  — the first derivative vanishes by the symmetry that kills the odd moments — and  $E''(0) = -1/8$ , and Taylor’s theorem in Mathlib’s `taylor_isLittle0.univ` form then yields `vonMisesSRatio_second_order`:  $E(a) - (\frac{1}{2} - a^2/16)$  is  $o(a^2)$  at the origin. The remainder is  $o(a^2)$  and not  $O(a^4)$ ; the stronger form would need a fourth derivative and nothing here requires it. `coherent_solution_critical`

`exponent` then combines the expansion with `coherent_iff_sRatio_eq`: writing  $a = Kr/D$  and rearranging  $E(a) = D/K$  gives  $K - 2D = Ka^2(\frac{1}{8} - 2q(a))$  with  $q$  the normalized remainder, and letting  $K \rightarrow 2D$  along any family of positive coherent solutions with  $r \rightarrow 0$  sends  $q \rightarrow 0$  and  $r^2/(K-2D) \rightarrow 1/D$ . The statement is deliberately filter-indexed rather than parameterized by a curve, so it constrains every such family and privileges no path. Because its hypotheses hold only eventually along the filter, and an eventual clause is vacuously true on the empty one, the file also exhibits a family satisfying it: `coherentBranch` selects the unique positive solution above threshold via `supercritical_fixed_point_existsUnique`, `coherentBranch.tendsto_zero` reads `coherent_branch_continuous_at_threshold` as a limit along  $K \downarrow 2D$ , and `coherentBranch.critical_exponent` discharges every hypothesis on it. Scope is inherited unchanged from the file above: this is asymptotics of the solutions of a stationary equation. It establishes no stability, no dynamical selection and no time course, so reading it as a square-root foot on a recovery curve requires the adiabatic separation and a threshold crossing at non-zero speed, both of which the main text states as assumptions. Every declaration in the file depends only on `propext`, `Classical.choice` and `Quot.sound`.

The gap between these stationary-density results and any statement about a finite trajectory is the dynamical mean-field limit — propagation of chaos for the Kuramoto system, which is a research programme rather than a lemma Sznitman (1991). No theorem in this development closes it.

## 4.1 Finite mean-field trajectory controls

Two finite- $N$  controls probe this gap without closing it. First, 500 seeded replicas of 1,000 identical-frequency oscillators were integrated at fixed coupling. Their steady order was 0.06266 below threshold, 0.16049 at threshold and 0.67735 above threshold. The last value lies -0.00535 from the stationary coherent value 0.68270. Refining the integration step by a factor of 4 at the same coupling, duration and tail sampling moves that steady order to 0.68168, so most of the gap is Euler–Maruyama bias. The last refinement still moves it by 0.00147, so the remaining -0.00102 bounds the finite- $N$  part rather than measuring it.

Early-growth fits are regressed over a window whose bounds follow the trace: it opens at  $2/\sqrt{N} = 0.06325$ , clear of the finite-size floor, and closes at half the static branch order, which runs from 0.20669 to 0.37050 across the swept couplings. The fits give  $\lambda = 0.41505K - 0.78619$ , against the infinite- $N$  linearization  $\lambda = 0.5K - 1$ . The shortfall from  $1/2$  is therefore not a property of the fit window; the sweep integrates at the step whose bias the control above measures, so finite  $N$  and the integrator both contribute to what remains. This exhibits finite sampled escape toward the coherent branch; it proves no stability or dynamical-selection theorem.

Second, 1,000 ensembles were swept from  $N = 10$  to  $N = 1000$  at  $K = 1$  and  $K = 3$ . In the subcritical control, mean order fell from 0.3726 to 0.0402 and absolute circular pair correlation from 0.0458 to 0.0010. The matched-sample sliced joint-product distance changed only from 0.0468 to 0.0419, exposing its finite-sample floor. Above threshold, mean order stabilized at 0.7208, but the unconditional lab-frame pair correlation remained between 0.459 and 0.525 rather than following the  $1/N$  reference. The unpinned collective orientation is a shared latent variable, so unconditional independence and independence conditional on that orientation are different claims. In the co-rotating frame, the  $N = 1000$  density at measured concentration  $a = 2.1624$  had  $L^1$  error 0.1291 from the von Mises target and Bessel self-consistency residual 0.00191. These finite, single-seed diagnostics prove neither form of propagation of chaos and discharge no Lean obligation.

## Dale-balanced rescue and calibration limits

Finite simulations separate synaptic balance from the magnitude of fluctuations. The network has no self-edges, an Erdős–Rényi edge probability of 0.2, and excitatory/inhibitory presynaptic fractions of 80/20. Each row’s positive and negative weights sum to  $g$  and  $-g$ , respectively. With identical frequencies,  $D = 1$ ,  $N = 500$ ,  $g = 1$ , step size 0.01 and duration 100, three seeded runs give zero-field mean order in  $[0.0508, 0.0574]$ , close to the  $1/\sqrt{N}$  reference. Controls at  $g \in \{0, 1, 2, 4\}$  and  $N \in \{250, 500, 1000\}$  show why row balance alone is insufficient to prescribe that reference: signed coupling can amplify finite-size fluctuations. Order decreases with size in both the weak and strong controls while  $N\langle r^2 \rangle$  remains above the independent-phase value of unity; these finite sizes do not establish an asymptotic scaling law. The  $g = 1$  regime is also below the positive-only mean-field threshold, so its low order does not establish a causal role for inhibition in preventing coherence.

A uniform per-pair term  $\epsilon$  adds mean-field strength  $K_{\text{eff}} = N\epsilon$ . The operational crossing, defined by second-half mean  $r = 0.2$  with all larger sampled couplings above it, occurs at  $K_{\text{eff}}/D \in [1.96, 1.99]$  across




Figure S1: Finite- $N$  convergence diagnostics for identical frequencies. The left panel compares branch-cut-invariant pair correlation with a  $1/N$  reference; the centre uses matched sliced-Wasserstein joint-product samples; the right compares the co-rotating  $N = 1000$  histogram with the von Mises density at measured concentration. The persistent supercritical lab-frame dependence is a negative result for the unconditional protocol, not a theorem about the conditional mean-field limit.

the three seeds (Fig. S2). Each value interpolates across one step of the  $\epsilon$  grid. That grid is geometric with ratio 1.136, and the one step of it the crossing falls in,  $[1.857, 2.110]$ , is subdivided linearly into intervals of width 0.025 and rerun on the same three seeds, which places the crossing in  $[1.958, 2.009]$  — the union of the refined intervals the three seeds cross in. That bracket contains  $K_c/D = 2$ , so refining the grid tenfold locates the crossing without separating it from the mean-field threshold at  $N = 500$ . This range describes finite-run variability, not a confidence interval or an infinite-size critical point. Equal-duration half-step controls probe numerical sensitivity near the transition; they do not establish convergence of the integrator.

The field conversion requires an explicit calibration rate  $\gamma$ :

$$E_{\text{req}} = \frac{K_{\text{eff}}}{\gamma N_{\text{cortex}} \text{shift } f}.$$

Since  $\text{shift } f$  is dimensionless,  $\gamma$  supplies the inverse-time units of the dynamical coupling. Under the numerical convention  $\gamma = 1$  in the simulation's inverse-time unit, the shared Fermi arithmetic gives field ranges  $[0.5856, 0.5930]$ ,  $[0.0732, 0.0741]$  and  $[0.0217, 0.0220]$  mV/mm at decay lengths 0.1, 0.2 and 0.3 mm. These conditional numbers fall below the 1–5 mV/mm comparison band and scale inversely with  $\gamma$ . The simulation does not measure that calibration, so it cannot independently validate physical-unit consistency or a cortical field threshold. Neither rescue nor these controls discharge a Lean or gluing obligation.

No clinical mechanism for perceptual anomalies is inferred from the threshold calculation.

*Lean 4 Implementation (Phase4RotatingFrame.lean):* The potential above omits the natural-frequency term  $-\sum_i \omega_i \theta_i$  that appears in the Lyapunov descent result `dV_dt.le.zero` in the combinatorial thermodynamics module. The omission is not cosmetic. With any  $\omega_i \neq 0$  the full potential is unbounded below (`kuramoto_potential_unbounded_below`), so it has no minimum for a phase-locked state to attain, and the two results concerned different functionals. This file chains them by the standard rotating-frame reduction  $\theta_i \mapsto \theta_i - \Omega t$ . For a system of identical natural frequencies  $\omega \equiv \Omega$ : `is_kuramoto_trajectory_rotate` sends a trajectory to a trajectory of the zero-frequency system `sys.reduced` — the same system Theorem 5's `ThermodynamicCover` works with; `kuramoto_potential_reduced` shows the two potentials coincide on it; `dynamic_potential_descent` transports the descent identity  $\dot{V} = -\sum_i \dot{\theta}_i^2$  onto the potential above; and `rotating_frame_chain` runs this through the minimum characterization to phase-locking and  $r^2 = 1$ . Phase differences, phase-locking and the order-parameter magnitude are proved frame-invariant (`rotate_sub`, `is_phase_locked_rotate`, `order_parameter_r_sq_shift`), so no observable is lost. `Examples.lean` §7 runs the chain end to end on a two-oscillator system, discharging every hypothesis, and exhibits a system whose full potential has no minimum. Scope: the reduction is exact only for identical frequencies. A spread of frequencies leaves residual detunings  $\omega_i - \Omega$  in the reduced system, to which the unboundedness result still applies, and phase-locking then depends on  $K_c$ , which this file does not treat — see the `Phase8SelfConsistency.lean` note above for what is proved about the threshold, and note that it concerns a stationary density rather than the trajectories this file transports.




Figure S2: Finite- $N$  rescue in the declared weak-synapse regime. Points are three-seed mean order; blue shading gives the seed range. Orange shading marks the union of sampled crossing brackets. The top axis uses the conditional  $\gamma = 1$  field conversion at decay length 0.2 mm; green shading is the 1–5 mV/mm comparison band. The dashed horizontal line is  $1/\sqrt{N}$ . The balanced strong-synapse control remains a distinct regime and does not satisfy the floor-level baseline gate.

*Trajectories that run.* The deterministic initial-value problem admits a unique global solution through each initial state. `kuramotoField.lipschitz` (`Phase3.CombinatorialThermodynamics.lean`) proves the Kuramoto vector field globally Lipschitz with constant  $2 \sum_{ij} |A_{ij}|$  —  $\sin$  is 1-Lipschitz and the index set is finite, so there is no smallness caveat and the natural frequencies drop out. `is_kuramoto_trajectory_unique` turns this into Picard–Lindelöf uniqueness on all of  $\mathbb{R}$ : two solutions agreeing at one instant agree forever, in both time directions. Existence follows by constructing solutions on arbitrary finite windows and gluing them using uniqueness. Mathlib’s `IsPicardLindelof` caps a local solution’s lifetime by the radius of the ball on which the field’s bounds hold, divided by the field’s sup norm there, and carries no global-in-time theorem; but `kuramotoField.norm.le` bounds this field by  $\sum_i |\omega_i| + \sum_{ij} |A_{ij}|$  on the *whole* state space, so the cap can be defeated by enlarging the ball after choosing the window, and `kuramoto_exists_on_window` produces a solution on every  $(t_0 - T, t_0 + T)$ . Uniqueness on an interval makes the family coherent — solutions on nested windows agree where both are defined — so  $\theta(t) := \alpha_{[t-t_0]}(t)$  is a solution on all of  $\mathbb{R}$  (`is_kuramoto_trajectory_exists`), and `kuramoto_cauchy_problem` states existence and uniqueness together as a single `ExistsUnique`: the trajectory predicate is inhabited exactly once for each initial state, rather than only by the solutions we write down by hand. What existence does not supply is behaviour, and for that a concrete solution is still what we have. On two oscillators with unit coupling and zero natural frequency the phase difference obeys  $\dot{\Delta} = -2 \sin \Delta$ , whose solution  $\Delta(t) = 2 \arctan(c e^{-2t})$  splits symmetrically into an exact trajectory of the full system (`Examples.lean` §15, `pairRelax.is_trajectory`). At  $c = 1$  it starts a quarter turn out of phase (`pairRelax_not_locked_at_zero`); the cosine of the phase difference tends to 1, the order parameter tends to 1 through the two-oscillator identity  $r^2 = (1 + \cos \Delta)/2$ , and the dynamic potential tends to its global minimum. By the uniqueness result it is the only solution through its initial state. Two general statements describe the long run, both for a zero-frequency system, which by `kuramoto_potential_unbounded_below` is the only case in which the potential is bounded below at all. `dynamic_potential_antitone` upgrades the instantaneous descent to monotonicity in  $t$  — `dynamic_potential_deriv_nonpos` constrains each instant separately and does not by itself say the trajectory settles — and with the uniform lower bound  $-\frac{1}{2} \sum_{ij} |A_{ij}|$  monotone convergence gives `dynamic_potential_tendsto`: the potential converges along *every* trajectory. Integrating the Lyapunov identity gives the quantitative form, `dissipation_integral_eq`: energy dissipated over  $[0, T]$  is exactly the potential dropped, so by `dissipation_integral_tendsto` the improper integral  $\int_0^\infty \sum_i \dot{\theta}_i^2 dt$  converges, to the total drop. This needed the Lyapunov identity in `HasDerivAt` form rather than as a statement about `deriv`, which carries

no differentiability information and cannot be fed to the fundamental theorem of calculus. `rotating-frame-dissipation` transports both to a uniform-frequency system. Finiteness of that integral is not the same as  $\sum_i \dot{\theta}_i^2 \rightarrow 0$ : an integrable function can spike forever on ever narrower intervals. The classical bridge is Barbalat's lemma, which Mathlib does not have, so it is proved here as `tendsto_zero_of_lipschitz_of_integral_le`: a non-negative Lipschitz function with  $\int_0^T g$  bounded above tends to zero, because a spike of height  $\varepsilon$  drags a window of width  $\varepsilon/2K$  up with it and only finitely many such windows fit under a finite integral. Its hypothesis is then checked for the dissipation rate without differentiating it: the velocities are bounded by  $\sum_i |\omega_i| + \sum_{ij} |A_{ij}|$ , the field is Lipschitz in the configuration, and the configuration is Lipschitz in time because its derivative *is* the bounded field (`trajectory_dist_le`, by the mean value theorem, with no invariant region needed since the bound holds on the whole state space). The conclusion is `velocity_sq_tendsto_zero`: along *every* trajectory of a zero-frequency system,  $\sum_i \dot{\theta}_i^2 \rightarrow 0$ ; equivalently each  $\dot{\theta}_i \rightarrow 0$  and the vector field along the trajectory tends to 0 in the state space. This is the first asymptotic statement here about an arbitrary trajectory rather than about a witness, and `Examples.lean` §15 fires it on the relaxing pair, whose velocity starts at  $-1$  and is therefore genuinely decaying rather than identically zero. It does not locate the limit. Locating it looks as though it needs the LaSalle invariance principle, which Mathlib does not carry, and a compact invariant set, which the state space  $V \rightarrow \mathbb{R}$  does not supply; neither is in fact required. LaSalle is what one needs when a decreasing functional is all one has, and here the flow is a gradient flow of the dynamic potential —  $\nabla_i$  of `kuramoto_potential_dynamic` is exactly  $-\text{kuramoto\_velocity}$  at zero frequencies, which is why `dynamic_potential_hasDerivAt` gives  $\dot{V} = -\sum_i \dot{\theta}_i^2$  as an identity — whose minimisers are already characterised. `lojasiewicz_estimate` adds the missing ingredient: on the region where every pairwise phase difference lies within a quarter turn,  $2a^2W \leq (\sum_{ij} A_{ij}) \sum_i \dot{\theta}_i^2$ , where  $W$  is the excess of the potential over its minimum and  $a > 0$  is a lower bound for the coupling. Its proof uses no spectral gap and no convexity, only two extremal indices: at the site  $m$  carrying the largest phase every coupling term has  $\theta_j - \theta_m \in [-D, 0]$  with  $D \leq \pi/2$  the spread, so every term is non-positive, and the term from the site carrying the smallest phase is at most  $-a \sin D$  (`velocity_max_le`); against  $W \leq \frac{1}{2} (\sum_{ij} A_{ij}) (1 - \cos D)$  the estimate follows from  $\sin^2 D = (1 - \cos D)(1 + \cos D)$  and  $\cos D \geq 0$ , which is the only place the quarter-turn hypothesis is used. The compactness objection is met without any invariant region put in by hand. Because all terms of the excess are non-negative, a single pair satisfies  $\frac{1}{2}a(1 - \cos(\theta_j - \theta_i)) \leq W$ , so an initial excess below  $a/2$  — which `dynamic_potential_antitone` preserves — keeps every pairwise cosine strictly positive for all time. A difference escaping  $[-\pi/2, \pi/2]$  would have to cross  $\pi/2$  in absolute value by the intermediate value theorem, and there the cosine vanishes; so the quarter-turn region is forward invariant (`spread_le_of_small`). This is branch-pinning by continuity on a connected half-line, and it replaces the quotient by the phase-shift symmetry we had thought necessary. The differential inequality  $\dot{W} \leq -cW$  with  $c = 2a^2/\sum_{ij} A_{ij}$  then integrates, in the elementary form “ $W(t)e^{ct}$  has non-positive derivative”, to  $W(t) \leq W(0)e^{-ct}$  (`excess_decay`). The rate transfers to the velocity through  $|\sin d| \leq \sqrt{2(1 - \cos d)}$ , giving  $|\dot{\theta}_i(t)| \leq Be^{-ct/2}$  (`velocity_abs_le_exp`), which is integrable on  $(0, \infty)$ ; since  $\theta_i(t) = \theta_i(0) + \int_0^t \dot{\theta}_i$ , the trajectory converges (`phase_tendsto`), and  $W \rightarrow 0$  forces the limit to be phase-locked, hence a global minimiser with  $r^2 = 1$  (`kuramoto_tendsto_global_minimum`). `rotating_frame_tendsto_global_minimum` transports it to uniform frequencies, where what converges is the relative configuration. The two hypotheses are on the initial data and neither can be removed: splay and twisted configurations are equilibria too, so “every trajectory reaches the phase-locked state” is false, and what is proved is the arc condition version — which the two-oscillator witness satisfies for every  $c$ , since  $2 \arctan$  lands in  $(-\pi, \pi)$ . `Examples.lean` §17 runs the theorem on three sites, where the method of §15 is unavailable because the system has no closed-form solution: the trajectory is produced by `is_kuramoto_trajectory_exists` from the initial data  $(0, 0, \frac{1}{2})$ , which is not phase-locked, and its limit by the convergence theorem as  $\theta(0) + \int_0^\infty \dot{\theta}$ ; the numerical content of checking the hypotheses is  $\cos \frac{1}{2} > \frac{3}{4}$ , supplied by `Real.cos_bound`.

*From continuous to discrete coupling:* The coupling matrix  $A_{ij}$  above is not postulated. A triangulation of the substrate assigns to each edge  $(u, v)$  a region of the manifold, and the coupling weight is the Lebesgue–Bochner integral of the scalar magnitude of the stress-energy tensor over that region (`edge_weight`); symmetry and non-negativity of  $A$  are then proved rather than assumed. The bridge back is mesh refinement: as the regions shrink, the discrete energy  $\frac{1}{2} \sum_{u,v} \mu(R_{uv}) f(x_u)$  should converge to the continuous energy  $\int_S f$ . Stated as a conjecture over the bare triangulation class, that is false: nothing there forces the regions to cover the manifold, empty regions have diameter zero, and the everywhere-empty triangulation refutes it. Stated correctly it is a theorem (`mesh_refinement`).

convergence), quantified over a sequence of triangulations of a region  $S$  whose regions are measurable, pairwise disjoint across distinct edges, covering, and anchored to the embedded vertices, with mesh size tending to zero. The fixed-mesh error is bounded by  $\varepsilon \mu(S)$ , with  $\varepsilon$  a modulus of continuity of the integrand, and the symmetric double sum is proved equal to the midpoint rule over unordered edges. Two limitations follow. The theorem gives convergence, not the  $O(1/N^2)$  rate reported by `simulations/mesh_refinement.py`. And `TriangulatedManifold` itself is bare data: it pairs a complex, an embedding and a family of edge regions without relating them, so the geometric conditions live in a separate predicate `IsRegularTriangulation`. Because those conditions are in tension — cells disjoint yet covering, every vertex anchored to each of its edges — `Examples.lean` §6 exhibits an instance: the uniform partition of  $[0, 1)$  into  $N$  half-open cells, on which the convergence theorem is instantiated and the approximations are checked to genuinely change with  $N$ .

The discretization structure `DiscreteThermodynamics` carries that predicate as a field, which is what makes its edge weights integrals over a geometrically constrained family rather than over arbitrary sets, and three statements follow that were unavailable before. `weight_self`: a vertex has no edge to itself, so the coupling matrix has zero diagonal. `face_of_weight_ne_zero`: a nonzero weight  $w(u, v)$  forces  $\{u, v\}$  to be an actual face of the complex, so the discretization cannot invent couplings between vertices the triangulation does not join. And `total_weight_eq_setIntegral`:  $\frac{1}{2} \sum_u \sum_v w(u, v) = \int_S |T|$ , i.e. the total coupling strength is exactly the scalar stress-energy of the region covered — nothing counted twice (disjointness), nothing dropped (covering). That last identity is the precise sense in which the coupling matrix is derived from the continuum rather than posited, and it is exactly what the everywhere-empty triangulation violates. `Examples.lean` §6 witnesses the strengthened structure on the same grid: each consecutive edge carries weight  $1/N$ , and the weights sum to 1 for every  $N$ . The witness is one-dimensional and its stress-energy tensor is constant, so it establishes satisfiability, not a physically interesting discretization. One field of the structure is deliberately left free and is marked `[MODELLING]`: `scalar_magnitude`, the scalar read off the stress-energy tensor, constrained only to be non-negative. A  $(0, 2)$  tensor has no canonical scalar — the trace, an observer’s energy density and the tensor norm are different functions, and the structure carries neither a metric nor a frame — so imposing one would pin a choice rather than derive it, and the three statements above hold for every non-negative choice, so it would prove nothing new. That the choice is nonetheless a physical input is stated as a theorem rather than left as a remark: §6 exhibits a second instance over the same triangulation, tensor and regions, differing only in that field, for which the partition identity again holds while every edge weight differs (`gridThermo_magnitude_matters`). The remaining constraint, non-negativity, is the discrete shadow of an energy condition and is what makes the induced couplings a legal Kuramoto coupling.

## 4.2 The fold-and-cross proof of strict decrease

Uniqueness of the coherent branch is the injectivity of  $E(a) = \mathbb{E}_a[\sin^2 \theta]$  on  $(0, \infty)$ , and what is proved is the stronger statement that  $E$  is strictly decreasing on  $[0, \infty)$  (`vonMisesSRatio_strictAnti0n`). The main text states the conclusion and the shape of the argument; this is the argument, and it avoids the differentiation route entirely. Differentiating the exponential family in its natural parameter gives  $E'(a) = -\text{Cov}_a(\cos^2 \theta, \cos \theta)$ , which costs differentiation under the integral sign and leaves a covariance that is not sign-definite pointwise: symmetrising with an independent copy gives  $\frac{1}{2} \mathbb{E}[(X - X')^2(X + X')]$  with  $X = \cos \theta$ , an integrand changing sign with  $X + X'$ . Nothing in the proof differentiates anything. The target is kept as the two-point inequality  $C_2(a)Z(b) < C_2(b)Z(a)$  for  $a < b$ , and two moves settle it. *Fold first*: the two reflections that proved  $I_2 \geq 0$  carry both integrals onto  $[0, \pi/2]$ , replacing  $e^{a \cos \theta}$  by  $W(a, \theta) = 2 \cosh(a \cos \theta)$  and confining  $\cos \theta$  to  $[0, 1]$ , after which  $\cos 2x - \cos 2y$  and  $W(b, x)W(a, y) - W(a, x)W(b, y)$  have the same sign, both being governed by  $\cos x$  against  $\cos y$ . This is the symmetrisation  $X \mapsto -X$  that the covariance form needed the tilt to perform, done once at the start; its analytic content is one identity,  $2 \cosh X \cosh Y = \cosh(X+Y) + \cosh(X-Y)$  with  $\cosh$  increasing in  $|\cdot|$ . *Then cross*: instead of integrating that inequality over a square, let  $\alpha$  be the mean of  $\cos 2x$  under  $W(a, \cdot)$  and  $x_0 \in [0, \pi/2]$  a point where  $\cos 2x_0 = \alpha$ , supplied by the intermediate value theorem; evaluating at  $y = x_0$  and integrating in  $x$  alone gives the result, since the term carrying  $W(b, x_0)$  vanishes by the choice of  $\alpha$ . No product measure and no Fubini appear. Run with  $\cos \theta$  in place of  $\sin^2 \theta$  the same argument needs no fold —  $R(a) = \mathbb{E}_a[\cos \theta]$  is the mean of the variable the family is tilted by — and gives `besselRatio_strictMono`:  $R$  is strictly increasing on all of  $\mathbb{R}$ , the monotonicity of the Bessel ratio, which Mathlib carries in no form.

## 5 Global Sections and the Unity of Consciousness

Phase synchronization does not imply overlap compatibility. The model represents global consistency by gluing compatible local descriptions; its identification with conscious unity remains a separate stipulation.

**Theorem 5 (Sheaf-Theoretic Unity):** Given a sheaf of probability measures  $\mathcal{F}$  over a continuous biological substrate  $X$ , if local states  $P_U \in \mathcal{F}(U)$  agree on all spatial intersections  $U \cap V$ , they uniquely glue into a Global Section  $P_X \in \mathcal{F}(X)$ . The overlap agreement is an independent physical hypothesis; phase-locking does not imply it.

*Lean 4 Implementation (Phase5\_GlobalSection.lean):* The headline gluing theorem separates two inputs. The class field `ThermodynamicCover.thermodynamic_equilibrium` places the phase field at the potential minimum, while `LocalSectionSynchronization.section_agrees_of_phase_eq` independently requires equal-phase patches to agree on overlaps. Only the second supplies the compatible family consumed by the sheaf theorem; no dynamics in the development derives it. What the second one says can be stated without the sheaf: on a discrete substrate two patch sections have the same restriction to their overlap exactly when they assign the same mass to every site they share (`restrict_eq_iff_densityOn_eqOn`), so the hypothesis is that patches at a common phase carry a common local density where they meet. It excludes data the rest of the structure admits: a two-patch cover of the three-site substrate at one phase, with profiles differing at the shared site, satisfies every other requirement and fails this one (`overlap_agreement_fails`). It does not require patches to agree away from their overlaps — the three-patch witness carries three pairwise-different profiles. The equilibrium field is discharged one file over on a class of initial data: `Phase5_EquilibriumBridge.lean` supplies `ThermodynamicCover.ofConvergentTrajectory` from `kuramoto_tendsto_global_minimum`, and `Examples.lean` §17.1 is a three-patch cover built over the resulting phase field. The cover class also has an independent instance (`Examples.lean`, §4), so the theorem is not conditional on a structure of unknown realizability. Its three-site substrate has two overlapping patches meeting at one site; local sections are germ families of the phase-indexed measure  $\mu_\theta = (1 + \cos \theta)^+ \delta_{\text{mid}}$  in the sheafification of the probability presheaf. The family is provably non-constant in  $\theta$  (`phaseMeasure_not_const`), and `thermodynamic_equilibrium` follows from `phase_locked_minimizes_potential`. This is a formalized gluing mechanism under an explicit compatibility postulate, not a derivation of compatibility from synchronization. Three results sharpen its scope. First, `potential_min_iff_phase_locked` (`Phase4_KuramotoDynamics.lean`) characterizes the minimizers of the reduced potential in both directions, from which `ThermodynamicCover.phase_locked` follows for every instance. Second, `Examples.lean` §13 carries a non-constant phase field and a non-uniform invariant measure, with patch-local densities read from the sections through `densityOn`; agreement at the shared site is an equation between separately computed numbers. The limit is structural: phase-locking equates phases, while agreement of the associated sections remains the independent class field named above.

The gluing construction assumes local overlap agreement and derives both existence and uniqueness of the global section. The lemma `probability_glue_unique` takes a compatible family to its unique global extension; `ThermodynamicCover.invariantMeasure` denotes that section. The construction uses local data, not a global state supplied as a premise.

## 6 Reflexive Topology and the Self

To minimize internal thermodynamic friction, the unified Global Section must model its own internal state changes.

**Theorem 6 (Conditional Auto-Resonance):** On a complete metric space of global sections, a self-prediction map has a unique fixed point if it is Lipschitz at the independently assumed rate  $e^{-(K-K_c)\tau/2} < 1$ . Coherence supplies  $K > K_c$ ; it does not supply the Lipschitz law.

*Lean 4 Implementation (Phase6\_ReflexiveTopology.lean):* We formalize the reflexive projection as a contraction on a complete metric space and apply the Banach fixed-point theorem to obtain a unique fixed point — the “Self” as a topological feature rather than a metaphysical addition. The entropy production functional does not appear in this file: the identification of the Banach fixed point with the minimizer of  $\sigma$  is an interpretive step made in the main text, and the Lipschitz bound at the mean-field rate is an instance obligation rather than a consequence of field dynamics. Without an instance of `ReflexiveBoundary` and a metric on `GlobalSection`, the theorem would be conditional on structures of unknown realizability — the same failure mode an uninhabited class has. `Examples.lean` §10 supplies both on the three-site cortex of §4, and the metric is built from the measures rather than imposed. Five stages. (i) The germ-measure dictionary: on a discrete substrate every point has a

smallest open neighbourhood, so “mass carried at  $x$ ” is a well-defined map out of the stalk at  $x$  (constructed as the colimit map of a cocone whose cocone law is precisely that restriction preserves point masses); reading a section’s germ at each site of an open  $U$  gives a map to mass profiles on  $U$ , and `massEquivOn` proves it a bijection — injective because a germ on a discrete space is determined by its restriction to the point, surjective because every prescribed profile is realized by a finite measure on  $U$ . `massEquiv` is the case  $U = \mathbb{T}$ ; the general case over an arbitrary open  $U$  is what lets Derivation 5’s §14 witness build its patch-local sections over their own patches. Mathlib supplies nothing here: `TopCat.Presheaf.sheafify` has no adjunction and no `isIso.toSheafify`, so the inverse had to be built. (ii) The metric: the distance between two sections is the largest difference between the masses their measures put on a site, i.e. the uniform distance between their densities (equivalent to the total variation of the glued measures on a finite substrate, within a factor of the number of sites, but not equal to it). It is complete, proved by exhibiting the limit section of a Cauchy sequence of densities, and `exists_dist_lt_one` exhibits distinct sections at distance  $1/2$ . (iii) The Self, built through the avatar. Carrying the predictive map as free data beside the avatar would let a blinded boundary—one whose avatar is replaced by a constant—retain the same predictive model and the same Self, since nothing would tie the map to the avatar. The structure instead carries the avatar’s `write (auto_resonance)` and its `read-out (readout)`, and `predict` is their composite rather than a field, so the self-model runs through the avatar region or it does not exist. On the cortex the map is `avatarRead` — the single mass a section over the one-site avatar region carries — followed by `avatarReadout`, which returns that mass averaged with the baseline; `cortexPredict_dist` gives the exact contraction factor (one half of *what the avatar sees*: the read-out is blind off the avatar region, and that loss is the fold, stated rather than hidden), `cortexPredict_not_const` that it is not constant, `cortexPredict_fixed_unique` that its fixed point is unique. A map such as `relax`, averaging the whole field with the baseline, is unavailable as a witness here: it reads every site and so does not factor through a one-site avatar. The 0/1 metric is retained as `gsDiscreteMetric` together with `contracting_implies_const`, which proves that under that metric every contraction is constant, so that metric supports no informative fixed point. (iv) The avatar as a constraint. Following the rule that data classes stay bare, the resonance condition is a predicate and not a field: `IsRestrictionResonance rb` says `rb.auto_resonance s` is the presheaf restriction of  $s$  to the avatar region. Four theorems fix its status. `eq_of_avatar_eq`: if the avatar regions of a family of resonant boundaries cover the substrate, two global states with the same reading at every avatar are equal — the readings are local, and the global state is reconstructed from them by `probability_glue_unique`, the same sheaf condition Derivation 5 uses. `self_mem_range_readout`: a fixed point lies in the range of the read-out, so the Self is reconstructed from a single section over the avatar region — the low-dimensional fold, as a theorem. `self_eq_readout_restrict`: under resonance the Self satisfies  $s = \text{readout}(s|_{\text{avatar}})$ , the field being what its own localized self-encoding reconstructs; this is the statement a structure carrying `predict` as free data cannot make, since there `predict` is unrelated to the restriction. And `not_self_of_constResonance` is the negative half: the blinded boundary’s map is constant, `constResonance_existsUnique_self` names its only fixed point as the state its stipulated contents extend to, and the original Self is not one. The opposite statement, that blinding leaves the Self where it was, does not typecheck against this structure — a contraction hypothesis about a boundary’s own map says nothing about the blinded one’s — and `constResonance_not_isRestrictionResonance` shows the blinded boundary fails the predicate as soon as two global states differ on the avatar region. On the cortex all of it is exercised: `cortexReflexive` is built from the restriction and discharges the predicate by `rfl`; `cortexReflexive_restrict_ne` checks the avatar region separates the silent field from the baseline, so the predicate is not vacuous there; `siteReflexive` gives one avatar per site, and since the singletons cover, `cortex_eq_of_avatar_eq` reconstructs the global state from three local readings; `cortexBlind.self_ne` proves the baseline is *not* a fixed point of the blinded map, and `cortexBlind_not_determined` exhibits two distinct states with identical blinded readings. Banach’s uniqueness half is `self.unique`. (v) The contraction rate, from  $K$  and  $D$ . `resonanceRate K D`  $\tau = \exp(-(K - K_c)\tau/2)$  with  $K_c = \text{critical_coupling}$   $D = 2D$  is the factor the self-model is assumed to be Lipschitz at over an interval  $\tau$ . The rate in its exponent is the growth rate of the incoherent state’s instability: near  $r = 0$  the mean-field order parameter obeys  $\dot{r} = \frac{1}{2}(K - K_c)r + O(r^3)$ . Relaxation about the *coherent* branch is twice as fast — linearizing  $\dot{r} = \frac{1}{2}(K - 2D)r - cr^3$  about its stable fixed point gives  $(K - K_c)$  — so the assumed Lipschitz constant is larger than the coherent branch’s relaxation factor, which makes the hypothesis weaker than the dynamics would support rather than stronger. The threshold is unaffected:  $\exp(-(K - K_c)\tau/2) < 1$  and  $\exp(-(K - K_c)\tau) < 1$  are both the condition  $K > K_c$ , so no theorem below depends on which of the two rates is used. `resonanceRate_lt_one` and `one_le_resonanceRate` split the parameter plane at Sakaguchi’s threshold; `self_of_supercritical` derives the unique fixed point from

$K > K_c$ , and `Chain.lean`'s `self_of_coherent_order_parameter` derives it instead from the *existence* of the coherent order parameter — a distinction that matters, since  $K_c$  is a definition unfolding to  $2D$ , so the first statement is insensitive to whether the bifurcation is proved and the second is not, and `not_contractingWith_resonanceRate` records that at or below threshold `ContractingWith` at that rate is false for every self-map of every metric space, so the argument is unavailable rather than merely unproved. The witness sits at  $K = 3$ ,  $D = 1$ ,  $\tau = 2 \ln 2$ , where the rate is exactly the one half the read-out achieves (`cortexResonanceRate`), and `cortexSubcritical_not_contracting` runs the negative half on the same map at  $K = 1$ . Three limitations remain and we state them. The read-out `avatarReadout` is a modelling choice: no field dynamics in the development produces it, and the Lipschitz bound at the mean-field rate is an instance obligation, not a theorem — the obligation is stated in the substrate's parameters and its consequence is conditional on the threshold. The identification of this three-site substrate's coupling with  $K = 3$  is a choice of units for the witness and not a measurement. Nothing derives the resonance condition either: that a physical avatar performs the restriction is an assumption on an instance, in the same class as the overlap-agreement obligation of Derivation 5. And the Lévy–Prokhorov metric named in the file header is *not* what is constructed: on a discrete substrate the  $\varepsilon$ -thickening of a set is the set itself for  $\varepsilon < 1$ , so  $\mu \mapsto \frac{1}{2}\mu + \frac{1}{2}\mu_0$  is not a Lévy–Prokhorov contraction there, and Mathlib carries no completeness result for that metric. The uniform distance between densities is the metric this substrate supports; we make no claim that it and Lévy–Prokhorov are equivalent.

## 7 Coupling support and continuum observables

Two mathematical comparisons distinguish a restriction on wiring support from a restriction on the measure used to evaluate a kernel.

**Theorem 7 (Resource-matched coupling reallocation):** If a fixed wiring support excludes a pair whose phase correlation exceeds the maximum over its wired pairs, a coupling using that pair can achieve greater total correlation at the same resource budget. This is a conditional support comparison, not a general inequality between biological and silicon hardware.

*Coupling reallocation (Phase7):* The lemma `exists_better_coupling_allocation`, wrapped as the theorem `continuous_beats_rigid_topology`, proves the reallocation step: given a coupling matrix satisfying `is_strictly_suboptimal` (positive weight on a less phase-correlated pair while a more correlated pair exists), there is another valid coupling with identical total resources and strictly greater total phase correlation.

*Lean 4 Implementation (Phase7.Rigidity.lean):* Two results specify the scope of the comparison. (i) *Rigidity as a support constraint.* `RealizableIn S R A` says  $A$  is a valid coupling meeting resource budget  $R$  and vanishing on every pair the architecture has no wire for — weights tunable, wiring fixed. `totalCorrelation_le_of_realizable` proves no realizable coupling exceeds  $R \cdot M$  when  $M$  bounds phase correlation over the wired pairs, and `exists_realizable_pair` attains it. `rigid_is_strictly_suboptimal` then *derives* `is_strictly_suboptimal` for every realizable coupling, whenever some pair  $(k, l)$  the wiring misses has correlation exceeding  $M$ : the hypothesis Theorem 7 would otherwise have to assume is a consequence of the wiring pattern. `rigid_gap` gives the quantitative form, a gain of at least  $R(\cos(\theta_l - \theta_k) - M) > 0$  at equal budget. `Examples.lean` §8 witnesses this on three sites wired only to the maximally anti-correlated pair, where the gap is exactly 2, and runs `exists_better_coupling_allocation` on it with no assumed hypothesis. (ii) *Continuity as a measure-theoretic distinction.* `fieldCorrelation` integrates a kernel against  $\mu \otimes \mu$  as in Derivation 4. A finitely-sited architecture is a `Finset X`; `fieldCorrelation_sited_eq_zero` proves that on an atomless substrate such a kernel contributes exactly zero, for every phase field and every choice of weights, while `fieldCorrelation_patchKernel` gives  $c\mu(U)^2$  for a patch. `sited_architecture_below_field_optimum` combines them; `Examples.lean` §8 instantiates it on  $\mathbb{R}$  with a thousand units of weight at the origin against the unit interval.

**Scope caveat.** Result (i) is resource-matched and not question-begging, but it turns on fixed wiring support, not on continuity: an architecture wired to the best pair is not beaten. Result (ii) is a genuine type-level separation but is *not* resource-matched, since a measure-zero substrate carries zero resource in the  $\mu \otimes \mu$  sense as well; what it establishes is that the continuum coupling energy cannot register a finitely-sited architecture at all, which is a claim about the model's observable rather than about computational power. Neither result formalizes manifold structure or lattice geometry. Neither result establishes that digital hardware cannot approximate field dynamics or support consciousness.

## 8 Claims and their Lean identifiers

The table below is the claim-by-claim map, kept here rather than in the main text. It is here because it is a reference rather than an argument: each row names the Lean identifier that carries a claim and states the scope limitation of that identifier, and no step of the main text depends on the reader working through it. Figure 4 and Table 1 in the main text give the eight connecting hypotheses, their kinds and their statements; the present table records node-level results and their implementation scope.

Table S1: Mapping of physical claims to Lean 4 formalization status. The development declares **no axioms**: `#print axioms` on any result above reports only `propext`, `Classical.choice` and `Quot.sound`. A physical postulate is an obligation on each instance rather than a global assertion, which is the shape a standalone axiom cannot have (Section 9); two of the postulates a reader might expect to find are not postulates at all, but consequences of the local data (Derivation 5). “Theorem” still does not mean “physically established”: the postulates live in hypotheses and instance obligations, and every such class is inhabited by a witness in `Examples.lean`, so no row is vacuous.

| Conceptual Claim                             | Lean 4 Implementation                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            | Status                       |
|----------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------|
| Capacity of a finite phase space             | The first premise, discharged by theorem rather than by footnote ( <code>Phase1.PhaseSpaceCapacity</code> ). $H(p) \leq \log  X $ for any distribution on a finite state space, with equality exactly at the uniform law ( <code>shannon_entropy.eq_log_card_iff</code> ); a stream with $ P ^k >  X $ histories cannot be recorded injectively ( <code>absorb_not_injective</code> ) and its entropy exceeds anything the system can hold; and on a <i>finite</i> space an unreachable state is an erasure, hence heat ( <code>finite_phase_space_dissipates</code> ). Finiteness is assumed, not derived, and is fenced — <code>Nat.succ</code> refutes the last step on an infinite space ( <code>Examples §19</code> ). The symmetry group and the metric are still consumed by no theorem                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Theorem                      |
| Pointwise Vacuum Minimization                | <code>pointwise_vacuum_of_global_min</code> ( <code>Phase1</code> ), potential energy defined as an integral; holds a.e., and a double-well witness proves the a.e. cannot be strengthened                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       | Theorem                      |
| Landauer’s Principle                         | Derived from statistical mechanics ( <code>Phase3</code> ); the axiom-free reversibility core becomes a heat bound via the field <code>StatisticalMechanics.heat_eq</code>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       | Theorem + instance postulate |
| Symmetry Breaking implies Boundaries         | Derived from pointwise potential minimization ( <code>Phase1</code> ); witnessed on a double well with a $\mathbb{Z}_2$ symmetry whose minimisation hypothesis is proved, not assumed                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            | Theorem                      |
| Inevitability of Boundaries ( $\pi_0$ )      | The domain-wall case of the Kibble mechanism ( <code>Phase1 §3</code> ): on a connected substrate, a continuous field taking values in two different connected components of the vacuum manifold must leave the vacuum somewhere ( <code>exists_notMem_of_connectedComponentIn</code> ). Witnessed on the double well, whose vacuum manifold is computed to be exactly $\{-1, 1\}$ and proved disconnected ( <code>Examples §11.1</code> ); the field class is non-empty and the forced wall is located. Higher homotopy ( $\pi_1$ vortices, $\pi_2$ monopoles) available and not taken                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          | Theorem                      |
| Nonpredictive information bounds dissipation | Still, Sivak, Bell and Crooks’ bound is the field <code>PredictiveDissipation.still_bound</code> ( <code>Phase3.PredictiveThermodynamics</code> ). Everything downstream is derived: mutual information is defined from Mathlib’s <code>klDiv</code> ; the data processing inequality for the chain $X_t \rightarrow S_t \rightarrow S_{t+1}$ proves the nonpredictive information non-negative ( <code>predictiveInfo.le_mutualInfo</code> ), from which $W_{\text{diss}} \geq 0$ follows rather than being assumed; a dissipation budget bounds wasted memory, and zero dissipation forces every retained bit to be predictive. Witnessed in <code>Examples §18</code> by two systems over one two-bit law, one permitted to dissipate nothing and one forced to dissipate its whole memory, with a counterexample showing the no-back-action hypothesis is load-bearing. Descent of $\sigma$ alone does not imply $D_{KL} \rightarrow 0$ . On a finite register with a bipartite environment the postulate is discharged rather than assumed: <code>PredictiveDissipation.ofLandauer</code> ( <code>Phase3.LandauerBridge</code> ) builds a predictive structure whose thermal scale and dissipated work are the register’s own and derives <code>still_bound</code> from <code>landauer_bound</code> , given the identification of wasted memory with erased entropy. That identification is the physical content and is not free — a reversible update erases nothing, so it forces the waste to zero ( <code>nonpredictive.eq_zero_of_injective</code> ) — and on the one-bit eraser the bound is attained ( <code>Examples §18.5</code> ) | Theorem + instance postulate |
| Structural Resonance KL Bound                | Gibbs’ inequality proved; the bound is the field <code>StructuralResonance.kl_bound</code> ( <code>Phase3.KLBound</code> ). The postulate is consistent and yields $\sigma \geq 0$ . No limit may be read off it                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 | Theorem + instance postulate |
| Fixed-phase kernel descent                   | The bounded drift operator and its Hilbert-space adjoint give the squared-drift gradient. The theorem <code>structural_resonance_decreases_sigmaContinuum</code> ( <code>Phase8</code> ) proves descent for the specified flow with phases fixed. An atomless witness is supplied in <code>Examples §20</code> . Neither joint dynamics nor thermodynamic entropy production is established                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      | Theorem                      |

| Conceptual Claim                        | Lean 4 Implementation                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              | Status                                                  |
|-----------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------|
| Phase-locking minimizes squared drift   | For a <i>reciprocal</i> coupling kernel, the phase-locked field has the least squared-drift objective of <i>any</i> field ( <code>phase.locked.minimizes_entropy_of_symm</code> , Phase8 §2a). The mean-matching condition is derived: the coupling term is antisymmetric under exchanging its two sites, so total drift is conserved across phase fields ( <code>total.drift.eq.of_symm</code> ). Reciprocity is necessary — a one-way coupling on two sites strictly beats its own phase-locked state (Examples §16)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             | Theorem                                                 |
| Electrodynamic Discretization           | Lebesgue-Bochner edge weights (Phase2) sitting on faces of the complex and partitioning the region's stress-energy; mesh refinement converges, witnessed on $[0, 1)$ . The scalar read off the stress-energy tensor is a marked modelling choice: the partition holds for any non-negative choice, and two instances differing only in it have different couplings (Examples §6)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Theorem + instance postulate                            |
| Coarse-graining of a symmetric coupling | <code>mesh.refinement.convergence</code> (Phase2) proves convergence of the scalar coupling energy along regular refining triangulations. On the formalized $[0, 1)$ grid the energy is the composite trapezoidal rule, so a $C^2$ integrand has error at most $\zeta/(12N^2)$ . The result produces neither a kernel nor dynamics: <code>edge.region</code> supplies no product-structured data on $M \times M$ . A genuine kernel limit requires new architecture before the separate dynamical mean-field problem can be stated                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 | Theorem (scalar energy only)                            |
| The field is the endogenous EM field    | <b>Conditions formalized; the identification itself not derived.</b> A continuum coupling kernel is posited after the scalar convergence result; identifying that kernel with cortical ephaptic coupling is the framework's central empirical commitment (the physical-calibration subsection of the main article). <code>IsEMFieldCoupling</code> (Phase9. <code>EMIdentification.lean</code> ) is the predicate E56 carries: a jointly continuous kernel, a normalised substrate, a coupling constant $K$ equal to that kernel's double average, no single site contributing more than half of it, and a noise $D$ that is the field's own. No object in this development denotes cortex, so whether cortex satisfies them is not a Lean question. Supported as <i>modulatory and integrative</i> , not as primary: measured endogenous fields of 1–5 mV/mm shift spike timing by 1–3 ms. Calibration and observational tests are stated in the main article                                                                                                                                                                                                                                                                                                                                     | Modelling assumption + empirical commitment             |
| Frustration as memory capacity          | <b>Outside the regime the chain covers.</b> <code>ThermodynamicCover.A.pos</code> excludes frustrated couplings at the gluing node, and circle-valued offsets from patch phases are coboundaries. Cortical phase singularities supply integer transitions between local real lifts instead: <code>loopWinding.eq.zero.of.hasGlobalLift</code> proves non-zero winding obstructs a global lift, with a winding-one witness. No theorem connects winding sectors to content                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          | Theorem (topological obstruction) + empirical candidate |
| Phase synchronization to Unity          | Sheaf-theoretic gluing (Phase5) over a <code>ThermodynamicCover</code> , whose fields are all local: the global section is constructed by <code>probability.glue.unique</code> , and the statement that every local section is its restriction, <code>sync.to.section.eq</code> , is a theorem about it. Witnessed by two patches built from different mass profiles, each section constructed over its own patch, whose glued state is neither of them (Examples §14). The cover's equilibrium hypothesis is not an assumption on every instance: <code>ThermodynamicCover.ofConvergentTrajectory</code> (Phase5. <code>EquilibriumBridge</code> ) discharges it from <code>kuramoto.tendsto.global.minimum</code> , and Examples §17.1 is a three-patch cover whose phase field is the limit of a trajectory started <i>unsynchronized</i> . The overlap-agreement hypothesis is untouched by any dynamics and is read pointwise: two patch sections have the same restriction to their overlap exactly when they assign the same mass to every shared site ( <code>restrict.eq.iff.densityOn.eqOn</code> ), and a two-patch cover at one phase with profiles differing at the shared site fails it while meeting every other requirement ( <code>overlap.agreement.fails</code> , Examples §20) | Theorem + instance postulate                            |
| Reflexive Topology (the Self)           | Banach fixed point for a self-prediction map (Phase6); metric, completeness and the Lipschitz law at $e^{-(K-K_c)\tau/2}$ are hypotheses, discharged on the three-site witness. Coherence yields $K > K_c$ and hence makes the stipulated rate less than one; it does not prove that the avatar read-out contracts at that rate. E89 carries both the independent Lipschitz law and the identification of the glued state with that fixed point                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    | Conditional theorem + modelling postulate (witnessed)   |
| Kuramoto Lyapunov descent               | Descent chained to the minimum characterization by the rotating frame (Phase4. <code>RotatingFrame</code> ); identical frequencies                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 | Theorem                                                 |

| Conceptual Claim                           | Lean 4 Implementation                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        | Status                                          |
|--------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------|
| Kuramoto dynamics run                      | The initial value problem is well posed on all of $\mathbb{R}$ : the field is globally Lipschitz and globally bounded ( <code>kuramotoField.lipschitz</code> , <code>kuramotoField.norm.le</code> ), and through every state at every instant passes exactly one trajectory, defined for all time ( <code>kuramoto-cauchy-problem</code> ). Along <i>every</i> trajectory of a zero-frequency system the dynamic potential converges ( <code>dynamic.potential.tendsto</code> ) and the total dissipation $\int_0^\infty \sum_i \dot{\theta}_i^2$ is finite, equal to the potential's total drop ( <code>dissipation.integral.tendsto</code> ). The motion stops: Barbalat's lemma is proved here, Mathlib lacking it, and applied, so $\sum_i \dot{\theta}_i^2 \rightarrow 0$ along every trajectory ( <code>velocity-sq.tendsto_zero</code> ). Where it stops is proved under an initial-data condition: if the phases start within a quarter turn and the potential excess starts below $a/2$ , where $a > 0$ bounds the coupling below, the trajectory converges to a phase-locked configuration that is a <i>global</i> minimiser, with $r^2 = 1$ ( <code>kuramoto.tendsto.global.minimum</code> ), by a Polyak–Łojasiewicz estimate rather than a LaSalle principle. Unconditionally this is false — splay and twisted states are equilibria — so the condition cannot be removed. Witnessed on three sites, where no closed form exists ( <code>Examples §17</code> ) | Theorem (conditional)                           |
| Continuous vs. Rigid Topology              | Fixed wiring support at matched resource ( <code>Phase7.Rigidity</code> ), deriving the suboptimality Phase 7 assumed; typed continuum/discrete separation                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Theorem                                         |
| Critical Coupling Threshold ( $K_c = 2D$ ) | Self-consistency equation ( <code>Phase8.SelfConsistency</code> ): at or below $K_c$ , threshold included, only $r = 0$ solves it; above $K_c$ <i>exactly one</i> $r > 0$ does ( <code>supercritical.fixed.point.existsUnique</code> ), so the non-negative solutions are $\{0\}$ below and $\{0, r\}$ above; that $r$ tends to 0 as $K \downarrow K_c$ and increases strictly with $K$ ( <code>coherent.branch.strictMono</code> ); that it does so at the mean-field rate $r^2/(K - K_c) \rightarrow 1/D$ , i.e. $\beta = 1/2$ ( <code>Phase8.CriticalExponent: coherent.solution.critical.exponent</code> , witnessed by <code>coherentBranch.critical.exponent</code> ); a solution is a density reproducing its own order parameter. Conditional on the von Mises stationary density, which is an input, and on no link to a trajectory                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 | Theorem (conditional)                           |
| Conditional composition                    | <code>Chain.lean</code> : <code>chain</code> takes eight named hypotheses, together with a <code>ThermodynamicCover</code> and a <code>ReflexiveBoundary</code> whose class fields carry physical content of their own, and concludes <code>UnifiedSelf</code> , meaning that the particular section glued by a thermodynamic cover is the unique fixed point of self-prediction. E89 explicitly supplies both the Lipschitz law and the identification of the glued state with a fixed point; Banach supplies uniqueness, not that identification. <code>chain.hypotheses.jointly.satisfiable</code> discharges all eight on one toy substrate and proves only that the hypotheses are simultaneously consistent, not that one mechanism realizes them. E12 remains equivalent to its conclusion, so the first two links are independent roots. Two of the eight ask for a structure tied to the objects of their premise rather than for bare existence: E34 for a predictive structure carrying the register's own temperature and Landauer heat and a non-zero waste, E78 for a cover coupled at least at $K$ whose equilibrium is the limit of a relaxation on that coupling ( <code>ThermodynamicCover.IsReachedByRelaxation</code> ); each is discharged in the joint witness and each has a regression rejecting a structure that would satisfy the bare form                                                                                                        | Conditional theorem with eight named hypotheses |
| Coarse-graining produces no kernel         | <b>Negative result</b> ( <code>Chain.lean §9</code> ). <code>ContinuousNeuralField</code> carries no conditions, so a constructor from a mesh limit would be free of content ( <code>continuousNeuralField.free</code> ). <code>mesh.refinement.convergence</code> converges a <i>scalar</i> , and <code>TriangulatedManifold.edge.region</code> is a subset of $M$ rather than of $M \times M$ , so the discrete side has no product structure to pass to a limit. The vertex-sited candidate does carry the weights ( <code>vertexKernel.apply.embedding</code> ) and still contributes exactly zero continuum coupling energy on an atomless substrate ( <code>vertexKernel.fieldCorrelation.eq.zero</code> ) — Derivation 8's hardware no-go, cutting against the framework's own coarse-graining step                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Theorem (negative)                              |

## 9 Formal soundness: why a global axiom is the wrong shape for a physical postulate

Formalization is only as trustworthy as the axioms it rests on, and an axiom that is inconsistent silently makes every theorem downstream of it vacuous while leaving the build green. The failure mode below is easy to reproduce: a declaration can constrain a symbol it does not itself bind, making the resulting class inconsistent.

The failure has one shape, and naming it is more useful than cataloguing its instances: *an axiom that constrains a symbol it does not itself bind*. Such a declaration is not a postulate about one system. It is a constraint on every model of the class the symbol belongs to, including models built to violate it, and a class that has such a model is thereby inconsistent. Three concrete cases, each of which proves **False** on its own:

1. `spontaneous_symmetry_breaking_pointwise_min` took the potential-energy functional as an *implicit* argument occurring only in its hypothesis `PotentialEnergy phi = V v0`. Instantiating that functional with the constant map  $\lambda \_. V(v_0)$  discharges the hypothesis by `rfl` and yields the conclusion for an arbitrary field. With  $V(v) = v^2$ ,  $v_0 = 0$  and  $\phi \equiv 1$  this proves  $1 = 0$ .
2. `landauer_heat_eq` equated the free class field `heat_dissipation` with a bath-entropy difference *for every instance* of the class. An instance with zero heat dissipation and a bath that gains one bit gives  $0 = \log 2$ .
3. `kl_bound_axiom` asserted  $\sigma \geq D_{KL}(P\|Q)/\Delta t$  for *all*  $P$  and  $Q$ , while  $\sigma$  is a single real fixed by a class field. Since  $D_{KL}$  is unbounded above, a point mass against a uniform distribution refutes it.

Two further declarations of the same shape are worth recording because their refutation is different in kind. `phase_invariant_periodic` and `sync_to_section_eq`, as global assertions pinning free fields of `LocalSectionSynchronization`, are refutable as soon as the probability presheaf admits two distinct global sections; the refutation needs no arithmetic at all, only that the sheaf is not trivial. Neither is a postulate in the current development: the periodicity condition is an obligation of the constructor that builds a cover from an invariant measure, and `sync_to_section_eq` is a theorem about the global section the sheaf condition produces.

The rule that prevents it is uniform: *a physical postulate that mentions a class field must be a field of that class, never a standalone axiom quantified over all instances*. A postulate is then an obligation each model discharges, not a global claim about every model. Under that rule the development declares no axioms at all.

## 9.1 What the witnesses establish, and what they do not

Explicit witnesses establish that the class assumptions can be satisfied. `Examples.lean` therefore constructs witnesses for every class that carries a physical postulate: a one-bit erasure system with a genuine bath that discharges `heat_eq` and reproduces  $\Delta Q = k_B T \ln 2$ ; a perfectly resonant system discharging `kl_bound`; and a scalar field on a one-point spacetime discharging the action principles. The gluing structures are the most demanding case and are witnessed by a `ThermodynamicCover` over a three-site substrate carrying the discrete topology and its Borel  $\sigma$ -algebra, covered by two overlapping open patches whose intersection is a single shared site.

What the witnesses establish is that the classes are inhabitable, so the theorems about them are not vacuous. What they do not establish is that a cortex is such an instance. In the cover above the phase field is already constant — every patch locked at  $\theta = 0$  — which is precisely what makes the equilibrium obligation dischargeable; getting a physical system into that configuration is the content of the informal dynamical argument.

## 9.2 Joint phase and coupling plasticity control

A finite driven control co-evolves phases and symmetric nonnegative couplings at fixed total coupling. It uses  $N = 100$ ,  $D = 0.1$ ,  $dt = 0.01$ , mean row coupling 4, and three nearly equal frequency clusters at 0.5, 1 and 1.5 per unit time. Couplings update every 50 fast steps with learning rate 0.2 for 400 time units, across three seeds. The symmetric-edge derivative of  $\frac{1}{2} \sum_i v_i^2$  is  $(v_i - v_j) \sin(\theta_j - \theta_i)$ ; each update clips negative weights, hollows the diagonal and restores the total resource. This constrained update does not guarantee monotonic descent during joint noisy dynamics. Each seed runs four arms on a single phase-noise stream: the gradient step, a symmetric hollow random step rescaled to the gradient's Frobenius norm, that same gradient step composed with a random node relabelling drawn afresh at every update, and a frozen kernel. At zero learning rate all four arms reproduce one trajectory, which is the contract the shared stream is tested against.

The nonzero grand-mean frequency drives the collective coordinate, since symmetric interactions cancel in its mean drift. That cancellation fixes  $\sum_i v_i$  across phase fields, so the objective cannot fall below  $N\langle\omega\rangle^2/D = 990$  and the update can only equalize drift across nodes rather than remove it.

The frozen arm sits at 1380.6063 to 1382.3448; the gradient arm’s second-half mean is 1110.5659 to 1116.6277, removing a fraction 0.6773 to 0.6914 of the headroom between the frozen arm and that floor. The matched random arm removes  $-0.0173$  to  $-0.0103$  of it, so the descent is specific to the gradient direction and not to displacing a kernel of that step size at fixed resource.

A matched step size is not a matched deformation. Fresh isotropic steps cancel across updates, so the random arm ends at a cumulative kernel norm growth of 1.2879 to 1.3056 against the gradient arm’s 12.4094 to 14.6958, and hardly enters the channel by which clipping and rescaling concentrate a fixed resource onto fewer edges. The permuted arm is matched on that channel instead. A node relabelling is an isometry of the matrix and a bijection of the edges, so it preserves the step’s Frobenius norm and the whole multiset of its entries and changes only which edges receive them; the arm therefore takes a step of the gradient’s size and of its shape, uncorrelated with the frequency clusters by construction. It reaches a norm growth of 11.4185 to 16.3296, a range bracketing the gradient arm’s, and a second-half mean order of 0.9184 to 0.9189, which lies inside the gradient arm’s range. Its descent fraction is  $-0.3319$  to  $-0.1607$ : a step of the gradient’s magnitude delivered to the wrong edges raises the objective. The comparison the two controls jointly support is therefore between arms of equal step size, equal cumulative deformation and equal coherence.

The learning rate is selected by a sweep from 0.001 to 0.4 on one seed, reported in the run’s `REPORT.md`: the objective’s second-half mean falls from 1377.9075 at the smallest rate to a minimum at 0.2 and rises again to 1248.2272 at the largest. Rates past the minimum lose descent while still meeting the plateau criterion, so the criterion does not by itself select a rate and no convergence in the rate is claimed.

Second-half mean order ranges from 0.9167 to 0.9222, so coherence survives the descent. All seeds meet the declared plateau criterion: the last-quarter mean differs from the preceding quarter by at most 5%. The environmental template is the covariance of unit-variance latent signals shared within each frequency cluster, rescaled and hollowed to match the coupling resource; it is a diagnostic template, not an input to plasticity. Alignment is reported by two statistics unchanged under rescaling the kernel. The ratio of mean within-cluster to mean between-cluster coupling starts at 0.9824 to 0.9948 and ends at 0.3507 to 0.6135, against 0.9965 to 0.9997 for the random arm and 0.7290 to 1.2935 for the permuted one. The permuted range is wide and contains the no-preference line. Concentrating a fixed resource onto few edges makes a ratio of block means fluctuate, and with three clusters two thirds of the node pairs are between-cluster, so a dominant edge lands between more often than within and the typical value sits a little under one. The three permuted values nonetheless lie above the gradient arm’s three without overlap, at kernel deformations that overlap almost exactly. Against 2000 node relabellings drawn per run, the final kernel is further from the true template than from a fraction 0.9510 to 1.0000 of them. The frozen arm’s initial kernel already sits at 0.6435 to 0.9110, so the gradient arm is read against the frozen arm and not against the null median; the permuted arm, at the gradient arm’s deformation, reaches 0.2680 to 0.8410.

A second partition of the same nodes settles what the loss of within-cluster coupling is a loss against. The frequency partition is the structure the environment carries; an interleaved partition of the same three group sizes, built so that each group draws equally from all three clusters, carries no frequency information, and a kernel with only frequency structure is neutral under it up to a correction of order  $1/N$ . Neither partition enters the update. The gradient arm’s interleaved ratio runs from 0.8071 to 3.0223 across the three seeds, on both sides of the no-preference line, where its frequency ratio is 0.3507 to 0.6135 on all three. The interleaved partition is one member of a family, so the calibrated form of the same reading is to score the frequency partition against 2000 relabellings of itself, each keeping the group sizes and losing the correspondence with frequency: at most a fraction 0.0495 of them give a lower ratio than the frequency partition does, on every seed, against 0.0755 to 0.3480 for the frozen arm’s untouched kernel. Concentrating a fixed resource onto few edges makes any such ratio fluctuate widely, which is why the interleaved value alone is uninformative and the percentile against the family is not.

Two readouts are reported and not relied on. The initial-minus-final Frobenius distance to the template lies in  $[-58.0962, -48.1760]$ , but the kernel norm grows by a factor 12.4094 to 14.6958 over the same interval, and a fixed total resource constrains the sum of the entries rather than their norm; most of the distance growth is therefore concentration of mass onto fewer edges rather than loss of alignment, and the distance to the shuffled template grows with it. The off-diagonal correlation with the template ends at  $-0.0362$  to  $-0.0208$ , close to its starting value: at the sparsity the descent produces, edge-level variance dominates a block-mean contrast that the within-over-between ratio resolves.

Minimizing squared drift under a conserved mean drift equalizes drift across nodes, which rewards coupling between clusters of unlike frequency, whereas the template rewards coupling within clusters of like frequency. The disagreement is a property of this pairing of objective and template. The blind-

partition reading is what separates that from the weaker statement it would otherwise be confused with: the descent does not degrade whatever partition a reader scores it against, but locates the one partition the environment carries and moves against it. Which structure a different objective would move against, or whether any objective of this family aligns with an environment, is not settled by one pairing and is not claimed. The squared-drift objective includes the drive and fluctuating relative drift; it is not a thermodynamic entropy-production estimate, which would require probability-current information. The four panels of Figure 3 in the main article are these trajectories: the dissipation function, smoothed over 21 saved samples, against the  $N\langle\omega\rangle^2/D$  floor; the order parameter against the  $1/\sqrt{N}$  reference; the within-over-between coupling ratio against the no-preference line at unity; and the distances to the cluster template and to its node-label shuffle, which grow together as the kernel norm grows. Only the gradient arm descends the objective, and only it loses cluster alignment, the permuted arm holding the same deformation without either; these diagnostics provide no evidence of structural resonance or a thermodynamic entropy-production rate.

These finite trajectories establish neither stationarity of the adaptive process nor noiseless phase locking, and do not discharge the E45 modelling assumption or extend the fixed-phase descent theorem.

## 10 The cortical field identification

The framework identifies its posited continuum coupling kernel with the endogenous electromagnetic field of mammalian cortex. This is a claim about biology, not a consequence of any theorem above it.

### 10.1 Supporting evidence

Neurons do not only signal through synapses: transmembrane currents generate extracellular fields that shift membrane potentials of neighbouring cells, a mechanism known as ephaptic coupling. Endogenous fields of 1–5 mV/mm entrain spike timing Fröhlich & McCormick (2010); Anastassiou & Koch (2015), and stable electric fields carry information that survives representational drift in the underlying spiking population Pinotsis & Miller (2022, 2023a,b). The field is spatially continuous and superposes contributions from large populations — a medium in which spatially separated activity is combined without further mechanism. This is the integrative role the CEMI field theory assigns it McFadden (2002, 2020); Pockett (2012), physically supported by astrocytic syncytia joined by gap junctions Verkhratsky & Nedergaard (2018).

These observations support a modulatory interaction, but do not determine its integrated strength relative to phase diffusion. Continuity does not itself establish compatible local states or conscious unity. The model requires independent calibration of the coupling kernel, its spatial aggregation and its rate scale.

### 10.2 Why the rate calibration is not derivable

The conversion  $K_{\text{eff}} = \gamma N s E f$  carries its whole physical content in  $\gamma$ , and no rearrangement of the remaining factors supplies it. The frustration control above records the dimensional half of this:  $s$  is a time shift per unit field, so  $sE$  is a time, and multiplying by an oscillation frequency leaves a pure number — the fractional phase advance a cell receives per cycle from the population field. Multiplying by a population size leaves it a pure number. A coupling in the sense of the main text’s stochastic model is an inverse time, and  $\gamma$  is the entire gap between the two.

The substantive half is that  $\gamma$  is a hypothesis about persistence, not a unit conversion to be fixed by convention. It states how a per-cycle phase advance accumulates into a drift rate, which depends on how long a field-driven timing shift survives. Three candidates are available and they are not interchangeable:  $\gamma = f$  takes each cycle’s advance to be retained in full and independently;  $\gamma = 1/\tau$  for a phase relaxation time takes it to decay as the medium relaxes;  $\gamma = D$  takes it to accumulate at the rate phase itself diffuses. These need not agree, in cortex or anywhere else, and the field measurement chooses between them only if the persistence of induced timing shifts is measured alongside it.

The third candidate must be declined for a further reason. Under  $\gamma = D$  the threshold test  $K_{\text{eff}}/D > 2$  reduces to  $N s E f > 2$ , a comparison of the dimensionless field estimate against the number 2 with the diffusion parameter cancelled out of a test whose content was the ratio of coupling to diffusion. The economy of that substitution is what makes it unsafe: it converts an unmeasured rate into an apparent confirmation while removing the quantity the confirmation was about. The same objection applies to any  $\gamma$  recovered by requiring the model to cross its own threshold, or read off the Bessel collapse used

to test the stationary density. Identification requires coupling and diffusion estimated in one cortical regime by a route independent of the field arithmetic, which is item 1 of the main text’s protocol and the reason the saved field-strength values are reported as conditional throughout.

### 10.3 Finite-sheet spatial-decay control

A finite stochastic control asks whether fixing total coupling while shortening its spatial reach fragments the phase field. On a periodic  $128 \times 128$  sheet spanning 2.0 mm, an exponential kernel was row-normalised to total coupling  $K_0 = 8$  at every decay length; the dynamics used thermodynamic diffusion  $0.5 \text{ rad}^2/\text{s}$ , quenched Gaussian frequency spread  $1.5 \text{ rad/s}$ , and 20,000 Euler–Maruyama steps per length. With steady  $r > 0.2$  as the declared operational coherence criterion, the smallest interpolated decay length above which every larger sampled length remained coherent was 0.0140 mm. On this grid that is 0.90 lattice spacings, and the shortest sampled length from which coherence is sustained is 0.94 spacings, so the sheet places the boundary below its own resolution. Rerunning the transition band and the lengths above it on a  $256 \times 256$  sheet of the same extent — half the spacing, so a given physical length spans twice as many cells — leaves the boundary at 0.01413 mm, 1.01 times its coarse value and 1.81 spacings on the finer grid. The boundary is fixed in millimetres rather than in cells: it is a property of the dynamics at this coupling, diffusion and frequency spread rather than of the discretisation, and the finer sheet resolves it. The transition-local trajectories are metastable and non-monotone on both sheets, so that number is an operational boundary rather than a critical-point estimate with statistical precision. At the empirical controls 0.1, 0.2 and 0.3 mm, steady  $r$  was respectively 0.9424, 0.9427 and 0.9427, and winding-defect density was at most  $2.4 \times 10^{-6}$ . From 0.1 mm out to the sweep’s longest range 8.0 mm, which is 4.0 times the sheet’s extent, steady order spans only 0.00038, so the empirical band sits inside a saturated plateau rather than near a boundary. Thus this finite, single-seed parameter sweep does not show fragmentation in the empirical decay-length band. It neither measures cortex nor proves a transition, and it leaves seed dependence and the near-boundary metastability unresolved.




Figure S3: Finite-sheet heuristic evidence under a row-normalised spatial kernel. The left panel shows global order traces and the  $1/\sqrt{N}$  finite-size floor; the right compares steady order and wrapped plaquette winding-defect density across physical decay lengths. The shaded band is 0.1–0.3 mm. This sampled simulation establishes no general transition or biological identification.

### 10.4 Critics

The longest-standing objection is that the fields are too weak to matter, pressed from within the field-theoretic camp by Pockett herself Pockett (2012, 2002). Vöröslakos and colleagues found that applied current is shunted through scalp, delivering fields at or below 1 V/m in cortex Vöröslakos et al. (2018). That is a bound on *applied* rather than endogenous fields, but it constrains how much causal work a field of that magnitude can do, and is the reason we do not claim the field drives spiking.

## 10.5 Falsification conditions

The physical identification would be challenged if independently calibrated field-mediated coupling were insufficient for the proposed regime, or if a selective perturbation of that interaction left the proposed neural and behavioral endpoints unchanged. Such a perturbation requires controls for direct effects on transmembrane currents and other coupling pathways; selectivity is not assumed experimentally available. Recovery tests must distinguish the coupling-gated account from other slow mechanisms under the same observation model. Coherence alone is not equated with report or experience, so dissociations must be evaluated against independently specified endpoints and the additional compatibility and reflexivity hypotheses.

## 11 Relationship to other theories

Integrated information and global-workspace theories supply different descriptions of integration and access Tononi et al. (2016); Baars (1997); Dehaene et al. (2014). The present model specifies a candidate physical coupling and a consistency condition on local states. Gluing does not derive an exclusion principle, and continuity alone does not force compatible local descriptions. A digital architecture is not excluded by the sheaf construction. These mathematical differences do not settle which systems have conscious experience.

## 12 The sleep-inertia prediction: assumptions, alternatives, and falsifiability

### 12.1 Assumptions

The effective coupling must depend on extracellular geometry through a separately calibrated relation. The stationary ansatz assumes identical frequencies in a rotating frame and positive phase diffusion. The square-root temporal onset additionally requires a nonzero crossing speed and adiabatic branch tracking, which critical slowing down limits near threshold. The theorem does not prove dynamical selection. Finite-size fluctuations require an operational escape criterion above the fluctuation floor; neither an exactly flat observable foot nor an infinite initial slope is guaranteed.

### 12.2 Finite-ramp numerical check

A finite- $N$  Euler–Maruyama ensemble quantifies the range over which the adiabatic reading is accurate. With identical frequencies,  $D = 1$ ,  $N = 2000$ , 32 seeded replicas and  $\dot{K} = v \in \{10^{-1}, 10^{-2}, 10^{-3}, 10^{-4}\}$ , escape is defined as the first of three consecutive sampled points at  $r \geq 0.2$  after  $K_c$ . The fastest leg is right-censored: only 5/32 replicas escape before  $K = 2.5$ . All replicas escape at the three slower speeds, with mean delays  $\Delta K = 0.1947, 0.0684$  and  $0.0253$ . Their log–log fit gives  $\Delta K \propto v^{0.443}$ , near the predicted exponent  $1/2$  at this resolution. At fixed  $v = 10^{-2}$  the delay is monotonic in  $N$ , increasing from 0.0619 through 0.1947 to 0.2341 as  $N$  increases from 500 through 2000 to 8000. The corresponding ratios  $\Delta K/\sqrt{\ln N}$  are 0.0248, 0.0706 and 0.0781, which do not collapse to a size-independent constant. At  $N = 500$ , individual replicas reach  $r = 0.436$  before  $K_c$ , while at  $N = 8000$  they reach only 0.106; the absolute criterion  $r \geq 0.2$  therefore fires almost immediately at the smallest size. The control demonstrates the protocol hazard warned of in the main text’s measurement protocol: a fixed absolute criterion meeting a size-dependent critical fluctuation floor.

The early-foot fit  $r = A(K - K_0)^{\beta_{\text{eff}}}$  is declared over the order window  $0.1 \leq r \leq 0.4$ . Its exponents are read against the same estimator applied to the exact stationary branch on each leg’s own coupling grid, which returns  $\beta_{\text{eff}} = 0.443, 0.443, 0.443, 0.448$  from fastest to slowest: the window reaches far enough past threshold that the estimator falls short of  $1/2$  even on a branch whose exponent is exactly  $1/2$ , and that shortfall belongs to the window rather than to the drive. Against that reference the drive still dominates at  $v = 10^{-2}$ , where  $\beta_{\text{eff}} = 0.898$  stands well above the adiabatic value; the two slower legs return 0.471 and 0.417, which fall on either side of their references rather than approaching them from one direction. The legs realise different windows and are comparable only so far as those agree:  $K - K_c \in [0.04, 0.31]$ ,  $[0.01, 0.19]$  and  $[0.00, 0.20]$  over 28, 19 and 200 samples. At the fastest ramp  $v = 10^{-1}$  the ensemble mean never reaches the declared window within the simulated interval: the fit truncates to  $r \in [0.100, 0.151]$  at  $K - K_c \in [0.32, 0.50]$  over 19 samples, so  $\beta_{\text{eff}} = 0.964$  is the local

log-slope of a strongly delayed trajectory rather than an onset exponent. The shift  $K_0$  is returned at its lower bound on 4 of the 4 legs, so pinning describes the estimator on this family of traces and does not single out the fastest fit. A mapping from these dimensionless rates to human recovery times requires an independently calibrated phase diffusion and coupling trajectory. The independently estimated  $(a, r)$  traces remain close to  $I_1(a)/I_0(a)$  even where the trajectory lags the stationary branch. Post-critical RMS deviations are 0.00245, 0.00599, 0.00689, 0.00705; using the slowest run as the finite- $N$  floor  $\text{Dev}_0$ , none reaches  $2\text{Dev}_0 = 0.01410$ . The specified critical speed  $v^*$  is therefore not bracketed by this sweep. This numerical check exhibits finite sampled dynamics only: it proves neither stability nor dynamical selection and does not discharge the adiabatic assumption.

### 12.3 Alternative models and failure conditions

Exponential relaxation is one comparator, not the entire class of electrical models. Slow latent variables, observation filtering and finite-size effects can also generate delayed or sigmoidal traces. Comparisons should use common preprocessing and observation models, with onset criteria specified before inspecting outcomes. A linear apparent onset counts against the adiabatic branch interpretation only in a regime where its assumptions hold.

The concentration-coherence relation tests the stationary-density model. Its failure over an informative concentration range constrains that model; its success does not identify the coupling medium. The ratio  $a/r$  is a consequence of the same density ansatz and is unreliable near the fluctuation floor. A coupling-geometry association requires an independently validated geometry proxy and an independently defined recovery endpoint. Stationary absence of a fold does not rule out finite-rate hysteresis.

## 13 Interpretive scope

The formal endpoint is a globally consistent section fixed by a specified reflexive map. Its identification with phenomenal unity is a philosophical commitment. Neither the mathematical constructions nor the simulations establish which biological or artificial systems have experience. The experimentally assessable content concerns coupling, phase organization, and recovery dynamics.

## 14 What a measurement of overlap compatibility would require

The commitment of the main text's final interpretive section turns on a property for which this article supplies no observable. Coherence has one: the resultant length of a phase distribution. Compatibility does not, and the dissociation tests that would bear on the commitment cannot be run until it does. This section states what such an observable must satisfy. It is a specification and not a result: nothing here is implemented, and the construction is left to separate work.

### 14.1 Three components

*A restriction map.* Two regions with overlapping receptive territories each carry a distribution over what occupies the shared part, decoded from the recorded activity of that region alone. Compatibility is a relation between the two restrictions to the shared territory, so the map taking a region's decoded distribution to a distribution over the overlap has to be specified before any statistic is computed. On a discrete substrate this is the marginalization the formalization already uses: two patch sections have the same restriction to their overlap exactly when they assign the same mass to every site they share (`restrict_eq_iff_density0n_eq0n`).

*A statistic with a null.* A distance between the two restricted distributions is not yet a measurement, because decoding error alone makes that distance positive. The statistic needs the distribution of the distance under compatibility — what two decoders of the same underlying content produce given their own noise — so that incompatibility is a departure from a stated null and not a nonzero number. Overlapping territories supply dependent samples, and the uncertainty has to be computed accordingly.

*Validation where the answer is built rather than inferred.* The estimator must be exercised on data in which compatible and incompatible overlaps are constructed by design, so that its sensitivity is known before it is applied where the answer is not. The formalization supplies the configuration to target: local states at a single common phase whose densities disagree where the patches meet, which satisfies every other requirement of the gluing theorem and fails compatibility alone (`overlap_agreement_fails`). An

estimator that cannot separate that configuration from its compatible counterpart at realistic decoding quality is not yet a measurement of the property the framework needs.

## 14.2 Three ways such a measurement fails silently

A decoder that shrinks both regional posteriors toward a shared prior reports compatibility whatever the regions encode, having replaced two local descriptions with one global one. A decoder carrying independent per-region bias reports incompatibility whatever the regions encode. These are opposite failures of the same component, neither is visible in the compatibility statistic itself, and both are why the constructed-answer validation above is a requirement rather than a refinement.

The third is circularity. A compatibility statistic computed from phases measures phase agreement, and the distinction between phase agreement and overlap compatibility is the entire content of the hypothesis labelled E78; the configuration named above exists to show that the two come apart. An observable that read compatibility off synchrony would confirm the framework’s central distinction by assuming its negation.

## 14.3 Scope

Such an observable would not establish E78, which asserts that coherence drives a cover the theorem is given to its relaxed equilibrium, nor the cover’s own compatibility field, which is assumed alongside the eight edges rather than by one of them. It would make the consequent measurable and leave the implication a hypothesis. It bears on neither the coupling-gated recovery prediction nor the identification of the glued section with experience. What it supplies is the missing term in the dissociation test: a state in which local descriptions are demonstrably compatible and the field supercritical while the subject reports no unified experience is not at present a describable observation, and with such an observable it becomes one. A null outcome — no reachable data separating compatible from incompatible overlaps at achievable decoding quality — is itself informative, and would place the interpretive commitment outside experimental reach rather than merely outside this article.

# 15 Empirical test of the (a, r) collapse

This section reports the methods and results of the empirical test of the parameter-free collapse prediction  $r = I_1(a)/I_0(a)$  described in the main text. The analysis was performed on a publicly available dataset of human scalp EEG recorded during propofol sedation (OpenNeuro ds005620 Bajwa et al. (2025)), comprising 21 subjects with 65-channel recordings at 5 kHz.

## 15.1 Preprocessing

For each recording block, 20–30 seconds of data were downloaded from the public S3 bucket. An additional 3 seconds of padding on each side was downloaded and trimmed after filtering to suppress edge artifacts. Each channel was detrended and bandpass filtered with a 4th-order Butterworth filter in the 4–40 Hz band using second-order sections (`sosfiltfilt`). The instantaneous phase  $\theta_j(t)$  was obtained from the Hilbert transform. Only the 62 scalp EEG channels were used; EOG and EMG channels were discarded.

## 15.2 (a, r) estimation

Each recording was divided into non-overlapping 100 ms bins (500 samples at 5 kHz). The estimand is the pooled spatiotemporal phase distribution over 62 sites and 100 ms, not an instantaneous 62-site spatial distribution. Pooling is required because the independent log-density estimator requires at least 100 phase samples. For each bin  $t$ :

- **Order parameter  $r$ :**  $r_b = \left| \frac{1}{NM} \sum_{j=1}^N \sum_{m=1}^M e^{i\theta_j(t_{bm})} \right|$ , where  $N = 62$  channels and  $M = 500$  samples per bin. This is the resultant length of the pooled sample, not an average of instantaneous spatial resultant lengths.
- **Concentration  $a$ :** The slope of the log-density of the same phase distribution regressed on  $\cos \theta$ , as specified in the main text. Phases are centred on their circular mean, binned into 40 bins on

$[-\pi, \pi]$ , and log of the binned density (with a 0.5 pseudocount) is regressed on  $\cos \theta$  by *unweighted* ordinary least squares. Under the von Mises model  $\log \rho(\theta) = a \cos \theta - \log(2\pi I_0(a))$ , so the slope is  $a$ .

The estimator must not be a maximum-likelihood one. The von Mises family is a one-parameter exponential family whose sufficient statistic is  $\sum_j \cos \theta_j$  — that is,  $r$  itself — so any ML estimate of the concentration solves its score equation  $I_1(\hat{a})/I_0(\hat{a}) = r$  and returns  $\hat{a} = A^{-1}(r)$  identically. The pair  $(\hat{a}, r)$  would then lie on  $r = I_1(a)/I_0(a)$  for any input whatsoever, including phases drawn from a distribution that is not von Mises, and the collapse would be an algebraic identity rather than a test. The implemented regression is unweighted, so it uses the log-density shape rather than defining concentration by inversion of the first trigonometric moment. Other weighting rules are different estimators and are not assumed equivalent to maximum likelihood.

The pseudocount biases the estimator. Against 200 synthetic replicates at this analysis’s per-bin sample size (62 channels  $\times$  500 samples = 31,000 phases, 40 bins), the bias is  $\leq 0.007$  for  $a \leq 3$  and reaches  $-0.375$  only at  $a = 5$ , where the troughs of the distribution empty out and the pseudocount dominates them; the sampling standard deviation over the same range is 0.008–0.048. These synthetic checks use independent samples; temporal and spatial dependence in EEG can change effective sample size and uncertainty.

The discriminating power of the estimator was verified directly on distributions that are not von Mises. At the same sample size, a wrapped Cauchy ( $\rho = 0.6$ ), a top-hat arc, and a 50/50 mixture of a von Mises with a uniform give residuals  $r - I_1(\hat{a})/I_0(\hat{a})$  of  $+0.088$ ,  $-0.117$  and  $+0.075$  respectively, against a worst case of 0.0024 on genuine von Mises samples — a separation of  $31.7\times$ . A uniform distribution passes, correctly, since it is von Mises at  $a = 0$ .

### 15.3 Bootstrap confidence intervals

95% confidence intervals on the mean ( $a, r$ ) per block were obtained by bootstrap resampling of the time bins (2000 resamples, percentile method). This procedure does not preserve dependence between neighboring bins; its intervals are exploratory rather than a calibrated time-series uncertainty estimate.

### 15.4 Narrowband analysis

The pipeline was run separately on four narrow bands (theta 4–8 Hz, alpha 8–12 Hz, beta 13–30 Hz, gamma 30–40 Hz) and the broad band (4–40 Hz) to test whether the collapse is band-specific.

### 15.5 Results

Residuals are reported per 100 ms bin and then averaged, rather than computed from block-mean  $a$  and block-mean  $r$ ; because  $I_1/I_0$  is nonlinear, the latter is not the mean residual. The estimator’s own noise floor at this sample size is 0.0024. Window sensitivity gave mean residuals  $-0.0261$ ,  $-0.0079$ ,  $+0.0017$ ,  $-0.0019$  and  $+0.0001$  at 5, 10, 20, 50 and 100 ms respectively (Figure S4); the result is stable over 20–100 ms, but not at 5–10 ms.

- **Single subject (sub-1016, 7 blocks).** In the bipolar montage the mean residual  $r - I_1(a)/I_0(a)$  lies between  $-0.0025$  and  $+0.0033$  in the 4–40 Hz band, between  $-0.0046$  and  $+0.0046$  in alpha, and between  $-0.0016$  and  $+0.0043$  in beta. These lie within  $\pm 0.005$ , of the order of the estimator’s 0.0024 noise floor and up to about twice it, and more than an order of magnitude below the 0.075–0.117 residuals the same estimator returns on the non-von-Mises alternatives above. The floor is a scale for reading the residuals, not a bound they all respect. In the common-reference montage the residual is systematically negative, between  $-0.041$  and  $-0.003$  (4–40 Hz) and between  $-0.052$  and  $-0.036$  (alpha).
- **Cross-subject (8 subjects, sed run-1).** Bipolar: mean residual  $-0.0002 \pm 0.0002$  (SEM), RMSE 0.0005, max  $|r - I_1/I_0(a)| = 0.0013$ . Common-reference: mean residual  $-0.0057 \pm 0.0028$  (SEM), RMSE 0.0093, max 0.0246. Group  $a = 0.373$  (range 0.302–0.542) and  $r = 0.173$  (range 0.145–0.222).
- **Narrowband.** Theta, alpha, beta, gamma and the broad band all collapse, with mean residuals of  $-0.0018$ ,  $+0.0001$ ,  $+0.0001$ ,  $+0.0000$  and  $+0.0000$  respectively (sub-1016, sed run-1, bipolar). The collapse is not band-specific.




Figure S4: Sensitivity of the mean collapse residual to the spatiotemporal pooling window. This assesses the pooled estimator and does not turn 62-channel EEG into an instantaneous spatial sample.

The gap between the two montages is the one systematic feature of the results. The common-reference residual is negative in every block and every subject, so the fitted concentration is high relative to the resultant length it accompanies:  $r < I_1(a)/I_0(a)$ . A common-mode component that simply raised  $r$  would push the residual the other way, so that is not the mechanism. The sign is consistent with the top-hat synthetic control above, whose concentrated support gives a negative residual of  $-0.117$ ; a shared reference could impose a sharp common-mode peak on channels that remain broadly distributed around it, steepening the log-density slope near the mode, and hence the fitted  $a$ , more than it raises the pooled resultant length. We report this as a candidate mechanism whose direction is consistent with the data, not as a demonstrated one. The bipolar derivation removes the shared component, and the residual falls by roughly an order of magnitude.

## 15.6 Interpretation

The  $(a, r)$  collapse to  $I_1/I_0$  is consistent with human scalp EEG during propofol sedation, in a bipolar montage and in every frequency band tested. This is a necessary condition for the framework’s phase-distribution claim: the von Mises stationary density is compatible with the data.

The strength of that confirmation is limited by the range of  $a$  the data occupy, and we state the limit rather than leave it to a referee. Across the eight subjects  $a$  spans only 0.302–0.542. Over so short an arc the Bessel ratio is nearly straight:  $I_1(a)/I_0(a)$  departs from the naive linear form  $a/2$  by at most 0.0095 across that interval, and from  $\tanh(a/2)$  by at most 0.0030, while the per-bin residual scatter is 0.02–0.03. The data are therefore compatible with the proposed relation between  $r$  and  $a$ , but they do not discriminate  $I_1/I_0$  from a straight line, and no claim of having done so should be read into them. Discriminating the curve from its tangent needs recordings that reach  $a \gtrsim 1.5$ , where the two separate by 0.154, and a separation of 0.17 is first reached at  $a = 1.56$ ; an awakening experiment would need to demonstrate that its sampled recovery spans such an informative range.

The result does not test the empirical commitment that the endogenous electromagnetic field is the coupling medium, nor the stipulation that the global section is experience. The coupling-gated awakening prediction remains untested here; these recordings concern propofol sedation rather than recovery from sleep.

## References

- Pinotsis, D. A., & Miller, E. K. (2022). Beyond dimension reduction: Stable electric fields emerge from and allow representational drift. *NeuroImage*, 253, 119058.
- Tononi, G., Boly, M., Massimini, M., & Koch, C. (2016). Integrated information theory: from consciousness to its physical substrate. *Nature Reviews Neuroscience*, 17(7), 450–461.

- Baars, B. J. (1997). In the Theater of Consciousness: The Workspace of the Mind. *Oxford University Press*.
- Dehaene, S., Charles, L., King, J. R., & Marti, S. (2014). Toward a computational theory of conscious processing. *Current Opinion in Neurobiology*, 25, 76-84.
- Friston, K. (2010). The free-energy principle: a unified brain theory?. *Nature Reviews Neuroscience*, 11(2), 127-138.
- Pinotsis, D. A., & Miller, E. K. (2023). In vivo ephaptic coupling allows memory network formation. *Cerebral Cortex*, 33(17), 9877-9890.
- Pinotsis, D. A., Fridman, G., & Miller, E. K. (2023). Cytoelectric coupling: Electric fields sculpt neural activity and "tune" the brain's infrastructure. *Progress in Neurobiology*, 226, 102465.
- McFadden, J. (2002). The conscious electromagnetic information (cemi) field theory: the hard problem made easy?. *Journal of Consciousness Studies*, 9(8), 45-60.
- Pockett, S. (2012). The electromagnetic field theory of consciousness: a testable hypothesis about the characteristics of conscious as opposed to non-conscious fields. *Journal of Consciousness Studies*, 19(11-12), 191-223.
- McFadden, J. (2020). Integrating information in the brain's EM field: the cemi field theory of consciousness. *Neuroscience of Consciousness*, 2020(1), niaa016.
- Fröhlich, F., & McCormick, D. A. (2010). Endogenous electric fields may guide neocortical network activity. *Neuron*, 67(1), 129-143.
- Anastassiou, C. A., & Koch, C. (2015). Ephaptic coupling to endogenous electric field activity: why bother?. *Current Opinion in Neurobiology*, 31, 95-103.
- Pockett, S. (2002). Difficulties with the electromagnetic field theory of consciousness. *Journal of Consciousness Studies*, 9(4), 51-56.
- Vöröslakos, M., et al. (2018). Direct effects of transcranial electric stimulation on brain circuits in rats and humans. *Nature Communications*, 9, 483.
- Verkhratsky, A., & Nedergaard, M. (2018). Physiology of Astroglia. *Physiological Reviews*, 98(1), 239-389.
- Bastos, A. M., Donoghue, J. A., Brincat, S. L., Mahnke, M., Yanar, J., Correa, J., ... & Miller, E. K. (2021). Neural effects of propofol-induced unconsciousness and its reversal using thalamic stimulation. *eLife*, 10, e60824.
- Bajwa, I. J., Nilsen, A. S., Skukies, R., Aamodt, A., Ernst, G., Storm, J. F., & Juel, B. E. (2025). A repeated awakening study exploring the capacity of complexity measures to capture dreaming during propofol sedation. *Scientific Reports*, 15, 32012.
- Xie, L., Kang, H., Xu, Q., Chen, M. J., Liao, Y., Thiyagarajan, M., ... & Nedergaard, M. (2013). Sleep drives metabolite clearance from the adult brain. *Science*, 342(6156), 373-377.
- Syková, E., & Nicholson, C. (2008). Diffusion in brain extracellular space. *Physiological Reviews*, 88(4), 1277-1340.
- Barbour, B. (2017). Analysis of claims that the brain extracellular impedance is high and non-resistive. *Biophysical Journal*, 113(7), 1636-1638.
- Norton, J. D. (2005). Eaters of the lotus: Landauer's principle and the return of Maxwell's demon. *Studies in History and Philosophy of Science Part B: Studies in History and Philosophy of Modern Physics*, 36(2), 375-411.
- Landauer, R. (1961). Irreversibility and heat generation in the computing process. *IBM Journal of Research and Development*, 5(3), 183-191.
- Kibble, T. W. B. (1976). Topology of cosmic domains and strings. *Journal of Physics A: Mathematical and General*, 9(8), 1387-1398.

- Sakaguchi, H. (1988). Cooperative phenomena in coupled oscillator systems under external fields. *Progress of Theoretical Physics*, 79(1), 39-46.
- Still, S., Sivak, D. A., Bell, A. J., & Crooks, G. E. (2012). Thermodynamics of prediction. *Physical Review Letters*, 109(12), 120604.
- Kawai, R., Parrondo, J. M. R., & Van den Broeck, C. (2007). Dissipation: The phase-space perspective. *Physical Review Letters*, 98(8), 080602.
- Corberi, F., dello Russo, S., Messuti, G., Scarpetta, S., & Smaldone, L. (2026). Thermodynamic learning. *arXiv preprint*. arXiv:2609.04732.
- Banach, S. (1922). Sur les opérations dans les ensembles abstraits et leur application aux équations intégrales. *Fundamenta Mathematicae*, 3(1), 133-181.
- Shenker, O. R. (2000). Logic and entropy. *PhilSci-Archive preprint*. philsci-archive.pitt.edu/115.
- Rosas, F. E., Mediano, P. A. M., Luppi, A. I., Boyd, A. B., Berta, M., Jensen, H. J., Gastpar, M., & Carroll, S. M. (2026). Four faces of information in natural systems. *Nature Reviews Physics*. doi:10.1038/s42254-026-00971-4.
- Sznitman, A.-S. (1991). Topics in propagation of chaos. In *École d'Été de Probabilités de Saint-Flour XIX—1989*, Lecture Notes in Mathematics, vol. 1464, pp. 165–251. Springer. doi:10.1007/BFb0085169.
- Townsend, R. G., Solomon, S. S., Chen, S. C., Pietersen, A. N. J., Martin, P. R., Solomon, S. G., & Gong, P. (2015). Emergence of complex wave patterns in primate cerebral cortex. *The Journal of Neuroscience*, 35(11), 4657-4662. doi:10.1523/JNEUROSCI.4509-14.2015.
- Xu, Y., Long, X., Feng, J., & Gong, P. (2023). Interacting spiral wave patterns underlie complex brain dynamics and are related to cognitive processing. *Nature Human Behaviour*, 7(7), 1196-1215. doi:10.1038/s41562-023-01626-5.