

Unity requires physically enforced agreement

Sebastian Bobadilla-Suarez

Independent Researcher

`sebastian.bobadilla.s@gmail.com`

ORCID: 0000-0002-5951-0772

Abstract

Human conscious experience is unified, and we argue that unity requires physically enforced agreement: within every window of a content's time scale, some graded carrier of each region must register the graded state of every region it overlaps, above its own noise, through a coupling that pulls their contents together. The premise is a functional condition stated at the grain of graded physical dynamics, finer than the grain of an algorithm: at that grain cortex and digital chips come apart, while at the grain of an algorithm they do not. It makes a prediction that an account stated on the information regions exchange does not: when two cortical regions are bridged by an analog and by a quantized loop that carry the same information, in total and about the task, unity follows the analog loop. We prove that under noise, content read from bits responds to a change of any region by at most its error rate, a bound that is attained, so a digital machine, clocked or clockless, fails the premise whatever algorithm it runs. Measured by one protocol over the states each occupies, a noisy cortical membrane passes near its threshold and reaches such a state within a percept's time scale, while the value a clocked latch reads passes at none of its states. Enforced agreement must arrive before the content changes,

which bounds how slowly cortical coupling may relax, and on a noisy sheet of coupled oscillators the coupling's reach decides how far phase agreement holds. We state the preparation that runs the bridge and the cortical outcomes that would refute the premise.

1 Introduction

A conscious scene is experienced as one scene. Many parts of the brain contribute to it, each describing part of the same situation, and their descriptions agree. This paper asks what kind of agreement that has to be.

Many parts can agree in two ways. In the first, the parts are tied together physically, like weights joined by springs. Each part's state is a graded physical quantity that pulls on its neighbours continuously, so no part can move, by however little, without pulling its neighbours after it. The parts are then coordinates of a single physical state, and they cannot vary independently at any scale. We call this agreement *physically enforced*. In the second, each part keeps its own record, like separate ledgers kept consistent by a clerk. A rule compares the records at discrete steps and corrects them where they differ. Each record is held in place by a margin, and below that margin the records vary independently of each other and of everything beneath them. We call this agreement *computed*. The first is one system in the sense physics uses, a single state space with coupled coordinates. The second is many systems kept consistent by a protocol. A clocked digital chip is of the second kind: its bits are separate records, and whatever agreement its program produces is computed. Our premise is that the unity of an experience is oneness of the first kind.

Premise P (physical unity). A system's contents are unified only if each region's contents are read from graded carriers and, within every window of the contents' own time scale, some carrier of each region depends on the

graded physical state of every region it overlaps, above the noise floor and through a coupling that contracts their disagreement.

P is a necessary condition for unity. It is itself a functional condition, since it picks out unity by what carriers do to one another, and it fixes that role at the grain of graded physical dynamics instead of at the grain of an algorithm. Functionalism leaves that grain open. At a coarse enough grain a nation whose citizens exchange messages by radio realizes a brain's program (Block, 1978), and the biological basis of cognition has been argued to tie mental states to a finer one (Godfrey-Smith, 2016). P and computational functionalism about unity, on which only the computation matters, are therefore two grains of one kind of claim, and they classify cortex alike. They part where an algorithm is carried by separate records: a perfect simulation of a coupled field reproduces the field's content-level dynamics and is unified at the algorithm's grain, and it fails P, because its agreement is computed. At P's grain cortex and mainstream digital chips come apart; at the algorithm's they do not. The bridge of Section 6 holds fixed the information two regions exchange and varies the grain of their coupling, so it asks whether unity tracks P's grain or the information exchanged.

P rests on two premises, argued in Section 3. The first is an account of unity: unified contents are parts of one total experience, so they need one physical system to carry them. The second is how physics counts systems: parts are one system when they interact at every scale between their own noise and their own thresholds. Everything else in the argument is a theorem. Content read through a margin responds to a region by at most its error rate, a rule that computes agreement among such records only adds another record and joins none of them into one system, and a graded reader joins the regions it reads.

Physical systems are noisy, so P is read on the probability laws of contents. Graded dependence then counts above the noise floor: every sub-threshold change larger than the noise of what a carrier depends on must move the

carrier’s next content by at least as much as that content fluctuates. P asks this of some carrier of each region within each window, and that is what one total state amounts to for regions made of noisy parts: two regions are one system when, within the window, some carrier of each registers the other above its noise, so that no change of either region above its noise is sealed off from the other. A region is many carriers, and a content is held over its own time scale, not at an instant. A noisy carrier registers its inputs near its threshold and not between spikes, so asking every carrier at every instant would split a single cortical neuron from its own inputs for most of each interspike interval and count no network of spiking cells as one system; the window form is the condition on the region and its time scale that the strong form states for a single carrier at a single step. P is neutral on where unified contents live: it constrains whichever regions carry them, whether they span the cortex or are confined to a posterior zone of it (Koch et al., 2016).

The paper makes three claims:

1. When two cortical regions are bridged by an analog and by a quantized loop that carry the same information, in total and about the task, P predicts that unity follows the loop that couples them gradedly, which tests whether unity tracks the graded grain or the information the regions exchange; a weak, sub-threshold perturbation of one region should also shift the graded content decoded in an overlapping region toward agreement (Section 6).
2. Under noise, content read from bits responds to a change of any region below its margin by at most its error rate, a bound that is attained, so content read from bits is not physically enforced, whatever algorithm the bits run. A region held by a margin drops out of every chain of couplings, corrections it writes back included, while a graded reader joins the regions it reads. Measured by one protocol over the states each occupies, a cortical membrane passes the noise-floor test near its threshold and visits such states within every window of a percept’s time

scale, while the value a clocked latch reads passes at none of its states, although the chip’s nodes cross their thresholds as often (Section 3).

3. Enforced agreement must arrive before the content changes, so P predicts that cortical coupling relaxes within 166 ms across the cortex; on a noisy coupled sheet, disagreement in phase is effective resistance, so the coupling’s reach decides how far phase agreement holds (Section 4).

Every formal result in the paper is machine-checked in Lean 4, and Appendix E maps each one to its declaration. The two premises are the steps that are argued rather than proved, and the bridge tests them together.

2 Unity as agreement over a cover of contents

Think of an atlas. Each page maps one patch of ground, neighbouring pages overlap, and where they overlap they show the same roads. Pages that agree on every overlap fit together into one map of the whole territory. We treat the brain’s regions the same way: each describes part of a shared situation, and unity is the property that these local descriptions agree wherever they overlap. When they do, they determine a single global state.

Synchrony is not unity. Two regions can oscillate in phase while one encodes a stimulus at one location and the other at another, so equal phases coexist with incompatible descriptions. Unity is a property of what overlapping regions describe, and synchrony is at most a route to it.

A region’s *content* is a variable that the system itself reads, not merely one an outside observer can decode. Two conditions define it: the regions this region drives respond to it, and it determines the quantity the region represents, such as a location or an orientation. A variable that meets both is a *carrier* of content. A variable that only an observer can decode is a trace. The supply current of a digital circuit carries enough of its logical state that a secret key can be recovered from it (Kocher et al., 1999), yet no gate of the

circuit reads the current, so the current is a trace. In cortex the membrane potential and the timing of spikes are carriers, because the postsynaptic cells integrate them.

The first condition fixes when a variable is a carrier as well as which one. A reader responds to a variable at the times it takes it, so the variable carries content at those times and is a trace between them. A cortical cell integrates a presynaptic spike as it arrives, so the time of each spike is taken when it happens. A clocked latch responds to its input only inside a short aperture at the clock edge, so the input's value there is a carrier, and the course of the voltage between edges, which no reader takes, is a trace, as the supply current is. A carrier's *next content* is therefore what its readers take from it next, at the time they take it: for a membrane the time of its next spike, and for a node that feeds a latch its value at the next clock edge.

Agreement is measured between regions whose contents represent a shared variable. The regions, and the content variable of each, are fixed by this causal role and by decodability before agreement is measured, so that nobody can make the measurement trivial by choosing the regions after the fact.

A running example. A bird flies across the visual field. A visual area represents its location on a retinotopic map, and a parietal area represents the same location while planning a reach toward it. The two areas overlap in one variable, the bird's location. Each area's content is its estimate of that location, and its carriers are the membrane potentials and spike times of the cells that encode it, which the cells each area drives integrate. An estimate an experimenter decodes from the area's summed activity, and that no cell reads, is a trace. Unity, here, is the two estimates agreeing. The sections below return to the bird: Section 3 asks how that agreement is produced, and Section 5 asks it of a chip that tracks the same bird.

The definition, stated precisely. Region i carries a content s_i , a description of the variables it represents, and regions i and j overlap in the

variables both represent. Restricting a content to an overlap keeps only the shared variables. The contents are *compatible* when every two agree on their overlap, $s_i|_{i \cap j} = s_j|_{i \cap j}$, and compatible contents extend to exactly one global description whose restriction to each region is that region’s content. Unity is compatibility, and its degree is the *overlap discrepancy*, the largest distance between two contents restricted to a common overlap. Approximate agreement still fixes one description to within its tolerance: for real-valued contents in the uniform distance, if every overlap discrepancy is at most ε , averaging the contents that cover each variable, with weights that sum to one, gives a global description within ε of every region’s content, and any two such descriptions lie within ε of each other. Agreement is measured by decoding. Decoders d_A and d_B fitted independently on two regions estimate a shared variable x with errors e_A and e_B , and $|d_A - d_B| \leq |e_A| + |e_B|$, so two accurate decoders agree whatever produced their agreement. The measure is therefore the agreement in excess of what the two errors imply, against a null that pairs decodes across trials with the same stimulus and the same errors. These results are proved in a companion paper (Bobadilla-Suarez, 2026), and P adds a condition on how that agreement is produced.

The companion proposes agreement on overlaps as a correlate of unity, to be tested, and P is conditional on that identification. The two papers therefore form a programme in two stages. The companion’s test asks whether decoded agreement in excess of single-region accuracy is larger on trials the subject perceives than on trials it does not, at matched coherence. A positive result makes the bridge of Section 6 the next question: whether that agreement has to be physically enforced. A negative one removes decoded agreement as a neural measure of unity and leaves the bridge its behavioural readout, so the second stage does not wait on the first.

3 Physical enforcement versus routed agreement

The argument in outline. This section argues for P from two premises and proves the rest. The first, *wholeness*, is an account of unity: unified contents are parts of one total experience, and being a part of a whole is not a further part of it (Bayne, 2010). The second, *counting*, is how physics counts systems once the parts fix the scale: two carriers are one system when every change of one between its own noise and its nearest threshold moves the other by at least the other’s own fluctuation. Three results are then proved from the definitions below. The margin lemma and its noisy form bound the response of content read from bits by its error rate. Counted by the second premise, a region held by a margin is one system only with itself and drops out of every chain of couplings, corrections written back through it included, while a graded reader joins the regions it reads. And graded dependence above the noise floor at every overlap holds exactly when every overlapping pair is one system. Wholeness says that a unified experience needs one physical system to carry it, counting says what one system is, and the theorems say that records never make one while graded coupling does. Where cortex and a chip hold their carriers is measured, by one protocol at the end of the section.

Definitions. The argument needs five definitions. Each is stated in words first.

A system has a physical state, its *micro state* $x = (x_i)_i$, made up of the states x_i of its regions. One step of its physics carries x to $f(x)$. A *content map* π reads each region’s content off the micro state. The content level is *closed* when the next content can be computed from the current content alone, without looking at the micro state: $\pi \circ f = g \circ \pi$ for some content dynamics g . A digital computer is built so that its logical level is closed in

this sense. The next content $\pi(f(x))$ of a carrier is what its readers take from it next, as the carrier condition of Section 2 fixes.

A readout has a *noise margin* on a set of states when every state in the set can be nudged a little without changing what is read, that is, when each has a neighbourhood on which the readout is constant. A bit is the standard example: any voltage well above the threshold reads as 1.

Content *depends gradedly* on a region when every change to that region, however small, can move the next content. Formally, no neighbourhood of the region's micro state leaves the next content unchanged while the other regions are held fixed.

The *overlap discrepancy* measures how much the contents of overlapping regions disagree. It *contracts* near a state when one step of f shrinks it by a factor below one throughout a neighbourhood of that state. Asking for a whole neighbourhood matters: perturbations too small to cross any threshold are then pulled back toward agreement along with larger ones.

Agreement is *physically enforced* at a state when content depends gradedly on each region there and the overlap discrepancy contracts near it. P requires this at the states the system occupies while its contents are unified.

The margin lemma. In words: if content is read through a margin, small enough physical changes never reach it. Precisely, if the micro dynamics is continuous and carries a state into a margin's operating set, the next content depends gradedly on no region's micro state. The proof is short: continuity carries a small enough neighbourhood of each region's micro state into the neighbourhood on which the readout is constant. A thresholded continuous quantity has a margin away from its threshold, and so does a finite word of such readouts, so the lemma covers any content read from bits. The lemma is elementary by design: once the definitions are fixed the separation follows in a line, so a reader who disputes it disputes a definition, and the rest of this section argues for each.

Noise. Real physics is noisy, and noise changes the question. A gate's chance of flipping is a smooth function of its input voltage and is never exactly zero, so the *probability* of a digital machine's next content depends, however slightly, on every region. Read on exact contents, the margin lemma would then pick out a difference lying below the noise floor of both substrates. P is therefore read on probabilities. The *response* of a content to a region at scale η is the largest change that moving the region by at most η makes to the probability of any outcome of the next content, a total-variation distance. As the yardstick for a large response we take the one that a shift of one standard deviation produces in a Gaussian carrier, $2\Phi(1/2) - 1 = 0.38$: a response of at least this size moves the content by as much as the content naturally fluctuates. Graded dependence above the noise floor asks the response to reach this fluctuation for every change of the region that is larger than its own noise and smaller than the change that crosses one of its thresholds. When the region moves a carrier linearly under additive noise whose density has a single peak, Gaussian or heavy-tailed, a larger change moves the content at least as much as a smaller one, so the criterion is met exactly when a change of one noise amplitude reaches the fluctuation; noise whose law changes with the state is treated in Appendix C. The constants are a Gaussian identity and the system's own scales, and none of them is fitted to either substrate.

The noisy margin theorem. The margin lemma has a noisy form, and it is this form that carries the argument. Let each realization of the noise fix an L -Lipschitz step of the micro dynamics, one that stretches no distance by more than a factor L . Let the content be constant within radius r of the noisy successor, except on a set of realizations of probability at most ε . Then a change of one region smaller than r/L moves the probability of every outcome of the next content by at most ε . For anything computed from bits, ε is the probability that some bit of the noisy successor lands within δ of its threshold, and the bound is attained.

Why the criterion sits below the content level. At the content level a computer can imitate a coupled field exactly. A digital system running a consensus algorithm, or integrating the equations of a coupled field, has a content map whose next step shrinks the disagreement between overlapping regions just as the field does. A criterion stated on the content dynamics would put that system in the same class as the field. P is stated instead on how content depends on graded physical state, and there the two differ. The field's content moves with every graded change in the regions it overlaps. The simulation's content is read from bits held inside their margins, and moves with none of them.

This is the simulation objection. A perfect simulation of a coupled cortical field reproduces the field's content-level trajectory and fails P, because its agreement is computed from contents and not enforced by the graded state that carries them. P denies the simulation unity, and the denial needs a reason.

Why unity tracks enforcement. Unity settles how many subjects there are. A unified experience is had by one subject and a divided one by two or more, which is what the split-brain debate of Section 6 asks of a single skull. Whatever makes contents unified therefore also decides whether they belong to one system or to several. Agreement at the content level cannot settle that count. Put the records of one program in different buildings, stretch the clerk's step to a year, or let one processor visit the records in turn: the content trajectory is unchanged while the physical system that realises it goes from one chip to several machines to one record at a time. Block's nation, whose citizens run a brain's functional organisation by radio for an hour (Block, 1978), is the same program spread over a country. A criterion stated on content-level agreement gives all of these the same number of subjects, without any physical fact to consult in doing so.

Wholeness: unity is not one more content. Wholeness takes experiences to be unified when they are parts of a single total experience, the subject’s experience at that time (Bayne, 2010). Being a part of a whole is not a further part of it, so the unity of two contents is not a third content alongside them. That fixes what can realise it. A *record* here is a content held by a margin: its identity survives every change of its carrier below the margin, so it is sealed off from the graded state of everything else. A rule that compares records and corrects them reads records and writes records, and whatever it produces, a flag, a merged value or a corrected copy, is one more record. Whether records can join anything is a theorem once systems are counted as the next paragraph counts them, by chains of couplings at the noise floor. A region whose content is held by a margin, with an error rate below the fluctuation yardstick, is coupled to no region that admits a change above its noise inside the margin, whatever the regions it reads carry, so it is one system only with itself. For any set of records, records that read records included, two regions are one system exactly when a chain of couplings that avoids every record joins them: records drop out of every chain, and the regress of rules on rules adds nothing at any depth. Writing corrections back changes nothing: when one region’s next content depends on another only through a record’s symbol, a change of the other inside the margin moves it by at most the record’s error rate, so the rule couples the two no more than the record does. A graded reader stays in the chain. A region whose content moves with two regions above their noise, and moves them in turn, is coupled to both and joins them into one system when they share no coupling of their own; under linear dynamics with uniform noise such a reader joins two regions that a record reading the same two leaves apart (Appendix E). The regress stops at a reader on the chain, not at a record beside it. In the running example, a register holding the bird’s location as a number is a record, and a cell in a third area that integrates spike times from the visual and the parietal area, and projects back to both, is a reader. A synapse,

and a membrane integrating spike times near its threshold, which a cortical membrane visits within every window as measured below, are readers in this sense. A workspace that makes every record available to a common reader is the case of the rule again: what the reader receives is a set of records, each held by its own margin. Computation therefore adds contents and never produces the whole they are parts of. What can make contents parts of one whole without adding a content is the physical state that carries them: carriers that are coordinates of a single state, whose agreement is a fact about where that state sits and is written nowhere. P asks for exactly this, which is why it is stated on how carriers depend on one another and not on what any rule computes about them.

Counting: systems at the noise floor. The second premise is argued in three steps from the criterion physics already uses. Physics counts systems by interaction: parts form one system when their joint state does not factor into independent coordinates. Applied at no particular scale, the criterion makes every chip one system, since wires, a shared supply and crosstalk couple all of its voltages, and it makes Block's nation one system through its radios. Of the voltages this is correct. Unity, though, is a property of contents, a chip reads its contents from its voltages through margins, and whether two coordinates factor depends on the resolution at which they are taken. The first step fixes the lower end of that resolution. A difference smaller than a part's own noise is not a state the part holds: its own fluctuations move its next state by more than the difference does, so the difference does not survive the part's next step as a distinguishable state, and a coordinate the part does not hold is not one along which it can be independent of anything. The second step fixes the upper end. A change that crosses one of a part's thresholds changes its symbol, and what another part registers of the crossing is the symbol, the same whatever graded state produced it. That registration is a message, the exchange of records whose agreement, by the previous paragraph, is one

more record, so crossings join nothing that records do not. The third step applies the criterion in the range that remains. There physics asks for a product decomposition: two parts are separate systems at a resolution when the dynamics of each is independent of the other's state at that resolution. A record meets it outright. Every change of another region inside the record's margin leaves the record's next content where it was, up to its error rate, so at that resolution the record's dynamics is independent of the other region's whole graded state and the joint state is a product. A graded reader is the opposite case: every change of the region it reads between noise and threshold moves it by at least its own fluctuation, so the region has no coordinate at that resolution that the reader's dynamics ignores. The count draws the line between these two cases at its strict end. Two carriers are one system when every change of one between its noise and its nearest threshold moves the other by at least the other's own fluctuation, and two when some such change leaves the other unmoved, a coordinate one of them holds and the other does not register. Records and graded readers are classed alike by the count and by a product decomposition, and they are the two cases the argument turns on. Both ends of the range are properties of each part, the size of its noise and the position of its thresholds, so nobody chooses the scale to separate brains from chips.

This is the noise-floor reading of graded dependence given above, and the identity is proved: for carriers whose contents overlap, each depends gradedly above the noise floor on every carrier it overlaps exactly when every overlapping pair is one system in this sense. A region holds many carriers, and a noisy carrier registers another only while it sits near its own threshold, as measured below. Two regions are therefore one system when some carrier of each registers some carrier of the other, and P asks this within every window of the contents' time scale. Asking it of some carriers is strictly weaker than asking it of all: a state exists at which the first holds and the second fails.

The chip, counted. A chip fails the count by design. On a logic node of 1 fF at room temperature the thermal noise is 2 mV (Johnson, 1928; Nyquist, 1928), and the textbook inverter’s margin is 506 times that. After a change of one noise amplitude, the probability of a flip, which bounds the response, is below 10^{-100} (Kish, 2002), and the response reaches the fluctuation only within 0.059% of the margin’s width from the threshold, where the bit is in the middle of switching. No other bit registers a change anywhere in the margin except through that error rate. By the noisy margin theorem, once that rate is under the fluctuation and each carrier has a change above its noise inside the margin, no carrier of any region registers any other, so the weaker form fails exactly as the stronger one does, and each part is a system of its own. Its wires pass on only the changes that cross a threshold, which is the rule applied at each clock step, so between thresholds, where a working chip spends nearly all of its time, its contents are as many systems as it has bits, carried on voltages that form one. Crosstalk and supply noise below the margin are real interactions, and they move the voltages; by the noisy margin theorem they move the contents by at most the error rate, which is the quantity a working chip is designed to keep small. The count looks strange because ordinary usage counts by program, and one program uses all of the bits. Counted by physics, each bit is restored to its own rail by its own gates, and what keeps the bits in step is the rule the program applies, a protocol among separate records. The citizens of Block’s nation likewise register each other only through the messages they exchange.

Where the line falls. A criterion built to exclude chips would draw its line between substrates. This one draws it inside each, at a carrier’s distance from its own threshold, and it discards threshold crossings in cortex as it does in silicon. A spike is all-or-none, so by the count above the occurrence of a spike is a message, and a spike count latched at a clock edge is a record. Cortex passes on what the count leaves it, the range below threshold in

which a membrane’s content moves with each input above its noise, and, as measured below, it passes there only near threshold. The verdict reverses within each substrate: an analog crossbar whose column currents drive the next array directly passes, and a neural circuit read only from latched counts fails (Section 5). What separates a working chip from cortex is where each holds its carriers, a fact about the system that the criterion reads and does not set.

The membrane, counted. A spike time passes the test at the moment the membrane rises through threshold, where one unitary cortical input shifts the spike by 1.5 times its jitter (Appendix C). Measured over every state a membrane occupies, by the protocol applied to the logic node, it passes near threshold and nowhere else. We model the membrane as a leaky integrate-and-fire cell driven by white noise, with time constants of 5 to 20 ms, spanning the in vivo high-conductance state (Destexhe & Paré, 1999; Destexhe et al., 2003), and stationary rates of 0.5 to 8 Hz, spanning the sparse rates of awake cortex (Hromádka et al., 2008). Its next content is the time of its first spike within 10 or 30 ms, and every change of one to 3 noise amplitudes, in either direction, must move that content by the fluctuation. Propagating the membrane’s probability exactly on a voltage grid (Appendix D), it sits on average 2.0 to 3.5 noise amplitudes below threshold and passes at 0.014 to 2.3% of the states it occupies. A bit restored to a rail at the logic node’s margin passes at none of its states, and its largest response over them is below 10^{-100} . No binary readout passes at any margin: under Gaussian noise, its response to a change of one noise amplitude, asked in both directions as the criterion asks, is at most $\Phi(1) - 1/2 = 0.34$. A graded content exceeds the fluctuation by a wide margin: at its best state the membrane responds by at least 0.64 in every regime and at every noise share.

A membrane passes at a given instant rarely and within a window often. A stationary membrane enters a passing state within 400 ms, the integra-

tion time that precedes a conscious percept, with probability 0.085 to 0.997, including the case in which the region it depends on supplies only 1/16 of its membrane’s noise variance. If two regions are joined by 10 carriers that pass independently, some carrier passes in such a window with probability at least 0.59, and with 100 within 40 ms with probability at least 0.59; for independent carriers the probability that none passes falls geometrically with their number. At a single instant, even 1000 such carriers link two regions with probability as low as 0.025, which is why P is stated per window. Reading the spike time through a clock of 1 or 5 ms leaves the passing fraction unchanged.

Neighbouring membranes share input, and in awake cortex their potentials are highly correlated during quiet wakefulness and less so during active behaviour (Poulet & Petersen, 2008). We model this as a shared component carrying a fraction c of every membrane’s noise variance, so that two membranes correlate by c and each keeps its own law. At least one of 100 such carriers enters a passing state within 400 ms with probability at least 0.90 for every share up to $c = 0.75$, in every regime, and at least one of 1000 with probability at least 0.99 (Appendix D). What separates the membrane from the restored bit is the distance from threshold at which each carrier is held when it is read, a few noise amplitudes against 506.

The switching chip, counted. A chip’s nodes also pass through their thresholds, many times in every window, and the protocol counts them too (Appendix D). Take a 1 GHz clock and nodes that switch in 0.02 to 0.5 of its cycles, each transition lasting 20 ps. On the rising half of a transition a change of one noise amplitude moves the time of the crossing by one jitter, as it moves a spike time. If a node’s next content were that time, the node would pass there, at 0.02 to 0.5% of its states, within the membrane’s range, and within every window of 40 ms except with probability below 10^{-100} . Counted at the voltage, a chip passes per window as a membrane does. By the carrier

condition of Section 2, its next content is what its latch reads at the clock edge, and timing closure keeps every transition on a clocked path outside the latch’s setup-and-hold window (Ginosar, 2011), so the latch reads the node at its rail: the restored bit above, which passes at none of its states and in no window. The crossing a latch discards is the content a membrane delivers, because a membrane’s readers integrate its spike times as they arrive. The one place a working chip latches a transition is a synchronizer between clock domains. With a 20 ps aperture and data changing every tenth cycle (Ginosar, 2011), its first stage samples inside the aperture at most 0.2% of the time, 8×10^5 times in a window of 400 ms, and which way it resolves moves with the timing of the transition; a second stage given one cycle to settle leaves it unresolved once in more than 10^{29} years. What that second stage takes is the bit the first resolves to, and no binary readout reaches the fluctuation: at the centre of the aperture a change of one noise amplitude in the transition’s timing moves the probability of either resolution by at most $\Phi(1) - 1/2 = 0.34$, and the stage’s own noise only lowers that. The synchronizer therefore fails graded dependence at every sample, those inside its aperture included, whichever protocol carries the data across the boundary, and so does a random source that reads its bits from how a metastable stage resolves. The value a chip reads from one bit passes nowhere, within a clock domain or across its boundaries.

The yardstick. The separation holds at every yardstick from 0.01 to 0.6 (Figure 1). Across that range the logic node passes at none of its states, and the membrane enters a passing state within 400 ms with probability at least 0.037 at the highest yardstick, rising as the yardstick is lowered. The formal results take the yardstick as a parameter and hold at every value of it. The fluctuation is the value that makes *large* mean what it means for the carrier itself: a change that moves the content by as much as the content moves of its own accord.




Figure 1: One protocol at every yardstick. Left: the share of occupied states at which every admissible change moves the next content by at least the yardstick. Right: the chance of entering such a state within 400 ms. The membrane is a band over every declared regime and, on the right, every noise share. A switching node's crossing time is a band over the declared activity factors; it passes exactly at the fluctuation (dotted) and at no higher yardstick. The synchronizer's share is the upper bound $T_W f_D$, at the yardsticks one bit can reach, up to $\Phi(1) - 1/2$. The value a clocked latch reads passes at no state and in no window.

Categories. Categorical contents have margins in every substrate, brains included. The dominant percept in binocular rivalry, a perceived category, or any class picked as the largest of several values stays the same on each side of its boundary. Away from the boundary, the margin lemma therefore denies it graded dependence in cortex exactly as in silicon. So P is stated on the graded variables from which categories are read, such as a location, an orientation, a time or the relative strength of competing percepts. A category meets P through the variable that carries it, and depends gradedly on a region only where that variable does. This opens no route for a digital machine. Every variable it computes, whether a number, a vector of numbers or a category, is a function of its bits, and after any number of steps that keep its bits away from their thresholds such a function depends gradedly

on no region. The voltages beneath the bits are not carriers in the sense of Section 2. An observer may decode the machine's contents from them, but the machine itself reads them only through the bits, at that step and at every later one.

Gradual replacement. The standard objection to a substrate condition replaces a brain's neurons one at a time by digital units that compute the same input-output function (Chalmers, 1995). If behaviour and report are preserved throughout while unity fades or divides, the subject is wrong about its own experience at the intermediate stages, which the argument takes to be absurd. P answers in two parts. First, exact replacement is unavailable at the grain P concerns: by the margin lemma, a unit that reads its inputs through margins reproduces a neuron's function to the resolution of its margins and discards the graded dependence below it. Whether discarding it changes behaviour is an empirical question, and the bridge of Section 6 asks it for one connection, which makes the bridge a local, testable instance of replacement. Second, where replacement preserves behaviour and report exactly, P accepts the conclusion the argument calls absurd and gives its reason: reports are contents computed from contents, and unity is not one more content, so reports can stay the same while unity changes. Behaviour is evidence about unity in a system whose contents are carried by its own coupling, and a system built to reproduce that evidence by computation is the simulation objection again. A bridged cortex is of the first kind, so at the bridge P stakes itself on behaviour (Section 6).

What the premise rests on. Cortex passes the count within every window, as shown above, and P asks exactly this of the regions whose contents are unified. P thus rests on wholeness and counting, and every step between them is proved: wholeness asks for one physical state, and counting says at which scale physics counts one. The bridge of Section 6 tests them together.

4 When and how far a coupling enforces agreement

P asks two things of the coupling: that it pull the disagreement of unified contents back before the content changes, and that it reach every pair of regions whose contents are unified. The first gives the prediction that the restoration test of Section 6 checks, and the second says over what distances that test should find restoration.

Agreement in time. A contraction that finishes after the content has changed enforces agreement on a content already replaced, so enforced agreement has a timing condition at both ends. It arrives no sooner than the coupling’s causal cone allows: a report at one site that tracks a variable at another to within ε , for two values more than 2ε apart, is possible after T rounds of the coupling only if the second site lies in the first site’s causal past for T rounds. It arrives at the coupling’s relaxation rate: if the coupling shrinks the overlap discrepancy by a factor $\kappa < 1$ at every step, on a set of states it keeps, a discrepancy that starts at most d_0 is at most ε from step $\log(d_0/\varepsilon)/\log(1/\kappa)$ on, and one that shrinks by at most a factor ρ per step is still above ε at every step before $\log(d_0/\varepsilon)/\log(1/\rho)$. Unity at the time scale of a content requires the later of the two edges to fall before the content changes. In cortex, signal travel uses at most 31 ms of the 400 ms that precede a conscious percept, so the condition falls on the coupling’s relaxation time (Appendix B). P therefore predicts a cortical relaxation time of at most 166 ms across the cortex and 171 ms within a posterior zone, and a slower restoration counts against it.

Reach. The pairs lie up to 15 cm apart if a scene’s contents span the cortex, and about 5 cm if they are confined to a posterior zone. On a noisy sheet of coupled oscillators sharing one natural frequency, in the harmonic

approximation, the variance of the phase difference between two sites is the noise strength times their *effective resistance*: the resistance a meter would read between them if each coupling were a wire whose conductance is the coupling’s strength. Classical results on resistance then decide how far phase agreement reaches (Appendix A). Under coupling of bounded range the resistance grows at least logarithmically with distance, the Gaussian form of the absence of true long-range order in two dimensions (Mermin & Wagner, 1966), and a kernel whose tail falls off more slowly than r^{-4} bounds it at every distance (Chung & Fuchs, 1951; Kunz & Pfister, 1976). A Lipschitz encoder, whose output changes by at most L times any change in its input, carries phase agreement to content agreement: at noise strength D the root-mean-square content discrepancy between two sites is at most $L\sqrt{D R_{\text{eff}}}$. In cortex the long-range part of the coupling is synaptic, since the field of a local source falls below detection within millimetres (Section 6), so at centimetre distances the restoration test measures synaptic coupling.

5 Current artificial systems

A clocked digital chip is an assortment of separate records. Each bit is held in place by its own noise margin, and below that margin it varies independently of every other bit and of the voltages beneath it. The records agree only when the program, applied at each clock step, makes them agree. A transformer running on such hardware reads every content from bits held within their margins, so the result of Section 3 applies to it: its agreement is computed, whatever it computes. Hardware whose contents are graded physical quantities coupled without barriers, such as analog, in-memory or neuromorphic substrates, can satisfy P.

A chip tracking the bird of Section 2 holds each area’s estimate as a word of bits. The two words agree when a step of the program copies or averages one into the other, and between steps neither moves with the voltages beneath

the other, however close the bird comes to a boundary between two of its values.

Symbols and physical contents. The line P draws is between symbolic and physical content, and the line between digital and analog hardware is its most common instance. A *symbol*, in the sense the margin lemma makes precise, is a content read through a noise margin: a physical state whose identity survives every small enough change of the state that carries it. Two symbols therefore agree only when a rule applied at discrete steps makes them agree. A *physical content* is a graded quantity that the system's own dynamics reads, so its agreement with the contents it overlaps is produced by the coupling between them. Symbols are physical objects too; what sets them apart is the margin through which they are read. A bit is a symbol, and so is a spike count latched at a clock edge. A spike time, a membrane potential and the phase of a coupled oscillator are physical contents. P requires physical contents, and a system whose contents are all symbols computes its agreement, whatever algorithm it runs.

Organisation, not material. The criterion concerns how a substrate is organised and says nothing of what it is made of. An analog crossbar whose column currents carry content and settle through shared conductances depends gradedly on every input voltage. An array of coupled electronic or optical oscillators shrinks phase disagreement through a physical coupling exactly as a cortical sheet does. A neuromorphic chip passes where its contents are read from membrane potentials coupled continuously, and fails where they are read from spike counts latched at a clock edge. Biological tissue passes or fails on the same terms: a neural circuit whose contents were read only from all-or-none events latched far from threshold would fall under the margin lemma as any digital circuit does.

Analog in-memory accelerators. In analog in-memory hardware a layer’s products are currents summed on the lines of a memory array, which is graded, and are then digitized. In a 64-core phase-change chip each row’s current is converted by its own analog-to-digital converter into integers, and activation functions and the traffic between cores run on those integers (Le Gallo et al., 2023). The next layer’s content is then read from bits, so the margin lemma applies to every content that crosses a converter, and such a chip fails P between layers as a digital chip does; its graded stage lies inside each product and couples no two contents. A design passes only where coupled graded contents meet with no restoring stage between them: column currents driving the next array directly, or oscillators and membranes coupled through shared conductances, as above.

Why spikes are not bits. Spikes are all-or-none events, and what makes them graded carriers is their timing. When a membrane potential rises through threshold, the moment it crosses moves continuously with any small input, and the receiving cells integrate spike times as they arrive rather than sampling them. Spike timing therefore carries graded dependence from one region to the next. A clocked latch compares the time of a transition with the clock edge, which puts a threshold on the transition time, so the margin lemma applies with that time as the thresholded quantity. The clock alone does not separate the substrates: a spike time read through a clock still passes near threshold, as measured above. What a digital circuit adds is restoration. Each gate drives its output back to a rail, so every transition the latch reads starts and ends far from any threshold, and the clock then removes the graded timing that spikes keep.

Clockless logic. Asynchronous circuits have no clock. In a quasi-delay-insensitive design each stage signals completion to the next through a handshake, and correctness holds under every assignment of gate and wire delays apart from a few forks assumed isochronic (Martin & Nyström, 2006). Their

transition times vary gradedly with every voltage, as a spike time does. The design makes the computed values independent of those times: a value is accepted once its code is complete, whatever delays produced it. The timing of such a circuit is therefore a trace in the sense of Section 2. It moves with the physics, and the circuit’s own contents are read from completed codes, each held at a rail by restoring gates, so they meet the margin lemma as a clocked circuit’s contents do.

Leakage. Real digital hardware leaks: supply noise, crosstalk and temperature move the analog quantities behind every bit. The margin lemma says where such leakage can reach content. If a content read as a word of thresholded continuous quantities depends gradedly on any region, some bit of the next state sits exactly on its threshold, and correct operation keeps every state off that set. Digital logic does pull perturbations back, but each gate does so alone: its gain pulls a disturbed voltage back toward its own rail. That restores each bit toward its own logical value, pulls no region toward agreement with another, and is what keeps the bits inside their margins. The chip’s records stay separate even in the one place where its physics contracts.

The margin also has a width. If every bit of the next state is at least δ from its threshold, each bit is read from an L_v -Lipschitz quantity and the microdynamics is L -Lipschitz, then a change to one region smaller than $\delta/(LL_v)$ changes no bit. Leakage below that scale reaches content only through the error rate (Section 3), and leakage above it flips a bit, which is a logical error. Digital design keeps ordinary noise well below that scale: a textbook inverter’s static noise margin is 41% of its supply voltage, against an expected supply variation of 10% (Rabaey et al., 2003), so the expected variation is 0.24 of the margin. The nondeterminism of floating-point reduction order is a choice among logical trajectories and lies outside both results: it changes which content is computed, not how content depends on graded state.

The margin's width. If P turned on how wide an engineer happened to make a noise margin, that would be a reductio of P: two machines running one program would differ in unity by a design parameter. The width is not free in that way. By the noisy margin theorem a content's response to a region is at most its error rate, so a machine whose contents reached the fluctuation at the noise scale would err with probability at least 0.38 at every step. Its bits would fail to hold their values from one step to the next, its program's trajectory would stop being reproducible, and whatever agreement survived would come from the physics beneath the bits and not from the program. The line P draws therefore coincides with the physical line between digital and analog operation, and working systems sit far from it on either side: a logic node's margin is 506 noise amplitudes wide, while a cortical membrane sits on average 2.0 to 3.5 noise amplitudes below its threshold and reaches a passing state within every window of a percept's time scale.

Attractor basins. Cortex has margins of its own. An attractor network's basin is a dynamical margin: every state in it flows to the same attractor, so which attractor the network settles in survives every small enough change. This is the categorical content of Section 3 in dynamical form. The basin gives the attractor's label a margin and leaves graded the state that carries the label: the position along a continuous attractor, such as the bump of activity that encodes a direction, the membrane potentials within the basin, and the spike times on the way to the attractor. Downstream cells integrate those. The voltage beneath a bit has no such reader, since the chip reads it only through the threshold. A cortical circuit whose contents were read only as attractor labels, far from the basins' boundaries, would fall under the margin lemma as a chip does, and the weak nudge of Section 6 tests which case cortex is in.

6 Tests

What counts as evidence about experience. Every test below measures decoded agreement: agreement of the contents available to readout. If a split in readout could always be dismissed as a split in access alone, with experience still unified behind it, no observation would bear on P. So the evidence about experience itself is fixed in advance. It consists of responses that combine stimuli across the regions in question when the task requires it, and reports of a single field of experience spanning them. A dissociation counts as evidence against P if such responses and reports split while graded coupling between the regions is intact, or stay integrated after graded coupling between them is removed.

This standard binds P. The answer to gradual replacement in Section 3 lets reports stay the same while unity changes only in a system built to reproduce them by computation. A bridged cortex is not such a system: each region still reads its contents from its own coupled carriers, only the connection between two regions changes, and the quantized loop is built to match the analog loop’s information, not its effect on behaviour. In such a system behaviour is evidence about unity, so if P holds, removing graded coupling between two regions shows in the responses and reports fixed above, and a split of unity behind unchanged reports is no answer P can give to a null result.

Separating coupling from information. This test asks whether unity tracks the graded grain or the information the regions exchange. In every manipulation of cortex available so far, graded coupling and information flow move together: anaesthesia and callosal section remove both. A test of the grain must change one while holding the other fixed. Stimulating the regions from outside does not separate them. Coherence between cortical regions arises largely as a consequence of communication between them (Schneider et al., 2021), and transcranial alternating current applied to two regions

drives both from a common source: where in-phase stimulation raised their synchrony, it raised directed information flow with it (Reinhart & Nguyen, 2019). The coupling has to be supplied instead.

The setup is a bridge (Figure 2). Two regions whose natural connection is interrupted are reconnected by stimulating each in proportion to the activity recorded at the other, through one of two loops. The *analog loop*’s output moves with every graded change in its input. It is analog end to end, an analog filter and amplifier driving the stimulator with no digitizing stage, because a loop that digitizes the recording before driving the stimulator would run the quantized condition in both arms. The *quantized loop* is the same filter and amplifier with a quantizer between them, so its output is read from bits and has a margin. The two loops share one filter shape, and they differ in the quantizer and in one cutoff frequency, set below.

The two loops must carry the same information and fall on opposite sides of P. Section 3 fixes what the second condition means: a change of the recorded region by its own noise amplitude must move the analog loop’s output by at least the fluctuation, and the quantized loop’s output by less. The analog loop passes exactly when its noise, in the region’s units, is at most the region’s own. The quantized loop fails at every state farther from a quantization edge than the region’s noise plus 0.3 times the recording noise, so its step must be many times both. Matching information by adding noise to the analog loop defeats this. At matched information the added noise is 0.24 of the quantizer’s step, so an analog loop that still passes caps the step at 4.1 noise amplitudes of the region, and the quantized loop then fails at no more than 52% of its states, taking states as uniform within a step. The analog loop’s bandwidth is lowered instead, with no noise added, until the transfer entropy it carries, estimated from the recorded activity, matches the quantized loop’s. Lowering the bandwidth slows the loop, and a loop that settles after the content has changed fails P’s timing condition (Section 4) however gradedly it couples. A first-order filter with cutoff f_c settles with

time constant $1/(2\pi f_c)$, and the bridged coupling relaxes no faster than its loop, so the analog loop's cutoff is kept at or above 0.99 Hz, whose time constant equals the shortest relaxation time that brings agreement in time in Appendix B. If matching would take the cutoff lower, the quantizer's step is coarsened instead, which lowers the quantized loop's transfer entropy and raises the share of its states that fail. A null result then cannot come from an analog loop too slow for the content's window. Above that floor the step is free: at a step of 10 noise amplitudes, with recording noise no larger than the region's own, the quantized loop fails at 74% of its states or more. Stepping the quantized loop's resolution checks that its transfer entropy rises with the resolution, so that the resolution, and not recording noise, sets the matched value.

Transfer entropy counts everything one region's activity tells about the other's, and the loops can match it while differing in how much of the task variable gets through, since the quantizer and the lower cutoff distort the waveform differently. A behavioural difference between the loops would then have an explanation at the level of information. So the task variable is also decoded from the current each loop delivers, by a decoder fitted on other blocks of the same loop, and the loops count as matched only when its decodability, as well as the transfer entropy, agrees between them within the confidence interval of the difference. The step and the cutoff are two free settings for these two conditions: a step coarse enough to fail P, and a cutoff above the floor, are chosen together on pilot blocks, and the check is repeated on the blocks that are compared. The two loops then carry the same information, in total and about the task, through the same filter shape, and only the analog one couples the regions gradedly. An account on which unity follows the information regions exchange predicts no difference between them.

The readout is unity, with the evidence fixed at the start of this section. The bridge holds information flow fixed and varies the grain of the

coupling. Unity that follows the analog loop at matched transfer entropy and task decodability is evidence that unity tracks the finer grain, and not only the information the regions exchange; unity that is the same under both loops counts against P. In animals the design measures unity by integration across the bridged connection in behaviour and by decoded agreement in the recordings; reported unity needs a bridged human preparation.

A first preparation. Closed loops that bridge two cortical sites already run for days in behaving animals. Spikes recorded at one site have triggered stimulation at another in freely moving monkeys (Jackson et al., 2006) and in rats after cortical injury (Guggenmos et al., 2013). Those loops are triggered by events, which is the quantized condition; the analog arm replaces the trigger with a continuous drive. A first preparation places chronic arrays in two areas that represent a shared variable, such as the visual and parietal areas that both represent the bird’s location in the running example, interrupts the natural connection between them, and alternates the two loops across blocks in the same animal. Transfer entropy (Schreiber, 2000) is estimated from each block’s recordings, and the analog loop’s cutoff and the quantizer’s step are set on pilot blocks so that the loops match in transfer entropy and in the decodability of the task variable from their delivered currents, with the cutoff above the floor that keeps the analog loop in time. The primary readout is behavioural: a task that requires the animal to combine the values the two areas represent, such as judging whether stimuli in their two fields match, whose accuracy is compared between the loops. The companion’s excess decoded agreement between the two areas is the neural readout, and the number of trials follows from the power of the weaker of the two comparisons.

Stimulation artefact. Event-triggered loops blank the recording around each pulse, and a continuous drive leaves no interval to blank. The loops therefore record and stimulate in separate frequency bands. Each loop reads multi-unit activity, above a few hundred hertz, takes its envelope, and drives




Figure 2: The bridge, one direction of two. Multi-unit activity recorded in region A is rectified and smoothed, passed through a filter both loops share, and then either kept analog or quantized, before an output stage both loops share drives the stimulator in region B. The quantizer’s step and the analog loop’s cutoff are the two settings that match the loops in transfer entropy and in task decodability. What leaks from the stimulator into a recording contact is removed by subtracting the loop’s own output, passed through the measured transfer. The loop from B to A is the same.

the other region with a current that follows the envelope and so lies in the band below it. Rectifying and smoothing are analog operations, so the analog loop stays analog end to end. What leaks across the bands is removed by subtracting the loop’s own output, which the loop knows exactly, passed through the transfer from stimulator to recording contact measured on open-loop calibration blocks. Both loops end in the same output stage, which smooths the quantizer’s steps into the stimulation band, and share the stimulator, the bands and the subtraction, so a residual artefact reaches both arms by one path. Open-loop blocks in which each stimulator follows activity recorded in another block measure that residual in each arm, and the comparison of the loops is made after it is removed.

A weak nudge should spread. Weak stimulation of one region, below the level that changes behaviour, should shift the content decoded in an overlapping region, in proportion to the stimulation’s strength and toward agreement with the stimulated region. Transcranial alternating current and weak intracortical fields supply the nudge. The decoded content must be a graded variable, such as a location, an orientation or the strength of evidence for a percept, read from the population activity that the overlapping region’s targets integrate, so that it is a carrier in the sense of Section 2. A decoded category would not do: it is constant away from its boundaries in any substrate, and would stay unchanged whether P holds or not. If the decoded content of the overlapping region stays unchanged at strengths that measurably move the stimulated one, that counts against P , since P then denies unity to the very system whose unity motivates it. A change alone confirms little, because most accounts expect stimulating one region to alter activity in regions connected to it. P predicts the change’s form. The decoded content of the overlapping region shifts with the sign of the shift induced in the stimulated region, as agreement requires, and in proportion to stimulation strength down to the weakest strength that measurably moves the stimulated region, with no strength below which the shift vanishes. A shift of the wrong sign, or one that appears only above a threshold strength, counts against P .

Agreement should be restored. Perturb one region and measure whether its agreement with its overlap partners comes back, and how fast. In cortex the perturbation is stimulation, and the measure is the disagreement between contents decoded from the perturbed region and from its partners. P predicts restoration within the relaxation time of Section 4, with a restoring pull proportional to the size of the perturbation, below any threshold as well as above. The same protocol run on an artificial system, by patching the activations of one layer, head or position, has an outcome fixed in advance by

the margin lemma. A digital system can restore agreement by computation, above the scale at which its logical trajectory changes, and below that scale nothing moves. Observing this checks that the hardware operates inside its margins. It is also the contrast that makes the cortical outcome informative: a restoring pull that stays graded below every threshold is the signature P attributes to enforcement.

The field’s causal cone. Brain tissue carries an electric field as well as synapses. At the frequencies of neural activity that field is quasi-static (Gratiy et al., 2017): it reaches every point of the tissue it spans with no propagation delay. Synaptic influence between distant regions, by contrast, arrives only after axonal conduction and synaptic transmission, no sooner than 8 ms across the longest cortico-cortical fibres. If two regions re-agree after a local perturbation faster than the synaptic route allows, the coupling has a non-synaptic part, and the timing measures the reach of the kernel tail of Section 4 directly.

The field’s strength sets the distance at which this test can decide anything. Endogenous fields in cortex peak near 2.4 mV/mm in vivo (Fröhlich & McCormick, 2010), networks of neurons detect fields from 0.14 mV/mm (Francis et al., 2003), and the potential of a local cortical source decays as $r^{-2.1}$ across millimetres (Rebollo et al., 2021). With the peak placed at the nearest distance of that measurement, the field of a local source falls to the detection threshold within 3.7 mm, and misses it by 5 orders of magnitude at 15 cm and by 3.5 at 5 cm. The field’s cone is therefore tested at millimetre distances, where activity has been observed to synchronize across a cut through cortex by the field alone (Rebollo et al., 2021). At centimetre distances P predicts that re-agreement arrives through synaptic coupling, which is graded as well.

The timing can be measured with intracortical recording, with care. The stimulation artefact saturates clinical amplifiers during the first millise-

onds, masking the earliest monosynaptic responses (Keller et al., 2014), and early responses at distant contacts carry a volume-conducted component that scales with distance as the artefact does (Prime et al., 2020). Three measures separate active re-agreement from the passive field that carries the stimulus: template subtraction of the artefact (Trebaul et al., 2016); multi-unit activity or current-source density at the distant site, which the passive spread of the local field potential does not reach (Kajikawa & Schroeder, 2011); and a deadline set for each subject from the tractography length of the connecting fibres (Caminiti et al., 2013).

Split brain. Cutting the corpus callosum removes the main route of both information and graded coupling between the hemispheres. Patients respond to stimuli in the two visual fields as if perception were divided, yet can report the presence of stimuli in either field with either hand. This has been read as divided perception with undivided consciousness (Pinto et al., 2017). Whether the field of experience divides is contested (de Haan et al., 2020), and on a rival switch model consciousness alternates between the hemispheres (Bayne, 2010). The remaining subcortical routes are synaptic, so they couple gradedly and carry information at once. The case therefore removes coupling and information together, which is exactly the confound the bridged preparation above is built to break.

7 Related work

Indicator properties. Butlin et al. (2023) assess current AI systems against indicator properties of consciousness drawn from scientific theories and stated computationally, and they adopt computational functionalism as a working hypothesis, which states every indicator at the grain of an algorithm. P adds to any list of indicators one stated at the grain of graded dynamics, a condition on how the substrate realises them, so a system can have every

computational indicator and fail P, as a faithful simulation of cortex does (Section 3). The bridge of Section 6 holds the information two regions exchange fixed and asks whether unity tracks the graded grain.

Computational closure. The closest formal neighbour is the account of Rosas et al. (2024). On it, a macroscopic level that is informationally, causally and computationally closed marks emergent software in a physical system. This closure is the property Section 3 places at the content level of a digital system, and P reads it the other way round: a level closed off from the graded state beneath it is one whose agreement is computed, not enforced.

Designed to suppress its physics. Kleiner & Ludwig (2024) reach a related exclusion from a different premise. If the physical dynamics of a system is relevant to its consciousness, processors designed and verified to suppress deviations from their computational dynamics cannot be conscious. The margin lemma states that suppression formally, and three things are added to it here. The criterion’s scale is fixed by the parts’ own noise, so it counts systems instead of asking whether physics is relevant at all. The coupling must contract disagreement, which a substrate meets or fails separately from how it suppresses noise. And the bridge of Section 6 asks in cortex whether unity tracks the physical dynamics or the information they carry.

Biology and the medium. Seth (2026) argues that consciousness may depend on biological properties that computation does not capture. P locates the relevant property in how the substrate is organised, and admits non-biological hardware that has it (Section 5). Mortal computation, whose parameters are useful only on the hardware instance that learned them (Hinton, 2022; Ororbia & Friston, 2023), shares P’s refusal to separate the computation from the medium.

Analog and digital. On the mechanistic account of Piccinini (2008), a digital computer manipulates strings of digits, states its components distinguish reliably, and which class a system belongs to is fixed by its mechanism. P’s symbol, a content read through a noise margin, is that reliability stated as a property of the readout, with the margin measured in the parts’ own noise; P then asks what the reliability costs, namely graded dependence below the margin. Maley (2011) separates analog from continuous and digital from discrete representation, and P’s line follows the second pair, not the first: a spike time is a physical content whether or not it represents by analogy, and a continuous quantity latched through a margin is a symbol. Maudlin (1989) argues that a computation’s counterfactual structure can be carried by machinery that is idle on the actual run, which sits badly with consciousness supervening on present physical activity. P states its counterfactual condition on present activity: graded dependence asks how the actual carriers respond to changes of the actual state, above their noise, and idle machinery responds to none.

Integrated information. Integrated information theory also denies experience to digital computers running faithful simulations of a brain (Tononi & Koch, 2015; Findlay et al., 2024). It does so through the causal structure of the hardware at the grain of its logic gates. P turns instead on how content depends on graded physical state, and on whether disagreement is pulled back. Its verdict on a machine is read from a measurement that runs on both substrates: the noise-floor test of Section 3, applied to a cortical membrane and to a logic node by one protocol.

Field theories. The electromagnetic field theory of McFadden (2020) identifies consciousness with the brain’s field, supported by evidence that endogenous fields entrain cortical activity (Fröhlich & McCormick, 2010). Miller et al. (2026) place cognition and consciousness in travelling waves that synaptic and ephaptic coupling sustain together, and take consciousness to arise

when those waves bring cortex into a globally integrated state. Their analog computation is set apart from digital computation by running in parallel. P treats the field, and the waves it carries, as ways among several to realize graded coupling, and draws the line between analog and digital at the noise margin, which a parallel digital machine keeps.

8 Limitations

- *The premise.* P is a premise and is not derived from the dynamics. The formal results show what follows from it, and the cortical tests of Section 6 bear on whether it holds. It is necessary for unity and not sufficient for consciousness: coupled pendulums satisfy it. The argument for it takes the mereological account of phenomenal unity, and an account on which unity is joint access to contents rejects its first step. The count makes two carriers two systems when a single change of one between its noise and its threshold leaves the other unmoved, which is stricter than a product decomposition; the two agree on records and graded readers, and a carrier that registers some changes in that range and not others is two systems by the count and may be one by a product decomposition. A functionalism that states its roles at the grain of graded physical state agrees with P on every test here, and it is the intended reading of P: the tests ask which grain unity tracks, not whether a functional account of unity is correct.
- *Hardware near its thresholds.* The noisy margin theorem bounds a digital content's response by its error rate. Hardware run so close to its thresholds that a change of one noise amplitude moves its contents by a fluctuation would escape the bound. One bit cannot, under Gaussian noise, since no binary readout reaches the fluctuation, so such hardware would read one graded quantity through several bits held near their thresholds together, a design the switching count does not cover.

The thermal estimate declares the node capacitance it uses, and the switching count declares its clock, activity factors, transition time and synchronizer.

- *The cortical estimates.* The single-input estimate of Appendix C combines in vitro measurements of one connection type and one noise recording, and reads the spike-time shift through a linear threshold crossing. The occupancy computation models one membrane as a leaky integrate-and-fire cell under white noise, with a threshold that has no dynamics of its own. Its time constants, rates, windows, noise shares and carrier counts are declared ranges, not fits. Correlation between carriers is modelled as one shared component with the membrane's own time constant, at declared shares; the measured correlations it is compared with are those of neighbouring cells, and the single spikes of awake cortex are driven by input that neighbouring cells do not share (Poulet & Petersen, 2008), which that model omits. The carrier counts stay far below the anatomical ones, and the per-window bound is proved formally in two forms: for independent carriers, and without independence, where the chance that no carrier passes is at most the variance of the number that pass over the square of its mean, which correlation between carriers keeps above a floor however many there are.
- *Reach.* The identification of phase variance with effective resistance rests on the harmonic approximation at identical natural frequencies. The logarithmic lower bound on resistance is proved for kernels of bounded range. The bound at every distance under a tail slower than r^{-4} is proved for a sheet that supplies the site counts and couplings of a dyadic chain of discs. That a kernel of finite second moment leaves the resistance unbounded is the classical recurrence criterion, for translation-invariant kernels on the lattice; that the growth is log-

arithmetic is established numerically on finite sheets. Above the order-disorder crossover of Appendix A the harmonic identity fails and agreement is lost at every distance. The step from phase agreement to content agreement assumes one oscillator population per content patch, which is a modelling choice about cortex, and across centimetres it is the phase-to-content assumption of Appendix A.

- *Which variables are contents.* Which variables a system's own dynamics reads, and so which are its contents, is settled system by system from its circuitry. The margin lemma takes that identification as given. For attractor networks the answer takes the unified contents to be carried by the graded state within a basin; a cortex whose unified contents were basin labels alone would fail P.
- *Artificial systems.* The predictions for artificial systems are consequences of the margin lemma. Confirming them checks that hardware runs inside its margins and leaves P untested; the tests of P are the cortical ones.
- *Agreement in time.* Its edges are proved for a given contraction factor; identifying that factor with the coupling's spectral gap is linear response. The relaxation time of cortical coupling is unmeasured, so the cortical estimate states the condition for agreement to arrive in time, and whether cortex meets it is the prediction the restoration test checks. It also reads agreement as a declared reduction factor of disagreement.
- *The field.* A long-range tail improves agreement only in proportion to its conductance against the short-range part. The field of a local cortical source falls below the detection threshold within millimetres, by an estimate that places the field's peak at a declared distance. Across centimetres graded coupling is carried chiefly by synapses, and the field's contribution there is a quantity to be measured.

- *The bridge's reach.* The bridge holds information flow fixed and varies the grain of the coupling, so it separates the graded grain from the grain of information exchanged. A functional account that counts the quantizer's nonlinearity, or the cutoff that matches information, as a difference between the loops can accommodate either outcome. Such an account individuates roles finely enough to see a quantizer, which is the grain P names, so the bridge tells P's grain from the grain of information and leaves that account on P's side of the line.
- *Human evidence.* The dissociation of coupling from information has, so far, no human preparation in which reported unity can be measured across a bridged connection. Human data can falsify P but cannot by themselves tell the graded grain from the grain of information exchanged, because graded coupling and information flow co-vary in every manipulation available so far.

References

- Bayne, T. (2010). *The Unity of Consciousness*. Oxford University Press. doi:10.1093/acprof:oso/9780199215386.001.0001.
- Block, N. (1978). Troubles with functionalism. In C. W. Savage (Ed.), *Perception and Cognition: Issues in the Foundations of Psychology*, Minnesota Studies in the Philosophy of Science, 9, 261–325. University of Minnesota Press.
- Bobadilla-Suarez, S. (2026). When does cortical coherence carry shared content? Compatibility and self-reconstruction as testable conditions. Manuscript. github.com/bobaseb/coupling_gated_cortical_coherence.
- Butlin, P., Long, R., Elmoznino, E., Bengio, Y., Birch, J., Constant, A.,

- Deane, G., Fleming, S. M., Frith, C., Ji, X., Kanai, R., Klein, C., Lindsay, G., Michel, M., Mudrik, L., Peters, M. A. K., Schwitzgebel, E., Simon, J., & VanRullen, R. (2023). Consciousness in artificial intelligence: Insights from the science of consciousness. *arXiv:2308.08708*. doi:10.48550/arXiv.2308.08708.
- Caminiti, R., Ghaziri, H., Galuske, R., Hof, P. R., & Innocenti, G. M. (2009). Evolution amplified processing with temporally dispersed slow neuronal connectivity in primates. *Proceedings of the National Academy of Sciences*, 106(46), 19551–19556. doi:10.1073/pnas.0907655106.
- Caminiti, R., Carducci, F., Piervincenzi, C., Battaglia-Mayer, A., Confalone, G., Visco-Comandini, F., Pantano, P., & Innocenti, G. M. (2013). Diameter, length, speed, and conduction delay of callosal axons in macaque monkeys and humans: Comparing data from histology and magnetic resonance imaging diffusion tractography. *Journal of Neuroscience*, 33(36), 14501–14511. doi:10.1523/JNEUROSCI.0761-13.2013.
- Chalmers, D. J. (1995). Absent qualia, fading qualia, dancing qualia. In T. Metzinger (Ed.), *Conscious Experience* (pp. 309–328). Imprint Academic.
- Chung, K. L., & Fuchs, W. H. J. (1951). On the distribution of values of sums of random variables. *Memoirs of the American Mathematical Society*, 6, 1–12.
- de Haan, E. H. F., Corballis, P. M., Hillyard, S. A., Marzi, C. A., Seth, A., Lamme, V. A. F., Volz, L., Fabri, M., Schechter, E., Bayne, T., Corballis, M. C., & Pinto, Y. (2020). Split-brain: What we know now and why this is important for understanding consciousness. *Neuropsychology Review*, 30(2), 224–233. doi:10.1007/s11065-020-09439-3.
- Destexhe, A., & Paré, D. (1999). Impact of network activity on the integrative properties of neocortical pyramidal neurons in vivo. *Journal of Neurophysiology*, 81(4), 1531–1547. doi:10.1152/jn.1999.81.4.1531.

- Destexhe, A., Rudolph, M., & Paré, D. (2003). The high-conductance state of neocortical neurons in vivo. *Nature Reviews Neuroscience*, 4(9), 739–751. doi:10.1038/nrn1198.
- Doyle, P. G., & Snell, J. L. (1984). *Random Walks and Electric Networks*. Carus Mathematical Monographs 22. Mathematical Association of America.
- Findlay, G., Marshall, W., Albantakis, L., David, I., Mayner, W. G. P., Koch, C., & Tononi, G. (2024). Dissociating artificial intelligence from artificial consciousness. *arXiv:2412.04571*. doi:10.48550/arXiv.2412.04571.
- Francis, J. T., Gluckman, B. J., & Schiff, S. J. (2003). Sensitivity of neurons to weak electric fields. *Journal of Neuroscience*, 23(19), 7255–7261. doi:10.1523/JNEUROSCI.23-19-07255.2003.
- Fröhlich, J., & Spencer, T. (1981). The Kosterlitz–Thouless transition in two-dimensional Abelian spin systems and the Coulomb gas. *Communications in Mathematical Physics*, 81(4), 527–602. doi:10.1007/BF01208273.
- Fröhlich, F., & McCormick, D. A. (2010). Endogenous electric fields may guide neocortical network activity. *Neuron*, 67(1), 129–143.
- Ginosar, R. (2011). Metastability and synchronizers: A tutorial. *IEEE Design & Test of Computers*, 28(5), 23–35. doi:10.1109/MDT.2011.113.
- Godfrey-Smith, P. (2016). Mind, matter, and metabolism. *The Journal of Philosophy*, 113(10), 481–506. doi:10.5840/jphil20161131034.
- Gratiy, S. L., Hanes, G., Denman, D., Hawrylycz, M. J., Koch, C., Einevoll, G. T., & Anastassiou, C. A. (2017). From Maxwell’s equations to the theory of current-source density analysis. *European Journal of Neuroscience*, 45(8), 1013–1023. doi:10.1111/ejn.13534.

- Guggenmos, D. J., Azin, M., Barbay, S., Mahnken, J. D., Dunham, C., Mohseni, P., & Nudo, R. J. (2013). Restoration of function after brain damage using a neural prosthesis. *Proceedings of the National Academy of Sciences*, 110(52), 21177–21182. doi:10.1073/pnas.1316885110.
- Herzog, M. H., Kammer, T., & Scharnowski, F. (2016). Time slices: What is the duration of a percept? *PLoS Biology*, 14(4), e1002433. doi:10.1371/journal.pbio.1002433.
- Hinton, G. (2022). The forward-forward algorithm: Some preliminary investigations. *arXiv:2212.13345*. doi:10.48550/arXiv.2212.13345.
- Hong, H., Park, H., & Choi, M. Y. (2005). Collective synchronization in spatially extended systems of coupled oscillators with random frequencies. *Physical Review E*, 72(3), 036217. doi:10.1103/PhysRevE.72.036217.
- Hromádka, T., DeWeese, M. R., & Zador, A. M. (2008). Sparse representation of sounds in the unanesthetized auditory cortex. *PLoS Biology*, 6(1), e16. doi:10.1371/journal.pbio.0060016.
- Jackson, A., Mavoori, J., & Fetz, E. E. (2006). Long-term motor cortex plasticity induced by an electronic neural implant. *Nature*, 444(7115), 56–60. doi:10.1038/nature05226.
- Jacobson, G. A., Diba, K., Yaron-Jakoubovitch, A., Oz, Y., Koch, C., Segev, I., & Yarom, Y. (2005). Subthreshold voltage noise of rat neocortical pyramidal neurones. *The Journal of Physiology*, 564(1), 145–160. doi:10.1113/jphysiol.2004.080903.
- Johnson, J. B. (1928). Thermal agitation of electricity in conductors. *Physical Review*, 32(1), 97–109. doi:10.1103/PhysRev.32.97.
- Kajikawa, Y., & Schroeder, C. E. (2011). How local is the local field potential? *Neuron*, 72(5), 847–858. doi:10.1016/j.neuron.2011.09.029.

- Katz, B., & Miledi, R. (1965). The measurement of synaptic delay, and the time course of acetylcholine release at the neuromuscular junction. *Proceedings of the Royal Society B*, 161(985), 483–495. doi:10.1098/rspb.1965.0016.
- Keller, C. J., Honey, C. J., Mégevand, P., Entz, L., Ulbert, I., & Mehta, A. D. (2014). Mapping human brain networks with cortico-cortical evoked potentials. *Philosophical Transactions of the Royal Society B*, 369(1653), 20130528. doi:10.1098/rstb.2013.0528.
- Kish, L. B. (2002). End of Moore’s law: thermal (noise) death of integration in micro and nano electronics. *Physics Letters A*, 305(3–4), 144–149. doi:10.1016/S0375-9601(02)01365-8.
- Kleiner, J., & Ludwig, T. (2024). The case for neurons: A no-go theorem for consciousness on a chip. *Neuroscience of Consciousness*, 2024(1), niae037. doi:10.1093/nc/niae037.
- Koch, C., Massimini, M., Boly, M., & Tononi, G. (2016). Neural correlates of consciousness: Progress and problems. *Nature Reviews Neuroscience*, 17(5), 307–321. doi:10.1038/nrn.2016.22.
- Kocher, P., Jaffe, J., & Jun, B. (1999). Differential power analysis. In *Advances in Cryptology — CRYPTO ’99*, Lecture Notes in Computer Science, vol. 1666, pp. 388–397. Springer. doi:10.1007/3-540-48405-1_25.
- Kunz, H., & Pfister, C.-E. (1976). First order phase transition in the plane rotator ferromagnetic model in two dimensions. *Communications in Mathematical Physics*, 46, 245–251. doi:10.1007/BF01609121.
- Le Gallo, M., Khaddam-Aljameh, R., Stanisavljevic, M., Vasilopoulos, A., Kersting, B., Dazzi, M., Karunaratne, G., Braendli, M., Singh, A., Mueller, S. M., Buechel, J., Timoneda, X., Joshi, V., Egger, U., Garofalo, A., Petropoulos, A., Antonakopoulos, T., Brew, K., Choi, S., Ok, I., Philip, T.,

- Chan, V., Silvestre, C., Ahsan, I., Saulnier, N., Narayanan, V., Francese, P. A., Eleftheriou, E., & Sebastian, A. (2023). A 64-core mixed-signal in-memory compute chip based on phase-change memory for deep neural network inference. *Nature Electronics*, 6, 680–693. doi:10.1038/s41928-023-01010-1.
- Mainen, Z. F., & Sejnowski, T. J. (1995). Reliability of spike timing in neocortical neurons. *Science*, 268(5216), 1503–1506. doi:10.1126/science.7770778.
- Maley, C. J. (2011). Analog and digital, continuous and discrete. *Philosophical Studies*, 155(1), 117–131. doi:10.1007/s11098-010-9562-8.
- Markram, H., Lübke, J., Frotscher, M., Roth, A., & Sakmann, B. (1997). Physiology and anatomy of synaptic connections between thick tufted pyramidal neurones in the developing rat neocortex. *The Journal of Physiology*, 500(2), 409–440. doi:10.1113/jphysiol.1997.sp022031.
- Martin, A. J., & Nyström, M. (2006). Asynchronous techniques for system-on-chip design. *Proceedings of the IEEE*, 94(6), 1089–1120. doi:10.1109/JPROC.2006.875789.
- Maudlin, T. (1989). Computation and consciousness. *The Journal of Philosophy*, 86(8), 407–432. doi:10.2307/2026650.
- McFadden, J. (2020). Integrating information in the brain’s EM field: the cemi field theory of consciousness. *Neuroscience of Consciousness*, 2020(1), niaa016.
- Miller, E. K., Brincat, S. L., & Roy, J. E. (2026). Analog cognition and consciousness. *The Journal of Neuroscience*, 46(33), e0711262026. doi:10.1523/JNEUROSCI.0711-26.2026.
- Mermin, N. D., & Wagner, H. (1966). Absence of ferromagnetism or antiferromagnetism in one- or two-dimensional isotropic Heisenberg models. *Physical Review Letters*, 17(22), 1133–1136. doi:10.1103/PhysRevLett.17.1133.

- Nyquist, H. (1928). Thermal agitation of electric charge in conductors. *Physical Review*, 32(1), 110–113. doi:10.1103/PhysRev.32.110.
- Nunez, P. L., & Srinivasan, R. (2014). Neocortical dynamics due to axon propagation delays in cortico-cortical fibers: EEG traveling and standing waves with implications for top-down influences on local networks and white matter disease. *Brain Research*, 1542, 138–166. doi:10.1016/j.brainres.2013.10.036.
- Ororbia, A., & Friston, K. (2023). Mortal computation: A foundation for biomimetic intelligence. *arXiv:2311.09589*. doi:10.48550/arXiv.2311.09589.
- Piccinini, G. (2008). Computers. *Pacific Philosophical Quarterly*, 89(1), 32–73. doi:10.1111/j.1468-0114.2008.00309.x.
- Pinto, Y., Neville, D. A., Otten, M., Corballis, P. M., Lamme, V. A. F., de Haan, E. H. F., Foschi, N., & Fabri, M. (2017). Split brain: Divided perception but undivided consciousness. *Brain*, 140(5), 1231–1237. doi:10.1093/brain/aww358.
- Poulet, J. F. A., & Petersen, C. C. H. (2008). Internal brain state regulates membrane potential synchrony in barrel cortex of behaving mice. *Nature*, 454(7206), 881–885. doi:10.1038/nature07150.
- Prime, D., Woolfe, M., O’Keefe, S., Rowlands, D., & Dionisio, S. (2020). Quantifying volume conducted potential using stimulation artefact in cortico-cortical evoked potentials. *Journal of Neuroscience Methods*, 337, 108639. doi:10.1016/j.jneumeth.2020.108639.
- Rabaey, J. M., Chandrakasan, A., & Nikolić, B. (2003). *Digital Integrated Circuits: A Design Perspective* (2nd ed.). Prentice Hall.
- Rebollo, B., Telenczuk, B., Navarro-Guzman, A., Destexhe, A., & Sanchez-Vives, M. V. (2021). Modulation of intercolumnar synchronization by

- endogenous electric fields in cerebral cortex. *Science Advances*, 7(10), eabc7772. doi:10.1126/sciadv.abc7772.
- Reinhart, R. M. G., & Nguyen, J. A. (2019). Working memory revived in older adults by synchronizing rhythmic brain circuits. *Nature Neuroscience*, 22(5), 820–827. doi:10.1038/s41593-019-0371-x.
- Rosas, F. E., Geiger, B. C., Luppi, A. I., Seth, A. K., Polani, D., Gastpar, M., & Mediano, P. A. M. (2024). Software in the natural world: A computational approach to hierarchical emergence. *arXiv:2402.09090*. doi:10.48550/arXiv.2402.09090.
- Sabatini, B. L., & Regehr, W. G. (1996). Timing of neurotransmission at fast synapses in the mammalian brain. *Nature*, 384(6605), 170–172. doi:10.1038/384170a0.
- Schneider, M., Broggin, A. C., Dann, B., Tzanou, A., Uran, C., Sheshadri, S., Scherberger, H., & Vinck, M. (2021). A mechanism for inter-areal coherence through communication based on connectivity and oscillatory power. *Neuron*, 109(24), 4050–4067. doi:10.1016/j.neuron.2021.09.037.
- Schreiber, T. (2000). Measuring information transfer. *Physical Review Letters*, 85(2), 461–464. doi:10.1103/PhysRevLett.85.461.
- Seth, A. K. (2026). Conscious artificial intelligence and biological naturalism. *Behavioral and Brain Sciences*, 49, e315. doi:10.1017/S0140525X25000032.
- Tononi, G., & Koch, C. (2015). Consciousness: here, there and everywhere? *Philosophical Transactions of the Royal Society B*, 370(1668), 20140167. doi:10.1098/rstb.2014.0167.
- Trebaul, L., Rudrauf, D., Job, A.-S., Măliia, M. D., Popa, I., Barborica, A., Minotti, L., Mîndruță, I., Kahane, P., & David, O. (2016).

Stimulation artifact correction method for estimation of early cortico-cortical evoked potentials. *Journal of Neuroscience Methods*, 264, 94–102. doi:10.1016/j.jneumeth.2016.03.002.

A Reach: resistance, encoders and simulations

The logarithmic growth is a lower bound for every kernel of bounded range. Group the sites into shells around a site a , so that the coupling joins only sites in the same or adjacent shells. The shells then act as conductors in series, and the effective conductance from a to any site beyond shell n is at most $(\sum_{k < n} 1/C_k)^{-1}$, where C_k is the coupling crossing from shell k to shell $k+1$. On a planar sheet the perimeter of shell k grows linearly in k , so $C_k \leq c(k+1)$ for a constant c set by the kernel’s strength and width, and the effective resistance is at least $\log(n+1)/c$. Widening the kernel enlarges its shells and raises c , which slows the growth and leaves it unbounded. Distance-independent agreement on a sheet therefore needs coupling of unbounded range.

A tail slower than r^{-4} bounds the resistance at every distance. Run a chain of site sets out from a through discs of radius 2^j and back in to b , the two widest discs spanning the separation. The phase difference across a and b telescopes through the discs’ mean phases, and the drop between the means of two consecutive discs is bounded by the coupling energy between them, shared among the pairs of their sites. A disc of radius 2^j on a planar sheet holds at least $\alpha 4^j$ sites, and a kernel $K(r) \gtrsim r^{-(2+\sigma)}$ couples two sites of consecutive discs by at least $\beta 2^{-(2+\sigma)j}$. The number of pairs grows as 16^j and their coupling falls as $2^{-(2+\sigma)j}$, so the drops shrink geometrically with ratio $q = 2^{-(2-\sigma)/2}$, which is below one for $\sigma < 2$, and the effective conductance is at least $(1-q)^2 \alpha^2 \beta / 8$ however far apart a and b lie.

Both results match classical potential theory. Imagine a random walker on the lattice that jumps between sites with probabilities proportional to the kernel, which is the same at every site. The resistance from a site to infinity is infinite exactly when that walker is sure to return to its start, that is, when the walk is recurrent (Doyle & Snell, 1984). Shorting every site outside a disc around a can only lower the resistance from a to a site b beyond the disc, so under recurrence that resistance grows without bound as b recedes. A symmetric kernel of finite second moment, which includes every tail steeper than r^{-4} , gives a recurrent walk in two dimensions (Chung & Fuchs, 1951). The bound above is the matching statement for tails slower than r^{-4} , whose walks escape. The same threshold holds beyond the harmonic approximation: the plane rotator in two dimensions with coupling falling as $r^{-\alpha}$, $2 < \alpha < 4$, has true long-range order at low noise (Kunz & Pfister, 1976), which short-range coupling forbids. The plane-rotator sheet also has a Kosterlitz–Thouless transition (Fröhlich & Spencer, 1981), an order–disorder crossover that the simulations below locate. When natural frequencies differ from site to site, short-range coupling fares worse still: phases then desynchronize in every dimension up to four (Hong et al., 2005).

Two further properties bound the resistance from above. Adding coupling never raises it, and a direct coupling caps it at its reciprocal. Across a finite cortex a wide short-range kernel and a slow tail differ by a slowly growing logarithm, so telling them apart takes a measurement of the kernel’s second moment as well as of its tail.

An encoder that reads the phases of a small neighbourhood relative to the phase at its centre does better than the Lipschitz bound of Section 4. It reports one content at every site of a travelling plane wave or a winding, whatever the gradient, and for a wave plus a departure from it its discrepancy between two sites is at most its Lipschitz constant times the departure’s relative readings there, so it pays for the local roughness of the field and for nothing else. Both bounds read one oscillator population per content

patch. Across centimetres, where two territories describing one quantity need not be neighbours on the sheet, that reading is an assumption, the *phase-to-content assumption*: each territory’s content is read, by encoders of Lipschitz constant L that are equal where territories overlap, from the phase of one oscillator population assigned to it, so that the phase agreement of the two populations bounds the agreement of the two contents. It is a modelling choice about cortex, since the coupling of a field is spatial while two territories describing one quantity need not be neighbours. Under it the coupling’s reach bounds phase disagreement between distant territories directly and content disagreement through it.

Simulated sheets at identical frequencies, of side 64 and 128, with nearest-neighbour, exponential and power-law kernels of tail exponent $\sigma = 1$ and $\sigma = 3$ at a common row sum, follow the harmonic identity in the ordered phase. The measured phase variance is at most 1.15 times $D R_{\text{eff}}$ at $D = 0.1$ and at most 1.26 times at $D = 0.4$, the excess rising with noise as the cosine coupling softens, and a logarithm fits its growth with distance better than a line in 16 of the 16 runs at those noise levels. At $D = 1.2$ the nearest-neighbour sheet disorders: the mean chord discrepancy $\langle 1 - \cos(\theta_a - \theta_b) \rangle$ at the largest separation is 1.00. The exponential and $\sigma = 1$ sheets stay ordered at that noise (0.46 and 0.50), and the $\sigma = 3$ sheet nearly disorders (0.92), so the kernel’s width sets the order–disorder crossover as well as the rate of decay. Exact resistance sums at sheet sides 128, 512 and 1024 separate width from tail. The exponential kernel, with 25 times the nearest-neighbour stiffness, gives resistance at half the side of 0.531, 0.549 and 0.558, still growing; the $\sigma = 1$ tail gives 0.580, 0.577 and 0.577, and the $\sigma = 3$ tail, steeper than r^{-4} , gives 1.198, 1.448 and 1.572, growing as a short-range kernel does.

B Agreement in time for cortex

The longest cortico-cortical fibres run about 15 cm (Nunez & Srinivasan, 2014). Callosal axons conduct at median velocities of 4.9 to 8.8 m/s, and the fastest of them at up to 20 m/s (Caminiti et al., 2009). A chemical synapse adds about 0.5 ms (Katz & Miledi, 1965; Sabatini & Regehr, 1996). The earliest a synaptic signal can cross the cortex is therefore 8 ms, along the fastest axons, and typical arrival takes 18 to 31 ms. For the time scale of a content, take the unconscious integration that precedes a conscious percept, which lasts up to 400 ms (Herzog et al., 2016). Count agreement as reached when disagreement has fallen 10-fold, decaying exponentially. Agreement then arrives in time when the coupling’s relaxation time is at most 160 ms with typical arrival along the slower fibres, and 166 ms along the faster. Over a localized extent of 5 cm the earliest arrival is 3 ms, typical arrival along the faster fibres is 6.2 ms, and the relaxation time may reach 171 ms. Signal travel thus uses only a small part of the time available, and whether agreement arrives in time depends on how fast the coupling settles.

C A single input and noise that changes with the state

A neuron in numbers. Suppose an input ΔV , too small to fire the cell on its own, arrives while the membrane potential is rising through threshold at rate $\dot{V}$. It moves the spike by $\Delta V/\dot{V}$. Voltage noise σ_V jitters the spike by $\sigma_V/\dot{V}$. The shift therefore exceeds the jitter exactly when ΔV exceeds σ_V , whatever the slope, and spike generation adds little noise of its own (Mainen & Sejnowski, 1995). A unitary excitatory potential between cortical pyramidal cells averages 1.3 mV, with a coefficient of variation of 0.52 (Markram et al., 1997), and the membrane noise near threshold is 0.54 mV (Jacobson et al., 2005). Counting the input’s own variability as noise, one such input

shifts the spike by 1.5 times its jitter, a response of 0.55. Every input larger than 0.63 mV passes, up to the input that fires the cell.

Noise whose law changes with the state. When the law of the noise changes with the state, as the variance of a spike count does, a larger change still moves the content at least as much as a smaller one on each side of the state, provided the laws have a monotone likelihood ratio, as exponential families do: such laws cross once along the parameter the region moves. The criterion of Section 3 is then met exactly when a change of one noise amplitude up and one down both reach the fluctuation.

D The occupancy computation

The membrane is an Ornstein–Uhlenbeck process in units of its own stationary standard deviation, with the threshold at zero. After a spike it is refractory and then reset to its mean, and the mean is calibrated so that the stationary rate equals the declared rate. Probability is propagated on a voltage grid through the exact Ornstein–Uhlenbeck step, and a crossing between two sampled times is counted with the Brownian-bridge probability $\exp(-2xy/s^2)$ for sub-threshold samples x and y and step deviation s . Nothing is sampled. From each grid state, the law of the first spike time in the window, with its absence as one more outcome, is propagated in one pass, and the response to a change is the total-variation distance between two such laws. A state passes when every admissible change moves the law by at least the fluctuation, where a change is admissible if it is at least one noise amplitude of the carrier and does not carry the membrane across threshold. The fraction of occupied states that pass is the stationary mass of the passing states. The probability of entering a passing state within a window is computed on the same chain with the passing states made absorbing. A carrier that supplies a share φ of the membrane’s noise variance has noise amplitude

$\sqrt{\varphi}$ in the membrane's units, and its changes are scaled accordingly. The bit is a carrier restored each step to a rail at a declared margin, read after one step of unit Gaussian noise, with occupancy given by that noise around the rail. The same computation is repeated at yardsticks from 0.01 to 0.6 in place of the fluctuation, at every regime and noise share. For correlated carriers, each membrane is the sum of a shared Ornstein–Uhlenbeck component of variance c and a private one of variance $1 - c$, both with the membrane's time constant. Given the shared path the carriers are independent, so the probability that none of K passes within a window is the average of $(1 - h)^K$ over shared paths, where h is one carrier's chance of passing given the path. For each of 256 shared paths from a fixed seed, the membrane's probability is propagated through the exact step with the private noise, its centre moved by the shared path's increment, after a burn-in of five time constants. Halving both the grid spacing and the time step changes the passing fraction by 10%, and the correlation of two carriers' passing by at most 0.012.

Correlation lowers the per-window link. The chance that 100 carriers leave a window of 400 ms without a link is at most 0.0015 at $c = 0.25$, 0.012 at $c = 0.5$ and 0.093 at $c = 0.75$. The shorter window is more sensitive: within 40 ms the link falls to 0.22 at $c = 0.75$, so P's per-window condition is met at the time scale of a percept rather than of its fastest components.

The switching chip is counted in closed form. On the rising half of a transition the node's voltage crosses its threshold at a slew $\dot{V}$, a change ΔV moves the crossing by $\Delta V/\dot{V}$ and the node's noise jitters it by $\sigma/\dot{V}$, so the response is $2\Phi(\Delta V/2\sigma) - 1$ whatever the slew. Every admissible change of at least one noise amplitude then reaches the fluctuation, exactly at one amplitude, so the node passes at the fluctuation and at no higher yardstick. A node that switches in a fraction α of the cycles of a clock of frequency f , with transitions of duration t , spends $\alpha ft/2$ of its time on the rising half of one, and a window of N cycles holds no transition with probability $(1 - \alpha)^N$. A synchronizer with aperture T_W , clock f_C and data changing at

rate f_D samples inside its aperture at the rate $T_W f_C f_D$, so at most a share $T_W f_D$ of its samples do. Each sample is read as one bit, so it passes at no yardstick above $\Phi(1) - 1/2$, and at lower yardsticks that share bounds its passing fraction. A second stage allowed a time S to resolve fails once in $e^{S/\tau}/(T_W f_C f_D)$ on average, with τ the stage’s resolution time constant (Ginosar, 2011); the declared values are that tutorial’s 28 nm example, $\tau = 10$ ps and $f_D = 100$ MHz, with S one cycle.

E Formal results

Each formal result the paper states is a theorem of the Lean 4 development in the project repository, in its four `Phase10` modules. Table 1 maps each result to its declaration and to the section that argues from it. Every theorem rests only on Lean’s foundational axioms `propext`, `Classical.choice` and `Quot.sound`, which the build verifies for every declaration of the development. Small concrete models in `Examples/` meet the hypotheses of the bounds and show where they are attained.

Table 1: The paper’s formal results and the Lean declarations that carry them.

Result	Section	Declaration
A readout with a margin, under continuous micro dynamics, has no graded dependence on any region	3	<code>not_physicallyEnforced_of_margin</code>
The same for any content read from a finite word of thresholded continuous quantities	3	<code>not_physicallyEnforced_of_bits</code>
A thresholded continuous quantity has a margin off its threshold	3	<code>hasMargin_threshold</code>

Result	Section	Declaration
A category read from a carrier depends gradedly on a region only where the carrier does	3	<code>GradedDependence.of_comp</code>
Every function of a bit word, after steps that keep the bits off their thresholds, depends gradedly on no region	3	<code>not_gradedDependence.iterate_of_bits</code>
Shells in series bound the effective conductance	A	<code>conductance_le_shell</code>
Crossing coupling at most $c(k+1)$ gives resistance at least $\log(n+1)/c$	A	<code>conductance_mul_log_le</code>
An averaging chain from a to b bounds the resistance by $2(\sum_k \rho_k)^2$	A	<code>one_le_conductance_of_chain</code>
A tail $K(r) \gtrsim r^{-(2+\sigma)}$ with $\sigma < 2$ bounds the conductance below at every distance	A	<code>le_conductance_of_powerLaw</code>
Adding coupling never raises the effective resistance	A	<code>conductance_mono</code>
A direct coupling caps the resistance at its reciprocal	A	<code>le_conductance_of_edge</code>
A Lipschitz encoder carries mean-square phase agreement to content agreement	4	<code>integral_dist_sq_le</code>
A relative encoder reports one content at every site of a wave	A	<code>relContent_eq_of_wave</code>
A relative encoder's discrepancy is bounded by the departure from a wave	A	<code>dist_relContent_le</code>
Tracking a remote variable requires it to lie in the causal past	4	<code>mem_ball_of_reconstructs</code>

Result	Section	Declaration
Contraction by κ reaches ε by step $\log(d_0/\varepsilon)/\log(1/\kappa)$	4	<code>disc_iterate_le_of_log_le</code>
A discrepancy shrinking by at most ρ per step is above ε before $\log(d_0/\varepsilon)/\log(1/\rho)$	4	<code>lt_disc_iterate_of_lt_log</code>
Graded dependence of a bit word requires some bit exactly on its threshold	5	<code>exists_threshold_of_gradedDependence</code>
A one-region change below $\delta/(LL_v)$ changes no bit	5	<code>bits_eq_of_lipschitz</code>
Under noise, a one-region change below r/L moves the probability of every event of the next content by at most the error rate	3	<code>law_le_of_marginRate</code>
The same for anything computed from noisy bits, with error rate the probability of landing within δ of a threshold	3	<code>bits_law_le_of_lipschitz</code>
The response of a content to one region grows with the scale	3	<code>response_mono</code>
Below r/L the response is at most the error rate	3	<code>response_le_of_marginRate</code>
Graded dependence above the noise floor at every overlap holds exactly when every overlapping pair is one system	3	<code>enforcedAtNoiseFloor_iff</code>
A region whose content has a margin with error rate below the fluctuation is coupled to no region, and is one system only with itself	3	<code>not_coupledAtNoiseFloor_of_record, sameSystem_iff_eq_of_record</code>

Result	Section	Declaration
For any set of such records, two regions are one system exactly when a chain of couplings avoiding every record joins them	3	<code>sameSystem_iff_of_records</code>
A content that depends on a region only through a record's symbol responds to it by at most the record's error rate, so write-back through a record couples nothing	3	<code>shiftResponse_le_of_factors, not_coupledAtNoiseFloor_of_factors</code>
A record that reads two graded regions joins neither, a write-back through it links neither, and a graded reader under linear dynamics with uniform noise joins two uncoupled regions	3	<code>tally_sameSystem_iff, writeBack_not_coupled, reader_sameSystem</code>
Asking graded dependence above the noise floor of every carrier implies asking it of some carriers of each region	3	<code>EnforcedAtNoiseFloor.enforcedByCarriers</code>
A state at which some carriers of each region pass and not every carrier does	3	<code>spare_enforcedByCarriers, spare_not_enforcedAtNoiseFloor</code>
Regions joined by carriers at every overlap are one system	3	<code>EnforcedByCarriers.regionSameSystem</code>
Bits with error rate below the fluctuation: no carrier of any region reaches any other, and each region is a system of its own	3	<code>not_reachesAtNoiseFloor_of_bits, regionSameSystem_iff_eq_of_bits</code>

Result	Section	Declaration
With K independent carriers each passing within a window with probability at least p , two regions go unlinked in the window with probability at most $(1 - p)^K$	3	<code>measure_not_reachesWithin_le</code>
Without independence, two regions go unlinked in the window with probability at most the variance of the number of passing carriers over the square of its mean	3	<code>measure_not_reachesWithin_le_variance</code>
Bits with error rate below the fluctuation couple no pair of regions, so each region is a system of its own	3	<code>sameSystem_iff_eq_of_bits</code>
The same bits fail graded dependence above the noise floor at any overlap of two regions	3	<code>not_enforcedAtNoiseFloor_of_bits</code>
A smaller shift of a Gaussian moves no event more than a larger one	3	<code>gaussian_shift_le</code>
Where the response grows with the change, graded dependence above the noise floor is decided by a change of one noise amplitude	3	<code>gradedAboveNoise_iff_of_mono</code>
Under linear dynamics with Gaussian noise, the response grows with the change, and graded dependence and coupling are decided at one noise amplitude	3	<code>gaussLin_coupledAtNoiseFloor_iff</code>
A shift of a density with a single peak moves no event more than a larger shift, in either direction	3	<code>lawRise_map_add_mono_of_unimodal</code>

Result	Section	Declaration
Under linear dynamics with additive noise whose density has a single peak, graded dependence and coupling are decided at one noise amplitude	3	<code>noisyLin_</code> <code>coupledAtNoiseFloor_iff</code>
Laws with a monotone likelihood ratio, exponential families among them, cross once	C	<code>singleCrossing_of_mlr,</code> <code>singleCrossing_</code> <code>expFamily</code>
Where the next content's laws cross once along a parameter the region moves monotonically, graded dependence above the noise floor is decided by a change of one noise amplitude up and one down	C	<code>gradedAboveNoise_iff_</code> <code>of_crossing</code>